

Type theory and language
From perception to linguistic communication

Robin Cooper

Draft, January 16, 2018
PLEASE QUOTE WITH CARE

Contents

Acknowledgements	v
I From perception and action to grammar	1
1 From perception to intensionality	3
1.1 Perception as type assignment	3
1.2 Modelling type systems in terms of mathematical objects	5
1.3 Situation types	7
1.4 The string theory of events	15
1.5 Doing things with types	21
1.6 Modal type systems	28
1.7 Intensionality: propositions as types	30
1.8 Summary	32
2 Information exchange	33
2.1 Speech events	33
2.2 Signs	38
2.3 Information exchange in dialogue	40
2.4 Resources	59
2.5 Summary	69
3 Grammar	71
3.1 Syntax	75
3.2 Semantics	82
3.3 Building a chart type	97
3.4 Summary	111
II Towards a dialogical view of semantics	113
4 Proper names, salience and accommodation	115
4.1 Montague's PTQ as a semantic benchmark	115
4.2 Montague's treatment of proper names and a sign-based approach	116

4.3	Proper names and communication	119
4.4	Proper names, salience and accommodation	130
4.5	Paderewski	139
4.6	Summary	145
5	Common nouns, intransitive verbs, frames, the Partee puzzle and passengers	147
5.1	Montague's treatment of common nouns and individual concepts	147
5.2	The Partee puzzle	148
5.3	Frames as records	152
5.4	Using frames in a compositional semantics for the Partee puzzle	160
5.5	Definite descriptions as dynamic generalized quantifiers	165
5.6	Individual vs. frame level nouns	174
5.7	Defining a compositional semantics for the Partee puzzle	176
5.8	Passengers and ships	183
5.9	Conclusion	195
6	Modality and intensionality without possible worlds	197
6.1	Possible worlds, modality and intensionality	197
6.2	Modality without possible worlds	200
6.3	Intensionality without possible worlds	214
6.4	Compositional semantics	247
6.5	Conclusion	247
7	Quantification, anaphora and underspecification	249
7.1	The interpretation of unbound pronouns	249
7.2	Quantifier scope and underspecification	250
7.3	Long distance dependencies	259
7.4	Conservativity and dynamic generalized quantifiers	261
7.5	Anaphora	261
7.6	Structured contexts, binding principles and projection	261
7.7	Compositional semantics and incremental processing	261
A	Type theory with records	263
A.1	Underlying set theory	263
A.2	Basic types	264
A.3	Complex types	265
A.4	Function types	267
A.5	List types	269
A.6	Set types	270
A.7	Singleton types	271
A.8	Join types	272
A.9	Meet types	272
A.10	Models and modal systems of types	273

A.11 The type <i>Type</i> and stratification	274
A.12 Records and Record types	276
A.13 Merges of record types	290
A.14 Relabelling of record types	294
A.15 Using records to restrict and specify record types	306
A.16 Strings and regular types	308
B Grammar rules	311
B.1 Universal resources	311
B.2 English resources	321
C Dialogue rules	325
C.1 Universal resources	325
C.2 English resources	330
Bibliography	331

Acknowledgements

I am grateful to many people for discussion which has led to significant changes in this material. Among them are: Ellen Breitholtz, Liz Coppock, Simon Dobnik, Tim Fernando, Jonathan Ginzburg, Staffan Larsson, Bengt Nordström, Aarne Ranta, Kwong-Cheong Wong. None of these people is responsible for what I have done with their ideas and suggestions.

This research was supported in part by the following projects: Records, types and computational dialogue semantics, Vetenskapsrådet, 2002-4879, Library-based grammar engineering, Vetenskapsrådet, 2005-4211, and Semantic analysis of interaction and coordination in dialogue (SAICD), Vetenskapsrådet, 2009-1569.

Part I

From perception and action to grammar

Chapter 1

From perception to intensionality

1.1 Perception as type assignment

Kim is out for a walk in the park and sees a tree. She knows that it is a tree immediately and does not really have to think anything particularly linguistic, such as “Aha, that’s a tree”. As a human being with normal visual perception, Kim is pretty good at recognizing something as a tree when she sees it, provided that it is a fairly standard exemplar, and the conditions are right: for example, there is enough light and she is not too far away or too close. We shall say that Kim’s perception of a certain object, a , as a tree involves the ascription of a type *Tree* to a . In terms of modern type theory (as in Martin-Löf (1984); Nordström *et al.* (1990)), we might say that Kim has made the *judgement* that a is of type *Tree* (in symbols $a : Tree$).

Objects can be of several types. An object a can be of type *Tree* but also of type *Oak* (a subtype of *Tree*, since all objects of type *Oak* are also of type *Tree*) and *Physical Object* (a supertype of *Tree*, since all objects of type *Tree* are of type *Physical Object*). It might also be of an intuitively more complicated type like *Objects Perceived by Kim* which is neither a subtype nor a supertype of *Tree* since not all objects perceived by Kim are trees and not all trees are perceived by Kim.

There is no perception without some kind of judgement with respect to types of the perceived object. When we say that we do not know what an object is, this normally means that we do not have a type for the object which is narrow enough for the purposes at hand. I trip over something in the dark, exclaiming “What’s that?”, but my painful physical interaction with it through my big toe tells me at least that it is a physical object, sufficiently hard and heavy to offer resistance to my toe. The act of perceiving an object is perceiving it *as* something. You cannot perceive something without ascribing some type to it, even if it is a very general type such as *thing* or *entity*.

Recognizing something as a tree may be immediate and not involve conscious reasoning. Recognizing a tree as an aspen, an elm or a tree with Dutch elm disease may involve closer inspection

and some conscious reasoning about the shape of the leaves or the state of the bark. For humans the relating of objects to certain types can be the result of a long chain of reasoning involving a great deal of conscious effort. But whether the perception is immediate and automatic or the result of a conscious reasoning process, from a logical point of view it still seems to involve the ascription of a type to an object.

The kind of types we are talking about here correspond to pretty much any useful way of classifying things and they correspond to what might be called properties in other theories. For example, in the classical approach to formal semantics developed by Montague (1974) and explicated by Dowty *et al.* (1981) among many others, properties are regarded not as types but as functions from possible worlds and times to (the characteristic functions of) sets of entities, that is, the property *tree* would be a function from possible worlds and times to the set of all entities which are trees at that world and time. Montague has types based on a version of Russell's (1903) simple theory of types but they were "abstract" types like *Entity* and *Truth Value* and types of functions based on these types rather than "contentful" types like *Tree*. Type theory for Montague was a way of providing basic mathematical structure to the semantic system in a way that would allow the generation of interpretations of infinitely many natural language expressions in an orderly fashion that would not get into problems with logical paradoxes. The development of type theory which we will undertake here can be regarded as an enrichment of an "abstract" type theory like Montague's with "contentful" types. We want to do this in a way that allows the types to account for content and relate to cognitive processing such as perception. We want our types to have psychological relevance and to correspond to what Gibson (1986) might call *invariants*, that is, aspects that we can perceive to be the same when confronted with similar objects or the same object from a different perspective. In this respect our types are similar to notions developed in situation theory and situation semantics (Barwise and Perry, 1983; Barwise, 1989).

Gibson's notion of attunement is adopted by Barwise and Perry. The idea is that certain organisms are attuned to certain invariants while others are not. Suppose that Kim perceives a cherry tree with flowers and that a bee alights on one of the flowers. One assumes that the bee's experience of the tree is very different from Kim's. It seems unlikely that the bee perceives the tree as a tree in the sense that Kim does and it is not at all obvious that the bee perceives the tree in its totality as an object. Different species are attuned to different types and even within a species different individuals may vary in the types to which they are attuned. This means that our perception is limited by our cognitive apparatus – not a very surprising fact, of course, but philosophically very important. If perception involves the assignment of types to objects and we are only able to perceive in terms of those types to which we are attuned, then as Kant (1781) pointed out we are not actually able to be aware of *das Ding an sich* ("the thing itself"), that is, we are not able to be aware of an object independently of the categories (or types) which are available to us through our cognitive apparatus.

1.2 Modelling type systems in terms of mathematical objects

In order to make our theory precise we are going to create models of the systems we propose as mathematical objects. This represents one of the two main strategies that have been employed in logic to create rigorous theories. The other approach is to create a formal language to describe the objects in the theory and define rigorous rules of inference which explicate the properties of the objects and the relations that hold between them. At a certain level of abstraction the two approaches are doing the same thing – in order to characterize a theory you need to say what objects are involved in the theory, which important properties they have and what relations they enter into. However, the two approaches tend to get associated with two different logical traditions: the *model theoretic* and *proof theoretic* traditions.

The philosophical foundation of type theory (as presented, for example, by Martin-Löf (1984)) is normally seen as related to intuitionism and constructive mathematics. It is, at bottom, a proof-theoretic discipline rather than a model-theoretic one (despite the fact that model theories have been provided for some type theories). However, it seems that many of the ideas in type theory that are important for the analysis of natural language can be adopted into the classical set theoretic framework familiar to linguists from the classical model-theoretic canon of formal semantics starting from Montague (1974).

Your theory is not very interesting if it does not make *predictions*, that is, by making certain assumptions you can infer some conclusions. This gives you one way to test your theory: see what you can conclude from premises that you know or believe to be true and then test whether the conclusion is actually true. If you can show that your theory allows you to predict some conclusion and its negation, then your theory is *inconsistent*, which means that it is not useful as a scientific theory. One way to discover whether a theory is consistent or not is to formulate it very carefully and explicitly so that you can show mathematical properties of the system and any inconsistencies will appear.

From the informal discussion of type theory that we have seen so far it is clear that it should involve two kinds of entity: the types and the objects which are of those types. (Here we use the word “entity” not in the sense that Montague did, that is, basic individuals, but as an informal notion which can include both objects and types.) This means that we should characterize a type theory with two domains: one domain for the objects of the types and another domain for the types to which these objects belong. Thus we see types as theoretical entities in their own right, not, for example, as collections of objects. Diagrammatically we can represent this as in Figure 1.1 where object a is of type T_1 .

A system of basic types consists of a set of types which are *basic* in the sense that they are not analyzed as *complex* objects composed of other objects in the theory. Each of these types is associated with a set of objects, that is, the objects which are of the type, that is the *witnesses* for the type. Thus if T is a type and A is the set of objects associated with T , then a is of type T (in symbols, $a : T$) just in case $a \in A$. We require that any object a which is a witness for a



Figure 1.1: System of basic types

basic type is not itself one of the types in the system. A type may be *empty* in the sense that it is associated with the empty set, that is, there is nothing of that type.

Notice that we are starting with the types and associating sets of objects with them. This means that while there can be types for which there are no witnesses, there cannot be objects which do not belong to a type. This relates back to our claim in Section 1.1 that we cannot perceive an object without assigning a type to it.

Notice also that the sets of objects associated with types may have members in common. Thus it is possible for objects to belong to more than one type. This is important if we want to have basic types *Elm*, *Tree* and *Physical Object* and say that a single object *a* belongs to all three types as discussed in Section 1.1.

An extremely important property of this kind of type system is that there is nothing which prevents two types from being associated with exactly the same set of objects. In standard set theory the notion of set is *extensional*, that is sets are defined by their membership. You cannot have two distinct sets with the same members. The choice of defining types as entities in their own right

rather than as the sets of their witnesses, means that they can be *intensional*, that is, you can have more than one type with the same set of witnesses. This can be important for the analysis of natural language words like *groundhog* and *woodchuck* which (as I have learned from the literature on natural language semantics) are the same animal. In this case one may wish to say that you have two different words which correspond to the same type, rather than two types with the same *extension* (that is, set of witnesses). Such an analysis is less appealing in the case of *unicorn* and *centaur*, both mythical animals corresponding to types which have an empty extension. If types were extensional, there would only be one empty type (just as there is only one empty set in set theory). In the kind of possible world semantics espoused by Montague the distinction between *unicorn* and *centaur* was made by considering their extension not only in the actual world (where both are empty) but also in all possible worlds, since there will be some worlds in which the extensions are not the same. However, this kind of possible worlds analysis of intensionality fails when you have types whose extensions cannot possibly be different. Consider *round square* and *positive number equal to 2 – 5*. The possible worlds analysis cannot distinguish between these since their extensions are both empty no matter which possible world you look at.

Finally, notice that there may be different systems of basic types, possibly with different types and different objects. One way of exploiting this would be to associate different systems with different organisms as discussed in Section 1.1. (Below we will see different uses of this for the analysis of types which model the cognitive system of a single agent.) Thus properly we should say that an object a is of type T with respect to a basic systems of types \mathbf{TYPE}_B , in symbols, $a :_{\mathbf{TYPE}_B} T$. However, we will continue to write $a : T$ in our informal discussion when there is no danger of confusion.

The definition of a system of basic types is made precise in Appendix A.2.

What counts as an object may vary from agent to agent (particularly if agents are of different species). Different agents have what Barwise (1989) would call different *schemes of individuation*. There appears to be a complex relationship between the types that an agent is attuned to and the parts of the world which the agent will perceive as an object. We model this in part by allowing different type systems to have different objects. In addition we will make extensive use in our systems of a basic type *Ind* for “individual” which corresponds to Montague’s notion of “entity”. The type *Ind* might be thought of as modelling a large part of an agent’s scheme of individuation in Barwise’s sense. However, this clearly still leaves a great deal to be explained and we do this in the hope that exploring the nature of the type systems involved will ultimately give us more insight into how individuation is achieved.

1.3 Situation types

Kim continues her walk in the park. She sees a boy playing with a dog and notices that the boy gives the dog a hug. In perceiving this event she is aware that two individuals are involved and that there is a relation holding between them, namely hugging. She also perceives that the boy is

hugging the dog and not the other way around. She sees that a certain action (hugging) is being performed by an agent (the boy) on a patient (the dog). This perception seems more complex than the classification of an individual object as a tree in the sense that it involves two individual participants and a relation between them as well as the roles those two individuals play in the relation. While it is undoubtedly more complex than the simple classification of an object as a tree, we want to say that it is still the assignment of a type to an object. The object is now an event and she classifies the event as a hugging event with the boy as agent and the dog as patient. We shall have complex types which can be assigned to such events.

Complex types are constructed out of other entities in the theory. As we have just seen, cognitive agents, in addition to being able to assign types to individual objects like trees, also perceive the world in terms of states and events where objects have properties and stand in relations to each other – what Davidson (1967) called events and Barwise and Perry (1983) called situations. We introduce types which are constructed from predicates (like ‘hug’) and objects which are arguments to this predicate like a and b . We will represent such a constructed type as $\text{hug}(a,b)$ and we will sometimes call it a *p*type to indicate that it is a type whose main constructor is a predicate. What would an object belonging to such a type be? According to the type-theoretic approach introduced by Martin-Löf it should be an object which constitutes a proof that a is hugging b . For Martin-Löf, who was considering mathematical predicates, such proof objects might be numbers with certain properties, ordered pairs and so on. Ranta (1994) points out that for non-mathematical predicates the objects could be events as conceived by Davidson (1967, 1980). Thus $\text{hug}(a,b)$ can be considered to be an event or a situation type. In some versions of situation theory Barwise (1989); Seligman and Moss (1997), objects (called *infons*) constructed from a relation and its arguments was considered to be one kind of situation type. Thus one view would be that ptypes are playing a similar role in type theory to the role that infons play in situation theory.

What kind of entity are predicates? The notion is made precise in Appendix A.3.1. The important thing about predicates is that they come along with an *arity*. The arity of a predicate tells you what kind of arguments the predicate takes and what order they come in. For us the arity of a predicate will be a sequence of types. The predicate ‘hug’ as discussed above we can think of as a two-place predicate both of whose arguments must be of type *Ind*, that is, an individual. Thus the arity of ‘hug’ will be $\langle \text{Ind}, \text{Ind} \rangle$. The idea is that if you combine a predicate with arguments of the appropriate types in the appropriate order indicated by the arity then you will have a type. Thus if $a : \text{Ind}$ and $b : \text{Ind}$ then $\text{hug}(a,b)$ will be a type, intuitively the type of situation where a hugs b .

It may be desirable to allow some predicates to combine with more than one assortment of argument types. Thus, for example, one might wish to say that the predicate ‘believe’ can combine with two individuals just like ‘hug’ (as in *Kim believes Sam*) or with an individual and a “proposition” (as in *Kim believes that Sam is telling the truth*). Similarly the predicate ‘want’ might be both a two-place predicate for individuals (as in *Kim wants the tree*) or a two-place predicate between individuals and “properties” (as in *Kim wants to own the tree*). We shall have more

to say about “propositions” and “properties” later. For now, we just note that we want to allow for the possibilities that predicates can be *polymorphic* in the sense that there may be more than one sequence of types which characterize the arguments they are allowed to combine with. The sequences need not even be of the same length (consider *Kim walked* and *Kim walked the dog*). We thus allow for the possibility that these pairs of natural language examples can be treated using the same polymorphic predicate. Another possibility, of course, is to say that the English verbs can correspond to different (though related) predicates in the example pairs and not allow this kind of predicate polymorphism in the type theory. We do not take a stand on this issue but merely note that both possibilities are available. If predicates are to be considered polymorphic then the arity of a predicate can be considered to be a set of sequences of types.

Predicates can be considered as functions from sequences of objects matching their arity to types. As such they would be a *dependent type*, that is, an entity which returns a type when provided with an appropriate object or sequence of objects. However, we have not made this explicit in Appendix A.3.1.

A system of complex types (made precise in Appendix A.3.2) adds to a system of basic types a collection of types constructed from a set of predicates with their arities, that is, it adds all the types which you can construct from the predicates by combining them with objects of the types corresponding to their arities according to the types in the rest of the system. The system also assigns a set of objects to all the types thus constructed from predicates. Many of these types will be assigned the empty set. Intuitively, if we have a type $\text{hug}(c,d)$ and there are no situations in which c hugs d then there will be nothing in the extension of $\text{hug}(c,d)$, that is, it will be assigned the empty set in the system of complex types. Notice that the intensionality of our type system becomes very important here. There may be many individuals x and y for which $\text{hug}(x,y)$ is empty but still we would want to say that the types resulting from the combination of ‘hug’ with the various different individuals corresponds to different types of situations. There are thus two important functions in a system of complex types: one, which we call A , which comes from the system of basic types embedded in the system and assigns extensions to basic types and the other, which we call F , which assigns extensions to types constructed from predicates and arguments corresponding to the arity of the predicates. We have chosen the letters A and F because they are used very often in the characterization of models of first order logic. A model for first order logic is often characterized as a pair $\langle A, F \rangle$ where A is the domain and F a function which assigns denotations to the basic expressions (constants and predicates) of the logic. In a slight variation on classical first order logic A may be a sorted domain, that is the domain is not a single set but a set divided into various subsets, corresponding to *sorts*. For us, A characterizes assignments to basic types and thus provides something like a sorted domain in first order model theory. In first order logic F gives us what we need to know to determine the truth of expressions like $\text{hug}(a,b)$ in first order logic. Thus F will assign to the predicate ‘hug’ a set of ordered pairs telling us who hugs whom. Our F also give us the information we need in order to tell who stands in a predicate relation. However, it does this, not by assigning a set of ordered n -tuples to each predicate, but by assigning sets of witnesses (or “proofs”) to each type constructed from a predicate with appropriate arguments. The set of ordered pairs assigned to ‘hug’ by the first

order logic F corresponds to the set of pairs of arguments $\langle x, y \rangle$ for which the F in a complex system of types assigns a non-empty set. For this reason we call the pair $\langle A, F \rangle$ a *model* within the type system, even though it is not technically a model in the sense of model theory for logic. The correspondence will become important below, however.

Kim sees this situation where a (the boy) hugs b (the dog) and perceives it to be of type $\text{hug}(a,b)$. However, there are intuitively other types which she could assign to this situation other than the type of situation where a hugs b which is represented here. For example, a more general type, which would be useful in characterizing all situations where hugging is going on between any individuals, is that of “situation where one individual hugs another individual”. Another type of situation she might use is that of “situation where a boy hugs a dog”. This is a more specific type than “situation where one individual hugs another individual” but still does not tie us down to the specific individuals a and b as $\text{hug}(a,b)$ does.

There are at least two different ways in type theory to approach these more general types. One is to use Σ -types such as (1).

- (1) a. $\Sigma x:\text{Ind}.\Sigma y:\text{Ind}.\text{hug}(x,y)$
 b. $\Sigma x:\text{Boy}.\Sigma y:\text{Dog}.\text{hug}(x,y)$

In general $\Sigma x:T_1.T_2(x)$ will have as witnesses any ordered pair the first member of which is a witness for T_1 and the second member of which is a witness for $T_2(x)$. Thus this type will be non-empty (“true”) just in case there is something a of type T_1 such that there is something of type $T_2(a)$. This means that Σ -types correspond to existential quantification. A witness for (1a) would be $\langle a, \langle b, s \rangle \rangle$ where $a:\text{Ind}$, $b:\text{Ind}$ and $s:\text{hug}(a,b)$. If there is such a witness then some individual hugs another individual and conversely if some individual hugs another individual there will be a witness for this type. Σ -types are exploited for the semantics of natural language by Ranta (1994) among others.

Another approach to these more general types is to use *record types* such as (2).

- (2) a. $\left[\begin{array}{l} x : \text{Ind} \\ y : \text{Ind} \\ c : \text{hug}(x,y) \end{array} \right]$
 b. $\left[\begin{array}{l} x : \text{Boy} \\ y : \text{Dog} \\ c : \text{hug}(x,y) \end{array} \right]$

We make the notion of record type precise in Appendix A.12. Record types consist of sets of fields such as $[x:\text{Ind}]$ and $[c:\text{hug}(x,y)]$. Fields themselves are pairs consisting of a *label* such as

‘x’ or ‘c’ in the first position (before the ‘:’ in our notation) and a type in the second position. You cannot have more than one field with the same label in a record type. The witnesses of record types are *records*. These are also sets of fields, but in this case the fields consist of a label and an object belonging to a type. A record, r , belongs to a record type, T , just in case r contains fields with the same labels as those in T and the objects in the fields in r are of the type with the corresponding label in T . The record may contain additional fields with labels not mentioned in the record type with the restriction there can only be one field within the record with a particular label. Thus both (3a) and (3b) are records of type (2a).

$$\begin{array}{ll}
 (3) \quad \text{a.} & \left[\begin{array}{lcl} x & = & a \\ y & = & b \\ c & = & s \end{array} \right] \quad \text{where } a:Ind, b:Ind \text{ and } s:hug(a,b) \\
 & \text{b.} \quad \left[\begin{array}{lcl} x & = & d \\ y & = & e \\ c & = & s' \\ z & = & f \\ w & = & g \end{array} \right] \quad \text{where } d:Ind, e:Ind, s':hug(d,e) \text{ and } f \text{ and } g \text{ are objects of} \\
 & & \text{some type}
 \end{array}$$

Note that in our notation for records we have ‘=’ between the two elements of the field whereas in record types we have ‘:’. Note also that when we have types constructed from predicates in our record types and the arguments are represented as labels as in (2a) this means that the type is *dependent* on what objects you choose for those labels in the object of the record type. Thus in (3a) the type of the object labelled ‘c’ is $hug(a,b)$ whereas in (3b) the type is $hug(d,e)$. Actually, the notation we are using here for the dependent types is a convenient simplification of what is needed as we explain in Appendix A.12.

Record types and Σ -types are very similar in an important respect. The type (2a) will be non-empty (“true”) just in case there are individuals x and y such that x hugs y . Thus both record types and Σ -types can be used to model existential quantification. In fact record types and Σ -types are so similar that you would probably not want to have both kinds of types in a single system and we will not use Σ -types. We have chosen to use record types for a number of reasons:

fields are unordered The Σ -types in (4) are distinct, although there is an obvious equivalence which holds between them.

$$\begin{array}{ll}
 (4) \quad \text{a.} & \Sigma x:Ind. \Sigma y:Ind. hug(x,y) \\
 & \text{b.} \quad \Sigma y:Ind. \Sigma x:Ind. hug(x,y)
 \end{array}$$

They are not only distinct types but they also have distinct sets of witnesses. The object $\langle a, \langle b, s \rangle \rangle$ will be of type (4a) just in case $\langle b, \langle a, s \rangle \rangle$ is of type (4b). In contrast, since we are regarding record types (and records) as *sets* of fields, (5a,b) are variant notations for the same type.

$$(5) \quad \begin{array}{ll} \text{a.} & \left[\begin{array}{l} x : Ind \\ y : Ind \\ c : \text{hug}(x,y) \end{array} \right] \\ \text{b.} & \left[\begin{array}{l} y : Ind \\ x : Ind \\ c : \text{hug}(x,y) \end{array} \right] \end{array}$$

labels Record types (and their witnesses) include labelled fields which can be used to access “components” of what is being modelled. This is useful, for example, when we want to analyze anaphoric phenomena in language where pronouns and other words refer back to parts of previous meanings in the discourse. They can also be exploited in other cases where we want to refer to “components” of utterances or their meanings as in clarification questions.

discourse representation The labels in record types can play the role of discourse referents in discourse representation structures (DRSs, Kamp and Reyle, 1993) and record types of the kind we are proposing can be used to model DRSs.

dialogue game boards Record types have been exploited to model dialogue game boards or information states (see in particular Ginzburg, 2012).

feature structures Record types can be used to model the kind of feature structures that linguists like to use (as, for example, in linguistic theories like Head Driven Phrase Structure Grammar, HPSG, Sag *et al.*, 2003). Here the labels in record types correspond to attributes in feature structures.

frames Record types can also be used to model something very like the kinds of frames discussed in frame semantics (Fillmore, 1982, 1985; Ruppenhofer *et al.*, 2006). Here the labels in record types correspond to roles (frame elements).

For discussion of some of the various uses to which record types can be put see Cooper (2005). We will take up all of the uses named here as we progress.

Another way of approaching these more general types in type theory is to use *contexts*. In (6) we take *True* to be the type of non-empty types.

$$(6) \quad \text{a. } x : Ind, y : Ind \vdash \text{hug}(x,y) : True$$

b. $x : \text{Boy}, y : \text{Dog} \vdash \text{hug}(x,y) : \text{True}$

(6a,b) mean in a context where x and y are individuals or a boy and a dog respectively the type $\text{hug}(x,y)$ is non-empty. This notation is normally taken to mean universal quantification over the parameters or variables in the context (i.e. sequence of parametric type judgements) to the left of ‘ \vdash ’. Thus they would mean that for any two individuals or pair of a boy and a dog, the first hugs the second. However, we can also devise ways for thinking of existential quantification over the variables of the context, e.g. for some boy, x , and some dog, y , the type $\text{hug}(x,y)$ is non-empty. We can also think of the contexts as being objects belonging to types in our type theory. Records and record types give us a way of doing this. Thus, for example, (2a) models the type of context which might be represented as the sequence of parametric type judgements given in (7).

(7) $x : \text{Ind}, y : \text{Ind}, c : \text{hug}(x,y)$

As in the comparison with Σ -types there is a difference in that the judgements in a standard type theory context are ordered whereas the fields in a record type are unordered. This means that technically (8) is a distinct context from (7) even though there is an obvious equivalence between them.

(8) $y : \text{Ind}, x : \text{Ind}, c : \text{hug}(x,y)$

They correspond to the same record type, however. Since we will use record types to model type theoretic contexts and records to model instantiations of contexts we will not introduce a separate notion of context.

Thus we use record types to replace both the Σ -types and contexts that one often finds in standard versions of type theory.

The introduction of predicates and ptypes raises some new questions. We said above that the arity of ‘hug’ is $\langle \text{Ind}, \text{Ind} \rangle$. However, when we look at (2b) where the types labelled with ‘ x ’ and ‘ y ’ are *Boy* and *Dog* we see that there is nothing explicit here that requires that the two arguments of ‘hug’ are of type *Ind*. One obvious way to achieve this would be to require that *Boy* and *Dog* are subtypes of *Ind*, that is, that any object of type *Boy* is also of type *Ind* and similarly for *Dog*. However, now that we have introduced predicates there is nothing to stop us having two predicates ‘boy’ and ‘dog’ with arity $\langle \text{Ind} \rangle$. Thus we could have the record type (9).

$$(9) \quad \left[\begin{array}{ll} x & : \text{Ind} \\ c_{\text{boy}} & : \text{boy}(x) \\ y & : \text{Ind} \\ c_{\text{dog}} & : \text{dog}(y) \\ c_{\text{hug}} & : \text{hug}(x,y) \end{array} \right]$$

How do we choose between a type like (9) where common nouns like *boy* and *dog* correspond to one-place predicates and a type like (2b) where common nouns correspond to basic types? One advantage is that (9) explicitly represents that the arity of ‘hug’ is fulfilled. Another advantage is that many, and on some analyses possibly all, nouns in natural languages will in more detailed treatments correspond to predicates of more than one argument. Consider, for example, the fact that boys grow into men. The same individual can be a boy at one time and a man at a later time. One way of treating this is to say that ‘boy’ is a predicate of two arguments with arity $\langle \text{Ind}, \text{Time} \rangle$. In fact if we are going to deal with tense and aspect in natural language in this way we will probably want to add time arguments to most if not all of our predicates and thus allow ourselves record types like (10).

$$(10) \quad \left[\begin{array}{ll} \text{e-time} & : \text{Time} \\ x & : \text{Ind} \\ c_{\text{boy}} & : \text{boy}(x, \text{e-time}) \\ y & : \text{Ind} \\ c_{\text{dog}} & : \text{dog}(y, \text{e-time}) \\ c_{\text{hug}} & : \text{hug}(x, y, \text{e-time}) \end{array} \right]$$

where ‘e-time’ stands for “event time”. Here we have required that the times in all the predicate fields be the event time but this is not always the case. Consider (11).

(11) The minister smoked pot in his youth

Here the time of the pot-smoking event most likely precedes the time of the pot-smoking individual being a minister. We will thus use our basic types for basic ontological categories like *individual* and *time* and use predicates for words that occur in natural language. Predicates can be n -ary whereas our types will always be unary. Note that a ptype like $\text{hug}(a, b, t)$ is constructed from a ternary predicate ‘hug’ but the type itself is a unary type of situations. Thus we might have the judgement $s : \text{hug}(a, b, t)$.

Below we will propose an alternative to this treatment of time as an argument. There is, however, another reason for allowing predicates corresponding to nouns to have more than one argument. This is the existence of relational nouns such as *friend* or *daughter*. (See Partee and Borschev, 2012 for recent discussion.)

In this book we will reserve basic types for two kinds of types: (i) those which correspond to intuitively fundamental ontological categories such as individual and (ii) those types which require a recursive definition to characterize the set of their witnesses. The latter is for a technical reason: defining recursive types as, for instance, record types could lead to the types themselves being a non-well-founded set of ordered pairs which contain themselves. We will discuss this more (???) when recursive types become relevant.

1.4 The string theory of events

Kim stands and watches the boy and the dog for a while. They start to play fetch.¹ This is a moderately complex game in that it consists of a number of components which are carried out in a certain order. The boy picks up a stick, attracts the attention of the dog (possibly shouting “Fetch!”), and throws the stick. The dog runs after the stick, picks it up in his mouth and brings it back to the boy. This sequence can be repeated arbitrarily many times. One thing that becomes clear from this is that events do not happen in a single moment but rather they are stretched out over intervals of time, characterized by the sub-events that constitute them. So if we were to have a type of event (that is, a kind of situation) `play_fetch(a,b,c)` where a is a human, b is a dog and c is a stick we can say something about the series of subevents that we have identified. So we might draw an informal diagram something like Figure 1.2.

In an important series of papers including Fernando (2004, 2006, 2008, 2009, 2011), Fernando introduces a finite state approach to event analysis where events are analyzed in terms of finite state automata something like what we have represented in Figure 1.3. Such an automaton will recognize a string of sub-events. The idea is that our perception of complex events can be seen as strings of punctual observations similar to the kind of sampling we are familiar with from audio technology and digitization processing in speech recognition. Thus events can be analyzed as strings of smaller events. What we mean by a string is made precise in Appendix A.16. Any object of any type can be part of a string. Any two objects (including strings themselves), s_1 and s_2 , can be *concatenated* to form a string $s_1 \frown s_2$. An important property of concatenation is *associativity*, that is if we concatenate s_1 with s_2 and then concatenate the result with s_3 we get the same string that we would obtain by concatenating s_2 with s_3 and then concatenating s_1 with the result. In symbols: $(s_1 \frown s_2) \frown s_3 = s_1 \frown (s_2 \frown s_3)$. For this reason we normally write $s_1 \frown s_2 \frown s_3$ (without the parentheses) or simply $s_1 s_2 s_3$ if it is clear from the context that we mean this to be string concatenation. Following Fernando we will use these strings to give us our notion of temporal order.

Although we will present strings in this way, we will model them as records with distinguished labels related to the natural numbers, t_0, t_1, \dots (‘t’ for “time”). The field labelled t_n will correspond to the n th place in the string. Thus a string of objects $a_1 a_2 a_3$ will be the record in (12).

¹[http://en.wikipedia.org/wiki/Fetch_\(game\)](http://en.wikipedia.org/wiki/Fetch_(game)), accessed 10th Oct 2011.

Figure 1.2: `play_fetch(a,b,c)`Figure 1.3: `play_fetch(a,b,c)` as a finite state machine

$$(12) \quad \begin{bmatrix} t_0 = a_1 \\ t_1 = a_2 \\ t_2 = a_3 \end{bmatrix}$$

The concatenation of (12) with the string a_4 , that is, (13a), will be (13b).

$$(13) \quad \begin{array}{l} \text{a. } \begin{bmatrix} t_0 = a_4 \end{bmatrix} \\ \text{b. } \begin{bmatrix} t_0 = a_1 \\ t_1 = a_2 \\ t_2 = a_3 \\ t_3 = a_4 \end{bmatrix} \end{array}$$

We will continue to represent strings for convenience in the traditional way but modelling strings as records will become important when following paths in records down to elements in strings. We will use $s[n]$ to represent the n th element in a string s (where the first element in the string is $s[0]$). But in terms of the record notation this is just a convenient abbreviation for $s.t_n$.

Now let us build further on the types that we have introduced so far to include string types. For any two types, T_1 and T_2 , we can form the type $T_1 \frown T_2$. This is the type of strings $a \frown b$ where $a : T_1$ and $b : T_2$. The concatenation operation on types (just like that on objects) is associative so we do not use parentheses when more than one type is involved, e.g. $T_1 \frown T_2 \frown T_3$.

Let us return to Kim watching the boy, a , playing fetch with the dog, b , using the stick, c . She perceives the event as being of type $\text{play_fetch}(a,b,c)$. But what does it mean to be an event of this type? Given our concatenation types we can build a type which corresponds to most of what we have sketched in Figure 1.2, namely (14).

$$(14) \quad \text{pick_up}(a,c) \frown \text{attract_attention}(a,b) \frown \text{throw}(a,c) \frown \text{run_after}(b,c) \frown \text{pick_up}(b,c) \frown \text{return}(b,c,a)$$

(14) is a type corresponding to everything we have represented in Figure 1.2 except for the arrow which loops back from the end state to the start state. In order to get the loop into the event type we will use a kind of type which introduces a Kleene-+. In standard notations for strings s^+ stands for a string consisting of one or more occurrences of s .² We will adopt this into types by saying that for any type T there is also a type T^+ which is the type of strings of objects of type T containing one or more members. (See Appendix A.16 for a more precise definition.) The type (15) will, then, give us a type corresponding to the complete Figure 1.2 since it will be the type consisting of strings of one or more events of the type (14).

²This notation was introduced by the mathematician Stephen Kleene.

$$(15) \text{ (pick_up}(a,c) \frown \text{attract_attention}(a,b) \frown \text{throw}(a,c) \frown \text{run_after}(b,c) \frown \text{pick_up}(b,c) \frown \text{return}(b,c,a))^+$$

We will complicate (15) slightly by substituting record types for the ptypes as in (16). We do this because we will want to allow for things happening simultaneously and record types will give us a straightforward way of allowing this.

$$(16) \text{ ([e:pick_up}(a,c)] \frown [e:\text{attract_attention}(a,b)] \frown [e:\text{throw}(a,c)] \frown [e:\text{run_after}(b,c)] \frown [e:\text{pick_up}(b,c)] \frown [e:\text{return}(b,c,a)])}^+$$

The label ‘e’ (“event”) occurs in each of the elements of the string type. In this case we will say that ‘e’ labels a *dimension* of events of this type. The ‘e’-dimension can be thought of as the dimension which characterizes what is happening at each stage of the event. If you want to think geometrically, you can think of the event-string as being located in a space of event types (that is, the ptypes).

What happens when Kim perceives an event as being of this type? She makes a series of observations of events, assigning them to types in the string type. Note that the ptypes in each of the types can be further broken down in a similar way. This gives us a whole hierarchy of perceived events which at some point have to bottom out in basic perceptions which are not further analyzed. In order to recognize an event as being of this type Kim does not need to perceive a string of events corresponding to each of the types in the string types. She may, for example, observe the boy waving the stick to attract the dog’s attention, get distracted by a bird flying overhead for a while, and then return to the fetch event at the point where the dog is running back to the boy with the stick. This still enables her to perceive the event as an event of fetch playing because she has seen such events before and learned that such events are of the string type in (16). It suffices for her to observe enough of the elements in the string to distinguish the event from other event types she may have available in her knowledge resources. Suppose, for example, that she has just two event string types available that begin with the picking up of a stick by a human in the company of a dog. One is (16). The other is one that leads to the human beating the dog with the stick. If she only observes the picking up of the stick she cannot be sure whether what she is observing is a game of fetch or a beating. However, as soon as she observes something in the event string which belongs only to the fetch type of string she can reasonably conclude that she is observing an event of the fetch type. She may, of course, be wrong. She may be observing an event of a type which she does not yet have available in her resource of event types, in which case she will need to learn about the new event type and add it to her resources. However, given the resources at her disposal she can make a prediction about the nature of the rest of the event. One could model her prediction making ability in terms of a function which maps a situation (modelled as a record) to a type of predicted situation, for example (17).

$$(17) \quad \lambda r: \left[\begin{array}{l} x:Ind \\ c_{human}:human(x) \\ y:Ind \\ c_{dog}:dog(y) \\ z:Ind \\ c_{stick}:stick(z) \\ e:[e:pick_up(x,z)] \frown [e:attract_attention(x,y)] \end{array} \right] \cdot \left[\begin{array}{l} e:play_fetch(r.x,r.y,r.z) \\ c_{init}:init(r.e,e) \end{array} \right]$$

Here the predicate ‘init’ has arity $\langle String, String \rangle$. The type $init(s_1, s_2)$ is non-empty just in case s_1 is an initial substring of s_2 . We achieve this by defining

If s_1 is a string of length n and s_2 is a string of any length, then $s : init(s_1, s_2)$ iff the length of s_2 is greater than or equal to n and for each i , $0 \leq i < n$, $s_1[i] = s_2[i]$ and $s = s_2$.

That is, if the initial substring condition holds then the second argument to the predicate (and nothing else) is of the ptype.

The kind of function of which (17) is an instance is a function of the general form (18).

$$(18) \quad \lambda a: T_1 . T_2(a)$$

where we use the notation $T_2(a)$ to represent the fact that T_2 depends on a . The nature of this dependence in (17) is seen in the occurrences of r in the body of the function, for example, ‘play_fetch($r.x, r.y, r.z$)’. Such a function maps an object of some type (represented by T_1) to a type (represented by $T_2(a)$). The type that results from an application of this function will depend on what object it is applied to – that is, we have the possibility of obtaining different types from different objects. In type theory such a function is often called a *dependent type*. These functions will play an important role in much of what is to come later in this book. They will show up many times in what appear at first blush to be totally unrelated phenomena. We want to suggest, however, that all of the phenomena we will describe using such functions have their origin in our basic cognitive ability to make predictions on the basis of partial observation of objects and events.

What happens when Kim does not observe enough of the event to be able to predict with any certainty that the complete event will be a game of fetch? One theory would be that she can only make categorical judgements, and that she has to wait until she has seen enough so that

there is only one type that matches in the collection of situation types in her resources. Another theory would be one where she predicts a disjunction of the available matching types when there is more than one that matches. One might refine this theory so that she can choose one of the available types but assign it a probability based on the number of matching types. If n is the number of matching types the probability of any one of them might be $\frac{1}{n}$. This assumes that each of the types is equally likely to be realized. It would be natural to assume, however, that the probability which Kim assigns to any one of the matching types would be dependent on her previous experience. Suppose, for example, that she has seen 100 events of a boy picking up a stick in the company of a dog, 99 of those events led to a game of fetch and only one led to the boy beating the dog. One might then assume that when she now sees the boy pick up the stick she would assign a 99% probability to the type of fetch events and only 1% probability to the boy beating the dog. That is, the probability she assigns to an event of a boy picking up a stick leading to a game of fetch is the result of dividing the number of instances of a game of fetch she has already observed by the sum of the number of instances she has observed of any types whose initial segment involves the picking up of a stick. In more general terms we can compute the probability which an agent A assigns on the basis of a string, ω , of previous observations to a predicted type T_{pr} given an observed type T_{obs} , $P_{A,\omega}(T_{pr} \mid T_{obs})$, in the case where T_{pr} is a member of the set of alternatives which can be predicted from T_{obs} according to A 's resources based on ω , $\text{alt}_{A,\omega}(T_{obs})$, by the following formula:

$$P_{A,\omega}(T_{pr} \mid T_{obs}) = \frac{|\{T_{pr}\}^{A,\omega}|}{\sum_{T_{alt} \in \text{alt}_{A,\omega}(T_{obs})} |\{T_{alt}\}^{A,\omega}|}$$

where $\{T\}^{A,\omega}$ is the set of objects of type T observed by A in ω . If T_{pr} is not a member of $\text{alt}_{A,\omega}(T_{obs})$, that is not one of the alternatives, we say that $P_{A,\omega}(T_{pr} \mid T_{obs}) = 0$.

While this is still a rather naive and simple view of how probabilities might be assigned it is not without interest, as shown by the following points:

Probability distributions It will always provide a probability distribution over sets of alternatives, that is,

$$\sum_{T_{pr} \in \text{alt}_{A,\omega}(T_{obs})} P_{A,\omega}(T_{pr} \mid T_{obs}) = 1$$

Alternatives We have assumed a notion of alternatives based on types of completed events for which the observed event is an initial segment but other notions of alternativeness could be considered and perhaps even combined.

Relativity of probability assignments The notion of probability is both agent and resource relative. It represents the probability which an agent will assign to a type when observing a given situation after a previous string of observations. Two agents may assign different probabilities depending on the resources they have available.

Learning Relevant observations will update the probability distributions an agent will assign to a given set of alternatives since the probability is computed on the basis of previous observations of the alternative types.

Kim is not alone in being able to draw conclusions based on partial observations of an event. The dog can do it too. As soon as the boy has raised the stick and attracted the dog's attention the dog is excitedly snapping at the stick and starting to run in the direction in which the boy seems to be about to throw. The dog also seems to be attuned to string types of events just as Kim is and also able to make predictions on the basis of partial observations. The types to which a dog is attuned will not be the same as those to which humans can be attuned and this can certainly lead to miscommunication between humans and dogs. For example, there may be many reasons why I would go to the place where outdoor clothes are hanging and where the dog's lead is kept. Many times it will be because I am planning to take the dog out for a walk, but not as often as the dog appears to think, judging from the excitement he shows any time I go near the lead. It is difficult to explain to the dog that I am just looking for a receipt that I think I might have left in my coat pocket. But the basic mechanism of being able to assemble types of events into string types of more complex events and make predictions on the basis of these types seems to be common to both humans and dogs and a good number of other animals too. Perhaps simple organisms do not have this ability and can only react to events that have already happened, but not to predicted outcomes.

This basic inferential ability is thus not parasitic on the ability to communicate using a human language. It is, however, an ability which appears to be exploited to a great extent in our use of language as we will see in later chapters. In the remaining sections of this chapter we will look at some aspects of the type theory which seem more likely to correspond to cognitive abilities which only humans have.

1.5 Doing things with types

The boy and the dog have to coordinate and interact in order to create an event of the game of fetch. This involves doing more with types than just making judgements. For example, when the dog observes the situation in which the boy raises the stick, it may not be clear to the dog whether this is part of a fetch-game situation or a stick-beating situation. The dog may be in a situation of entertaining these two types as possibilities prior to making the judgement that the situation is of the fetch type. We will call this act a query as opposed to a judgement. Once the dog has made the judgement that what it has observed so far is an initial segment of a fetch type situation it has to make its own contribution in order to realize the fetch type, that is, it has to run after the stick and bring it back. This involves the creation of a situation of a certain type. Thus creation acts are another kind of act related to types. Creating objects of a given type often has a *de se* (see, for example, Perry, 1979; Lewis, 1979a; Ninan, 2010; Schlenker, 2011) aspect. The dog has to know that it itself must run after the stick in order to make this a situation in which it and the boy are playing fetch. There is something akin to what Perry calls an essential indexical here,

though, of course, the dog does not have indexical linguistic expressions. It is nevertheless part of the basic competence that an agent needs in order to be able to coordinate its action with the rest of the world that it has a primitive sense of self which is distinct from being able to identify an object which has the same properties as itself. We will follow Lewis in modelling *de se* in terms of functional abstraction over the “self”. In our terms this will mean that *de se* type acts involve dependent types.

In standard type theory we have judgements such as $o : T$ “ o is of type T ” and $T \text{ true}$ “there is something of type T ”. We want to enhance this notion of judgement by including a reference to the agent A which makes the judgement, giving judgements such as $o :_A T$ “agent A judges that o is of type T ” and $:_A T$ “agent A judges that there is some object of type T ”. We will call the first of these a *specific* judgement and the second a *non-specific* judgement. Such judgements are one of the three kinds of acts represented in (19) that we want to include in our type act theory.

(19) *Type Acts*

judgements

specific $o :_A T$ “agent A judges object o to be of type T ”

non-specific $:_A T$ “agent A judges that there is some object of type T ”

queries

specific $o :_A T?$ “agent A wonders whether object o is of type T ”

non-specific $:_A T?$ “agent A wonders whether there is some object of type T ”

creations

non-specific $:_A T!$ “agent A creates something of type T ”

Note that creations only come in the non-specific variant. You cannot create an object which already exists.

Creations are also limited in that there are certain types which a given agent is not able to realize as the main actor. Consider for example the event type involved in the fetch game of the dog running after the stick. The human cannot be the main creator of such an event since it is the dog who is the actor. The most the human can do is wait until the dog has carried out the action and we will count this as a creation type act. This will become important when we discuss coordination in the fetch-game below. It is actually important that the human makes this passive contribution to the creation of the event of the dog running after the stick and does not, for example, get the game confused by immediately throwing another stick before the dog has had a chance to retrieve the first stick. There are other cases of event types which require a less passive contribution from an agent other than the main actor. Consider the type of event where the dog returns the stick to the human. The dog is clearly the main actor here but the human has also a role to play in making the event realized. For example, if the human turns her back on the dog and ignores

what is happening or runs away the event type will not be realized despite the dog's best efforts. Other event types, such as lifting a piano, involve more equal collaboration between two or more agents, where it is not intuitively clear that any one of the agents is the main actor. So when we say "agent A creates something of type T " perhaps it would be more accurate to phrase this as "agent A contributes to the creation of something of type T " where A 's contribution might be as little as not realizing any of the other types involved in the game until T has been realized.

De se type acts involve functions which have the agent in its domain and return a type, that is, they are dependent types which, given the agent, will yield a type. We will say that agents are of type *Ind* and that the relevant dependent types, \mathcal{T} , are functions of type $(Ind \rightarrow Type)$. We characterize *de se* type acts in a way parallel to (19), as given in (20).

(20) *De Se Type Acts*

judgements

specific $o :_A \mathcal{T}(A)$ "agent A judges object o to be of type $\mathcal{T}(A)$ "

non-specific $:_A \mathcal{T}(A)$ "agent A judges that there is some object of type $\mathcal{T}(A)$ "

queries

specific $o :_A \mathcal{T}(A)?$ "agent A wonders whether object o is of type $\mathcal{T}(A)$ "

non-specific $:_A \mathcal{T}(A)?$ "agent A wonders whether there is some object of type $\mathcal{T}(A)$ "

creations

non-specific $:_A \mathcal{T}(A)!$ "agent A creates something of type $\mathcal{T}(A)$ "

From the point of view of the type theory *de se* type acts seem more complex than non-*de se* type acts since they involve a dependent rather than a non-dependent type and a functional application of that dependent type to the agent. However, from a cognitive perspective one might expect *de se* type acts to be more basic. Agents which perform type acts using types directly related to themselves are behaving egocentrically and one could regard it as a more advanced level of abstraction to consider types which are independent of the agent. This seems a puzzling way in which our notions of type seem in conflict with our intuitions about cognition.

While these type acts are prelinguistic (we need them to account for the dog's behaviour in the game of fetch) we will try to argue later that they are the basis on which the notion of speech act (Austin, 1962; Searle, 1969) is built. Our notion of using types in query acts seems intuitively related to work on inquisitive semantics (Groenendijk and Roelofsen, 2012) where some propositions (in particular disjunctions) are regarded as inquisitive. However, this will still allow us to make a distinction between questions and assertions in natural language as argued for by Ginzburg (2012).

Let us now apply these notions to the kind of interaction that has to take place between the human and the dog in a game of fetch. First consider in more detail what is actually involved in playing a game of fetch, that is creating an event of type (16). Each agent has to keep track in some way of where they are in the game and in particular what needs to happen next. We analyze this by saying that each agent has an information state which we will model as a record. We need to keep track of the progression of types of information state for an agent during the course of the game. We will refer to the types of information states as gameboards.³ The idea is that as part of the event occurs then the agent's gameboard is updated so that an event of the next type in the string is expected. For now, we will consider gameboards which only place one requirement on information states, namely that there is an agenda which indicates the type of the next move in the game. Thus if the agent is playing fetch and observes an event of the type where the human throws the stick, then, according to (16), the next move in the game will be an event of the type where the dog runs after the stick. If the actor in the next move is the agent herself then the agent will need to create an event of the type of the next move if the game is to progress. If the actor in the next move is the other player in the game, then the agent will need to observe an event and judge it to be of the appropriate type in order for the game to progress. The type of information states, *InfoState*, will be (21a). (In Chapter 2, when we apply these ideas to dialogue, we will see more complex information states.) The type of the initial information state, *InitInfoState*, will be one where the agenda is required to be the empty list.

- (21) a. [agenda : [RecType]]
 b. [agenda=[] : [RecType]]

We can now see the rules of the game corresponding to the type (16) as a set of update functions which indicate for an information state of a given type what type the next information state may belong to if an event of a certain type occurs. These update functions correspond to the transitions in a finite state machine. This is given in (22).

³Our notions of *information state* and *gameboard* are taken from Larsson (2002) and Ginzburg (2012) respectively as well as a great deal of related literature on the gameboard or information state approach to dialogue analysis originating from Ginzburg (1994). We have adapted the notions somewhat to our own purposes and will take this up in more detail in Chapter 2.

$$(25) \quad \{(([\text{agenda}:[\text{RecType}]] \rightarrow (\text{Rec} \rightarrow \text{RecType})) \vee ([\text{agenda}:[\text{RecType}]] \rightarrow \text{RecType}))\}$$

Let us call the type in (25) *GameRules*. Sets of game rules of this type define the rules for specific participants as in (22). In order to characterize the game in general we need to abstract out the roles of the individual participants in the game. This we will do by defining a function from a record containing individuals appropriate to play the roles in the game thus revising (22) to (26).

$$(26) \quad \lambda r^*: \left[\begin{array}{ll} h & : \text{Ind} \\ c_{\text{human}} & : \text{human}(h) \\ d & : \text{Ind} \\ c_{\text{dog}} & : \text{dog}(d) \\ s & : \text{Ind} \\ c_{\text{stick}} & : \text{stick}(s) \end{array} \right] .$$

$$\{ \begin{array}{l} \lambda r: [\text{agenda}=[] : [\text{RecType}]] . \\ \quad [\text{agenda}=[[e:\text{pick_up}(r^*.h, r^*.s)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{pick_up}(r^*.h, r^*.s)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{pick_up}(r^*.h, r^*.s)] . \\ \quad \quad [\text{agenda}=[[e:\text{attract_attention}(r^*.h, r^*.d)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{attract_attention}(r^*.h, r^*.d)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{attract_attention}(r^*.h, r^*.d)] . \\ \quad \quad [\text{agenda}=[[e:\text{throw}(r^*.h, r^*.s)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{throw}(r^*.h, r^*.s)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{throw}(r^*.h, r^*.s)] . \\ \quad \quad [\text{agenda}=[[e:\text{run_after}(r^*.d, r^*.s)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{run_after}(r^*.d, r^*.s)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{run_after}(r^*.d, r^*.s)] . \\ \quad \quad [\text{agenda}=[[e:\text{pick_up}(r^*.d, r^*.s)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{pick_up}(r^*.d, r^*.s)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{pick_up}(r^*.d, r^*.s)] . \\ \quad \quad [\text{agenda}=[[e:\text{return}(r^*.d, r^*.s, r^*.h)]] : [\text{RecType}]], \\ \lambda r: [\text{agenda}=[[e:\text{return}(r^*.d, r^*.s, r^*.h)]] : [\text{RecType}]] \\ \quad \lambda e: [e:\text{return}(r^*.d, r^*.s, r^*.h)] . \\ \quad \quad [\text{agenda}=[] : [\text{RecType}]] \end{array} \}$$

(26) is of type $(\text{Rec} \rightarrow \text{GameRules})$ which we will call *Game*.

Specifying the rules of the game in terms of update functions in this way will not actually getting anything to happen, though. For that we need type acts of the kind we discussed. We link the update functions to type acts by means of *licensing conditions on type acts*. A basic licensing

condition is that an agent can create (or contribute to the creation of) a witness for the first type that occurs on the agenda in its information state. Such a licensing condition is expressed in (27).

- (27) If A is an agent, s_i is A 's current information state,
 $s_i :_A [\text{agenda} = T | R : [RecType]]$, then $:_A T!$ is licensed.

(Here we use the notation $T | R$ to represent a list whose first member is T and whose rest is R . For example, if the list is $[T_1, T_2, T_3]$ then T would correspond to T_1 and R would correspond to $[T_2, T_3]$. See Appendix A.5.)

Update functions of the kind we have discussed are handled by the licensing conditions in (28).

- (28) a. If $f : (T_1 \rightarrow (T_2 \rightarrow Type))$ is an update function, A is an agent, s_i is A 's current information state, $s_i :_A T_i, T_i \sqsubseteq T_1$ (and $s_i : T_1$), then an event $e :_A T_2$ (and $e : T_2$) licenses $s_{i+1} :_A f(s_i)(e)$.
 b. If $f : (T_1 \rightarrow (T_2 \rightarrow Type))$ is an update function, A is an agent, s_i is A 's current information state, $s_i :_A T_i, T_i \sqsubseteq T_1$ (and $s_i : T_1$), $s_{i+1} :_A f(s_i)$ is licensed.

(28a) is for the case where the update function requires an event in order to be triggered and (28b) is for the case where no event is required. There are two variants of licensing conditions which can be considered. One variant is where the licensing conditions rely only on the agent's judgement of information states and events occurring. The other variant is where in addition we require that the information states and events actually are of the types which the agent judges them to be of. (These conditions are represented in parentheses in (28).) In practical terms an agent has to rely on its own judgement, of course, and there is one sense in which any resulting action is licensed even if the agent's judgement was mistaken. There is another stricter sense of license which requires the agent's judgement to be correct. In the real world, though, the only way we have of judging a judgement to be correct is to look at judgements by other agents.

Licensing conditions will regulate the coordination of successfully realized games like fetch. They enable the agents to coordinate their activity when they both have access to the same objects of type *Game* and are both willing to play. The use of the word "license" is important, however. The agents have free will and may choose not to do what is licensed and also may perform acts that are not licensed. We cannot build a theory that will predict exactly what will happen but we can have a theory which tells us what kinds of actions belong to a game. It is up to the agents to decide whether they will play the game or not. At the same time, however, we might regard whatever is licensed at a given point in the game as an obligation. That is, if there is a general obligation to continue a game once you have embarked on it, then whatever type is placed on an agent's agenda as the result of a previous event in the game can be seen as an obligation on the agent to play its part in the creation of an event of that type.

1.6 Modal type systems

Kim continues her walk still thinking about the boy and the dog. She thinks, “Was the boy standing too close to the pond? Suppose he had fallen in. If he had been my son, I wouldn’t have let him play just there.” An important aspect of human cognition is that we are not only able to observe things as they are but also to conceive of alternatives which go beyond the completion of observed events in the way discussed in Section 1.4. We can not only observe objects and perceive them to be of certain types we can also consider possibilities in which they belong to different types and perhaps do not belong to the type we have observed. We have managed to unhook type judgements from direct perception. While the seeds of this ability can be seen in the kind of event perception and prediction discussed above in that it gives us a way to consider types which have not yet been realized, it is at least one step further in cognitive evolution to be able to consider alternative type assignments which do not correspond to completions of events already perceived.

This leads us to construct *modal type systems* with alternative assignments of objects to types.⁴ Figure 1.4 provides an example of a modal system of basic types with two possibilities, one where the extensions of types T_1 and T_2 overlap and another possibility where they do not.



Figure 1.4: Modal system of basic types

The object a is of type T_1 in the first possibility but not in the second possibility. There is an object, b , of type T_1 in the second possibility. b does not exist at all in the first possibility. In the

⁴The term *modal* is taken from modal logic. See Hughes and Cresswell (1968) for a classic introduction. A modern introduction is to be found in Blackburn *et al.* (2001).

figure we just show two possibilities but our general definition in Appendix A.2 allows for there to be any number of possibilities, including infinitely many.

Given this apparatus we define four simple modal notions:

(necessary) equivalence Two types are (necessarily) equivalent just in case the extension of one type is identical with that of the other type in all the possibilities. While the different possibilities may provide different extensions for the types, it will always be the case that in any given possibility the two types will have the same extension.

subtype One type is a subtype of another just in case whatever possibility you look at it is always the case that the extension of the first type is a subset of the extension of the second. We can also say that the first type “entails” the second, that is, any object which is of the first type will also be of the second type, no matter which possibility you are considering.

necessity The notion of necessity we characterize for a type could be glossed as “necessarily realized” or “necessarily instantiated”. A type will be necessary just in case there is something of the type in all the possibilities.

possibility This notion corresponds to “possibly realized” or “possibly instantiated”. A type will be possible just in case there is some possibility according to which it has a non-null extension.

These notions are made precise in Appendix A.2. Note that all of these notions are relativized to the modal system you are considering and the possibilities it offers. We may think of the family of assignments \mathcal{A} as providing a modal base (cf. Kratzer) or alternatives (in the sense of ???). For these kinds of applications we may wish to consider very small families of assignments corresponding to the knowledge we have. Alternatively, we may want to consider strong logical variants of these modal notions where we consider all the logical possibilities, for example, all possible assignments of extensions to types.

So far we have talked about modal systems of basic types. Modal systems of complex types, where we introduce ptypes, create a minor complication. What ptypes that are present in a system depends on what objects there are of the types that are used in the arities of the predicates. Thus if we have some predicate r with arity $\langle Ind, Ind \rangle$ and a possibility where the set assigned to Ind is $\{a, b\}$ then according to that possibility the ptypes formed with r will be $r(a, a)$, $r(a, b)$, $r(b, a)$ and $r(b, b)$. In a possibility where Ind is assigned a different set the set of available ptypes will be different. It is an important feature of type theories with types constructed from predicates that the collection of such types depends on what objects are available as arguments to the predicates. This makes type theory very different from a logical language such as predicate calculus where the notion of well-formedness of syntactic expressions containing predicates is defined independently of what is provided by the model as denotations of arguments to the predicate.

This leads us (in Appendix A.10) to define two variants of each of our modal notions: *restrictive* variants which are only defined for types which exist in all possibilities and *inclusive* variants which require that the modal definition holds for all the possibilities in which the types exist and disregards those in which the types do not exist. For example, a type is *necessary_r* (that is, “restrictively necessary”) just in case the type is available in all possibilities and has a non-empty set of witnesses in all possibilities. It is *necessary_i* (“inclusively necessary”) just in case in all the possibilities in which the type is provided it has a non-empty set of witnesses. It is clear that if a type is *necessary_r* it will also be *necessary_i* but there may be types which are *necessary_i* but not *necessary_r* (if the type is not provided in all possibilities). A similar relationship between the restrictive and inclusive notions holds for all the modal notions we have discussed.

There may be significant classes of modal type systems in which the types available in the different possibilities do not vary. This could be achieved by requiring that the types used in the arities of predicates always have the same witnesses in all the possibilities. This seems feasible if we restrict the types used in predicate arities to basic ontological categories such as individual or time point. It seems reasonable to consider modal systems in which an individual in one possibility will be an individual in any other possibility, for example. It seems reasonable to say that we wish to consider possibilities where, for example, Kim is a man rather than a woman, but not possibilities where Kim is a point in time rather than an individual. However, the notion “basic ontological category” is a slippery one and we do not want to be forced to make commitments about that.

In the definition of a system of complex types in section A.3.2 we call the pair of an assignment to basic types and assignment to ptypes, $\langle A, F \rangle$, a *model* because of its similarity to first order models.⁵ The model provides an interface between the type theoretical system and a domain external to the type theory. The natural domain to relate to the type theory is that of individuals and situations, that is the kind of things we can perceive or at least consider as possibilities. However, we may want to use models which relate to our perceptual apparatus, as in Larsson (2011), rather than directly to the world. This can also be the key for relating the type theory to a dynamically changing world where the models representing our perceived possibilities are not fixed.

1.7 Intensionality: propositions as types

Kim continues to think about the boy and the dog as she walks along. It was fun to see them playing together. They seemed so happy. The boy obviously thought that the dog was a good playmate. Kim is not only able to perceive events as being of certain types. She is able to recall and reflect on these types. She is able to form attitudes towards these types: it was fun that the boy and the dog were playing but a little worrying that they were so close to the pond. This means that the types themselves seem to be arguments to predicates like ‘fun’ and ‘worrying’.

⁵For a more detailed discussion of the relationship between this and first order models as used in the interpretation of first order logic see Cooper (fthc).

This seems to be an important human ability – not only to be able to take part in or observe an event and find it fun or worrying but to be able to reflect independently of the actual occurrence of the event that it or in general similar events are fun or worrying. This is a source of great richness in human cognition in that it enables us to consider situation types independently of their actual instantiation.⁶ This abstraction also enables us to consider what attitudes other individuals might have. For example, Kim believes that the boy thought that the dog was a good playmate. She is able to ascribe this belief to the boy. Furthermore, we are able to reflect on Kim’s state of mind where she has a belief concerning the type of situation where the boy thinks that the dog was a good playmate. And somebody else could consider of us that we have a certain belief about Kim concerning her belief about the boy’s belief. There is in principle no limit to the depth of recursion concerning our attitudes towards types.

We propose to capture this reflective nature of human cognition by making the type theory technically *reflective* in the sense that we allow types themselves to be objects which can belong to other types. In classical model theoretic semantics we think of *believe* as corresponding to a relation between individuals and propositions. In our type theory, however, we are subscribing to the “propositions as types” view which comes to us via Martin-Löf (1984) but has its origins in intuitionistic logic [????]. Propositions are true or false. Types of situations such as $\text{hug}(a,b)$ correspond to propositions in the sense that if they are non-empty then the proposition is true. If there is nothing of this type then it is false. The reasoning is thus that we do not need propositions in our system as separate semantic objects if we already have types. We can use the types to play the role of propositions. To believe a type is to believe it to be non-empty. From the point of view of a type theory for cognition in which we connect types to our basic perceptual ability, this provides a welcome link between our perceptual ability and our ability to entertain propositions (that is, to consider whether they are true or false).

A predicate like ‘believe’ which represents that an individual has an attitude (of belief) to a certain type should thus have an arity which requires its arguments to be an individual and a type. That is, we should be able to construct the type $\text{believe}(c, \text{hug}(a,b))$ corresponding to *c believes that a hugs b*. We thus create *intensional* type systems where types themselves can be treated as objects and belong to types. Care has to be taken in constructing such systems in order to avoid paradoxes. We use a standard technique known as stratification Turner (2005). We start with a basic type system and then add higher order levels of types. Each higher order includes the types of the order immediately below as objects. In each of these higher orders n there will be a type of all types of the order $n - 1$ but there is no ultimate “type of all types” – such a type would have to have itself as an object. This is made precise in Appendix A.11. For more detailed discussion see Cooper (fthc). Figure 1.5 represents an intensional modal type system where we indicate just the initial three orders of an infinite hierarchy of type orders.

⁶This richness also has its downside in that we often become so engaged in our internal cognitive abstraction that it can be difficult to be fully present and conscious of our direct perception of the world – for example, worrying about what might happen in the future rather than enjoying the present.



Figure 1.5: Intensional modal type system

1.8 Summary

Chapter 2

Information exchange

2.1 Speech events

In Chapter 1 we talked about the perception of events such as a boy and a dog playing fetch. We imagined Kim walking through the park and perceiving various kinds of events. Suppose that she meets a friend in the park and they start to have a conversation. A conversation is a kind of event involving language which seems to be uniquely human. The kind of dialogue involved in a conversation enables humans to exchange information in a way that is more complex and more abstracted from currently occurring events than other animals seem capable of. Nevertheless, we will argue that the basic mechanisms of dialogue involve assigning types to events in way that we discussed in Chapter 1. The events involved are *speech events*.

Consider the kind of event type prediction that we considered in Chapter 1. Suppose that Kim sees the boy playing fetch with the dog and the boy is standing close to the lake with his back to it. As the dog runs towards him with the stick he takes a step backwards. “No,” says Kim, seeing that the boy is about to fall in the lake. “Watch out,” she shouts to the boy who takes a step forward just in time and narrowly misses falling in the lake. Her utterance of *no* represents a negative attitude towards a predicted outcome. This kind of negation is discussed briefly in Cooper and Ginzburg (2011a,b) where examples are given of cases where *no* is a response to a completed event and where it is used as an attempt to prevent the predicted outcome. This latter exploits the fact that agents cannot only perceive and classify events according to the types to which they are attuned but can also intervene and prevent a predicted outcome. Kim’s linguistic utterance of *watch out* is used in this way. While Kim is using words of English this is not yet completely linguistic interaction. A dog, sensing danger, will begin to bark and this can have the effect of preventing a predicted outcome. It is a kind of inter-agent communication nevertheless in that it is an intervention in the flow of events which involves predicting and changing the behaviour of another agent. In this sense it is similar to human dialogue, although human dialogue is normally a much more abstract affair, involving predicting and influencing the other agent’s linguistic behaviour and the attitudes and beliefs which the other agent has concerning certain types.

- (1) John: Hello doctor.
 Anon 1: Hello.
 Well Mr [last or full name], what can I [do for you today]₁?
 John: [Er, it's]₁ a wee problem I've had for a ⟨pause⟩ say about a year now.
 Anon 1: Mhm.
 John: It's er my face.
 And my skin.
 I seem to get an awful lot of, it's like
 Anon 1: Aha.
 John: dry flaky skin.
 Anon 1: Yeah.
 John: And I get it on my forehead, [down here]₂
 Anon 1: [I can see]₂

BNC file G43, sentences 1–13

Dialogues themselves are events and, just like other events, can be regarded as strings of smaller events. Consider the dialogue excerpt (1) from the British National Corpus which is the beginning of a consultation between a patient (John) and a doctor (Anon 1).

We might assign the whole dialogue of which this is a part to a genre type for patient doctor consultation.¹ The genre type could be seen as an event type which, like the type for the game of fetch discussed in Chapter 1, can be broken down into a string of subevent types such as greeting (realized here by the exchange *Hello doctor./Hello*), establishing the patient's symptoms (realized here by the remainder of (1)), making a diagnosis, prescribing treatment and so on. Events belonging to these subtypes can be further broken down into strings of turns which further can be broken down into strings of utterances of phrases. In turn phrase utterances are constituted by strings of word utterances which in turn can be regarded as strings of phoneme events. Notice that the temporal relationships between the elements of these strings is more varied than we accounted for in Chapter 1. In dialogue utterances may temporally overlap each other (as indicated in (1) by the notation [...]_n). When we consider adjacent phoneme events in a string overlap becomes the norm (referred to as coarticulation in phonetics). Although we did not take it up in Chapter 1, temporal overlap in event strings is not restricted to speech events. For example, in the game of fetch it is quite often the case that the dog will start running after the stick before the human has finished throwing it. Perceiving temporally overlapping events is part of our basic perceptual apparatus.

¹For a recent discussion of genre in the kind of framework that we are describing see Ginzburg (2010, 2012).

We will work on developing a type for speech events, *SEvent*.² Crucial here is the type of phonological event, *Phon*, that is the type of event where certain speech sounds are produced. A field for events of this type will play a role corresponding to the phonology feature in HPSG (Sag *et al.*, 2003). For simplicity we might assume that *Phon* is an abbreviation for $[e:Word]^+$ that is a non-empty string of events where a word is uttered.³ Here *Word* is the type of event where word forms of the language are uttered. A more accurate proposal might be that *Phon* is $[e:Phoneme]^+$ ⁴ where *Phoneme* is the type of utterance event where a phoneme is uttered. This would still be a simplification and an abstraction from the actual events that are being classified, however. A phoneme type is rather to be regarded as a complex type of acoustic and articulatory event and what we regard as a string of phonemes is in fact a string of events where the phoneme types overlap (corresponding to what is known as *coarticulation* in phonological and phonetic theory). For example, the pronunciation of the phoneme /k/ in “kit” is distinct from its pronunciation in “cat” due to the influence of the following vowel. Suppose that the dimensions of phoneme utterance events are given by place, manner, rounding, voicing and nasal. Then we might represent the type of an utterance of /k/ as

$$(2) \quad \left[\begin{array}{ll} \text{place} & : \textit{Velar} \\ \text{manner} & : \textit{Stop} \\ \text{rounding} & : \textit{NonRound} \\ \text{voicing} & : \textit{NonVoiced} \\ \text{nasal} & : \textit{NonNasal} \end{array} \right]$$

the type of an utterance of /i/ by

$$(3) \quad \left[\begin{array}{ll} \text{place} & : \textit{FrontHigh} \\ \text{manner} & : \textit{Vocalic} \\ \text{rounding} & : \textit{NonRound} \\ \text{voicing} & : \textit{Voiced} \\ \text{nasal} & : \textit{NonNasal} \end{array} \right]$$

and the type of an utterance of /æ/ by

²This type will be different for different languages, dialects, even idiolects. Thus there will be a different type corresponding to what we think of as speech events of English as opposed to speech events of French. Similar remarks can be made about all the linguistic types that we introduce. We will ignore this in our grammatical types in order to avoid proliferation of subscripts.

³If we want to be more grammatically sophisticated we might want to allow silent speech events by allowing empty phonologies, that is, we say that *Phon* is the type $[e:Word]^*$.

⁴Or $[e:Phoneme]^*$.

$$(4) \left[\begin{array}{ll} \text{place} & : \text{BackHigh} \\ \text{manner} & : \text{Vocalic} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{Voiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right]$$

Naively, one might think that the type of the phoneme string /ki/ would be

$$(5) \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{Velar} \\ \text{manner} & : \text{Stop} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{NonVoiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right] \sim \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{FrontHigh} \\ \text{manner} & : \text{Vocalic} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{Voiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right]$$

However, the place of articulation of the /k/ will be influenced by the place of articulation of the following vowel as in (6)

$$(6) \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{Palatal} \\ \text{manner} & : \text{Stop} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{NonVoiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right] \sim \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{FrontHigh} \\ \text{manner} & : \text{Vocalic} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{Voiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right]$$

In addition to this the voice onset associated with the vowel will normally begin before the articulation of the stop is complete as in (7).

$$(7) \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{Palatal} \\ \text{manner} & : \text{Stop} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{NonVoiced} \sim \text{Voiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right] \sim \left[\begin{array}{ll} e & : \left[\begin{array}{ll} \text{place} & : \text{FrontHigh} \\ \text{manner} & : \text{Vocalic} \\ \text{rounding} & : \text{NonRound} \\ \text{voicing} & : \text{Voiced} \\ \text{nasal} & : \text{NonNasal} \end{array} \right] \end{array} \right]$$

This is not meant to be a serious phonological analysis. We include it here to show how the well-studied phenomenon of coarticulation could be included in the general framework and to show that the notion of overlapping events which we will need later for semantics and dialogue

is the same notion that is needed for phonology. We have no more to say about phonology and will limit our analysis of phonological events to strings of words.

We will keep the simplifying assumption that phonology is a string of words here (that is, that *Phon* is *Word*⁺ and we do not say more about what is of type *Word*) as we do not aim to give a detail account of phonology. Thus a proposal for the type *SEvent* might be (8).

$$(8) \quad \left[\begin{array}{ll} e & : \text{Phon} \end{array} \right]$$

To this we might usefully add the speech location as in (9).

$$(9) \quad \left[\begin{array}{ll} e\text{-loc} & : \text{Loc} \\ e & : \text{Phon} \\ c_{\text{loc}} & : \text{loc}(e, e\text{-loc}) \end{array} \right]$$

We will take *Loc* to be the type of regions in three dimensional space without specifying more detail. Further if *e* is an event and *l* a location we will say that the type *loc(e, l)* is non-empty just in case *e* is located at *l*, again without saying exactly what that means for now.

It might seem natural to add roles of speaker and audience, given what we know about speech act theory (Searle, 1969). Thus we might consider *SEvent* to be the type in (10).

$$(10) \quad \left[\begin{array}{ll} e\text{-loc} & : \text{Loc} \\ \text{sp} & : \text{Ind} \\ \text{au} & : \text{Ind} \\ e & : \text{Phon} \\ c_{\text{loc}} & : \text{loc}(e, e\text{-loc}) \\ c_{\text{sp}} & : \text{speaker}(e, \text{sp}) \\ c_{\text{au}} & : \text{audience}(e, \text{au}) \end{array} \right]$$

However, while many speech events may be considered to be of this type, not all will. Of course, some speech events are not addressed to any audience. An example might be an exclamation uttered after hitting one's thumb with a hammer. Longer speech events like dialogues will not have a single speaker or audience. Even shorter chunks corresponding perhaps to single speech acts do not always have a single speaker or audience. For example, consider split utterances as discussed by Purver *et al.* (2010) who give the example (11).

- (11) A: I heard a shout. Did you
B: Burn myself? No, luckily.

Here we probably want to consider the utterance of *Did you . . . burn myself?* as a speech event on which *A* and *B* collaborate. Otherwise it might be hard to explain how *you* can be interpreted as the subject of *burn*. We have a single predication split across two speakers. Similarly, speakers can address different audiences within the same predicate structure as in (12).

(12) You [pointing] work with you [pointing] and you [pointing] work on your own.

Nevertheless, we might consider that the majority of speech events would belong to the more restricted type (10).

Because we have taken a neo-Davidsonian (Dowty, 1989) approach to the more restricted speech-event types, where the objects playing the various roles in the speech events are introduced in separate fields, both (9) and (10) are subtypes of (8). We will use *SEvent* below to represent the most specific of the types, (10), while bearing in mind that many events we may want to call “speech events” will belong only to more general types such as (9) and (8).

2.2 Signs

We interpret many speech events as being associated with a semantic content, but not all. When John in (1) says *It’s a wee problem I’ve had for a, say, about a year now* he is using the speech event to refer to another situation - a situation in which he has dry skin for a period of a year. This is what Barwise and Perry (1983) would refer to as the *described situation* which is distinct from the speech situation. In contrast the doctor’s utterance of *Hello* in (1) does not tell us anything about a described situation external to the current conversation, although it does give us information about where we are in the conversation (the beginning) and indicate that the doctor is paying attention. We shall say that the former utterance with a type of described situation and call this the *content* of the utterance. A situation type is an appropriate content for a declarative sentence used to make an assertion.⁵ The contents of phrases within such a sentence such as *a wee problem* or *about a year* will be objects which can be combined to produce such a type. The contents of other kinds of speech acts, for example, associated with questions like the doctor’s utterance of *what can I do for you today?* will be objects based on situation types, in the case of this question a function which maps actions to a situation type. (See Ginzburg (2012) for a discussion of the kind of treatment of questions we have in mind.)

We can think of this association of content with a speech event in similar terms to prediction of event completion discussed in Section 1.4 of Chapter 1. At least in the case of declarative assertions it is a mapping from an observation of a situation to a type of situation. In the case of the event completion the result of the mapping was a type for the completion of the event so far

⁵We will discuss later that alternative proposed in Ginzburg (2012) that it should be a pairing of a situation type with a situation, that is an Austinian proposition as introduced by Barwise and Perry (1983) based on Austin (1961).

observed. In the case of the speech event we are relating the observation to a type of situation which is entirely distinct from the speech event. The association is less immediate and more abstract but the underlying mechanism, associating the observation of a situation of a given type with another type and drawing the conclusion that the second type must be non-empty, is the same. We could represent the association by a function of the form (13), corresponding to (18) in Chapter 1.

$$(13) \quad \lambda s : T_{SpEv} . T_{Cnt}(s)$$

This represents a mapping from a speech event s of a given type T_{SpEv} to a type T_{Cnt} which is the content of the speech event. The type T_{Cnt} can depend on s (for example, the type of the described situation may require that the described situation be related to the utterance situation temporally or spatially).

de Saussure (1916) called the association between speech and content a sign and this notion has been taken up in modern linguistics in Head Driven Phrase Structure Grammar (HPSG, Sag *et al.*, 2003). In HPSG a sign is regarded technically as a feature structure and our notion of record type corresponds to a feature structure. One way in which our type system differs from HPSG is that we have both records and record types where HPSG has just feature structures. We will consider a sign to be a record representing a pairing of a speech event and a type representing the content. One advantage of considering a sign as a record rather than a function as in (13) is that there is no directionality in a record as there is in a function. Thus the record can be associated with either interpretation (from speech event to content) or generation (from content to speech event). We can make a straightforward relationship between a function such as (13) and a record type (14).

$$(14) \quad \left[\begin{array}{ll} \text{s-event} & : T_{SpEv} \\ \text{cnt}=T_{Cnt} & : Cnt \end{array} \right]$$

(14) is a type of signs. Notice that the ‘cnt’-field in (14) is a manifest field corresponding to the fact that the function in (13) returns the type T_{Cnt} , not an object of type T_{Cnt} . This means that the ‘cnt’ field in (14) requires that the type itself is in the ‘cnt’ field in a record of the type, that is, in the sign. The type Cnt is the type of contents. For the moment we will say that Cnt is the type *RecType*, that is, that contents are record types. This is because, for the moment, we will restrict our attention to declarative sentences. When we come to look at constituents of sentences and speech acts other than assertions we will need to expand Cnt to include other kinds of entities as well. Restricting our attention first to complete declarative sentences is similar to starting with propositional logic before moving on to more complex analysis. The type *Sign* of signs in general is given in (15).

$$(15) \quad \left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{cnt} & : \text{Cnt} \end{array} \right]$$

A record of this type, a sign, will pair a speech event with a content. We will refine our definition of *Sign* as we progress.

2.3 Information exchange in dialogue

We start by considering simple dialogues such as (16) which might occur between two people one of whom is instructing the other about simple facts or between a user and a system where the user is adding simple facts to a database using a natural language interface.

- (16) User: Dudamel is a conductor
 System: Aha
 User: Beethoven is a composer
 System: OK

The job of the dialogue partner identified as “System” is to record the facts in memory and confirm to the dialogue partner identified as “User” that this has happened. It seems straightforward to think of the user’s utterances in (16) as corresponding to signs as described in Section 2.2. For example, the user’s first utterance could be regarded as corresponding to a sign of the type in (17).

$$(17) \quad \left[\begin{array}{l} \text{s-event:} \left[\begin{array}{ll} e & : \text{“Dudamel is a conductor”} \end{array} \right] \\ \text{cnt=} \left[\begin{array}{ll} e & : \text{conductor(dudamel)} \\ c_{\text{tns}} & : \text{final_align}(\uparrow \text{s-event.e}, e) \end{array} \right] \end{array} \right] : \text{RecType}$$

Here “Dudamel is a conductor” is a convenient abbreviation for (18).

$$(18) \quad [e:\text{“Dudamel”}] \cap [e:\text{“is”}] \cap [e:\text{“a”}] \cap [e:\text{“conductor”}]$$

where for any word w , “ w ” is the type of event where w is uttered. “Dudamel is a conductor” is thus a type of string of events of word utterances and is thus a subtype of *Phon*, given our assumptions in Section 2.1.

The content is that Dudamel is a conductor and that his being a conductor is aligned with the speech event in that the speech event occurs simultaneously with the end of the event of Dudamel being a conductor. This is not to say that Dudamel will not continue to be a conductor after the

speech event but rather to say that we are aligning the speech event with what has happened so far up to and including the speech event. (The simple present in English in contrast to the present progressive and the simple present in many other languages seems to require this.) How do we align events? We use the technique developed by Fernando (see, for example, Fernando, 2008) of creating a single event which includes both events as a part. We will exploit our record technology to keep track of the separate events in the larger event and to achieve something corresponding to what Fernando calls *superposition*. We might require that the event which is the coordination of the two events of type “Dudamel is a conductor” and ‘conductor(dudamel)’ is of the type in (19).

$$(19) \left[\begin{array}{l} e_1 : \text{“Dudamel is a conductor”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right]$$

Another option is to require that the coordinated event type explicitly allow for there to be events of the type ‘conductor(dudamel)’ prior to the utterance as in (20).

$$(20) \left[e : \text{conductor(dudamel)} \right] * \left[e : \left[\begin{array}{l} e_1 : \text{“Dudamel is a conductor”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right] \right]$$

Here the dimension ‘e’ splits into two subdimension ‘e.e₁’ and ‘e.e₂’. If we wish to be explicit about the fact that a situation of type “Dudamel is a conductor” is a string of word utterances we can give the more detail type in (21).

$$(21) \left[e : \text{conductor(dudamel)} \right] * \left[e : \left[\begin{array}{l} e_1 : \text{“Dudamel”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right] \right] \wedge \\ \left[e : \left[\begin{array}{l} e_1 : \text{“is”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right] \right] \wedge \\ \left[e : \left[\begin{array}{l} e_1 : \text{“a”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right] \right] \wedge \\ \left[e : \left[\begin{array}{l} e_1 : \text{“conductor”} \\ e_2 : \text{conductor(dudamel)} \end{array} \right] \right]$$

This explicitly requires that Dudamel is a conductor during the utterance of each individual word. Both the types (20–21) are facilitated by the fact that ‘conductor(dudamel)’ is a *state-type*, that is, given a situation $e : \text{conductor(dudamel)}$ we can regard it as a string of events of type $[e : \text{conductor(dudamel)}]^+$. We will return to aspectual types other than state below. The predicate ‘final_align’ in (17) requires alignment of the speech event and the described event in the way we have exemplified in (20) and (21). The definition of what counts as a witness for $\text{final_align}(e_1, e_2)$ given in Appendix A.16 requires that e is of this type just in case e is an event

where e_1 is aligned with a final segment of e_2 , that is in e there is a split in dimension in the final segment as illustrated in (21). The notation ‘ \uparrow ’ in (17) indicates that the path ‘ x ’ is not to be found in the local record type which is required to be the value of ‘ cnt ’ but in the next higher record type with the fields ‘ s-event ’ and ‘ cnt ’. This notation is explained in Appendix A.12.

This sign type (17) seems to give us what we need in order to explain how an utterance of *Dudamel is a conductor* can convey the information that Dudamel is a conductor. If both dialogue participants have this sign type among their resources then the User knows that in order to convey this content she has to make an utterance which witnesses the appropriate speech event type. The System knows that on observing a speech event of this type the corresponding content should be recorded.

Things are not as straightforward, however, for the acknowledgements *Aha* and *OK* expressed by the system. It is not obvious whether these utterances are to be regarded as signs at all. Certainly a speech event is involved but one might question what content they have. One suggestion would be that the content of *Aha* uttered after an assertion by the other dialogue partner would be the same as the content of that assertion. Thus the system is expressing the same content as the user. This may or may not be true. But such an analysis seems to be missing a central point about what is going on in this dialogue, namely that the user is making an assertion and the system is acknowledging that the content has been accepted and duly processed. In order to account for this kind of fact Ginzburg in a large body of work has developed the notion of a dialogue gameboard, most recently formulated in terms of TTR in Ginzburg (2012); Ginzburg and Fernández (2010). In the computational dialogue systems literature this has given rise to the Information State Update (ISU) approach (Larsson and Traum, 2001; Larsson, 2002) which is also described in Ginzburg and Fernández (2010). In Chapter 1 we introduced the notion of an information state as a record containing a field labelled ‘agenda’ and used the word “gameboard” to refer to a type of information state. Our aim there was to show that the kind of gameboard analysis introduced for dialogue in this literature is also important for the coordination of joint action by agents in general. The gameboards that have been used for dialogue analysis have a number of fields in addition to the agenda. Each dialogue participant will have among their resources a record type, their dialogue gameboard which represents their understanding of (what Larsson call their take on) their current information state. Following Larsson (2002) we place information which the agent assumes to be common with its interlocutors under the label ‘shared’ in the gameboard and also have a field with the label ‘private’ representing information about the state of the dialogue which is not shared with other dialogue participants. This will include, for example, plans for what should be said next represented in the agenda. In Figure 2.1 we give a schematic view of the gameboards associated with each of the dialogue participants in the first exchange in (16).

This assumes ideal communication. There is lots that could go wrong which could have the consequence that the two agents become misaligned and an important part of this framework is to provide a basis for the description of miscommunication as well as communication. (See Ginzburg (2012) for more discussion of this.)



Figure 2.1: Dialogue management:“Dudamel is a conductor”

We treat the dialogue information states represented by the square boxes as records as in (22).

$$(22) \quad \left[\begin{array}{lcl} \text{private} & = & \left[\begin{array}{lcl} \text{agenda} & = & \text{AGENDA} \end{array} \right] \\ \text{shared} & = & \left[\begin{array}{lcl} \text{latest-utterance} & = & \text{L-UTT} \\ \text{commitments} & = & \text{COMM} \end{array} \right] \end{array} \right]$$

What kinds of objects should *AGENDA*, *L-UTT* and *COMM* be? They will be defined with respect to the agent who owns the information state which, for convenience, we will refer to as *SELF*. We will see as we proceed with the discussion below that *SELF* is related to the notion of *de se* type act discussed in Chapter 1.5.

We will say that *AGENDA* is a list of dialogue move types, that is, the types of dialogue moves that *SELF* plans to realize by means of a creation type act. Recall from Chapter 1 that this does not necessarily mean that *SELF* is the main actor in the event realizing the move type. It can for example be a type of move to be carried out by an interlocutor which *SELF* should wait for. This will give us a mechanism for handling basic turn-taking in dialogue. (See Sacks *et al.*, 1974 for the classic work on turn-taking.) We will define a dependent type *Move* such that for any agent *A*, *Move*(*A*) is the type of dialogue moves in which *A* is involved. For now we will say that there are two ways in which an agent can be involved in a dialogue act: as speaker (or performer) or as hearer (part of the audience to whom the dialogue act is addressed).⁶ Performing a dialogue move is a *de se* type act of creation as discussed in Chapter 1.5. Being the hearer or audience of a move type involves a *de se* type act of judgement as discussed there. *AGENDA* should thus be a list of move types depending on *SELF* (that is, subtypes of *Move*(*SELF*)). We introduce a dependent type, *MoveType*, such that for any *a:Ind*, *T:MoveType*(*a*) iff *T* \sqsubseteq *Move*(*a*). *AGENDA* is thus a list of move types and will have type [*MoveType*(*SELF*)]. For any type *T*, [*T*] is the type of lists all of whose members are of type *T* (see Appendix A.5). We will come back to the details of *Move* below.

L-UTT should tell us what move (or moves⁷) has just been carried out. But we will need more information than this. We will need information about what (the agent *SELF* thinks) was actually said. For this we will use a chart, i.e. a set of edges between vertices representing hypotheses about parts of the utterance, that is, sign types associated with parts of the utterance. The move should be predictable from the chart by a process of move-interpretation for which we will use the predicate ‘m-interp’. Thus *L-UTT* should itself be of the type (23).

$$(23) \quad \left[\begin{array}{lcl} \text{move} & : & \text{Move}(\text{SELF}) \\ \text{chart} & : & \text{Chart} \\ \text{e} & : & \text{m-interp}(\text{chart}, \text{move}) \end{array} \right]$$

⁶A third way of being involved in a dialogue act which we will not take account of here is as an overhearer.

⁷See Larsson (2002) for a proposal where dialogue contributions involve several moves. For now we will make the simplifying assumption that utterances are associated with a single move.

The type *Chart* we will say more about in Chapter 3.

The commitments field has normally been considered as a set of facts or propositions (Ginzburg, 2012; Larsson, 2002). Here we will treat them as a single record type, i.e. a member of the type *RecType*. Using a single type will make it more straightforward to deal with issues like consistency and anaphora [????].

Thus information states can belong to the type (24).

$$(24) \quad \left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda:} [MoveType(SELF)] \\ \text{latest-utterance:} \left[\begin{array}{l} \text{move:} Move(SELF) \\ \text{chart:} Chart \\ \text{e:m-interp(chart,move)} \end{array} \right] \\ \text{commitments:} RecType \end{array} \right] \right]$$

(Here, by convention the labels ‘chart’ and ‘move’ in $[e:m\text{-interp}(\text{chart}, \text{move})]$ refer to the path down to the minimal record in which the e-field occurs, that is ‘shared.latest-utterance.chart’ and ‘shared.latest-utterance.move’ respectively.)

(24) is, however, not quite general enough. It requires that there always will be a latest utterance. At the beginning of a dialogue this will not be the case and we need a way of representing that there is no previous utterance. We will use a type whose only witness is the empty record for this. Records, it will be recalled, are sets of ordered pairs (see Appendix A.12). This will include the empty set, \emptyset which could also be notated as ‘[]’ if we are thinking of the empty set as the empty record. However, this latter notation is confusing since it could also be used to represent the empty record type, that is the type that does not place any constraints on which records it has as witnesses, that is, the type of all records which we represent as *Rec* in order to avoid confusion. The type of the empty record could be constructed as the singleton type Rec_{\emptyset} (or if you are using the bracket notation for the empty record, $Rec_{[]}$). In order to avoid notational confusion we will use *ERec* to represent the type whose only witness is the empty record, that is, the empty set. Thus (25) will hold.

$$(25) \quad a : ERec \text{ iff } a = \emptyset$$

At the beginning of a dialogue there will not be any shared commitments either. Therefore, it will be natural to use *Rec* for the commitments at the beginning of a dialogue. *Rec* is the type of all records. If we think of records as modelling situations then a commitment represented by *Rec* is a commitment to the existence of a situation but not to a situation of any particular type. Thus it corresponds to “there is a situation” or “the world is not empty”. It plays a similar role in our theory to the set of all possible worlds in a system based on possible worlds. It represents

a state where no constraints have been placed on the nature of the world. The adjustment we need to make to (24) in order to include dialogue initial information states is to the shared.latest-utterance field as in (26). The ‘commitments’-field does not need to be adjusted as the type *Rec* is one of the witnesses of *RecType* (see Appendix A.12).

$$(26) \quad \left[\begin{array}{l} \text{private:} [\text{agenda:} [\text{MoveType}(\text{SELF})]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move:} \text{Move}(\text{SELF}) \\ \text{chart:} \text{Chart} \\ \text{e:m-interp}(\text{chart}, \text{move}) \end{array} \right] \vee \text{ERec} \\ \text{commitments:} \text{RecType} \end{array} \right] \end{array} \right]$$

(26) uses a join type (Appendix A.8). For any two types T_1 and T_2 you can form the join (or disjunction) $T_1 \vee T_2$. $a : T_1 \vee T_2$ just in case either $a : T_1$ or $a : T_2$.

We will use notation including ‘*SELF*’ as in (26) to represent types which are derived from dependent types by applying them to the argument represented by *SELF*. This notational convention will save us a good deal of complication in presentation and it is always possible to recover the dependent type from which the type is derived by creating a function which maps an individual to the appropriate type. Thus in the case of (26) the dependent type would be (27).

$$(27) \quad \lambda a: \text{Ind} . \left[\begin{array}{l} \text{private:} [\text{agenda:} [\text{MoveType}(a)]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move:} \text{Move}(a) \\ \text{chart:} \text{Chart} \\ \text{e:m-interp}(\text{chart}, \text{move}) \end{array} \right] \vee \text{ERec} \\ \text{commitments:} \text{RecType} \end{array} \right] \end{array} \right]$$

[???? This needs revising in order to include moves by agents other than *SELF*!]

Dialogue moves are a type of event in which an actor (normally speaker) is related to an intended audience, an illocutionary force (such as ‘assert’) and a content (that is, for our present purposes, a record type such as $[\text{e:conductor}(\text{Dudamel})]$).

We will take dialogue moves to be a pairing of speech acts and content. The type of speech acts (*SpeechAct*) will be taken to be a subtype of the type of speech events (*SEvent*) as defined in (10) on p. 37. In particular this will mean that there is a field in a speech act for the speaker (labelled by ‘sp’) and another for the audience (labelled by ‘au’). More specifically we will take the type *Move(a)* to be an abbreviation for (28).

$$(28) \quad [\text{e:SpeechAct}] \wedge ([\text{e:}[\text{sp}=a:\text{Ind}]] \vee [\text{e:}[\text{au}=a:\text{Ind}]]) \wedge \text{MoveContent}$$

The type in (28) is a meet type (Appendix A.9).⁸ If T_1 and T_2 are types, then an object a is of type $T_1 \wedge T_2$ just in case $a : T_1$ and $a : T_2$.

Note that (28) requires a to be either the speaker or the audience of the speech act and does not rule out the possibility that a is both speaker and audience (i.e. a is talking to herself).

We will not attempt a complete inventory of speech act types here. Preliminarily, we could define *SpeechAct* to be

$$\textit{Assertion} \vee \textit{Query} \vee \textit{Command} \vee \textit{Acknowledgement},$$

that is, a join type (Appendix A.8) of all the available speech act types.⁹ Something will be of this type just in case it is of at least one of the types of the join. Each of the speech act types are subtypes of *SEvent* and can be defined as in (29).

$$\begin{array}{lll} (29) & \textit{Assertion} & - \textit{SEvent} \wedge \left[\begin{array}{l} \text{e:Phon} \\ \text{c}_{\text{illoc}}:\text{assertion}(\text{e}) \end{array} \right] \\ & \textit{Query} & - \textit{SEvent} \wedge \left[\begin{array}{l} \text{e:Phon} \\ \text{c}_{\text{illoc}}:\text{query}(\text{e}) \end{array} \right] \\ & \textit{Command} & - \textit{SEvent} \wedge \left[\begin{array}{l} \text{e:Phon} \\ \text{c}_{\text{illoc}}:\text{command}(\text{e}) \end{array} \right] \\ & \textit{Acknowledgement} & - \textit{SEvent} \wedge \left[\begin{array}{l} \text{e:Phon} \\ \text{c}_{\text{illoc}}:\text{acknowledgement}(\text{e}) \end{array} \right] \end{array}$$

Here the subscript ‘illoc’ stands for “illocutionary” indicating that the condition provides information about the illocutionary force of the speech act. The symbol \wedge represents the merge operation defined in Appendix A.13. In (29) the relevant merges will be the unions of the sets of fields represented by *SEvent* and the type consisting of the ‘e’ and ‘illoc’ fields. This is illustrated in (30) for *Assertion*. (30a) (where *SEvent* is spelled out) is identical with (30b).

⁸Strictly speaking, this type should be written with parentheses since we are assuming a binary meet operation: $([\text{e:SpeechAct}] \wedge (([\text{e:sp=a:Ind}] \vee [\text{e:au=a:Ind}]) \wedge \textit{MoveContent}))$ but we will often omit parentheses for clarity.

⁹Strictly speaking, this type should be written with parentheses since we are assuming a binary join operation: $(\textit{Assertion} \vee (\textit{Query} \vee (\textit{Command} \vee \textit{Acknowledgement})))$ but we will often omit parentheses for clarity.

$$\begin{array}{lcl}
 (30) \text{ a. } & \left[\begin{array}{lcl} \text{e-loc} & : & \text{Loc} \\ \text{sp} & : & \text{Ind} \\ \text{au} & : & \text{Ind} \\ \text{e} & : & \text{Phon} \\ \text{c}_{\text{loc}} & : & \text{loc(e,e-loc)} \\ \text{c}_{\text{sp}} & : & \text{speaker(e,sp)} \\ \text{c}_{\text{au}} & : & \text{audience(e,au)} \end{array} \right] & \wedge \left[\begin{array}{l} \text{e:Phon} \\ \text{c}_{\text{illoc}}:\text{assertion(e)} \end{array} \right] \\
 \text{b.} & \left[\begin{array}{lcl} \text{e-loc} & : & \text{Loc} \\ \text{sp} & : & \text{Ind} \\ \text{au} & : & \text{Ind} \\ \text{e} & : & \text{Phon} \\ \text{c}_{\text{loc}} & : & \text{loc(e,e-loc)} \\ \text{c}_{\text{sp}} & : & \text{speaker(e,sp)} \\ \text{c}_{\text{au}} & : & \text{audience(e,au)} \\ \text{c}_{\text{illoc}} & : & \text{assertion(e)} \end{array} \right] &
 \end{array}$$

Finally, the type *MoveContent* in (28) relates the type of the content of the move to the type of the move. We define it preliminarily as the join type in (31).

$$\begin{array}{lcl}
 (31) & \left[\begin{array}{lcl} \text{e} & : & \text{Assertion} \\ \text{cnt} & : & \text{RecType} \\ \text{c}_{\text{cnt}} & : & \text{content(e,cnt)} \end{array} \right] \vee \left[\begin{array}{lcl} \text{e} & : & \text{Query} \\ \text{cnt} & : & \text{Question} \\ \text{c}_{\text{cnt}} & : & \text{content(e,cnt)} \end{array} \right] \vee \\
 & \left[\begin{array}{lcl} \text{e} & : & \text{Command} \\ \text{cnt} & : & \text{RecType} \\ \text{c}_{\text{cnt}} & : & \text{content(e,cnt)} \end{array} \right] \vee \left[\text{e} : \text{Acknowledgement} \right]
 \end{array}$$

Note that this allows for acknowledgements such as *ok* not to have any content (although it does not prevent them from having content). We will return later to discussion of whether this is a reasonable claim for acknowledgements, while noting that this would be one way of dealing with “phatic” communication such as greetings like *Hello*.

We will be able to read partial information about a dialogue move from certain aspects of a speech-event. For example, an utterance of *ok* may tell us that the dialogue move is of type (32).

$$(32) \left[\begin{array}{lcl} \text{e} & : & \text{Acknowledgement} \\ \text{sp=SELF} & : & \text{Ind} \end{array} \right]$$

If an utterance of *ok* is to have content after all, we will have to look to the previous utterance to find it. An utterance of *Dudamel is a conductor* may give us information about the type of

speech act (*Assertion*) and the content ($[e:conductor(Dudamel)]$). Note, however, that we only get such a fully specified content if we have a unique individual ‘Dudamel’ whom we associate with utterances of the name *Dudamel*. If the resources we have available do not give us such an individual associated with *Dudamel* then we only get the information that somebody named Dudamel is a conductor, that is, we may only get partial information about the content the speaker intended to communicate. Consider the example *Strauss is a composer*. There are at least two famous composers named Strauss (and also some more not so famous ones). If our available resources give us two people associated with the name *Strauss* we will not know which of them is being referred to. Representing contents as record types will enable us to handle this content underspecification. However, in order to do this we will need to abandon our current simplifying “propositional logic” assumption that sentences come as unanalyzed wholes associated with their contents. This we will do in Chapter 3.

Not surprisingly, when we are dealing with an agenda as in (26), a plan for future action, we have got ourselves into a situation where we need types rather than the objects. The things that are on the agenda list are not actual events, but rather *types* of events planned for the future. Normally the types occurring on the agenda will be subtypes of *Move(SELf)*, though we may wish to include types of events like looking something up in a database, i.e. non-speech events.

For the most part types on the agenda will not be completely specified types. That is they will not be types all of whose fields are manifest (restricted to particular objects of those types). Frequently it will be the case that we specify the content of the move but leave open the phonology, that is, the type will specify the content of what is to be said but not actually what is to be said or even perhaps which language it should. We want, for example, to be able to say that a speaker is carrying out the same type of move independently of which language they are speaking. Thus, if the user says to the system “Dudamel is a conductor” or (in Swedish) “Dudamel är dirigent” she will in both cases have carried out a move involving the assertion of the content $[e:conductor(Dudamel)]$. This abstraction will be important, for example, if we want to change language in the middle of a dialogue, as people sometimes do.¹⁰ At the same time it is the normal case to continue a dialogue in the same language and thus we need to note which language was used in the previous utterance, i.e. keep track of what was actually said. This information will be in the chart which is part of the latest move. In the chart there will be more information about what was actually said which will be important when it comes to dealing with parts of the utterance for things like clarification and anaphora. But this again requires us to abandon our current simplifying “propositional logic” assumption.

We will assume that agents do not have complete information about the information state, that is, they reason in terms of *types* of information state (that is, gameboards). The basic intuition behind our reasoning about information state updates can be expressed as in (33).

(33) If $r_i : T_i$, then $r_{i+1} : T_{i+1}(r_i)$

¹⁰This phenomenon is known as code-switching (Bullock and Toribio, 2009).

That is, given that we believe that the current information state is of type T_i (recall that we can come to this belief without having any belief about which specific information state is involved), then we can conclude that the next information state is of type T_{i+1} which can depend on the current information state. According to this, we can have a hypothesis about the type of the next information state even though we may not know exactly what the current information state is. Exactly which type the next information state belongs to depends, though, on the exact nature of the current information state. Thus the dependency in our types provides us with an additional means for representing underspecification.

This basic rule of inference corresponds to a function from records to record types, a function of type $(T_i \rightarrow \text{RecType})$, that is, one kind of update function we were using in Chapter 1. Such a function is of the form (34).

$$(34) \quad \lambda r : T_i . T_{i+1}(r)$$

Things are a little more complicated than this, however, because this only represents the change from one information state to another, whereas in fact this change is triggered by a speech event which bears an appropriate relation to the current information state represented by r . Thus we are actually interested in functions from the current information state to a function from events to the new information state, as in (35).

$$(35) \quad \lambda r : T_i . \lambda e : T_e(r) . T_{i+1}(r, e)$$

This is the other kind of update function we were using in Chapter 1.¹¹ Let us consider the update function which the user could use in order to update her information state after her own utterance of *Dudamel is a conductor*. This is modelled on the kind of integration rules discussed in Larsson (2002).

$$(36)$$

¹¹This is one of a number of ways of characterizing update in this kind of framework. One might for instance think of the type of the speech event as being part of the current information state. Also instead of using an update function one can use a record type with a ‘preconditions’-field and an ‘effect’-field. Both Ginzburg (2012) and Larsson (2002) have this kind of approach.

$$\begin{aligned}
& \lambda r: \left[\text{private} : \left[\text{agenda} : {}_{ne}[\text{MoveType}(\text{SELF})] \right] \right] \\
& \lambda u: \left[\begin{array}{l} \text{move} : \text{fst}(r.\text{private}.\text{agenda}) \wedge \left[\begin{array}{l} \text{sp}=\text{SELF:Ind} \\ \text{au:Ind} \end{array} \right] \wedge [\text{e:Assertion}] \\ \text{chart} : \text{Chart} \\ \text{e} : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] . \\
& \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} \text{e:Acknowledgement} \wedge \left[\begin{array}{l} \text{sp}=u.\text{move.e.au:Ind} \\ \text{au=SELF:Ind} \end{array} \right] \\ \text{cnt}=u.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(\text{e}, \text{cnt}) \end{array} \right] \\ \text{rst}(r.\text{private}.\text{agenda}) \end{array} \right] : [\text{MoveType}(\text{SELF})] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u.\text{move:Move}(\text{SELF}) \\ \text{chart}=u.\text{chart:Chart} \\ \text{e}=u.\text{e:m-interp}(\text{chart}, \text{move}) \end{array} \right] \end{array} \right] \end{array} \right]
\end{aligned}$$

This function maps information states (records), r , which have a non-empty agenda to a function that maps events to a type of information state. (See Appendix A.5 for an account of non-empty list types.) It thus requires that the current information state (the first argument to the function) have a non-empty agenda. The second argument to the function (represented by u) requires the move associated with the speech-event to be of the first type on the agenda in r , the current information state, and also to be an assertion with *SELF* as the speaker (see Appendix A.9 for a discussion of meet types, that is, conjunction). It also requires that the chart associated with this utterance can be interpreted as a move of that type. The requirements on the arguments to the function represent the preconditions. The type that results from applying the function to its arguments represents the effect of the update. This type requires the agenda to be result of replacing the first type on the agenda in r with an acknowledgement where the speaker is the audience of the assertion move and the audience of the acknowledgement is *SELF*. The content of the acknowledgement is the same as the content of the assertion. That is, what is being acknowledged is the content of the assertion. It furthermore requires the latest-utterance field to contain the move and chart of the utterance u . The idea is that this function should be used to predict the type of the next information state on the basis of the current information state and the observed event. That is, if we believe the current information state to be of the domain type of the update function and we observe an event of the required type then we reason that the updated information state should be of the type resulting from applying the function to the current information state. Thus this update function will be used in the same way as the update functions we discussed in Chapter 1. However, the gameboards involved are now more complex.

We will now examine how such an update function could be used to reason about an update. Let us suppose that the user considers the current information state to be of type:

$$(37) \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda} = \left[\begin{array}{l} e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e, cnt) \end{array} \right] : [\text{RecType}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} E\text{Rec} \\ \text{commitments} = \text{Rec}:\text{RecType} \end{array} \right] \end{array} \right]$$

This represents that the user intends to assert that Dudamel is a conductor represented by the record type $[e:\text{conductor}(\text{Dudamel})]$. The user also believes that there was no previous utterance and no commitments, i.e. that the planned utterance will be dialogue initial.

Suppose now that the user utters *Dudamel is a conductor* and judges this utterance event u_1 to be an event of type (38).

$$(38) \left[\begin{array}{ll} \text{move} & : \left[\begin{array}{l} e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{conductor}(\text{Dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e, cnt) \end{array} \right] \\ \text{chart} & : \text{Chart} \\ \text{e} & : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right]$$

The user will have more information about the nature of the chart (that is, about what was actually said and how it might be analyzed) than we have represented but we will leave this underspecified for now.

Clearly in the user's judgement the utterance u_1 fulfils the requirements placed on it by (36) since the move interpretation associated with it is of the type which occurs at the head of the agenda. Note that we are reasoning with this function without actually providing it with an argument since we only have a (hypothesized) type of the current information state, not the actual information state. The crucial judgement is that the type of the current information state is a subtype of the domain type of the function. This is sufficient to allow us to come to a conclusion about the type of the new information state.

According to the update function the next information state must be of the type (39).

$$(39) \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda} = \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} sp=u_1.\text{move}.e.\text{au}:Ind \\ au=SELF:Ind \end{array} \right] \\ cnt=u_1.\text{move}.cnt:\text{RecType} \\ c_{cnt}:\text{content}(e, cnt) \end{array} \right] : [\text{RecType}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move} = u_1.\text{move}:\text{Move} \\ \text{chart} = u_1.\text{chart}:\text{Chart} \\ e = u_1.e:\text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \end{array} \right] \end{array} \right]$$

But we know more about the new information state than what is expressed by the type which results from the update function. Everything we know about the current information state which remains unchanged by the function must be carried over from the current information state. This is related to the frame problem introduced by McCarthy and Hayes (1969).¹² We handle this performing an *asymmetric merge* (see Appendix A.13) of the type we have for the current information state with the type resulting from the update function. The asymmetric merge of two types T_1 and T_2 is represented by $T_1 \sqcup T_2$. If one or both of T_1 and T_2 are non-record types then $T_1 \sqcup T_2$ will be T_2 . If they are both record types, then for any label ℓ which occurs in both T_1 and T_2 , $T_1 \sqcup T_2$ will contain a field labelled ℓ with the type resulting from the asymmetric merge of the corresponding types in the ℓ -fields of the two types (in order). For labels which do not occur in both types, $T_1 \sqcup T_2$ will contain the fields from T_1 and T_2 unchanged. In this informal statement we have ignored complications that arise concerning dependent types in record types. This is discussed in Appendix A.13. Our notion of asymmetric merge is related to the notion of priority unification (Shieber, 1986).

Let us see how this works with our example. We have assumed that the type under consideration for the current information state, T_{curr} , is (37) and computed that the predicted type of the updated information state, T_{pr} , is (39). Therefore we need to compute $T_{curr} \sqcup T_{pr}$, that is, (40).

$$(40) \quad \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} \text{e:Assertion} \wedge \left[\text{sp=SELF:Ind} \right] \\ \text{cnt} = \left[\text{e:conductor(Dudamel)} \right] : \text{RecType} \\ \text{c}_{\text{cnt}} : \text{content(e,cnt)} \end{array} \right] : [\text{Move(SELF)}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:ERec} \\ \text{commitments=Rec:RecType} \end{array} \right] \end{array} \right] \sqcup \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} \text{e:Acknowledgement} \wedge \left[\begin{array}{l} \text{sp=u}_1.\text{move.e.au:Ind} \\ \text{au=SELF:Ind} \end{array} \right] \\ \text{cnt} = \text{u}_1.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}} : \text{content(e,cnt)} \end{array} \right] : [\text{RecType}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move=u}_1.\text{move:Move} \\ \text{chart=u}_1.\text{chart:Chart} \\ \text{e=u}_1.\text{e:m-interp(chart,move)} \end{array} \right] \end{array} \right] \end{array} \right]$$

A straightforward way to think of the asymmetric merge of two record types is in terms of the paths in each of them. Both T_{curr} and T_{pr} contain paths ‘private.agenda’. The types at the end of the respective paths, however, are distinct singleton types. (Recall that manifest fields $[\ell=a:T]$ are a convenient notation for $[\ell:T_a]$ where T_a is a restriction of the type T whose only witness is a .) Therefore we include the complete path from the second type in the result of the asymmetric merge. In the case of the path ‘shared.latest-utterance’ we have the type of the empty record $ERec$ compared with a record type of non-empty records in T_{pr} and since these cannot be merged we

¹²For a recent overview of the frame problem see Shanahan (2009).

choose the second record type in the result. Finally, the path ‘shared.commitments’ occurs in the first type but not in the second and therefore it occurs in its form from the first type in the result of the asymmetric merge. The result is given in (41) which represents the type of the new information state which has been computed as a result of the update.

$$(41) \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} \text{sp}=u_1.\text{move.e.au:Ind} \\ \text{au}=SELF:Ind \end{array} \right] \\ \text{cnt}=u_1.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u_1.\text{move:Move} \\ \text{chart}=u_1.\text{chart:Chart} \\ e=u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \text{commitments}=\text{Rec:RecType} \end{array} \right] \end{array} \right] \end{array} \right]$$

Why has the field ‘shared.commitments’ not been updated after the user has asserted that Dudamel is a conductor? This is because the audience has not yet confirmed that they have understood and accepted the move. We assume that our agents are *cautious* and do not assume that commitments are shared until the dialogue participant(s) they are addressing have confirmed acceptance. This interaction is known as grounding and is discussed (among other places) in Traum (1994) and Larsson (2002).

We shall call the update function (36) **IntegrateOwnAssertion** following the style of Larsson (2002) although this does not correspond exactly to any of Larsson’s particular update rules. This then can be used to account for the state that the user is in after asserting that Dudamel is a conductor.

We now need an update function that will account for the effect of this utterance on another dialogue participant. For this we will define a function **IntegrateOtherAssertion** which allows an agent to integrate a move which it perceives to be an assertion.

$$(42) \lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : [\text{RecType}] \end{array} \right] \\ \lambda u: \left[\begin{array}{l} \text{move:} \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp:Ind} \\ \text{au}=SELF:Ind \end{array} \right] \\ \text{cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart:Chart} \\ e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right] .$$

$$\left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} \text{sp}=\text{SELF:Ind} \\ \text{au}=u.\text{move.e.sp:Ind} \end{array} \right] \\ \text{cnt}=u.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] :[\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u.\text{move:Move} \\ \text{chart}=u.\text{chart:Chart} \\ e=u.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right] \mid r.\text{private.agenda} \end{array} \right]$$

If an agent uses (42) to update then the new information state will contain a move type on the agenda which involves acknowledging the content of the assertion by the other dialogue partner. This update function is also cautious in that it does not yet update the shared commitments since the acknowledgement is only scheduled on the agenda but has not yet been performed.

If an agent performs an acknowledge-event (“ok”) and it can integrate it with the update function **IntegrateOwnAcknowledgement** which will finally perform an update of shared commitments. Before we define this update function we will examine what needs to happen in order to update the commitments.

Suppose that in the dialogue so far it has been established that Dudamel is a conductor and that this is represented by the record type (43).

$$(43) \quad [e : \text{conductor}(\text{Dudamel})]$$

Suppose further that the latest move has the content that Beethoven is a composer, namely (44).

$$(44) \quad [e : \text{composer}(\text{Beethoven})]$$

One obvious way to combine them would be to merge them, that is, (45a) which is identical with (45b) which in turn is identical with (45c), given the definition in Appendix A.13 which requires that the merge of any two types which are not both record types is identical with the meet of the two types.

$$(45) \quad \begin{array}{ll} \text{a.} & [e : \text{conductor}(\text{Dudamel})] \wedge [e : \text{composer}(\text{Beethoven})] \\ \text{b.} & [e : \text{conductor}(\text{Dudamel}) \wedge \text{composer}(\text{Beethoven})] \\ \text{c.} & [e : \text{conductor}(\text{Dudamel}) \wedge \text{composer}(\text{Beethoven})] \end{array}$$

For the simple storing of information represented by predicates and names represented in (16) this might be sufficient. It makes the claim that all the information is collected into one eventuality. In more narrative dialogues referring to separate events which we may wish to be able to refer back to this would be an inadequate solution, however. It would be better if we have a way of keeping the labels ‘e’ separate so that they don’t clash, for example in (46a) which is identical with (46b)

$$(46) \text{ a. } [e_1 : \text{conductor}(\text{Dudamel})] \wedge [e_2 : \text{composer}(\text{Beethoven})]$$

$$\text{b. } \begin{bmatrix} e_1 : \text{conductor}(\text{Dudamel}) \\ e_2 : \text{composer}(\text{Beethoven}) \end{bmatrix}$$

The potential problems of label clash become very clear if we consider the types in (47a) corresponding to *a boy hugged a dog* and *a girl stroked a cat*. (47a) is identical with (47b) and has a single individual which is both a girl and a boy stroking another individual which is both a dog and a cat.

$$(47) \text{ a. } \begin{bmatrix} x : Ind \\ c_{\text{boy}} : \text{boy}(x) \\ y : Ind \\ c_{\text{dog}} : \text{dog}(y) \\ e : \text{hug}(x,y) \end{bmatrix} \wedge \begin{bmatrix} x : Ind \\ c_{\text{girl}} : \text{girl}(x) \\ y : Ind \\ c_{\text{cat}} : \text{cat}(y) \\ e : \text{stroke}(x,y) \end{bmatrix}$$

$$\text{b. } \begin{bmatrix} x : Ind \\ c_{\text{boy}} : \text{boy}(x) \\ c_{\text{girl}} : \text{girl}(x) \\ y : Ind \\ c_{\text{dog}} : \text{dog}(y) \\ c_{\text{cat}} : \text{cat}(y) \\ e : \text{hug}(x,y) \wedge \text{stroke}(x,y) \end{bmatrix}$$

One way to get around this problem is to ensure that whenever you introduce new types you always use fresh labels that have not been used before and then use explicit constraints to require identity in cases where it is required. However, when we come to examine compositional semantics in Chapter 3 we will see that it is quite important to refer to particular labels in our rules of combination. Instead of introducing unique *labels* we will use the power of records to introduce unique *paths* when contents are combined. We will use the label ‘prev’ (“previous”). If T_{old} is the content so far and T_{new} is the content we wish to add then the new combined content will be as in (48a). Thus adding the content of *a girl stroked a cat* to that of *a boy hugged a dog* will yield (48b).

$$(48) \text{ a. } [\text{prev} : T_{\text{old}}] \wedge T_{\text{new}}$$

$$b. \left[\begin{array}{l} \text{prev} : \left[\begin{array}{l} x : \text{Ind} \\ c_{\text{boy}} : \text{boy}(x) \\ y : \text{Ind} \\ c_{\text{dog}} : \text{dog}(y) \\ e : \text{hug}(x,y) \end{array} \right] \\ x : \text{Ind} \\ c_{\text{girl}} : \text{girl}(x) \\ y : \text{Ind} \\ c_{\text{cat}} : \text{cat}(y) \\ e : \text{stroke}(x,y) \end{array} \right]$$

In the case of our example with Dudamel and Beethoven the result will be (49).

$$(49) \left[\begin{array}{l} \text{prev} : \left[e : \text{conductor}(\text{Dudamel}) \right] \\ e : \text{composer}(\text{Beethoven}) \end{array} \right]$$

If we add a further fact to this, say, that Uchida is a pianist we would obtain (50)

$$(50) \left[\begin{array}{l} \text{prev} : \left[\begin{array}{l} \text{prev} : \left[e : \text{conductor}(\text{Dudamel}) \right] \\ e : \text{composer}(\text{Beethoven}) \end{array} \right] \\ e : \text{pianist}(\text{Uchida}) \end{array} \right]$$

This means that we now have to add additional information if we want to require identity, for example if we want the Beethoven and Uchida eventualities (prev.e and e in (50)) to be identical. We will return to these matters when we deal with anaphora in Chapter 3. Note that this strategy also gives us a straightforward record of the order in which content was added.

The update function **IntegrateOwnAcknowledgement** is given in (51).

$$(51) \lambda r: \left[\begin{array}{l} \text{private} : \left[\text{agenda} : {}_{ne}[\text{RecType}] \right] \\ \text{shared} : \left[\begin{array}{l} \text{latest-utterance} : \left[\text{move} : \left[\text{content} : \text{RecType} \right] \right] \\ \text{commitments} : \text{RecType} \end{array} \right] \end{array} \right] \\ \lambda u: \left[\begin{array}{l} \text{move} : \text{fst}(r.\text{private}.\text{agenda}) \wedge [e:\text{Acknowledgement}] \wedge [e:[\text{sp}=\text{SELF}:\text{Ind}]] \\ \text{chart} : \text{Chart} \\ e : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \\ \left[\begin{array}{l} \text{private:} [\text{agenda}=\text{rst}(r.\text{private}.\text{agenda}):[\text{RecType}]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u.\text{move}:\text{Move} \\ \text{chart}=u.\text{chart}:\text{Chart} \\ e=u.e:\text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \\ \text{commitments}=[\text{prev}:r.\text{commitments}] \wedge u.\text{move}.\text{cnt}:\text{RecType} \end{array} \right] \end{array} \right] \end{array} \right] .$$

This function will

1. update the agenda with the result of removing the first item on the agenda in r , the information state prior to update
2. update the latest utterance with the current utterance (e.g. the utterance of *ok*)
3. update the commitments to be the result of placing the commitments of r under the label ‘prev’ and merging with the content of the move in the acknowledgement, u , (which by the update function **IntegrateOtherAssertion** will be the content of the previous assertion, e.g. the utterance of *Dudamel is a conductor*)

We then need an update function **IntegrateOtherAcknowledgement** which is like **IntegrateOwnAcknowledgment** except that it requires that the move event is directed towards the agent doing the updating. This is given in (52).

$$(52) \quad \lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : {}_{ne}[\text{RecType}] \\ \text{shared} : \left[\begin{array}{l} \text{latest-utterance} : \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} \text{content} : \text{RecType} \end{array} \right] \\ \text{commitments} : \text{RecType} \end{array} \right] \end{array} \right] \end{array} \right] \\ \lambda u: \left[\begin{array}{l} \text{move} : \text{fst}(r.\text{private}.\text{agenda}) \wedge [e:\text{Acknowledgement}] \wedge [e:[\text{au}=\text{SELF}:\text{Ind}]] \\ \text{chart} : \text{Chart} \\ \text{e} : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \end{array} \right] . \\ \left[\begin{array}{l} \text{private:} [\text{agenda}=\text{rst}(r.\text{private}.\text{agenda}):[\text{RecType}]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u.\text{move}:\text{Move} \\ \text{chart}=u.\text{chart}:\text{Chart} \\ \text{e}=u.\text{e}:\text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \\ \text{commitments}=[\text{prev}:r.\text{commitments}] \wedge u.\text{move}.\text{cnt}:\text{RecType} \end{array} \right] \end{array} \right] \end{array} \right] .$$

We have so far talked of update functions in this chapter, functions which given an information state and an utterance will return a type for an updated information state. Update functions specify something about the state that an agent will be in after the occurrence of a certain type of event. We have not, however, specified what it is that will specify that an agent should carry out an action which gives rise to an event of the appropriate type. Formally, these will also be functions which map an information state of a given type to a new type, the type of event which the agent is to bring about. Thus they too will be functions from objects to types (or dependent functions). We will call such functions *action functions*. These are associated with creation type acts (Chapter 1.5). We will introduce one such function, **ExecTopAgenda**, which takes an information state with a non-empty agenda and returns a type for a move of that type and a chart which can be interpreted as that move. It is given in (53).

$$(53) \quad \lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : {}_{ne}[\text{RecType}] \end{array} \right] \end{array} \right] .$$

$$\left[\begin{array}{ll} \text{move} & : \text{fst}(r.\text{private.agenda}) \\ \text{chart} & : \text{Chart} \\ \text{e} & : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right]$$

2.4 Resources

While there is no formal distinction between update functions and action functions they are to be used in different ways. Update functions are to be used as instructions to conclude that there is something of the resulting type. Action functions are to be used as instructions to create something of the resulting type. We shall say that they are different kinds of *resources* that are available to an agent. The update and action functions we have discussed in this chapter belong to a general resource for *dialogue management*. We shall see that there are a number of resources which contain both update and action functions and that in general they can be viewed as the two kinds of enthymemes (inferential and imperative) discussed in Breitholtz' work in progress on Aristotelian enthymemes (Breitholtz and Villing, 2008; Breitholtz, 2010; Breitholtz and Cooper, 2011).

We need more resources: signs and move-interpretations of charts containing signs. In this chapter we are taking signs to be objects of type (15) and the sign corresponding to *Dudamel is a conductor* is (17). For compactness of representation we can define an operation which takes a speech event type and a content and constructs the corresponding sign. This can be defined as in (54).

(54) If σ is a type of speech event and κ is a type (of situation) then

$$\text{sign}(\sigma, \kappa) = \left[\begin{array}{l} \text{s-event: } [e:\sigma] \\ \text{cnt} = \left[\begin{array}{l} e:\kappa \\ \text{c}_{\text{tns}}:\text{final_align}(\uparrow\text{s-event.e}, e) \end{array} \right] : \text{RecType} \end{array} \right]$$

Note that the operation 'sign' introduces the interpretation of present tense (represented by the field 'c_{tns}'). This is only possible because the resources we are considering concern only simple present tense assertions such as *Dudamel is a conductor*. We will see already in the next chapter that things are not this simple. We can use (54) to create signs types for utterances with specific contents such as *Dudamel is a conductor* or *Beethoven is a composer*. We will use another operation 'sign_{uc}' to create signs with underspecified content as defined in (55).

(55) If σ is a type of speech event then

$$\text{sign}_{uc}(\sigma) = \left[\begin{array}{l} \text{s-event: } [e:\sigma] \\ \text{cnt:RecType} \end{array} \right]$$

Now we can characterize the sign types that an agent that can deal with the simple dialogues that we have been characterizing in this chapter as (56).

- (56) $\{\text{sign}(\text{"Dudamel is a conductor"}, \text{conductor}(\text{dudamel})),$
 $\text{sign}(\text{"Beethoven is a composer"}, \text{composer}(\text{beethoven})),$
 $\text{sign}(\text{"Uchida is a pianist"}, \text{pianist}(\text{uchida})),$
 $\text{sign}_{uc}(\text{"ok"}),$
 $\text{sign}_{uc}(\text{"aha"})\}$

Recall that “Dudamel is a conductor” etc. represent a type of a string of word utterance events. For any word w , “ w ” is the type of event where w is uttered. For present purposes we assume that the agent has basic types of word utterances as given in (57a). In order to cope with the content the agent must have a basic type *Ind* to which certain individuals belong as given in (57b). Finally in order to construct the ptypes used for the content the agent would have to have the predicates given in (57c).

- (57) a. “Dudamel”, “is”, “a”, “conductor”, “Beethoven”, “composer”, “Uchida”, “pianist”,
 “aha”, “ok”
 b. dudamel, beethoven, uchida : *Ind*
 c. predicates with arity $\langle \text{Ind} \rangle$: conductor, composer, pianist

The set of ptypes based on (57b,c) is thus (58).

- (58) $\{p(a) \mid p \in \{\text{conductor}, \text{composer}, \text{pianist}\} \text{ and}$
 $a \in \{\text{dudamel}, \text{beethoven}, \text{uchida}\}\}$

Of the ptypes in (58) we could say that ‘conductor(dudamel)’, ‘composer(beethoven)’ and ‘pianist(uchida)’ are non-empty (“true”) and the rest are empty, although that may not correspond to the actual facts of the world. (Beethoven was a pianist, for example.) Very often, we are mainly interested in whether a ptype has witnesses (something of the type) or not and not particularly what those witnesses are. In a complete formal treatment, of course, the type system would specify objects which belong to those types. For example, we could say s_1 : conductor(dudamel), s_2 : composer(beethoven) and s_3 : pianist(uchida). Informally, we can say s_1 is a situation where Dudamel is a conductor or which shows that Dudamel is a conductor and so on. The idea of saying that an agent has a certain type in its resources is not so much to say that it has complete information about what belongs to the type (although its memory will contain partial information about what belongs to what types) but rather that it has a way (possibly not entirely decidable) of recognizing an object of the type if it sees one. Thus since I am an agent with the type ‘composer(uchida)’ in my resources I know (sort of) what it would mean for a situation to be of this type, e.g. a situation in which Uchida has written original musical compositions, had them performed and so on. When we are using our type theory to give an analysis of certain

fragments of language we are sometimes interested in going into more detail concerning the criteria for belonging to a given type. Other times we just treat the type as basic and only need to assume that the agent has some way of recognizing objects of the type. It depends on the level of detail we are interested in for the particular analysis.

We have used predicates other than those given in (58) in the types that we have discussed in this chapter. There are “technical” predicates such as ‘m-interp’ (“move interpretation”) which takes as its arguments a chart and a move. If c is a chart and m is a move then $\text{m-interp}(c,m)$ will be a non-empty type just in case “ m is an interpretation of c ”. Clearly, this is a case where it is of theoretical interest to us to say more about what constraints this places on c and m .

For the purposes of this chapter, since the parsing involved is a trivial association of strings of words with signs without any constituent analysis, we will equate charts with signs, that is $a : \text{Chart}$ just in case $a : \text{Sign}$. Thus ‘m-interp’ will relate signs to moves. The definition is given in (59).

- (59) a. if $c : \text{Chart}$ and $m : \text{Move}$ and for some σ and κ , $c = \text{sign}(\sigma, \kappa)$, then $\text{m-interp}(c,m)$ is non-empty iff $m : \begin{bmatrix} \text{e:Assertion} \\ \text{cnt}=c.\text{cnt:RecType} \end{bmatrix}$
- b. if $c : \text{Chart}$ and $m : \text{Move}$ and for some σ , $c = \text{sign}_{uc}(\sigma)$, then $\text{m-interp}(c,m)$ is non-empty iff $m : \begin{bmatrix} \text{e:Acknowledgement} \\ \text{cnt:RecType} \end{bmatrix}$

Let us now check that we can characterize the types of information states of A and B in the dialogue (60), where we represent the information states associated with the two agents at various points in the dialogue as a_i and b_i , and the utterance events as u_i .

- (60)
- | | | |
|----|-------------------------|-------|
| | a_0, b_0 | |
| A: | Dudamel is a conductor | u_1 |
| | a_1, b_1 | |
| B: | Aha | u_2 |
| | a_2, b_2 | |
| A: | Beethoven is a composer | u_3 |
| | a_3, b_3 | |
| B: | ok | u_4 |
| | a_4, b_4 | |
| A: | Uchida is a pianist | u_5 |
| | a_5, b_5 | |
| B: | ok | u_6 |
| | a_6, b_6 | |

We will assume that a_0 and b_0 are initial states, essentially empty except for A 's agenda to make the three assertions. This is shown in (61).

$$(61) \quad \begin{array}{l} \text{a. } a_0 : \\ \text{b. } b_0 : \end{array} \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda} = \left[\begin{array}{l} [e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right], \\ \left[\begin{array}{l} [e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{composer}(\text{beethoven})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right], \\ \left[\begin{array}{l} [e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{pianist}(\text{uchida})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \end{array} \right] : [\text{RecType}] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance}:\text{ERec} \\ \text{commitments}=\text{Rec}:\text{RecType} \end{array} \right] \end{array} \right] \\ \left[\begin{array}{l} \text{private:} [\text{agenda}=\text{Rec}:[\text{RecType}]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance}:\text{ERec} \\ \text{commitments}=\text{Rec}:\text{RecType} \end{array} \right] \end{array} \right]$$

(61) indicates that a_0 is an appropriate argument to the function **ExecTopAgenda** given in (53) and repeated in Appendix C. The result of applying **ExecTopAgenda** to a_0 is given in (62).

$$(62) \quad \text{ExecTopAgenda}(a_0) = \left[\begin{array}{l} \text{move:} \left[\begin{array}{l} [e:\text{Assertion} \wedge [sp=SELF:Ind] \\ cnt = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{chart:Chart} \\ e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right]$$

Note that given our notational convention on using *SELF* (p. 46), (62) is actually a dependent type as in (63).

$$(63) \quad \lambda a:Ind. \left[\begin{array}{l} \text{move:} \left[\begin{array}{l} [e:\text{Assertion} \wedge [sp=a:Ind] \\ cnt = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{chart:Chart} \\ e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right]$$

This means that the appropriate licensing condition on type acts associated with **ExecTopAgenda** (given in Appendix C.1.2) is the *de se* variant in (64).

- (64) If $f : (T \rightarrow (Ind \rightarrow Type))$ is an action function then for any object a and agent A , $a :_A T$ licenses $:_A f(a)(A)!$

That is, A is licensed to create (or contribute to the creation of) something of the type (62) with A itself as *SELF*. We have not yet said anything about what is involved in creating something of this type. The procedure involves generating a chart (in this chapter conceived of as a sign) whose content corresponds to the content of the move. We will not make this formally precise here but will wait until Chapter 3 where we have developed a more serious approach to grammar. Suffice it to say that the agent's resources must include the sign types introduced in (56) and that there is just one sign type here involving the type 'conductor(dudamel)' which figures in the content of the move specified in (62), namely $\text{sign}(\text{"Dudamel is a conductor"}, \text{conductor(dudamel)})$. For convenience we will abbreviate the notation of this type as $\Sigma_{\text{"Dudamel is a conductor"}}$. Only a sign of this type will satisfy *m-interp* for a move of the move type. Thus in order to realize something of type (62) A must in fact create something of type (65), a subtype of (62).

$$(65) \left[\begin{array}{lcl} \text{move} & : & \left[\begin{array}{lcl} e & : & \text{Assertion} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt}=[e:\text{conductor(dudamel)}] & : & \text{RecType} \\ c_{\text{cnt}} & : & \text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} & : & \Sigma_{\text{"Dudamel is a conductor"}} \\ e & : & \text{m-interp}(\text{chart},\text{move}) \end{array} \right]$$

(Here it is important that $\Sigma_{\text{"Dudamel is a conductor"}}$ is an abbreviation for the *notation* of the sign type, since when the notation is interpreted *in situ* in (65) each local path-name occurring as an argument to a predicate will be prefixed by 'chart.' and thus what occurs in the 'chart'-field of (65) is actually a modified version of the original type. It is possible to develop a notation that is more explicit but it becomes cluttered and unwieldy.)

Thus we can conclude that u_1 is judged by A to be of type (65). We can now predict that a_1 is of type **IntegrateOwnAssertion**(a_0)(u_1) which, given the types we have hypothesized for a_0 and u_1 will be (66).

(66)

$$\left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} \text{sp}=u_1.\text{move.e.au:Ind} \\ \text{au}=SELF:Ind \end{array} \right] \\ \text{cnt}=u_1.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{composer}(\text{beethoven})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{pianist}(\text{uchida})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \end{array} \right] : [\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u_1.\text{move:} \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart}=u_1.\text{chart:}\Sigma \text{“Dudamel is a conductor”} \\ e=u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right]
\end{array}
\right]$$

We can now use (66) to update the type we had for a_0 , given in (61a), as in (67a) which is identical with (67b).

$$(67) \text{ a. } \left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{composer}(\text{beethoven})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{l} e:\text{Assertion} \wedge \left[\begin{array}{l} \text{sp}=SELF:Ind \\ \text{cnt}=[e:\text{pianist}(\text{uchida})]:\text{RecType} \end{array} \right] \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \end{array} \right] : [\text{RecType}] \\ \text{latest-utterance:ERec} \\ \text{commitments=Rec:RecType} \end{array} \right] \right] \quad \boxed{\wedge}$$

$$\begin{array}{c}
\left[\begin{array}{c} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{c} \text{agenda} = \left[\begin{array}{c} e:\text{Acknowledgement} \wedge \left[\begin{array}{c} \text{sp} = u_1.\text{move.e.au:Ind} \\ \text{au} = \text{SELF:Ind} \end{array} \right] \\ \text{cnt} = u_1.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{composer}(\text{beethoven})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{pianist}(\text{uchida})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \end{array} \right] : [\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{c} \text{move} = u_1.\text{move:} \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} = u_1.\text{chart:} \Sigma \text{“Dudamel is a conductor”} \\ e = u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right] \end{array} \right]
\end{array}$$

b.

$$\begin{array}{c}
\left[\begin{array}{c} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{c} \text{agenda} = \left[\begin{array}{c} e:\text{Acknowledgement} \wedge \left[\begin{array}{c} \text{sp} = u_1.\text{move.e.au:Ind} \\ \text{au} = \text{SELF:Ind} \end{array} \right] \\ \text{cnt} = u_1.\text{move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{composer}(\text{beethoven})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{pianist}(\text{uchida})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \end{array} \right] : [\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{c} \text{move} = u_1.\text{move:} \left[\begin{array}{c} e:\text{Assertion} \wedge \left[\begin{array}{c} \text{sp} = \text{SELF:Ind} \\ \text{cnt} = [e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} = u_1.\text{chart:} \Sigma \text{“Dudamel is a conductor”} \\ e = u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \text{commitments} = \text{Rec:RecType} \end{array} \right] \end{array} \right]
\end{array}$$

Thus we can conclude that a_1 is of type (67b). A type for b_1 can be obtained in a similar fashion using **IntegrateOtherAssertion**(b_0)(u_1). The type that B will assign to u_1 can be predicted by the perception function (Appendix C) in (68).

$$(68) \quad \lambda e: \left[\begin{array}{c} e:\text{“Dudamel is a conductor”} \\ \text{au} = \text{SELF:Ind} \end{array} \right] .$$

$$\left[\begin{array}{l} \text{move} : \left[\begin{array}{l} e : \text{SpeechAct} \wedge [\text{au}=\text{SELF:Ind}] \\ \text{cnt} : \text{Cnt} \\ c_{\text{cnt}} : \text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} : \Sigma \text{"Dudamel is a conductor"} \\ e : \text{m-interp}(\text{chart},\text{move}) \end{array} \right]$$

(68) together with the type we have for b_0 predicts that **IntegrateOtherAssertion**(b_0)(u_1) will be the type (69).

$$(69) \quad \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] :[\text{RecType}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u_1.\text{move:} \left[\begin{array}{l} e:\text{Assertion} \wedge [\text{au}=\text{SELF:Ind}] \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart}=u_1.\text{chart:} \Sigma \text{"Dudamel is a conductor"} \\ e=u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right]$$

We can now use (69) to update the type we had for b_0 , obtaining (70a) identical with (70b).

$$(70) \quad \text{a.} \quad \left[\begin{array}{l} \text{private:} [\text{agenda}=[]:[\text{RecType}]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:ERec} \\ \text{commitments=Rec:RecType} \end{array} \right] \end{array} \right] \boxed{\wedge} \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] :[\text{RecType}] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move}=u_1.\text{move:} \left[\begin{array}{l} e:\text{Assertion} \wedge [\text{au}=\text{SELF:Ind}] \\ \text{cnt}=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart}=u_1.\text{chart:} \Sigma \text{"Dudamel is a conductor"} \\ e=u_1.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right]$$

b.

$$\left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge [sp=SELF:Ind] \\ cnt=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \end{array} \right] :[\text{RecType}] \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} move=u_1.move: \left[\begin{array}{l} e:\text{Assertion} \wedge [au=SELF:Ind] \\ cnt=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \end{array} \right] \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ chart=u_1.chart: \Sigma \text{"Dudamel is a conductor"} \\ e=u_1.e:m\text{-interp}(chart,move) \end{array} \right] \\ \text{commitments}=\text{Rec}:\text{RecType} \end{array} \right]
\end{array}
\right]$$

Now we are in a situation where both A and B are in information states (a_1 and b_1) with non-empty agendas. But they are coordinated in that A has an acknowledgement to be spoken by B topmost on the agenda and B has an acknowledgement with the same content to be spoken by B with A as the audience. **ExecTopAgenda** is applicable to both a_1 and b_1 . (71a) is **ExecTopAgenda**(a_1) and (71b) is **ExecTopAgenda**(b_1).

$$\begin{array}{ll}
(71) \text{ a.} & \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} sp=u_1.move.e.au:Ind \\ au=SELF:Ind \end{array} \right] \\ cnt=u_1.move.cnt:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{chart} : \text{Chart} \\ \text{e} : m\text{-interp}(chart,move) \end{array} \right] \\
\text{b.} & \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} sp=SELF:Ind \\ au=u_1.move.e.sp:Ind \end{array} \right] \\ cnt=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{chart} : \text{Chart} \\ \text{e} : m\text{-interp}(chart,move) \end{array} \right]
\end{array}$$

By substituting values for $SELF$ and values from u_1 we can see the extent to which A and B are coordinated. Both (71a) and (71b) reduce to the type in (72).

$$(72) \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} e:\text{Acknowledgement} \wedge \left[\begin{array}{l} sp=B:Ind \\ au=A:Ind \end{array} \right] \\ cnt=[e:\text{conductor}(\text{dudamel})]:\text{RecType} \\ c_{cnt}:\text{content}(e,cnt) \end{array} \right] \\ \text{chart} : \text{Chart} \\ \text{e} : m\text{-interp}(chart,move) \end{array} \right]$$

Both A and B can now, in virtue of **ExecTopAgenda** play their respective roles in creating an event of type (72), A by waiting for and paying attention to the acknowledgement and B by uttering the acknowledgement. This is an elementary form of what is known as *turn-taking* in the dialogue literature (Sacks *et al.*, 1974). The acknowledgement is u_2 in (60). In virtue of what is on the top of their respective agendas both A and B can judge u_2 to be of the type (72). A and B can now update their gameboards in virtue of **IntegrateOtherAcknowledgement** and **IntegrateOwnAcknowledgement** respectively. A will thus judge a_2 to be of type (73a) as a result of updating (67b) and B will judge b_2 to be of type (73b) as a result of updating (70b).

(73) a.

$$\left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda} = \left[\begin{array}{l} e:\text{Assertion} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt} = [\text{e:composer(beethoven)}]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right], \\ \left[\begin{array}{l} e:\text{Assertion} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt} = [\text{e:pianist(uchida)}]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] : [\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{l} \text{move} = u_2.\text{move:} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge [\text{au}=\text{SELF:Ind}] \\ \text{cnt} = [\text{e:conductor(dudamel)}]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} = u_2.\text{chart:} \Sigma_{\text{"Aha"}} \\ e = u_2.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \text{commitments} = \left[\begin{array}{l} \text{prev:Rec} \\ e:\text{conductor(dudamel)} \end{array} \right] : \text{RecType} \end{array} \right] \right]$$

b.

$$\left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda} = [] : [\text{RecType}] \\ \text{latest-utterance:} \left[\begin{array}{l} \text{move} = u_2.\text{move:} \left[\begin{array}{l} e:\text{Acknowledgement} \wedge [\text{sp}=\text{SELF:Ind}] \\ \text{cnt} = [\text{e:conductor(dudamel)}]:\text{RecType} \\ c_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart} = u_2.\text{chart:} \Sigma_{\text{"Aha"}} \\ e = u_2.e:\text{m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \text{commitments} = \left[\begin{array}{l} \text{prev:Rec} \\ e:\text{conductor(dudamel)} \end{array} \right] : \text{RecType} \end{array} \right] \right]$$

A and B are coordinated in that they both hypothesize the same type for shared.commitments. Now the assertion-acknowledgement cycle can begin again and repeat until both agents have gameboards with empty agendas. The final gameboards for A and B respectively are given in (74).

(74) a.

$$\begin{array}{l}
\left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda} = [] : [RecType] \\ \text{latest-utterance:} \\ \text{commitments} = \end{array} \left[\begin{array}{l} \text{move} = u_6.\text{move:} \left[\begin{array}{l} e: Acknowledgement \wedge [au = SELF: Ind] \\ cnt = [e: pianist(uchida)] : RecType \\ c_{cnt}: content(e, cnt) \end{array} \right] \\ \text{chart} = u_6.\text{chart:} \Sigma_{\text{"ok"}} \\ e = u_6.e: m\text{-interp}(\text{chart}, \text{move}) \\ \text{prev:} \left[\begin{array}{l} \text{prev:} Rec \\ e: conductor(dudamel) \\ e: composer(beethoven) \end{array} \right] \\ e: pianist(uchida) \end{array} \right] : RecType \end{array} \right]
\end{array}$$

b.

$$\begin{array}{l}
\left[\begin{array}{l} \text{private:} \\ \text{shared:} \end{array} \left[\begin{array}{l} \text{agenda} = [] : [RecType] \\ \text{latest-utterance:} \\ \text{commitments} = \end{array} \left[\begin{array}{l} \text{move} = u_6.\text{move:} \left[\begin{array}{l} e: Acknowledgement \wedge [sp = SELF: Ind] \\ cnt = [e: pianist(uchida)] : RecType \\ c_{cnt}: content(e, cnt) \end{array} \right] \\ \text{chart} = u_6.\text{chart:} \Sigma_{\text{"ok"}} \\ e = u_6.e: m\text{-interp}(\text{chart}, \text{move}) \\ \text{prev:} \left[\begin{array}{l} \text{prev:} Rec \\ e: conductor(dudamel) \\ e: composer(beethoven) \end{array} \right] \\ e: pianist(uchida) \end{array} \right] : RecType \end{array} \right]
\end{array}$$

2.5 Summary

Chapter 3

Grammar

In Chapter 2 we made the simplifying assumption that sentences come as single unanalyzed units (something like the assumption that is made in propositional logic). In this chapter we will deal with the same simple examples but break the sentences down into their constituent parts. (This will be something like moving from propositional logic to predicate logic without quantifiers.) In order to do this we will need more complex signs.

We will first consider how linguistic constituent structure is related to our general perception of events. We have so far talked of events in terms of string types which we have related to finite state automata. Finite state automata are equivalent to regular grammars. We will now consider an example of how we perceive events which suggest a more complex structure in terms of strings of regular types. This gives us something which is equivalent to recursive transition networks (RTNs) which are in turn equivalent to context free grammars.¹ Consider an event type of bus trips, *BusTrip*. This could be defined as in (1).

$$(1) \quad \textit{BusTrip} \equiv \textit{GetBus} \frown \textit{TravelOnBus} \frown \textit{GetOffBus}$$

Each of the three event types which are concatenated in (1) could be further broken down into strings of events. For example, *GetBus* might be defined as in (2).

$$(2) \quad \textit{GetBus} \equiv \textit{WaitAtBusstop}^* \frown \textit{BusArrive} \frown \textit{GetOnBus}$$

The elements in (2) could be broken down further. For example, getting on the bus could be analyzed in terms of going towards a door on the bus, waiting for the door to open, placing one foot on the step into the bus and then the other, paying for your ticket and so on. There seems

¹For a general introduction to automata theory and its relation to the Chomsky hierarchy see, for example, Partee *et al.* (1990).

almost no limit to how finegrained an analysis of events we can give. Which muscles do you have to move in order to place your right foot inside the bus? What events are involved in the contraction of this muscle? However, there seems to be a limit on the level of detail we need to be conscious of (or even are capable of being conscious of) in order to carry out a high level action like getting on a bus. We can also build upwards from the type *BusTrip*. For example, many bus trips are not direct in that we have to change buses in order to reach our destination. Thus a bus trip can consist of a string of events where you get on a bus, travel on it and then get off it again. A return bus trip involves a bus trip from one place to another followed (after intervening events) by a bus trip from the second place back to the first. Both of those bus trips might involve several buses if the connection is not direct.

The notation we have used in (1) and (2) is used to mean that what occurs to the left of \equiv is a convenient notational abbreviation for what occurs on the right. That is, whenever we write the symbol on the left, that is just shorthand for the longer expression on the right. Given the two definitions in (1) and (2), *BusTrip* is just an abbreviation for the regular string type in (3).

$$(3) \quad \text{WaitAtBusstop}^* \frown \text{BusArrive} \frown \text{GetOnBus} \frown \text{TravelOnBus} \frown \text{GetOffBus}$$

Thus while our notation is giving us the beginnings of a hierarchical organization, the type that is represented by the notation is not hierarchically organized. We are still in the realm of a finite state system. Compare this with the statements in (4).

$$(4) \quad \begin{array}{ll} \text{a. } e : \text{BusTrip} \text{ iff } e : \text{GetBus} \frown \text{TravelOnBus} \frown \text{GetOffBus} \\ \text{b. } e : \text{GetBus} \text{ iff } e : \text{WaitAtBusstop}^* \frown \text{BusArrive} \frown \text{GetOnBus} \end{array}$$

The statements in (4) claim that there are distinct types *BusTrip* and *GetBus* in addition to the regular types used on the right-hand side of ‘iff’. These types are *equivalent* to the regular types in the sense that anything of the one type will be of the other type. Now the actual type system (not just the notation) is hierarchically organized and includes two additional “higher” types *BusTrip* and *GetBus*. On the face of it one might think that the type system with the additional higher types would be just a more complicated way of achieving the same result and would be less efficient than a system which just includes the regular types. However, there seems to be good reason to suppose that an organism that organizes its event perception in terms of such a hierarchical type system would have serious advantages over an organism that lacks the hierarchical organization. These advantages include at least the following:

access and compact representation Recall from Chapter 1 that we want to consider the types that an agent has available as resources as being represented in the brain states of the agent. Having higher types means that something corresponding to a complex type can be stored

as a single element. In a complex reasoning task this can give considerable advantage in that the task can be represented in a more compact fashion and it can be easier to access (search and find) something which is a single element rather than something which is represented in terms of a complex string each element of which has to be checked in order to be sure that you have found the right element.

planning Having a compact representation facilitates planning. It is feasible to plan to take a bus trip given that we can conceive of it as such without having to plan for all the small subevents that make it up, for example, all that is involved in lifting your legs in the right way in order get on the bus. The ability to plan actions seems based on an ability to classify events in a hierarchical way.

reuse A hierarchical organization of event types means that certain event types can be reused in other event types. For example, getting on a bus (waiting for the doors to open, putting one foot inside and so on) can be very much like getting on a train. Similarly, paying for a ticket on a bus trip involves an exchange of money for a ticket in much the same way for a bus, a tram, a train, a theatre performance and so on. An agent which is not able to perceive this kind of generalization would at best use up a lot of memory coding the same event types over and over as parts of different larger event types.

learning The hierarchical organization of event types and the reuse capabilities it offers also facilitates learning of new event types. In learning to take the tram it can be useful to reuse what you have learnt about buying tickets on buses and insert it ready made into your type for tram trips. If it turns out that the procedure for buying tickets for trams is slightly different from for buses (for example, you can buy a ticket on the bus but you have to pay before you get on the tram) you nevertheless have a buying ticket type which you can modify. This might involve creating more types corresponding to those strings which the two ticket buying procedures have in common to separate out the differences between the two procedures.

Related observations about the importance of hierarchical structure for behaviour and its relationship to hierarchical reinforcement learning and neurological structure have been made for example by Botvinick (2008); Botvinick *et al.* (2009); Ribas-Fernandes *et al.* (2011).

Introducing hierarchical types in this way is an important step in our cognitive processing of events because of the computational and learning processes indicated above even if the class of events we are formally able to recognize is the same as what could be recognized by non-hierarchical regular string types, that is, technically, finite state languages. An organism with hierarchically organized types will have important advantages in acquiring new finite state event patterns. An evolutionary step from non-hierarchically organized string types to hierarchically organized types is a significant development and organisms with hierarchical types will have clear evolutionary advantages over those that do not.

However, hierarchical organization brings with it, almost as a kind of side effect, something which means that the organism could recognize classes of events that are not finite state. This is known as *recursion*. Hierarchical organization means that we can give type definitions of the form in (5).

$$(5) \quad a : T \text{ iff } a : T_1 \frown \dots \frown T_n$$

If we do not explicitly rule it out there is nothing to say that one of the T_i is not T itself. Of course, things will go badly wrong if we have a definition such as (6).

$$(6) \quad a : T \text{ iff } a : T_1 \frown T \frown T_2$$

If we try to perceive or create something of this type we will not be able to terminate and get into an endless string of objects of type T_1 and never be able to move on to T_2 . However, if we define T in terms of a join type where at least one of the types in the join does not contain T , things will work fine. For example, (7):

$$(7) \quad a : T \text{ iff } a : (T_1 \frown T \frown T_2 \vee T_1 \frown T_2)$$

According to (7) anything of type T will be a string of objects of type T_1 followed by a string of equal length of objects of type T_2 . It is the requirement “of equal length” which means that this type is not a regular type. For example, we could have the regular type $T_1^+ \frown T_2^+$ but this only expresses that we require a non-empty string of objects of type T_1 followed by a non-empty string of objects of type T_2 without the equal length requirement. What we have done here is restate a basic result from formal language theory in terms of our types. In formal language theory one talks of languages of the form $a^n b^m$ (the set of strings of n a ’s followed by a string of m b ’s, for any n and m greater than 0) which is a regular or finite state language and $a^n b^n$ (the set of strings of n a ’s followed by n b ’s, for any n greater than 0) which is context free. While this possibility of recursion is offered as soon as we allow the hierarchical typing of events in this way, it is not clear that it is exploited to a great extent in non-linguistic events. The clear examples that seem to exist are examples like opening and closing Chinese boxes, that is, boxes within boxes. The type of opening and closing (reassembling) a Chinese box could be characterized as the $a^n b^n$ -type in (8).

$$(8) \quad \begin{aligned} e : \text{OpenClose} & \text{ iff} \\ e : (\text{Open} \frown \text{OpenClose} \frown \text{Close} \vee \text{Open} \frown \text{Close}) \end{aligned}$$

It is significant in this kind of example that the ordering of the events is forced on the agent by the physical reality of the boxes. There is only one order in which you can open all the boxes and only one order (the reverse order) in which you can close them if you are going to assemble all the boxes within a single box. It is unclear that such ordering is required in non-linguistic event types when it is not dictated by physical reality.

3.1 Syntax

We now turn our attention to how this hierarchical organization is reflected in the nature of linguistic events. In Chapter 2 we used (9) as our sign type.

$$(9) \quad \left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{cont} & : \text{Cont} \end{array} \right]$$

This represents the pairing of a speech event with content in a Saussurean sign. It does not, however, require the presence of any hierarchical information in the sign corresponding to what in linguistic theory is normally referred to as the *constituent* (or *phrase*) structure of the utterance. To some extent it is arbitrary where we add this information. We could, for example, add it under the label ‘s-event’, perhaps by dividing ‘s-event.e’ into two fields ‘phon’ and ‘syn’ (“syntax”). However, it will be more convenient (in terms of keeping paths that we need to refer to often shorter) to add a third field labelled ‘syn’ at the top level of the sign type as in (10).

$$(10) \quad \left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{syn} & : \text{Syn} \\ \text{cont} & : \text{Cont} \end{array} \right]$$

However, as we will see below, *Syn* will require a ‘daughters’-field for a string of signs. This means that *Sign* becomes a recursive type. It will be a *basic* type with its witnesses defined by (11).

$$(11) \quad \sigma : \text{Sign} \text{ iff } \sigma : \left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{syn} & : \text{Syn} \\ \text{cont} & : \text{Cont} \end{array} \right]$$

We shall take *Syn* to be the type (12).²

²One might think that *Syn* should also be defined as a recursive type since it can contain *Sign* which in its turn can contain *Syn*. However, in the types we are currently proposing the only way for *Syn* to recur is through *Sign* and it is sufficient for *Sign* to be defined recursively to ensure that we do not introduce record types that are non-well founded sets of ordered pairs. That is, we want to avoid the mathematical object which is the type being a set which contains itself. In contrast the set of witnesses for a recursive type, while it will be infinite, will be well-founded.

$$(12) \begin{bmatrix} \text{cat} & : & \text{Cat} \\ \text{daughters} & : & \text{Sign}^* \end{bmatrix}$$

The type *Sign*, as so far defined, can be seen as a *universal resource*. By this we mean that it is a type which is available for all languages. *Cat* is the type of names of syntactic categories. In this chapter we will take the witnesses of *Cat* to be: *s* (“sentence”), *np* (“noun phrase”), *det* (“determiner”), *n* (“noun”), *v* (“verb”) and *vp* (“verb phrase”). These correspond to the categories we will use to cover the expressions of the fragment of English we introduced in Chapter 2. We will use capitalized versions of these category names to represent types of signs with the appropriate path in a sign type as in (13).

$$(13) \begin{aligned} \text{a. } S &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{s}: \text{Cat}]] \\ \text{b. } NP &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{np}: \text{Cat}]] \\ \text{c. } Det &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{det}: \text{Cat}]] \\ \text{d. } N &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{n}: \text{Cat}]] \\ \text{e. } V &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{v}: \text{Cat}]] \\ \text{f. } VP &\equiv \text{Sign} \wedge [\text{syn}: [\text{cat}=\text{vp}: \text{Cat}]] \end{aligned}$$

Recall that the symbol \wedge represents the merge operation on types as defined in Appendix A.13. This means that, for example, (13a) is the type in (14).

$$(14) \begin{bmatrix} \text{s-event} & : & \begin{bmatrix} \text{e-loc} & : & \text{Loc} \\ \text{sp} & : & \text{Ind} \\ \text{au} & : & \text{Ind} \\ \text{e} & : & \text{Phon} \\ \text{c}_{\text{loc}} & : & \text{loc}(\text{e}, \text{e-loc}) \\ \text{c}_{\text{sp}} & : & \text{speaker}(\text{e}, \text{sp}) \\ \text{c}_{\text{au}} & : & \text{audience}(\text{e}, \text{au}) \end{bmatrix} \\ \text{syn} & : & \begin{bmatrix} \text{cat}=\text{s} & : & \text{Cat} \\ \text{daughters} & : & \text{Sign}^* \end{bmatrix} \\ \text{cont} & : & \text{Cont} \end{bmatrix}$$

We might think that the type *Cat* is a language specific resource and indeed if we were being more precise we might introduce separate types for different languages such as Cat_{eng} , Cat_{swe} and Cat_{tag} for the type of category names of English, Swedish and Tagalog respectively. However, there is a strong intuition that categories in different languages are more or less related. For example, we would not be surprised to find that the categories available for English and Swedish

closely overlap (despite the fact that their internal syntactic structure differs) whereas the categories of English and Tagalog have less overlap. (See Gil, 2000 for discussion.) For this reason we assume that there is a universal resource *Cat* and that each language will have a subtype of *Cat* which specifies which of the categories are used in that particular language. This is related to the kind of view of linguistic universals as a kind of toolbox from which languages can choose which is put forward by Jackendoff (2002).

The ontological status of objects of type *Cat* as we have presented them is a little suspicious. Intuitively, categories should be subtypes of *Sign*, that is, like the types such as *S*, *NP* and so on in (13). We have identified signs belonging to these types as containing a particular object in *Cat* in their ‘cat’-field. But one might try to characterize such signs in a different way, for example, as fulfilling certain conditions such as having certain kinds of daughters. However, this is not quite enough, for example, for lexical categories, which do not have daughters. We have to have a way of assigning categories to words and we need to create something in the sign-type that will indicate the arbitrary assignment of a category to a word. For want of a better solution we will introduce the category names which belong to the type *Cat* as a kind of “book-keeping” device that will identify a sign-type as being one whose witnesses belong to category bearing that name.

The ‘daughters’-field is required to be a string of signs, possibly the empty string, since the type *Sign** uses the Kleene-*, that is the type of strings of signs including the empty string, ε . (See Appendix A.16.) Lexical items, that is words and phrases which are entered in the lexicon, will be related to signs which have the empty string of daughters. We will use *NoDaughters* to represent the type $[\text{syn}:[\text{daughters}=\varepsilon:\text{Sign}^*]]$.

@@Eliminate “normally” here.

If T_{phon} is a type (normally a phonological type, that is, $T_{\text{phon}} \sqsubseteq \text{Phon}$) and T_{sign} is a type (normally a sign type, that is, $T_{\text{sign}} \sqsubseteq \text{Sign}$), then we shall use $\text{Lex}(T_{\text{phon}}, T_{\text{sign}})$ to represent (15)

$$(15) \quad T_{\text{sign}} \wedge [\text{s-event}:[\text{e}:T_{\text{phon}}]] \wedge \text{NoDaughters}$$

This means, for example, that (16a) represents the type in (16b) which, after spelling out the abbreviations, can be seen to be the type in (16c).

- (16) a. $\text{Lex}(\text{“Dudamel”}, \text{NP})$
 b. $\text{NP} \wedge [\text{s-event}:[\text{e}:\text{“Dudamel”}]] \wedge \text{NoDaughters}$

$$c. \left[\begin{array}{lcl} & & \left[\begin{array}{ll} e\text{-loc} & : \text{Loc} \\ sp & : \text{Ind} \\ au & : \text{Ind} \\ e & : \text{"Dudamel"} \\ c_{loc} & : \text{loc}(e, e\text{-loc}) \\ c_{sp} & : \text{speaker}(e, sp) \\ c_{au} & : \text{audience}(e, au) \end{array} \right] \\ s\text{-event} & : & \\ \left[\begin{array}{ll} cat=np & : \text{Cat} \\ daughters=\varepsilon & : \text{Sign}^* \end{array} \right] \\ syn & : & \\ cont & : & \text{Cont} \end{array} \right]$$

We can think of ‘Lex’ as the function in (17)³

$$(17) \lambda T_1:Type \\ \lambda T_2:Type . \\ T_1 \wedge [s\text{-event}: [e:T_2]] \wedge NoDaughters$$

This function, which creates sign types for lexical items in a language, associating types with a syntactic category, can be seen as a universal resource. We can think of it as representing a (somewhat uninteresting, but nevertheless true) linguistic universal: “There can be speech events of given types which have no daughters (lexical items)”.

@ @ Make clear that NP and S for aha etc. is arbitrary.

The lexical resources needed to cover our example fragment is given in (18).

$$(18) \text{Lex}(\text{"Dudamel"}, NP) \\ \text{Lex}(\text{"Beethoven"}, NP) \\ \text{Lex}(\text{"a"}, Det) \\ \text{Lex}(\text{"composer"}, N) \\ \text{Lex}(\text{"conductor"}, N) \\ \text{Lex}(\text{"is"}, V) \\ \text{Lex}(\text{"ok"}, S) \\ \text{Lex}(\text{"aha"}, S)$$

The types in (18) belong to the specific resources required for English. This is not to say that these resources cannot be shared with other languages. Proper names like *Dudamel* and *Beethoven*

³We are using the notational convention for function application as used, for example, by Montague (1973) that if f is a function $f(a, b)$ is $f(b)(a)$.

have a special status in that they can be reused in any language, though often in modified form, at least in terms of the phonological type with which they are associated without this being perceived as quotation, code-switching or simply showing off that you know another language.

Resources like (18) can be exploited by update rules. If $\text{Lex}(T_w, C)$ is one of the lexical resources available to an agent A and A judges an event e to be of type T_w , then A is licensed to update their gameboard with the type $\text{Lex}(T_w, C)$. Intuitively, this means that if the agent hears an utterance of the word “composer”, then they can conclude that they have heard a sign which has the category noun. This is the beginning of *parsing*, which we will regard as the same kind of update involved in event perception as discussed in the previous chapters. The licensing condition corresponding to lexical resources like (18) is given in (19). We will return below to how this relates to gameboard update.

- (19) If $\text{Lex}(T, C)$ is a resource available to agent A , then for any u , $u :_A T$ licenses $:_A \text{Lex}(T, C) \wedge [\text{s-event}: [\text{e}=u:T_1]]$

(19) says that an agent with lexical resource $\text{Lex}(T, C)$ who judges a speech event, u , to be of type T is licensed to judge that there is a sign of type $\text{Lex}(T, C)$ whose ‘s-event.e’-field contains u .

Strings of utterances of words can be classified as utterances of phrases. That is, speech events are hierarchically organized into types of speech events in the way that we discussed at the beginning of this chapter. Agents have resources which allow them to reclassify a string of signs of certain types (“the daughters”) into a single sign of another type (“the mother”). So for example a string of type $\text{Det} \cap N$ can lead us to the conclusion that we have observed a sign of type NP whose daughters are of the type $\text{Det} \cap N$. The resource that allows us to do this is a rule which we will model as the function in (20a) which we will represent as (20b).

- (20) a. $\lambda u : \text{Det} \cap N .$
 $NP \wedge [\text{syn}: [\text{daughters}=u:\text{Det} \cap N]]$
 b. $\text{RuleDaughters}(NP, \text{Det} \cap N)$

‘RuleDaughters’ is to be the function in (21).

- (21) $\lambda T_1 : \text{Type}$
 $\lambda T_2 : \text{Type} .$
 $\lambda u : T_1 . T_2 \wedge [\text{syn}: [\text{daughters}=u:T_1]]$

Thus ‘RuleDaughters’, if provided with a subtype of Sign^+ and a subtype of Sign as arguments, will return a function which maps a string of signs of the first type to the second type with

the restriction that the daughters field is filled by the string of signs. ‘RuleDaughters’ is one of a number of sign type construction operations which we will introduce as universal resources which have the property of returning what we will call a sign combination function. The licencing conditions associated with sign combination functions are as characterized in (22).

@@ Give names to such principles

- (22) If $f : (T_1 \rightarrow Type)$ is a sign combination function available to agent A , then for any u ,
 $u :_A T_1$ licenses $:_A f(u)$

This means, for example, that if you categorize a string of signs as being of type $Det \cap N$ then you can conclude that there is a sign of type NP with the additional restriction that its daughters are u .

‘RuleDaughters’ takes care of the ‘daughters’-field but it says nothing about the ‘s-event.e’-field, that is the phonological type associated with the new sign. This should be required to be the concatenation of all the ‘s-event.e’-fields in the daughters. If $u : T^+$ where T is a record type containing the path π , we will use $concat_i(u[i].\pi)$, the concatenation of all the values $u[i].\pi$ for each element in the string u in the order in which they occur in the string. (This notation is made precise in Appendix A.16.) We can now formulate the function ConcatPhon as in (23)

- (23) $\lambda u : [s\text{-event} : [e : Phon]]^+ .$
 $\quad [s\text{-event} : [e = concat_i(u[i].s\text{-event}.e) : Phon]]$

ConcatPhon will map any string of speech events to the type of a single speech event whose phonology (that is the value of ‘s-event.e’) is the concatenation of the phonologies of the individual speech events in the string.

We want to combine the function (23) with a function like that in (20). We do this by merging the domain types of the two functions and also merging the types that they return. This is shown in (24a) which in deference to standard linguistic notation for phrase structure rules could be represented as (24b).⁴

- (24) a. $\lambda u : Det \cap N \wedge [s\text{-event} : [e : Phon]]^+ .$
 $\quad NP \wedge [syn : [daughters = u : Det \cap N]]$
 $\quad \wedge [s\text{-event} : [e = concat_i(u[i].s\text{-event}.e) : Phon]]$

⁴Note that ‘ \longrightarrow ’ used in the phrase structure rule in (24b) is not the same arrow as ‘ \rightarrow ’ which is used in our notation for function types. We trust that the different contexts in which they occur will help to distinguish them.

$$\text{b. } NP \longrightarrow Det\ N$$

In general we say that if C, C_1, \dots, C_n are category sign types as in (13) then $C \longrightarrow C_1 \dots C_n$ represents $\text{RuleDaughters}(C, C_1 \frown \dots \frown C_n) \frown \text{ConcatPhon}$ where for any type returning functions $\lambda r : T_1 \cdot T_2(r)$ and $\lambda r : T_3 \cdot T_4(r)$, $\lambda r : T_1 \cdot T_2(r) \frown \lambda r : T_3 \cdot T_4(r)$ denotes the function $\lambda r : T_1 \frown T_3 \cdot T_2(r) \frown T_4(r)$.

@@Make the definition of function merging displayed.

Thus the function in (24) can be represented in a third way as in (25).

$$(25) \text{ RuleDaughters}(NP, Det \frown N) \frown \text{ConcatPhon}$$

The hope is that the ability to factorize rules into “bite-size” components will enable us to build a theory of resources that will allow us to study them in isolation and will also facilitate the development of theories of learning. It gives us a clue to how agents can build new rules by combining existing components in novel ways. It has implications for universality as well. For example, while the rule $NP \longrightarrow Det\ N$ is not universal (though it may be shared by a large number of languages), ConcatPhon is a universally available rule component, albeit a trivial universal which says that you can have concatenations of speech events to make a larger speech event.

The rules associated with our small grammar are given by (26)

$$(26) \begin{array}{l} S \longrightarrow NP\ VP \\ NP \longrightarrow Det\ N \\ VP \longrightarrow V\ NP \end{array}$$

It may seem that we have done an awful of work to arrive at simple phrase structure rules. Some readers might wonder why it is worth all this trouble to ground the rules in a theory of events and action when what we come up with in the end is something that can be expressed in a standard notation which is one of the first things that a student of syntax learns. One reason has to do with our desire to explore the relationship between the perception and processing of non-linguistics events and speech events as discussed at the beginning of this chapter. Another reason has to do with placing natural constraints on syntax. By grounding syntactic structure in types of events we provide a motivation for the kind of discussion in Cooper (1982). An abstract syntax which proposes constituent structure which does not correspond to speech events is not grounded in the same way and thus presents a different kind of theory.

3.2 Semantics

We have so far specified our sign types in terms of phonology and syntax. Now we need to specify the content in the ‘cont’-field. We shall start by accounting for the contents of the lexical items specified in (18). We consider first the common nouns *composer* and *conductor*. For each of these we introduce a predicate of arity $\langle Ind \rangle$. (See Appendix A.3.2 for discussion of predicates and arity.) Our universal resources will include a function, ‘SemCommonNoun’ which will construct a common noun content from such a predicate, p . This is defined as in (27).

$$(27) \text{ SemCommonNoun}(p) = \lambda r: [x:Ind] . [e : p(r.x)]$$

The function in (27) is of type $([x:Ind] \rightarrow RecType)$. That is, it is a function which maps any record containing a field labelled ‘x’ with an individual as value to a record type. We will abbreviate this type as *Ppty* (for “property”) and we will call functions of this type *properties*. In our compositional semantics, properties will play a similar role as functions from individuals to truth values $(\langle e, t \rangle)$ in Montague semantics. In place of individuals, we use records with an ‘x’-field containing an individual. The motivation for this will become apparent in Chapter 5 when we discuss the temperature puzzle. In place of Montague’s truth-values (that is, objects of Montague’s type t) we use record types. Record types play the role of “propositions” in our system. Types, thought of as types of situations, can be considered as truth-bearing objects. They are true just in case there is something of the type and false otherwise, that is, if there is nothing of the type. The fact that we use the “proposition-like” objects as the results that our properties return is an essential ingredient in our intensional treatment of properties. In this way it follows in the tradition of property theory (Chierchia and Turner, 1988; Fox and Lappin, 2005) and Thomason’s intensional approach to propositional attitudes (Thomason, 1980).

We can now combine the ‘Lex’-function which builds sign types excluding content information with our new way of constructing common noun content. We define a function $\text{Lex}_{\text{CommonNoun}}$ which takes a phonological type and a predicate and returns a sign type. This is defined in (28).

$$(28) \text{ Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p) = \text{Lex}(T_{\text{phon}}, N) \wedge [\text{cont} = \text{SemCommonNoun}(p):Ppty]$$

Note that the type of the content required here is *Ppty*. In Chapter 2 we defined the content type *Cont* to be identical with *RecType*. Now we have to revise the definition of *Cont* to be $(RecType \vee Ppty)$. We will add further disjuncts to allow for more possibilities as we progress.

In order to cover the two common nouns *conductor* and *composer* we can include the sign types in (29) among our resources.

- (29) a. $\text{Lex}_{\text{CommonNoun}}(\text{"composer"}, \text{composer})$
 b. $\text{Lex}_{\text{CommonNoun}}(\text{"conductor"}, \text{conductor})$

Following Montague's (1973) original strategy we shall treat the contents of noun-phrases such as *Dudamel* or *a conductor* as being functions from properties to truth-bearing elements, that is, in our terms, record types. That is, noun-phrase contents will be of type $(P\text{pty} \rightarrow \text{RecType})$ which we will abbreviate as *Quant* (for "quantifier"). This means that we should now redefine the type of contents, *Cont*, as $\text{RecType} \vee P\text{pty} \vee \text{Quant}$.⁵

Dudamel and *Beethoven* will receive proper name contents. The recipe for constructing a proper name content based on a particular individual *a* is given by $\text{SemPropName}(a)$ as defined in (30).

- (30) $\text{SemPropName}(a) =$
 $\lambda P:P\text{pty} . P([x=a])$

We define $\text{Lex}_{\text{PropName}}$ which takes a phonological type (a name) and an individual (the referent of the name) and returns a sign type as in (31).

- (31) $\text{Lex}_{\text{PropName}}(T_{\text{Phon}}, a) =$
 $\text{Lex}(T_{\text{Phon}}, NP) \wedge [\text{cont}=\text{SemPropName}(a):\text{Quant}]$

Resources to cover the proper names in our grammar could be as in (32) where $d, b : \text{Ind}$ (two individuals, Dudamel and Beethoven).

- (32) a. $\text{Lex}_{\text{PropName}}(\text{"Dudamel"}, d)$
 b. $\text{Lex}_{\text{PropName}}(\text{"Beethoven"}, b)$

Note that there is nothing to prevent us from constructing sign types with the same phonological type but different contents. Thus proper names are not required to be "logically proper" in the sense that there is one and only one individual which can be referred to by an utterance belonging to the phonological type. Names can be ambiguous. For example, there are many composers named Bach and Strauss. We have the means to construct sign types for all of them on an as needed basis.

Now that we have both properties and quantifiers let us check at this point that we are on the right track for combining them in something like the kind of way that we will need for compositional

⁵Omitting parentheses for clarity.

semantics. Suppose we want to combine a proper name content for *Dudamel* (33a) with the property of being a conductor (33b). The obvious way to do this is by applying the function in (33a) to the argument (33b) as represented in (33c). According to the definition of functional application in Appendix A.4, (33c) is identical to (33d) which in turn is identical to (33e). In turn the dot notation for record path values defined in Appendix A.12 shows (33e) to be identical to (33f).

- (33) a. $\lambda P:Ppty . P([x=d])$
 b. $\lambda r:[x:Ind] . [e : conductor(r.x)]$
 c. $\lambda P:Ppty . P([x=d])$
 $(\lambda r:[x:Ind] . [e : conductor(r.x)])$
 d. $\lambda r:[x:Ind] . [e : conductor(r.x)]([x=d])$
 e. $[e : conductor([x=d].x)]$
 f. $[e : conductor(d)]$

This means that if we were dealing with a language like Russian where *Dudamel* is a *conductor* corresponds to a proper name followed by a common noun we would have a good way of combining the two contents by applying the content of the proper name to the content of the common noun.⁶ However, things are not quite so straightforward in English. Here we use an indefinite article to form the noun phrase *a conductor*. We shall treat the content of indefinite articles as a function that maps properties to quantifiers involving the existential relation between properties. That is, it will be a function of type $(Ppty \rightarrow Quant)$, a type which should be added to our definition of *Cont* which now becomes $RecType \vee Ppty \vee Quant \vee (Ppty \rightarrow Quant)$. As part of our universal resources we introduce a function ‘SemIndefArt’ which is defined as the function in (34).

- (34) $\lambda Q:Ppty .$
 $\lambda P:Ppty . \left[\begin{array}{ll} \text{restr}=Q & : Ppty \\ \text{scope}=P & : Ppty \\ e & : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$

We can also define a universal resource, $\text{Lex}_{\text{IndefArt}}$, which associates a phonological type (corresponding to an indefinite article in the language) with this content, as defined in (35).

- (35) $\text{Lex}_{\text{IndefArt}}(T_{\text{Phon}}) =$
 $\frac{\text{Lex}(T_{\text{Phon}}, Det) \wedge [\text{cont}=\text{SemIndefArt}:(Ppty \rightarrow Quant)]}{}$

⁶An alternative would be to treat the content of a proper name as a record rather than a quantifier and apply the property to the record as in (33d). This would correspond to the treatment of proper names as individual denoting as discussed, for example, by Partee (1986).

The local resource for the English indefinite article would thus be (36).

$$(36) \text{ Lex}_{\text{IndefArt}}(\text{“a”})$$

The compositional semantics of a noun-phrase consisting of a determiner followed by a noun will be the content of the determiner applied to the content of the noun. This is a case of *content forward application*. We define a function ‘ContForwardApp’, which is part of the universal resources, as in (37).

$$(37) \lambda T_1:\text{Type} \lambda T_2:\text{Type} . \\ \lambda u: [\text{cont}:(T_2 \rightarrow T_1)] \frown [\text{cont}:T_2] . \\ [\text{cont}=u[0].\text{cont}(u[1].\text{cont}):T_1]$$

The intuition behind this function is that if you observe a string of two utterances, the first of which has a content of type $(T_2 \rightarrow T_1)$ and the second of which has a content of type T_2 then you are licensed to conclude that there is an utterance whose content is the result of applying the content of the first element in the string to the content of the second element of the string. (For the notation $s[n]$ representing the n th element of a string s see Appendix A.16.) We can use ‘ContForwardApp’ to add constraints on content to a phrase structure rule as in the example in (38).

$$(38) NP \longrightarrow Det N \frown \text{ContForwardApp}(Ppty, Quant)$$

Recall from (24a) that $NP \longrightarrow Det N$ is the function (39a). $\text{ContForwardApp}(Ppty, Quant)$ is the function (39b). Merging these two functions yields (39c).

$$(39) \begin{array}{ll} \text{a. } \lambda u : Det \frown N \frown [s\text{-event}:[e:Phon]]^+ . \\ \quad NP \frown [syn:[daughters=u:Det \frown N]] \\ \quad \frown [s\text{-event}:[e=\text{concat}_i(u[i].s\text{-event}.e):Phon]] \\ \text{b. } \lambda u: [\text{cont}:(Ppty \rightarrow Quant)] \frown [\text{cont}:Ppty] . \\ \quad [\text{cont}=u[0].\text{cont}(u[1].\text{cont}):Quant] \\ \text{c. } \lambda u : Det \frown N \frown [s\text{-event}:[e:Phon]]^+ \\ \quad \frown [\text{cont}:(Ppty \rightarrow Quant)] \frown [\text{cont}:Ppty] . \\ \quad NP \frown [syn:[daughters=u:Det \frown N]] \\ \quad \frown [s\text{-event}:[e=\text{concat}_i(u[i].s\text{-event}.e):Phon]] \\ \quad \frown [\text{cont}=u[0].\text{cont}(u[1].\text{cont}):Quant] \end{array}$$

A convenient abbreviatory notation for this interpreted phrase structure rule is given in (40).

$$(40) \quad NP \longrightarrow Det\ N \mid Det'(N')$$

Here Det' and N' represent the contents of the determiner and noun.

We can represent the type (41a) using an informal diagrammatic tree notation which is common in linguistics as in (41b).⁷

$$(41) \quad a. \quad NP \wedge \left[\begin{array}{l} \text{s-event: } [e=\text{syn.daughters}[0].\text{s-event.e} \frown \text{syn.daughters}[0].\text{s-event.e:Phon}] \\ \text{syn: } [\text{daughters: } Det \frown N] \\ \text{cont}=\text{syn.daughters}[0].\text{cont}(\text{syn.daughters}[1].\text{cont}):Quant \end{array} \right]$$

b.

$$\begin{array}{c} NP \\ \alpha(\beta) \\ \swarrow \quad \searrow \\ Det \quad N \\ \alpha \quad \beta \end{array}$$

Here what is written under the category type (e.g. α, β) represents the value in the ‘cont’-field.

The content of an utterance of *a conductor* will be (42a) applied to (42b), that is (42c).

$$(42) \quad a. \quad \lambda Q:Ppty . \left[\begin{array}{ll} \text{restr}=Q & : Ppty \\ \text{scope}=P & : Ppty \\ e & : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$$

b. $\lambda r:[x:Ind] . [e : \text{conductor}(r.x)]$

c. $\lambda P:Ppty . \left[\begin{array}{ll} \text{restr}=\lambda r:[x:Ind] . [e : \text{conductor}(r.x)] & : Ppty \\ \text{scope}=P & : Ppty \\ e & : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$

We will now look in more detail at the nature of the generalized quantifier in (42c). ‘exist’ is a predicate with arity $\langle Ppty, Ppty \rangle$, that is, it corresponds to a relation between two properties. The classical account of generalized quantifiers (Barwise and Cooper, 1981; Peters and Westerståhl, 2006, and much other literature) treats such quantifier relations as relations between sets. Here we will follow Cooper (2011, 2013a) in relating our treatment directly to the classical relation

⁷A similar use of tree notation, though relating to typed feature structures rather than types, is used in HPSG (see, for example, Ginzburg and Sag, 2000, Chapter 2).

between sets, although, as argued in Cooper (2012a) based on earlier work by Keenan and Stavi (1986), there are ultimately good reasons for exploiting the intensionality of properties. If P is a property the relevant set is the set of individuals which have the property, which we will represent as $\lfloor \downarrow P \rfloor$. This is defined as in (43) where we use the notation $\lceil T \rceil$ to represent $\{a \mid a : T\}$.

$$(43) \quad \lfloor \downarrow P \rfloor = \{a \mid \exists r[r : \lceil x=a:Ind \rceil \text{ and } \lceil P(r) \rceil \neq \emptyset]\}$$

Following the terminology of Cooper (2011, 2013a) we will call $\lfloor \downarrow P \rfloor$ the property extension, or *P-extension*, of property P . Intuitively the property extension of P is the set of objects which have the property in some situation. Let us compute this for a particular example of a property, the property of being a dog given in (44).

$$(44) \quad \lambda r : \lceil x:Ind \rceil . \lceil e : \text{dog}(r.x) \rceil$$

The property extension of (44) is given in (45).

$$(45) \quad \{a \mid \exists r[r : \lceil x=a:Ind \rceil \text{ and } \lceil \lambda r : \lceil x:Ind \rceil . \lceil e : \text{dog}(r.x) \rceil (r) \rceil \neq \emptyset\}$$

By β -reduction (45) is the same set as (46).

$$(46) \quad \{a \mid \exists r[r : \lceil x=a:Ind \rceil \text{ and } \lceil \lceil e : \text{dog}(r.x) \rceil \rceil \neq \emptyset\}$$

Since r is required to be of the type $\lceil x=a:Ind \rceil$ we know that $r.x$ must be a . Therefore (46) is identical to (47).

$$(47) \quad \{a \mid \exists r[r : \lceil x=a:Ind \rceil \text{ and } \lceil \lceil e : \text{dog}(a) \rceil \rceil \neq \emptyset\}$$

By the definition of record types, a record $\lceil e=s \rceil$ is of type $\lceil e:\text{dog}(a) \rceil$ just in case $s : \text{dog}(a)$. Therefore this type is non-empty just in case there is such an s . For this reason (47) is the same set as (48).

$$(48) \quad \{a \mid \exists r[r : \lceil x=a:Ind \rceil \text{ and } \exists s[s : \text{dog}(a)]]\}$$

Since r is no longer bound in the second conjunct of (48), (49) also defines the same set.

$$(49) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } \exists s[s : \text{dog}(a)]\}$$

Given the nature of records, there will be an r of the required type just in case $a:Ind$. Therefore we can characterize the same set as in (50).

$$(50) \quad \{a \mid a:Ind \text{ and } \exists s[s : \text{dog}(a)]\}$$

Finally, since the existence of a situation of type $\text{dog}(a)$ requires the a is an individual given that the arity of ‘dog’ is $\langle Ind \rangle$ we can eliminate the first conjunct altogether so the minimal characterization of this set is (51).

$$(51) \quad \{a \mid \exists s[s : \text{dog}(a)]\}$$

If P and Q are properties we want $\text{exist}(P, Q)$ to be a type of situations which will be non-empty (that is, “true”) just in case the P-extensions of P and Q have a non-empty overlap, that is there is some individual which has both property P and property Q . In symbols we can express this as (52).

$$(52) \quad [\text{exist}(P, Q)] \neq \emptyset \text{ iff } [\downarrow P] \cap [\downarrow Q] \neq \emptyset$$

This places a requirement on objects which are assigned to the type ‘ $\text{exist}(P, Q)$ ’ without actually tying down what kind of object they have to be. That is, it leaves it open as to which objects get assigned to the type, as long as they respect this requirement. It places a constraint on F in the models discussed on p. 9. We can, however, go a step further and make precise exactly which objects these should be. The intuition is that a situation e should be of type $\text{exist}(P, Q)$ just in case it is a witness (or “proof”) of the fact that the “exist”-relation holds between P and Q (that is, that the P-extensions of P and Q have a non-empty overlap). We will say that a situation is such a witness just in case the P-extensions of the properties restricted to the situation in question stand in the required relation, that is, intuitively that the set of objects in the situation which have P overlaps with the set of objects in the situation which have Q . We will get at this notion by restricting properties to a particular situation (what we have called a resource situation in previous literature such as Barwise and Perry, 1983; Cooper, 1996). We will represent the restriction of property P to situation s as $P \upharpoonright s$. We take our previous example of the property of being a dog, repeated in (53a). Its restriction to the situation s is given in (53b).

$$(53) \quad \begin{array}{ll} \text{a. } \lambda r:[x:Ind] . [e : \text{dog}(r.x)] \\ \text{b. } \lambda r:[x:Ind] . [e \in s : \text{dog}(r.x)] \end{array}$$

In (53b) the restricted field $[e \underline{\varepsilon} s : \text{dog}(r.x)]$ requires that the object in ‘e’-field is not only of type ‘ $\text{dog}(r.x)$ ’ but also that it is either s itself or a component of s , that is, for some path π in s it is the object $s.\pi$. See Appendix A.12 for the definition of components. A definition of restriction for properties in general is given in Appendix B.1. We will not be concerned with the general definition here.

Now we can compute the property extension of (53b) in a similar fashion to the calculation for the non-restricted property. The property extension of (53b) is given in (54).

$$(54) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } [\lambda r : [x:Ind]. [e \underline{\varepsilon} s : \text{dog}(r.x)](r)] \neq \emptyset\}$$

By β -reduction (54) is the same set as (55).

$$(55) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } [\lambda [e \underline{\varepsilon} s : \text{dog}(r.x)]] \neq \emptyset\}$$

Since r is required to be of the type $[x=a:Ind]$ we know that $r.x$ must be a . Therefore (55) is identical to (56).

$$(56) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } [\lambda [e \underline{\varepsilon} s : \text{dog}(a)]] \neq \emptyset\}$$

By the definition of record types, a record $[e=s']$ is of type $[e \underline{\varepsilon} s : \text{dog}(a)]$ just in case $s' \underline{\varepsilon} s$ and $s' : \text{dog}(a)$. Therefore this type is non-empty just in case there is some s' such that $s' \underline{\varepsilon} s$ and $s' : \text{dog}(a)$. For this reason (56) is the same set as (57).

$$(57) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } \exists s'[s' \underline{\varepsilon} s \text{ and } s' : \text{dog}(a)]\}$$

Since r is no longer bound in the second conjunct of (57), (58) also defines the same set.

$$(58) \quad \{a \mid \exists r[r : [x=a:Ind]] \text{ and } \exists s'[s' \underline{\varepsilon} s \text{ and } s' : \text{dog}(a)]\}$$

Given the nature of records, there will be an r of the required type just in case $a:Ind$. Therefore we can characterize the same set as in (59).

$$(59) \quad \{a \mid a:Ind \text{ and } \exists s'[s' \underline{\varepsilon} s \text{ and } s' : \text{dog}(a)]\}$$

Finally, since the existence of a situation of type $\text{dog}(a)$ requires the a is an individual given that the arity of ‘dog’ is $\langle \text{Ind} \rangle$ we can eliminate the first conjunct altogether so the minimal characterization of this set is (60).

$$(60) \quad \{a \mid \exists s' [s' \underline{\leq} s \text{ and } s' : \text{dog}(a)]\}$$

Now we can use the notion of property restriction to characterize the witness condition for ptypes constructed with ‘exist’, as in (61).

$$(61) \quad e : \text{exist}(P, Q) \text{ iff } [\downarrow P \upharpoonright e] \cap [\downarrow Q \upharpoonright e] \neq \emptyset$$

This will have the consequence that any record of the type (62a) will be of type (62b).

$$(62) \quad \begin{array}{l} \text{a. } \left[\begin{array}{l} x : \text{Ind} \\ c : \text{dog}(x) \\ e : \text{run}(x) \end{array} \right] \\ \text{b. } \text{exist}(\lambda r : [x : \text{Ind}] . [e : \text{dog}(r.x)], \lambda r : [x : \text{Ind}] . [e : \text{run}(r.x)]) \end{array}$$

In other words, (62a) is a subtype of (62b). Consider an arbitrary record of type (62a) as given in (63).

$$(63) \quad \left[\begin{array}{l} x = a \\ c = s_1 \\ e = s_2 \end{array} \right]$$

where $s_1 : \text{dog}(a)$ and $s_2 : \text{run}(a)$

The two relevant property extensions to be considered will both be the set $\{a\}$ and thus the conditions for (63) being of type (62b) will be fulfilled since the two property extensions have a non-empty overlap.

Let us consider what we get when we apply the content we have for *a conductor*, (64a) (repeated from (42c)), to the property of composing, (64b). The result which would correspond to *a conductor composes* if we were to introduce *composes* as an intransitive verb in our resources, is given in (64c).

$$(64) \quad \text{a. } \lambda P : \text{Ppty} . \left[\begin{array}{l} \text{restr} = \lambda r : [x : \text{Ind}] . [e : \text{conductor}(r.x)] : \text{Ppty} \\ \text{scope} = P : \text{Ppty} \\ e : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$$

- b. $\lambda r: [x:Ind] . [e : compose(r.x)]$
- c.
$$\left[\begin{array}{lcl} \text{restr} = \lambda r: [x:Ind] . [e : conductor(r.x)] & : & Ppty \\ \text{scope} = \lambda r: [x:Ind] . [e : compose(r.x)] & : & Ppty \\ e & : & exist(\text{restr}, \text{scope}) \end{array} \right]$$

What would it mean for there to be something of type (64c)? In other words, what would be required to make the sentence *a conductor composes* true? There would have to be a record, r^* , which contains the three fields in the record in (65) and which meets the condition indicated.

$$(65) \quad r^* = \left[\begin{array}{lcl} \text{restr} & = & \lambda r: [x:Ind] . [e : conductor(r.x)] \\ \text{scope} & = & \lambda r: [x:Ind] . [e : compose(r.x)] \\ e & = & s \end{array} \right]$$

where $\lfloor \downarrow r^*.restr \rfloor s$ and $\lfloor \downarrow r^*.scope \rfloor s$ have a non-empty overlap.

This gives us a version of the classical treatment of indefinite articles as involving the existential quantifier, expressed in terms of a generalized quantifier which compares sets. There is, of course, a real and important question whether this is an appropriate content for the sentence *a conductor composes* which tends to get a generic reading something like “conductors, in general, compose”. We will return to this issue in Chapter 7 where we will deal with the indefinite article in more detail. For now, we will ignore the sentence *a conductor composes* since we are not considering syntactic resources for it anyway.

We are concerned with finding a way to interpret the verb phrase *is a conductor*. Can we find a content for *is* which could be combined with the content for *a conductor* given in (42c) to produce an appropriate interpretation for the verb-phrase? Montague’s (1973) strategy for assigning a content to *is* is reproduced in our terms in (66).

$$(66) \quad \lambda Q: Quant . \\ \lambda r_1: [x:Ind] . \\ Q(\lambda r_2: [x:Ind] . \left[\begin{array}{lcl} x=r_2.x, r_1.x & : & Ind \\ e & : & be(x) \end{array} \right])$$

Here we use a manifest field based on a multiple singleton type (a singleton type formed from a singleton type, see Appendices A.7 and A.12) to require the identity of $r_1.x$ and $r_2.x$. In the ‘e’-field of the type with the manifest field we use the predicate ‘be’ which we will take to be polymorphic with the set of arities as given in (67a). The witness condition associated with types constructed with ‘be’ is given in (67b).

$$(67) \quad \text{a. } \text{arity}(\text{be}) = \{ \langle T \rangle \mid T \text{ is a type} \}$$

$$\text{b. } e : \text{be}(a) \text{ iff } a \in e$$

The intuition behind (67b) could be expressed as “To be is to be a component of a situation”, that is, more technically, a “is” just in case there is a path, π , in some record, r , such that $r.\pi = a$.⁸

We will call (66) ‘SemBe’. It will be included among the universal resources, together with the ‘Lex_{be}’ as defined in (68).

$$(68) \text{ If } T_{\text{Phon}} \text{ is a phonological type, then } \text{Lex}_{\text{be}}(T_{\text{Phon}}) \text{ is } \text{Lex}(T_{\text{Phon}}, V) \wedge [\text{cont}=\text{SemBe}:(\text{Quant} \rightarrow \text{Ppty})]$$

Among the lexical resources for English we have Lex_{be}(“is”).

Now let us see what we get when we combine (66) with the content of *a conductor*. This involves applying (66), repeated as (69a), to (42c), repeated as (69b). The result of this application is (69c).

$$(69) \text{ a. } \lambda Q:\text{Quant} . \lambda r_1:[x:\text{Ind}] . Q(\lambda r_2:[x:\text{Ind}] . \left[\begin{array}{ll} x=r_2.X, r_1.X & : \text{Ind} \\ e & : \text{be}(x) \end{array} \right])$$

$$\text{b. } \lambda P:\text{Ppty} . \left[\begin{array}{ll} \text{restr}=\lambda r:[x:\text{Ind}] . \left[e : \text{conductor}(r.x) \right] & : \text{Ppty} \\ \text{scope}=P & : \text{Ppty} \\ e & : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$$

$$\text{c. } \lambda r_1:[x:\text{Ind}] . \left[\begin{array}{ll} \text{restr}=\lambda r:[x:\text{Ind}] . \left[e : \text{conductor}(r.x) \right] & : \text{Ppty} \\ \text{scope}=\lambda r_2:[x:\text{Ind}] . \left[\begin{array}{ll} x=r_2.X, r_1.X & : \text{Ind} \\ e & : \text{be}(x) \end{array} \right] & : \text{Ppty} \\ e & : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$$

In order to obtain a content for *Dudamel is a conductor* we apply the content of *Dudamel*, (33a), repeated as (70a), to (69c), repeated as (70b), with result (70c).

$$(70) \text{ a. } \lambda P:\text{Ppty} . P([x=d])$$

$$\text{b. } \lambda r_1:[x:\text{Ind}] .$$

⁸This might be compared with Quine’s (1948) dictum: “To be is to be the value of a variable”.

$$\begin{array}{c}
\left[\begin{array}{l} \text{restr} = \lambda r: [x:Ind] . \left[\begin{array}{l} e : \text{conductor}(r.x) \end{array} \right] \\ \text{scope} = \lambda r_2: [x:Ind] . \left[\begin{array}{l} x=r_2.x, r_1.x : Ind \\ e : \text{be}(x) \end{array} \right] \\ e \end{array} \right] \begin{array}{l} : Ppty \\ : Ppty \\ : \text{exist}(\text{restr}, \text{scope}) \end{array} \\
\text{c. } \left[\begin{array}{l} \text{restr} = \lambda r: [x:Ind] . \left[\begin{array}{l} e : \text{conductor}(r.x) \end{array} \right] \\ \text{scope} = \lambda r_2: [x:Ind] . \left[\begin{array}{l} x=r_2.x, d : Ind \\ e : \text{be}(x) \end{array} \right] \\ e \end{array} \right] \begin{array}{l} : Ppty \\ : Ppty \\ : \text{exist}(\text{restr}, \text{scope}) \end{array}
\end{array}$$

The type (70c) is distinct from the type (33f), repeated as (71), which we obtained by applying the content of *Dudamel* directly to the content of *conductor*.

$$(71) \left[\begin{array}{l} e : \text{conductor}(d) \end{array} \right]$$

There is, however, an equivalence that holds between (70c) and (71). The equivalence is not that they share the same set of witnesses. We can characterize the set of witnesses of (70c) and (72a) and the witnesses of (71) as (72b).

$$\begin{array}{l}
(72) \text{ a. } \left\{ \left[\begin{array}{l} \text{restr} = P \\ \text{scope} = Q \\ e = s \end{array} \right] \mid \begin{array}{l} P = \lambda r: [x:Ind] . \left[\begin{array}{l} e : \text{conductor}(r.x) \end{array} \right] \\ \text{and } Q = \lambda r: [x:Ind] . \left[\begin{array}{l} x=r.x, d : Ind \\ e : \text{be}(x) \end{array} \right] \\ \text{and } [\downarrow P \upharpoonright s] \cap [\downarrow Q \upharpoonright s] \neq \emptyset \end{array} \right\} \\
\text{b. } \left\{ \left[\begin{array}{l} e = s \end{array} \right] \mid s : \text{conductor}(d) \right\}
\end{array}$$

The sets in (72) do not have any members in common. The equivalence is a weaker “truth-conditional” equivalence. (70c) has a witness (“is true”) if and only if (71) has a witness. This is because the P-extensions of the property of being a conductor and the property of being identical with *Dudamel* can have a non-empty overlap if and only if *Dudamel* is a conductor. We might try to characterize the difference between the property associated with *conductor* and the property associated with *is a conductor* as “the property of being an x such that $\text{conductor}(x)$ ” and “the property of being an x such that there is a y such that $\text{conductor}(y)$ and $y = x$ ”. The two are truth-conditionally equivalent and for this reason in Montague’s system they turn out to be the same property. For us, since we are taking a more intensional approach than Montague, they are distinct properties but they are nevertheless truth-conditionally equivalent.

Since we have two distinct properties, the question is raised whether the property that is associated with the verb-phrase should be the same as the property associated with the common noun or whether it should be the property proposed here involving existential quantification. One way to do this is to create a type corresponding to the tree in (73).

(73)



This is not compositional in the standard sense because the content of the verb phrase is not defined as some operation applied to the contents of the verb and the noun phrase, but rather it makes the content of the verb phrase be the content of the noun. Furthermore, it requires the verb and determiner utterances be of the specific types “is” and “a” respectively. This gives (73) the flavour of representing a construction type as discussed in a variety of approaches to Construction Grammar, see, for example, Boas and Sag (2012). We can allow the type corresponding to (73) by introducing the update function (74).

$$(74) \quad \lambda u: V \wedge [s\text{-event}: [e: \text{"is"}]] \frown NP \wedge \left[\text{syn}: \left[\text{daughters}: Det \wedge [s\text{-event}: [e: \text{"a"}]] \right] \right] \frown N \wedge [cont: Ppty] \Big] \\
 VP \wedge [cont = u[2].\text{syn}.daughters[2].cont: Ppty]$$

We can call this function *CnstrIsA* (“is-a construction”) and merge it with $VP \longrightarrow V NP$. Thus one of the resources available for English is (75).

$$(75) \quad VP \longrightarrow V NP \frown \text{CnstrIsA}$$

This suggests that a phrase structure and construction based approach can be combined within a single framework. Since we are working with a toy fragment where the only verb is *is* and the only determiner is *a*, we can make do with (75) as the only resource for assigning content to verb-phrases. In a more general grammar we would, of course, require in addition a rule that applies the content of the verb to the content of the object noun-phrase as in (76).

$$(76) \quad VP \longrightarrow V NP \frown \text{ContForwardApp}(Quant, Ppty)$$

Allowing both resources (75) and (76) simultaneously raises the issue of what the relationship should be between them. Should the more specific rule (75) take precedence and guarantee that the only content associated with the verb phrase *is a conductor* is the property which is the

content of *conductor*? Or should the verb phrase be ambiguous between this interpretation and the property obtained by applying the content of *is* to the content of *a conductor*?

A more pressing issue, perhaps, is what to do about the sentence in (77).

(77) #A conductor is Dudamel

We have used the marking ‘#’ in (77) to indicate that an utterance of this sentence would under most, if not all, circumstances be considered to be odd, though it is difficult to rule it out as ungrammatical, particularly if we are to use something corresponding to context-free phrase structure rules as we are. The oddness of (77) may have something to do with the tendency to interpret noun phrases with indefinite articles in subject position as generic as in (78).

(78) A conductor is a high-ranking individual in the musical hierarchy

(77) can be improved without becoming generic. Examples are given in (79).

- (79) a. A conductor to reckon with is Dudamel
 b. A conductor to consider is Dudamel
 c. A conductor who impresses me as a leader in his generation is Dudamel
 d. A conductor I would like to see more often in Gothenburg is Dudamel

This raises a lot of issues which we do not currently have tools to deal with. There is, however, something we can say, if we choose to allow the is-a construction interpretation of *is a conductor*. The content of the sentence *Dudamel is a conductor* on the construction analysis becomes (71), repeated as (80a), rather than (70c), repeated as (80b).

- (80) a. $[e : \text{conductor}(d)]$
 b.
$$\left[\begin{array}{l} \text{restr} = \lambda r : [x : \text{Ind}] . [e : \text{conductor}(r.x)] \\ \text{scope} = \lambda r_2 : [x : \text{Ind}] . \left[\begin{array}{ll} x = r_2.x, d & : \text{Ind} \\ e & : \text{be}(x) \end{array} \right] \\ e \end{array} \right] \begin{array}{l} : \text{Ppty} \\ : \text{Ppty} \\ : \text{exist}(\text{restr}, \text{scope}) \end{array}$$

If we include the resource (76) then the content of *is Dudamel* is (81a) applied to (81b), that is, (81c).

- (81) a. $\lambda Q:Quant .$
 $\lambda r_1:[x:Ind] .$
 $Q(\lambda r_2:[x:Ind] . \left[\begin{array}{ll} x=r_2.x, r_1.x & : Ind \\ e & : be(x) \end{array} \right])$
- b. $\lambda P:Ppty . P([x=d])$
- c. $\lambda r_1:[x:Ind] . \left[\begin{array}{ll} x=d, r_1.x & : Ind \\ e & : be(x) \end{array} \right])$

The content of *A conductor is Dudamel* is (82a) applied to (82b) (identical with (81c)), which is (82c).

- (82) a. $\lambda P:Ppty . \left[\begin{array}{ll} restr=\lambda r:[x:Ind] . [e : conductor(r.x)] & : Ppty \\ scope=P & : Ppty \\ e & : exist(restr, scope) \end{array} \right]$
- b. $\lambda r_1:[x:Ind] . [x=d, r_1.x : Ind])$
- c. $\left[\begin{array}{ll} restr=\lambda r:[x:Ind] . [e : conductor(r.x)] & : Ppty \\ scope=\lambda r_1:[x:Ind] . \left[\begin{array}{ll} x=d, r_1.x & : Ind \\ e & : be(x) \end{array} \right]) & : Ppty \\ e & : exist(restr, scope) \end{array} \right]$

(82c) is almost exactly the same type as (80b). The difference between them is that d is the first restrictor of Ind in (82c) whereas in (80b) it is the second restrictor. In (82c) we have, for some conductor, c , $Ind_{d,c}$ whereas in (80b) we have $Ind_{c,d}$. Thus while an analysis that only uses the content of *is* that is based on Montague's original interpretation does predict different contents for *Dudamel is a conductor* and *a conductor is Dudamel*, the difference between the types hardly seems enough to explain the difference in reaction we have to the two sentences. Given the construction analysis for *Dudamel is a conductor* we get a markedly different type (80a) which does not involve existential quantification (even though it is truth conditionally equivalent to both the types with existential quantification). The only way that (80a) can be expressed according to the resources that we have developed in this chapter is by the sentence *Dudamel is a conductor*, using the non-compositional construction 'CnstrIsA'. Thus if (80a) is the target content and we do not wish to express a content involving existential quantification, *a conductor is Dudamel* is not an option.

We thus have the beginnings of an explanation of the difference in acceptability between the two sentences. It is not the whole story since we have not explained why the quantificational readings appear odd in these cases. Note that the distinction we are making between a non-quantified reading and a reading involving an existential quantification is not available on Montague's 1973 original approach since the fact that the two contents are truth-conditionally equivalent means

for Montague that they are identical. The same holds for the kind of analysis discussed in Partee (1986) where even though the content may not be built up using existential quantification the final result is still the same content that would be expressed by using existential quantification because of the truth-conditional equivalence. One might try to introduce the distinction we are making by relating utterances to an expression in an artificial logical language in addition to the content. This would correspond to the notion of logical form as discussed for example by Heim and Kratzer (1998) and much current work in linguistic semantics. The idea might be that there are two distinct logical forms such as (83) which correspond to identical contents in Montague's terms.

- (83) a. $\text{conductor}(\text{dudamel})$
 b. $\exists x [\text{conductor}(x) \wedge x = \text{dudamel}]$

Here the challenge would be to give an explanatory account of why one expression in an artificial language (83a) should be preferred over another (83b) when they both express the same content. An alternative is to follow Lewis (1972) (further developed by Cresswell, 1985). The idea here is that we keep a record not only of the final content but the way in which that content is constructed – that is we keep a record of the content of each of the syntactic constituents of the English sentence and the way these contents are combined. This idea, which goes back to the notion of intensional isomorphism introduced by Carnap (1956), provides enough structure to make the distinction required here. However, there are other problems with the proposal which we will take up in Chapter 6 when we discuss intensionality.

@@Do we really want to say that “ok” and “aha” don't have content?

Since acknowledgements like *aha* and *ok* do not have a specified content (*cf.* the function ‘ sign_{uc} ’ we used for these words in Chapter 2), we do not need a function that specifies their content but can make do with the function ‘Lex’ which associates them with a sign type in which the content is unspecified.

3.3 Building a chart type

In Chapter 2 we made the simplifying assumption that a chart was a sign. Now we have a grammar we need to complicate this picture. We will present here a version of chart parsing as it is used in computational linguistics. For a recent textbook introduction to chart parsing see Jurafsky and Martin (2009), Chap. 13. The idea of a chart is that it should store all the hypotheses that we make during the processing of an utterance and allow us to compute new hypotheses to be added to the chart on the basis of what is already present in the chart. We will say that a chart is a record and we will use our resources to compute a chart type on the basis of utterance events. We will first go through an example of the incremental construction of a chart type for an agent processing an utterance of the sentence *Dudamel is a conductor*. Then we will consider

what kind of update functions are needed in order to achieve this. We will, as usual, make the simplifying assumption that what we have at bottom is a string of word utterances as we are not dealing with the details of phonology. Thus we are giving a simplified view of incremental processing at the word level.

Suppose that we have so far heard an utterance of the word *Dudamel*. At this point we will say that the type of the chart is (84).

$$(84) \left[\begin{array}{ll} e_1 & : \text{“Dudamel”} \\ e & : [e_1:\text{start}(\uparrow^2 e_1)] \cap [e_1:\text{end}(\uparrow^2 e_1)] \end{array} \right]$$

The main event of the chart type (represented by the e -field) breaks the phonological event of type “Dudamel” down into a string of two events, the start and the end of the “Dudamel”-event.⁹

Why are the arguments to ‘start’ and ‘end’ in the string type prefixed by ‘ \uparrow^2 ’? Recall from the discussion on p. 15 that a string of type (85a) will be a record of type (85b).

$$(85) \begin{array}{ll} \text{a.} & [e_1:T_1] \cap [e_1:T_2] \\ \text{b.} & \left[\begin{array}{ll} t_0 & : [e_1:T_1] \\ t_1 & : [e_1:T_2] \end{array} \right] \end{array}$$

Thus a record of type (84) will be of the type (86).

$$(86) \left[\begin{array}{ll} e_1 & : \text{“Dudamel”} \\ e & : \left[\begin{array}{l} t_0 : [e_1:\text{start}(\uparrow^2 e_1)] \\ t_1 : [e_1:\text{end}(\uparrow^2 e_1)] \end{array} \right] \end{array} \right]$$

Thus the arguments to the ‘start’ and ‘end’ predicates are to be found two levels up.

Thus (84) records that we have observed an event of the phonological type “Dudamel” and an event consisting of the start of that event followed by the end of that event. Given that we have the resource the resource $\text{Lex}_{\text{PropName}}(\text{“Dudamel”, } d)$ available (see Appendix B), we can update (86) to (87).

⁹These starting and ending events correspond to what are standardly called *vertices* in the chart parsing literature.

$$(87) \left[\begin{array}{ll} e_1 & : \text{“Dudamel”} \\ e_2 & : \text{Lex}_{\text{PropName}}(\text{“Dudamel”, } d) \wedge [\text{s-event: } [e=e_1:\textit{Phon}]] \\ e & : \left[e_1:\text{start}(\uparrow^2 e_1) \right] \cap \left[e_1:\text{end}(\uparrow^2 e_1) \right] \end{array} \right]$$

That is, we add the information to the chart that there is an event (labelled ‘ e_2 ’) of the type which is the sign type corresponding to “Dudamel” and that the event which is the speech event referred to in that sign type is the utterance event, labelled by ‘ e_1 ’. Furthermore the duration of the event labelled ‘ e_2 ’ is the same as that labelled ‘ e_1 ’. One could discuss where there are two events which are contemporaneous or whether there is a single utterance event which is of both types. The fact that we have presented two fields labelled ‘ e_1 ’ and ‘ e_2 ’ does not of itself prevent the two fields containing the same event. However, the fact that we have analyzed the sign as containing the speech event as a part (corresponding to the basic intuition that signs are pairings of utterances and contents) decides the issue for us. A sign is a record (a labelled set) which models a situation and we are not allowing sets to be members of themselves. Thus records cannot be a part of themselves.¹⁰

The type $\text{Lex}_{\text{PropName}}(\text{“Dudamel”, } d)$ is a subtype of NP . Thus the event labelled ‘ e_2 ’ could be the first item in a string that would be appropriate for the function which we have abbreviated as (88a) (see Appendix B) which has the type (88b).

$$(88) \text{ a. } S \longrightarrow NP \ VP \mid NP'(VP') \\ \text{ b. } (NP \cap VP \rightarrow \textit{Type})$$

Thus in a way that is similar to the prediction by the dog in Chapter 1 that it should run after the stick which is help up and the kind of event that this will contribute to is a game of fetch so on observing a noun-phrase event we can predict that it might be followed by a verb phrase event thus creating a sentence event. We will add a hypothesis event to our chart which takes place at the end of the noun-phrase event as in (89).¹¹

¹⁰‘ e_1 ’ and ‘ e_2 ’ correspond to what are known as *passive edges* in the chart parsing literature. They represent information about potential constituents that have been found.

¹¹In terms of the traditional chart parsing terminology this corresponds to an *active edge* involving a *dotted rule*. The fact that the addition of this type to the chart type is triggered by finding something of an appropriate type to be the leftmost element in a string the would be an appropriate argument to the rule corresponds to what is called a *left-corner* parsing strategy.

$$(89) \left[\begin{array}{l} e_1 : \text{"Dudamel"} \\ e_2 : \text{Lex}_{\text{PropName}}(\text{"Dudamel"}, d) \wedge [s\text{-event}: [e=\uparrow^2 e_1 : \text{Phon}]] \\ e_3 : \left[\begin{array}{l} \text{rule} = S \rightarrow NP VP \mid NP'(VP') : (NP \frown VP \rightarrow \text{Type}) \\ \text{fnd} = \uparrow e_2 : \text{Sign} \\ \text{req} = VP : \text{Type} \\ e : \text{required}(\text{req}, \text{rule}) \end{array} \right] \\ e : \left[\begin{array}{l} [e_1 : \text{start}(\uparrow^2 e_1)] \\ [e_2 : \text{start}(\uparrow^2 e_2)] \end{array} \right] \frown \left[\begin{array}{l} e_1 : \text{end}(\uparrow^2 e_1) \\ e_2 : \text{end}(\uparrow^2 e_2) \\ e_3 : \text{start}(\uparrow^3 e_3) \frown \text{end}(\uparrow^3 e_3) \end{array} \right] \end{array} \right]$$

In the e_3 -field the ‘rule’-field is for a syntactic rule, that is, a function from a string of signs of a given type to a type. The ‘fnd’-field is for a sign or string of signs so far found which match an initial segment of a string of the type required by the rule. The ‘req’-field is the type of the remaining string required to satisfy the rule as expressed in the ‘e’-field. This hypothesis event both starts and ends at the end of the event of the noun-phrase event e_2 .¹²

We can now progress to the next word in the input string as shown in (90).

$$(90) \left[\begin{array}{l} e_1 : \text{"Dudamel"} \\ e_2 : \text{Lex}_{\text{PropName}}(\text{"Dudamel"}, d) \wedge [s\text{-event}: [e=\uparrow^2 e_1 : \text{Phon}]] \\ e_3 : \left[\begin{array}{l} \text{rule} = S \rightarrow NP VP \mid NP'(VP') : (NP \frown VP \rightarrow \text{Type}) \\ \text{fnd} = \uparrow e_2 : \text{Sign} \\ \text{req} = VP : \text{Type} \\ e : \text{required}(\text{req}, \text{rule}) \end{array} \right] \\ e_4 : \text{"is"} \\ e : \left[\begin{array}{l} [e_1 : \text{start}(\uparrow^2 e_1)] \\ [e_2 : \text{start}(\uparrow^2 e_2)] \end{array} \right] \frown \left[\begin{array}{l} e_1 : \text{end}(\uparrow^2 e_1) \\ e_2 : \text{end}(\uparrow^2 e_2) \\ e_3 : \text{start}(\uparrow^3 e_3) \frown \text{end}(\uparrow^3 e_3) \\ e_4 : \text{start}(\uparrow^2 e_4) \end{array} \right] \frown [e_4 : \text{end}(\uparrow^2 e_4)] \end{array} \right]$$

Note that the start of the “is”-event is aligned with the end of “Dudamel”-event. This allows for the fact that there is no break between the words and that the exact pronunciation of the final /l/ in “Dudamel” is influenced by the pronunciation of the initial /i/ in “is” through coarticulation.¹³

We can now go through similar procedures as we did for *Dudamel* adding both a lexical event based on our lexical resources and a hypothesis event based on the only rule for strings beginning with a *V* that we have in our resources. The result of these two steps is given in (91).

¹²With respect to the word string event labelled by ‘e’, it is a *punctual* event.

¹³It also means that the number of elements in the string labelled ‘e’ is the same as the number of vertices in a standard chart.

(91)

$$\begin{array}{l}
 \begin{array}{l}
 e_1 : \text{"Dudamel"} \\
 e_2 : \text{Lex}_{\text{PropName}}(\text{"Dudamel"}, d) \wedge [s\text{-event}: [e = \uparrow^2 e_1 : \text{Phon}]] \\
 e_3 : \left[\begin{array}{l} \text{rule} = S \rightarrow NP VP \mid NP'(VP') : (NP \frown VP \rightarrow Type) \\ \text{fnd} = \uparrow e_2 : \text{Sign} \\ \text{req} = VP : Type \\ e : \text{required}(\text{req}, \text{rule}) \end{array} \right] \\
 e_4 : \text{"is"} \\
 e_5 : \text{Lex}_{\text{be}}(\text{"is"}) \wedge [s\text{-event}: [e = \uparrow^2 e_4 : \text{Phon}]] \\
 e_6 : \left[\begin{array}{l} \text{rule} = VP \rightarrow [V \text{"is"}] [NP [Det \text{"a"}] N] \mid N' : \\ \quad (V \wedge [s\text{-event}: [e : \text{"is"}]] \frown \\ \quad \quad NP \wedge \left[\begin{array}{l} \text{syn}: \left[\begin{array}{l} \text{daughters}: Det \wedge [s\text{-event}: [e : \text{"a"}]] \\ \frown N \wedge [\text{cont}: Ppty] \end{array} \right] \\ \rightarrow Type \end{array} \right] \\ \text{fnd} = \uparrow e_5 : \text{Sign} \\ \text{req} = NP : Type \\ e : \text{required}(\text{req}, \text{rule}) \end{array} \right] \\
 e : \left[\begin{array}{l} [e_1 : \text{start}(\uparrow^2 e_1)] \frown [e_2 : \text{start}(\uparrow^2 e_2)] \frown \left[\begin{array}{l} e_1 : \text{end}(\uparrow^2 e_1) \\ e_2 : \text{end}(\uparrow^2 e_2) \\ e_3 : \text{start}(\uparrow^3 e_3) \frown \text{end}(\uparrow^3 e_3) \\ e_4 : \text{start}(\uparrow^2 e_4) \\ e_5 : \text{start}(\uparrow^2 e_5) \end{array} \right] \frown \left[\begin{array}{l} e_4 : \text{end}(\uparrow^2 e_4) \\ e_5 : \text{end}(\uparrow^2 e_5) \\ e_6 : \text{start}(\uparrow^3 e_6) \frown \text{end}(\uparrow^3 e_6) \end{array} \right] \end{array} \right]
 \end{array}
 \end{array}$$

Now we can add *a* and *conductor* in a similar way with the result shown in (92).

(92)

e_1	:	“Dudamel”
e_2	:	$\text{LexPropName}(\text{“Dudamel”, } d) \wedge [s\text{-event}: [e=\uparrow^2 e_1: \text{Phon}]]$
e_3	:	$\left[\begin{array}{l} \text{rule}=S \rightarrow NP VP \mid NP'(VP'): (NP \cap VP \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_2: \text{Sign} \\ \text{req}=VP: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_4	:	“is”
e_5	:	$\text{Lex}_{be}(\text{“is”}) \wedge [s\text{-event}: [e=\uparrow^2 e_4: \text{Phon}]]$
e_6	:	$\left[\begin{array}{l} \text{rule}=VP \rightarrow [V \text{ “is”}] [NP [Det \text{“a”}] N] \mid N': \\ \quad (V \wedge [s\text{-event}: [e: \text{“is”}]] \cap \\ \quad \quad NP \wedge [syn: [daughters: Det \wedge [s\text{-event}: [e: \text{“a”}]]] \cap \\ \quad \quad \quad N \wedge [cont: Ppty]]]) \cap \\ \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_5: \text{Sign} \\ \text{req}=NP: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_7	:	“a”
e_8	:	$\text{Lex}_{IndefArt}(\text{“a”}) \wedge [s\text{-event}: [e=\uparrow^2 e_7: \text{Phon}]]$
e_9	:	$\left[\begin{array}{l} \text{rule}=NP \rightarrow Det N \mid Det'(N'): (Det \cap N \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_8: \text{Sign} \\ \text{req}=N: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_{10}	:	“conductor”
e_{11}	:	$\text{Lex}_{CommonNoun}(\text{“conductor”, conductor}) \wedge [s\text{-event}: [e=\uparrow^2 e_{10}: \text{Phon}]]$
e	:	$\left[\begin{array}{l} [e_1: \text{start}(\uparrow^2 e_1)] \cap [e_2: \text{start}(\uparrow^2 e_2)] \cap \left[\begin{array}{l} e_1: \text{end}(\uparrow^2 e_1) \\ e_2: \text{end}(\uparrow^2 e_2) \\ e_3: \text{start}(\uparrow^3 e_3) \cap \text{end}(\uparrow^3 e_3) \\ e_4: \text{start}(\uparrow^2 e_4) \\ e_5: \text{start}(\uparrow^2 e_5) \end{array} \right] \cap \left[\begin{array}{l} e_4: \text{end}(\uparrow^2 e_4) \\ e_5: \text{end}(\uparrow^2 e_5) \\ e_6: \text{start}(\uparrow^3 e_6) \cap \text{end}(\uparrow^3 e_6) \\ e_7: \text{start}(\uparrow^2 e_7) \\ e_8: \text{start}(\uparrow^2 e_8) \end{array} \right] \cap \\ \left[\begin{array}{l} e_7: \text{end}(\uparrow^2 e_7) \\ e_8: \text{end}(\uparrow^2 e_8) \\ e_9: \text{start}(\uparrow^3 e_9) \cap \text{end}(\uparrow^3 e_9) \\ e_{10}: \text{start}(\uparrow^2 e_{10}) \\ e_{11}: \text{start}(\uparrow^2 e_{11}) \end{array} \right] \cap [e_{10}: \text{end}(\uparrow^2 e_{10})] \\ [e_{11}: \text{end}(\uparrow^2 e_{11})] \end{array} \right]$

Note that there is no possibility of adding a hypothesis event based on the utterance of *conductor* given the resources we have since our small grammar does not include a phrase structure rule for strings whose first element is of type N . However, now for the first time we have found something which fulfills one of our hypotheses. The hypothesis event labelled ‘ e_9 ’ has the type N in its ‘req’-field. The event labelled ‘ e_{11} ’ is required to be of a subtype of N and thus fulfils the requirement of ‘ e_9 ’. Furthermore, the start of e_{11} is aligned with the end (and also the start) of ‘ e_9 ’. This means that we can update the chart-type by adding a new field for an event of the

The event labelled ‘ e_{12} ’ will be of type *NP* and thus satisfy the requirement of e_6 . By carrying out the same procedure as before we will obtain a new event (labelled ‘ e_{13} ’) of type *VP* which will satisfy the requirement of ‘ e_3 ’ which will allow us to add a new event (labelled ‘ e_{14} ’) of type *S* whose start is at the beginning of the string labelled ‘*e*’ and whose end is at the end of that string. The final chart type is given in (94).

(94)

e_1	:	“Dudamel”
e_2	:	$\text{LexPropName}(\text{“Dudamel”, } d) \wedge [s\text{-event: } [e=\uparrow^2 e_1: \text{Phon}]]$
e_3	:	$\left[\begin{array}{l} \text{rule}=S \rightarrow NP VP \mid NP'(VP'): (NP \cap VP \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_2: \text{Sign} \\ \text{req}=VP: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_4	:	“is”
e_5	:	$\text{Lex}_{be}(\text{“is”}) \wedge [s\text{-event: } [e=\uparrow^2 e_4: \text{Phon}]]$
e_6	:	$\left[\begin{array}{l} \text{rule}=VP \rightarrow [V \text{ “is”}] [NP [Det \text{“a”}] N] \mid N': \\ \quad (V \wedge [s\text{-event: } [e:\text{“is”}]] \cap \\ \quad \quad NP \wedge \left[\text{syn: } \left[\begin{array}{l} \text{daughters: } Det \wedge [s\text{-event: } [e:\text{“a”}]] \\ \cap N \wedge [\text{cont: } Ppty] \end{array} \right] \right] \\ \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_5: \text{Sign} \\ \text{req}=NP: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_7	:	“a”
e_8	:	$\text{Lex}_{\text{IndefArt}}(\text{“a”}) \wedge [s\text{-event: } [e=\uparrow^2 e_7: \text{Phon}]]$
e_9	:	$\left[\begin{array}{l} \text{rule}=NP \rightarrow Det N \mid Det'(N'): (Det \cap N \rightarrow \text{Type}) \\ \text{fnd}=\uparrow e_8: \text{Sign} \\ \text{req}=N: \text{Type} \\ e: \text{required}(\text{req}, \text{rule}) \end{array} \right]$
e_{10}	:	“conductor”
e_{11}	:	$\text{Lex}_{\text{CommonNoun}}(\text{“conductor”, conductor}) \wedge [s\text{-event: } [e=\uparrow^2 e_{10}: \text{Phon}]]$
e_{12}	:	$e_9.\text{rule}(e_9.\text{fnd} \cap e_{11})$
e_{13}	:	$e_6.\text{rule}(e_6.\text{fnd} \cap e_{12})$
e_{14}	:	$e_3.\text{rule}(e_3.\text{fnd} \cap e_{13})$
e	:	$\left[\begin{array}{l} \left[\begin{array}{l} e_1: \text{start}(\uparrow^2 e_1) \\ e_2: \text{start}(\uparrow^2 e_2) \\ e_{14}: \text{start}(\uparrow^2 e_{14}) \end{array} \right] \cap \left[\begin{array}{l} e_1: \text{end}(\uparrow^2 e_1) \\ e_2: \text{end}(\uparrow^2 e_2) \\ e_3: \text{start}(\uparrow^3 e_3) \cap \text{end}(\uparrow^3 e_3) \\ e_4: \text{start}(\uparrow^2 e_4) \\ e_5: \text{start}(\uparrow^2 e_5) \\ e_{13}: \text{start}(\uparrow^2 e_{13}) \end{array} \right] \cap \left[\begin{array}{l} e_4: \text{end}(\uparrow^2 e_4) \\ e_5: \text{end}(\uparrow^2 e_5) \\ e_6: \text{start}(\uparrow^3 e_6) \cap \text{end}(\uparrow^3 e_6) \\ e_7: \text{start}(\uparrow^2 e_7) \\ e_8: \text{start}(\uparrow^2 e_8) \\ e_{12}: \text{start}(\uparrow^2 e_{12}) \end{array} \right] \cap \left[\begin{array}{l} e_7: \text{end}(\uparrow^2 e_7) \\ e_8: \text{end}(\uparrow^2 e_8) \\ e_9: \text{start}(\uparrow^3 e_9) \cap \text{end}(\uparrow^3 e_9) \\ e_{10}: \text{start}(\uparrow^2 e_{10}) \\ e_{11}: \text{start}(\uparrow^2 e_{11}) \end{array} \right] \cap \left[\begin{array}{l} e_{10}: \text{end}(\uparrow^2 e_{10}) \\ e_{11}: \text{end}(\uparrow^2 e_{11}) \\ e_{12}: \text{end}(\uparrow^2 e_{12}) \\ e_{13}: \text{end}(\uparrow^2 e_{13}) \\ e_{14}: \text{end}(\uparrow^2 e_{14}) \end{array} \right] \end{array} \right]$

We now need to turn our attention to the update functions that will achieve this building of the chart type. We will introduce a field ‘current-utterance’ into the field ‘shared’ on the game-board. This field will be used for the incremental construction of a chart during the course of

an utterance. We will not at this point include a ‘move’-field here but reserve that for the field ‘latest-utterance’, though one could, of course, consider an alternative with incremental hypotheses about moves that have or about to be made formed on the basis of the utterance so far. Here, however, we will restrict ourselves to the mechanisms involved in the construction of the chart. We will add a field to the gameboard ‘shared.current-utterance’ which will be used to store the chart during the course of processing an utterance. The new type *InfoState* is given in (95).

$$(95) \left[\begin{array}{l} \text{private:} [\text{agenda:} [RecType]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move:} Move \\ \text{chart:} RecType \\ \text{e:} m\text{-interp}(\text{chart}, \text{move}) \end{array} \right] \vee ERec \\ \text{current-utterance:} [\text{chart:} RecType] \\ \text{commitments:} RecType \end{array} \right] \end{array} \right]$$

The initial type *InitInfoState* is now (96).

$$(96) \left[\begin{array}{l} \text{private:} [\text{agenda}=[]:[RecType]] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} ERec \\ \text{current-utterance:} [\text{chart}=Rec:RecType] \\ \text{commitments}=Rec:RecType \end{array} \right] \end{array} \right]$$

We first address update functions for integrating lexical events into the chart. We introduce update functions defined by **IntegrateLexicalEvent**(T_{phon} , T_{chart}) where T_{phon} is the type the agent assigns to the phonological event perceived and T_{chart} is the type the agent assigns to the current chart. This is governed by the clauses in (97).

(97) a. If T_{phon} is a lexical phonological resource and T_{chart} is *Rec*, then

IntegrateLexicalEvent(T_{phon} , T_{chart}) is

$\lambda r: [\text{shared:} [\text{current-move:} [\text{chart:} T_{\text{chart}}]]]$

$\lambda u: T_{\text{phon}} .$

$\left[\text{shared:} \left[\text{current-move:} \left[\text{chart:} \left[\begin{array}{l} e_1: T_{\text{phon}} \\ e: [e_1:\text{start}(e_1)] \frown [e_1:\text{end}(e_1)] \end{array} \right] \right] \right] \right]$

b. If T_{phon} is a lexical phonological resource, T_{chart} is a record type such that ‘ e_n ’ is the maximal distinguished label ‘ e_i ’ in T_{chart} and T_{chart} is $T_1 \wedge [e: T_2 \frown T_3]$ where T_1 is a record type, T_2 is a string type and $T_3 \sqsubseteq [e_n: \text{end}(e_n)]$,

then **IntegrateLexicalEvent**(T_{phon} , T_{chart}) is

$\lambda r: [\text{shared:} [\text{current-move:} [\text{chart:} T_{\text{chart}}]]]$

$\lambda u: T_{\text{phon}} .$

$$\left[\text{shared:} \left[\text{current-move:} \left[\text{chart:} \left[\begin{array}{l} e_{n+1}:T_{\text{phon}} \\ e:T_2 \frown (T_3 \wedge [e_{n+1}:\text{start}(e_{n+1})]) \frown \\ [e_{n+1}:\text{end}(e_{n+1})] \end{array} \right] \right] \right] \right]$$

The licensing condition associated with chart update functions is the same as for other update functions (see Appendix C.1.4).

We now need update rules that will add signs to the chart which are derived from the lexical resources for signs associated with phonological types. For the lexical resources associated with this chapter in Appendix B.2.1 we will define the notion of a resource lexical sign type based on a phonological type as in (98).

- (98) T_{lex} is a *resource lexical sign type based on phonological type* T_{phon} according to a collection of resources R just in case T_{lex} is in R and is identical with either $\text{Lex}_{\text{PropName}}(T_{\text{phon}}, a)$, for some $a:\text{Ind}$, $\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p)$, for some predicate p , $\text{Lex}_{\text{IndefArt}}(T_{\text{phon}})$ or $\text{Lex}_{\text{be}}(T_{\text{phon}})$

We introduce update functions for integrating such lexical resources into the chart. These update functions are defined by **IntegrateLexicalResources**($T_{\text{phon}}, T_{\text{chart}}$) where T_{phon} is the type the agent assigns to the phonological event perceived and T_{chart} is the type the agent assigns to the current chart. This is governed by the clause in (99).

- (99) If
1. T_{phon} is a resource phonological type
 2. T_{event} is either $T_{\text{start}} \frown T_{\text{end}}$ or $T_{\text{evpref}} \frown T_{\text{start}} \frown T_{\text{end}}$ (where $T_{\text{start}} \sqsubseteq [e_k:\text{start}(e_k)]$, $T_{\text{end}} \sqsubseteq [e_k:\text{end}(e_k)]$ and T_{evpref} , “event prefix”, is a type)
 3. $T_{\text{chart}} \sqsubseteq \left[\begin{array}{l} e_k:T_{\text{phon}} \\ e:T_{\text{event}} \end{array} \right]$ whose maximal ‘ e_i ’ label is ‘ e_n ’
 4. T_{sign} is a resource lexical sign type based on T_{phon} such that for no j
 - a) $T_{\text{chart}} \sqsubseteq [e_j:T_{\text{sign}}]$,
 - b) $T_{\text{start}} \sqsubseteq [e_j:\text{start}(e_j)]$ and
 - c) $T_{\text{end}} \sqsubseteq [e_j:\text{end}(e_j)]$

then **IntegrateLexicalResources**($T_{\text{phon}}, T_{\text{chart}}$) is

$$\lambda r : \left[\text{shared:} \left[\text{current-move:} \left[\text{chart:} T_{\text{chart}} \right] \right] \right] . \\ \left[\text{shared:} \left[\text{current-move:} \left[\text{chart:} \left[\begin{array}{l} e_{n+1}:T_{\text{sign}} \\ e:T_{\text{newevent}} \end{array} \right] \right] \right] \right]$$

where T_{newevent} is

$$\begin{aligned}
& \textbf{either } (T_{\text{start}} \wedge [e_{n+1}:\text{start}(e_{n+1})]) \frown (T_{\text{end}} \wedge [e_{n+1}:\text{end}(e_{n+1})]) \\
& \textbf{or } T_{\text{evpref}} \frown (T_{\text{start}} \wedge [e_{n+1}:\text{start}(e_{n+1})]) \frown (T_{\text{end}} \wedge [e_{n+1}:\text{end}(e_{n+1})]) \\
& \text{depending on whether } T_{\text{event}} \text{ is } T_{\text{start}} \frown T_{\text{end}} \text{ or } T_{\text{evpref}} \frown T_{\text{start}} \frown T_{\text{end}}
\end{aligned}$$

There are several complexities in (99) which need some explanation. Firstly, notice that the update functions generated by **IntegrateLexicalResources** are of the form (100).

$$(100) \lambda r:T_1 . T_2$$

That is, they are *tacit* update functions which do not require a second event argument. They map directly from an information state of a certain type to a type for the new information state. An update using this update function is thus not driven by an agent-external event, merely by the state that the agent is currently in. An important issue in the design of tacit update functions is to develop mechanisms to prevent them from applying indefinitely many times adding the same information repeatedly and getting the agent carrying out the updates into a infinite loop. We will discuss how this has been avoided here below.

Condition 2 in (99) allows the ‘e’-field in the current chart to contain either a concatenation of just two events or to be a string of events ending in two events. The two final events of the event string are required to include the starting and ending respectively of some particular event labelled by ‘ e_k ’ (for some natural number k). Note that other things can also be going on in these two final events as indicated by the use of subtyping to characterize T_{start} and T_{end} here.

Condition 3 in (99) requires that the event labelled ‘ e_k ’ in the current information state is of the phonological type which we are going to use to construct the sign which we are going to add to the chart. We might have required ‘ e_k ’ to be the maximal ‘ e_i ’, that is, we might have required that the field labelled ‘ e_k ’ was the last to have been added. This would have been one way of avoiding an infinite loop, since once we have added the new field with the sign type, ‘ e_k ’ would no longer be maximal and **IntegrateLexicalResources** would not become applicable again until a further lexical event was entered into the chart. This would in fact have worked given the restricted collection of resources we are considering in this chapter, since they only allow for one sign type to be associated with any phonological type corresponding to a word. In general, this will not be the case since we want to allow for ambiguous words like *bank* and *can* to be associated with different sign types and we want to allow for all of the alternative sign types to be added to the same chart. For this reason we want **IntegrateLexicalResources** to apply even if ‘ e_k ’ is not maximal.

As condition 3 does not prevent looping we introduce a mechanism that will prevent it in condition 4 in (99) instead. This introduces a sign type based on the relevant phonological type which is going to be used for the update. But it requires that the sign type has not already be introduced

on the chart with the start and end of the sign at the end of the event string we are considering. (We do not wish to prevent it having been associated with a previous part of the event, of course, since the same word can occur more than once in an utterance.)

After integrating lexical sign types into the chart, the next step is to integrate rules from our resources that apply to strings which could begin with a lexical sign of this type. We will use **IntegrateRule**($f_{\text{rule}}, T_{\text{chart}}$) to generate such update rules. This is governed by the clause in (101).

(101) If

1. $T_{\text{chart}} \sqsubseteq \left[\begin{array}{l} \mathbf{e}_k : T_{\text{sign}} \\ \mathbf{e} : T_{\text{evpref}} \frown T_{\text{end}} \end{array} \right]$
 where:
 $T_{\text{sign}} \sqsubseteq T_{\text{cat}}$ (T_{cat} is one of *NP*, *VP*, ...)
 $T_{\text{end}} \sqsubseteq [\mathbf{e}_k : \text{end}(\mathbf{e}_k)]$
2. ' \mathbf{e}_i ' max in T_{chart} is ' \mathbf{e}_n '
3. $f_{\text{rule}} : ((T_1 \frown T_2 \frown \dots \frown T_m) \rightarrow \text{Type})$ where $T_{\text{sign}} \sqsubseteq T_1$
4. there is no l such that

$$T_{\text{chart}} \sqsubseteq \left[\begin{array}{l} \mathbf{e}_l : [\text{rule} = f_{\text{rule}} : ((T_1 \frown T_2 \frown \dots \frown T_m) \rightarrow \text{Type})] \\ \mathbf{e} : T_{\text{evpref}} \frown [\mathbf{e}_l : \text{start}(\uparrow^3 \mathbf{e}_l) \frown \text{end}(\uparrow^3 \mathbf{e}_l)] \end{array} \right]$$

then **IntegrateRule**($f_{\text{rule}}, T_{\text{chart}}$) is

$$\lambda r : T_{\text{chart}} . \left[\begin{array}{l} \mathbf{e}_k = r . \mathbf{e}_k : T_{\text{sign}} \\ \mathbf{e}_{n+1} : \left[\begin{array}{l} \text{rule} = f_{\text{rule}} : ((T_1 \frown T_2 \frown \dots \frown T_m) \rightarrow \text{Type}) \\ \text{fnd} = \uparrow \mathbf{e}_k : \text{Sign} \\ \text{req} = T_2 \frown \dots \frown T_m : \text{Type} \\ \mathbf{e} : \text{required}(\text{req}, \text{rule}) \end{array} \right] \\ \mathbf{e} : T_{\text{evpref}} \frown (T_{\text{end}} \frown [\mathbf{e}_{n+1} : \text{start}(\uparrow^3 \mathbf{e}_{n+1}) \frown \text{end}(\uparrow^3 \mathbf{e}_{n+1})]) \end{array} \right]$$

Condition 1 in (101) identifies a category field in the chart whose event is the latest event in the event string which has been processed. Condition 2 identifies the label of the latest addition to the chart so that it can be incremented for what is now going to be added. Condition 3 identifies a rule whose “left corner” (T_1) is a supertype of the type in the category field and Condition 4 requires that this rule has not already been added to the chart and related to the current final event in the event string — this in order to prevent an infinite loop. The result of **IntegrateRule**($f_{\text{rule}}, T_{\text{chart}}$) is then a function which adds a new field to the chart which contains a record of the rule, what has so far been found matching the “left corner” of the rule (that is, the category field that has been identified by Condition 1), and what is still required in order for the rule to be fully satisfied (that is, that is the type of strings required by the rule minus the “left corner”). Finally, the new

event is added as both starting and ending at the current end of the event string (that is, it does not extend the length of the event string, but the new event starts and ends simultaneously with the end of the event matching the “left corner” of the rule.¹⁴

The final kind of update functions that we need in order to build charts involves combining an event with a non-empty requirement with an event of a type matching the requirement whose start coincides with the end of the first event. In general there are two variants of such update functions that we need: one for the case where what is required is a string of category signs and one for the case where what is required is a single category sign. In the first case we need to create a new event with a requirement which is the remainder of the requirement after removing the left corner of the original requirement and a found string which concatenates the found event at the end of the original found event string. In the second case we need to add an event of the type which results from applying the rule to the concatenation of the found event to original found event string. As we only have binary rules in our small grammar the first case will not be necessary as we will only introduce a rule onto the chart when we have found an event matching its first element and the requirement result from this addition will thus be a single event of a given category type. We will thus only introduce update functions for the second case. We will use **Combine**($T_{\text{chart}}, \ell_1, \ell_2$) to generate such update rules. This is governed by the clause in (102).

(102) If

$$1. T_{\text{chart}} \sqsubseteq \left[\begin{array}{l} e_f : T_{\text{sign}_1} \\ \left[\begin{array}{l} \text{rule} = f_{\text{rule}} : (T \rightarrow \text{Type}) \\ e_k : \left[\begin{array}{l} \text{fnd} = \uparrow e_f : \text{Sign} \\ \text{req} = T_{\text{sign}_2} : \text{Type} \end{array} \right] \end{array} \right] \\ e_l : T_{\text{sign}_3} \\ e : T_1 \frown \left[\begin{array}{l} e_f : \text{start}(\uparrow^3 e_f) \\ e_f : \text{end}(\uparrow^3 e_f) \\ e_k : \text{start}(\uparrow^3 e_k) \frown \text{end}(\uparrow^3 e_k) \\ e_l : \text{start}(\uparrow^3 e_l) \\ e_l : \text{end}(\uparrow^3 e_l) \end{array} \right] \frown T_2 \frown T_3 \frown T_4 \end{array} \right]$$

where

T_{sign_1} , T_{sign_2} and T_{sign_3} are subtypes of one of NP , VP , \dots , that is, they are types of category signs.

$T_{\text{sign}_3} \sqsubseteq T_{\text{sign}_2}$

T is a type of strings of category signs

2. ‘ e_i ’ max in T_{chart} is ‘ e_n ’

3. There is no i such that

¹⁴The fact that the new event has a non-empty requirement for future events means that it corresponds to what is known in the chart parsing literature as an active edge and the new event encodes a dotted rule.

$$T_{\text{chart}} \sqsubseteq \left[\begin{array}{l} e_f : T_{\text{sign}_1} \\ e_l : T_{\text{sign}_3} \\ e_i : f_{\text{rule}}(e_f \frown e_l) \\ e : \text{Rec}^* \frown \left[\begin{array}{l} e_i : \text{start}(\uparrow e_i) \\ e_f : \text{start}(\uparrow e_f) \end{array} \right] \frown \text{Rec}^* \end{array} \right]$$

then **Compose**($T_{\text{chart}}, e_k, e_l$) is

$$\lambda r : T_{\text{chart}} . \left[\begin{array}{l} e_f : T_{\text{sign}_1} \\ e_k : \left[\begin{array}{l} \text{rule} = f_{\text{rule}} : (T \rightarrow \text{Type}) \\ \text{fnd} = \uparrow e_f : \text{Sign} \\ \text{req} = T_{\text{sign}_2} : \text{Type} \end{array} \right] \\ e_l : T_{\text{sign}_3} \\ e_{n+1} : r.e_k.\text{rule}(r.e_f \frown r.e_l) \\ e : T_1 \frown \left[\begin{array}{l} e_f : \text{start}(\uparrow^3 e_f) \\ e_{n+1} : \text{start}(\uparrow^3 e_{n+1}) \end{array} \right] \frown T_2 \frown \left[\begin{array}{l} e_k : \text{start}(\uparrow^3 e_k) \frown \text{end}(\uparrow^3 e_k) \\ e_l : \text{start}(\uparrow^3 e_l) \end{array} \right] \\ \frown T_3 \frown \left[\begin{array}{l} e_l : \text{end}(\uparrow^3 e_l) \\ e_{n+1} : \text{end}(\uparrow^3 e_{n+1}) \end{array} \right] \frown T_4 \end{array} \right]$$

Condition 1 in (102) requires that the chart to be updated has a rule event labelled ' e_k ' where the found event is labelled ' e_f ' and that there is an event ' e_l ', starting at the end of ' e_f ', simultaneously with ' e_k '. The type specified for ' e_l ' must be a subtype of the type identified as the required type in ' e_k ', that is T_{sign_2} .

Condition 2 identifies the maximum event index in the chart as n .

Condition 3 requires that the result of applying the rule to the event string $e_f \frown e_l$ has not already been added to the chart. (This will prevent the creation of an infinite loop.)

The resulting update function **Compose**($T_{\text{chart}}, e_k, e_l$) is a function which adds a new event field labelled ' e_{n+1} ' for an event of the type returned by applying the rule to the event string consisting of the found event, ' e_f ', followed by the required event, ' e_l '. Event ' e_{n+1} ' starts at the beginning of ' e_f ' and ends at the end of ' e_l '.

3.4 Summary

In this chapter we have explored how the type theoretical apparatus developed in Chapters 1 and 2 can be applied to the notion of grammar, viewing grammatical phenomena in terms of event perception and information state update. While we have included both syntax and semantics in this framework and taken a fairly detailed look at how incremental parsing can be incorporated in this approach, the actual grammatical phenomena that we have looked at are linguistically trivial.

In Part II we will look at a variety of linguistic phenomena and argue that this approach provides theoretically interesting insights into the way that they function in dialogue.

Part II

Towards a dialogical view of semantics

Chapter 4

Proper names, salience and accommodation

4.1 Montague's PTQ as a semantic benchmark

In this chapter and the following chapters we will extend the linguistic coverage of the toy grammar we presented in Chapter 3. We will take Montague's PTQ (Montague, 1973, 1974) as providing a benchmark of linguistic phenomena that need to be covered and try to cover a sizeable part of what Montague covered, although we will add a few things which are obviously closely related to Montague's original benchmark and which have been treated subsequently in the literature.

For many of the phenomena we discuss we will first present a treatment which is as close as possible to Montague's original treatment and then present a treatment which exploits the advantages of the approach we are proposing in this book as well as more recent developments since Montague's original work. Our aim is to show that we have something to say about all these phenomena in an overall consistent framework, that is, to show that we can cover a significant part of the benchmark using the tools we are proposing and in many cases say something new concerning a dialogical approach to these phenomena. In doing this within the space of a single book we will not be able to cover all the aspects of these phenomena which have been studied in the literature following after Montague. We hope, however, to show that it is a fruitful line of research to add a rich type theoretic perspective and a dialogical approach to current work in linguistic semantics.

4.2 Montague's treatment of proper names and a sign-based approach

The treatment of proper names that we presented in Chapter 3, encapsulated in the definition of SemPropName and $\text{Lex}_{\text{PropName}}$ in Appendix B, is an adaptation of Montague's original treatment in that it has the content of a proper name utterance as a quantifier generated from an individual. The essence of Montague's treatment was that if we have a proper name *Sam* whose denotation is based on an individual 'sam', then the denotation of *Sam* is the characteristic function of the set of properties possessed by the individual concept of 'sam'. Montague modelled individual concepts as functions from possible worlds to individuals. Using more or less Montague's logical notation, the denotation of *Sam* would be represented by (1).

$$(1) \quad \lambda P.P\{[\hat{\text{sam}}]\}$$

Here $[\hat{\text{sam}}]$ represents the individual concept of 'sam', that is, that function, f , on the set of possible worlds such that for any world w , $f(w) = \text{sam}$. The reason that Montague used the individual concept (and the associated special notion of application involved in applying a property to an individual concept represented by the ' $\{\}$ '-brackets) was to treat what is known as the Partee-puzzle concerning temperature and price which we will discuss in Chapter 5. Many subsequent researchers came to the conclusion that Montague's treatment of this puzzle was not the correct one and that the individual concept was not necessary in the treatment of proper names. Thus (1) could be simplified to (2).

$$(2) \quad \lambda P.P(\text{sam})$$

The content that we assigned to an utterance of *Sam* in Chapter 3 is represented in (3).

$$(3) \quad \lambda P:P_{\text{pty}}.P([\text{x}=\text{sam}])$$

The reason that we have chosen to characterize properties as having records as their domain rather than individuals, has to do with our treatment of the Partee puzzle as we will explain in Chapter 5. Thus the reason that we have the record $[\text{x}=\text{sam}]$ as the argument to the property rather than an individual as in (4) is for the same reason as Montague introduced an individual concept.

$$(4) \quad \lambda P:P_{\text{pty}}.P(\text{sam})$$

The treatment of proper names we presented in Chapter 3 has an important advantage over Montague's original. For Montague, (1) is the result of applying an interpretation function to the linguistic expression *Sam* and a number of indices for the interpretation, \mathfrak{A} , a possible world, i , a time, j , and an assignment to variables, g . This is represented in (5).

$$(5) \quad \llbracket \text{Sam} \rrbracket^{\mathfrak{A}, i, j, g} = \lambda P. P\{\hat{\text{sam}}\}$$

This requires that the English expression *Sam* is always associated with the same individual 'sam' with respect to \mathfrak{A} and any i, j, g related to \mathfrak{A} . This seems to go against the obvious fact that more than one individual can have the name *Sam*. It does not work to say that a different individual can be associated with *Sam* when it is evaluated with respect to different parameters. g is irrelevant since it is defined as an assignment to variables and the English expression *Sam* is not (associated with) a variable — it cannot be bound by a quantifier.¹ A strategy which involves varying the possible world and time to get a different individual associated with *Sam* would be defeated by the fact that there are many people called Sam in the actual world right now as well as having the unintuitive consequence that *Sam might be Sam* would be true if it is true that Sam might be somebody else called Sam and *Sam will be Sam* could be true if somebody called Sam now is somebody else called Sam at a future time. We might try saying that associating a different individual with Sam involves a different interpretation, \mathfrak{A}' , of the language. This has some intuitive appeal and we will discuss a variant of it in Section 4.5 in relation to a recent proposal by Ludlow (2014). But it will come to grief when we need to talk about two people named Sam in the same sentence unless we allow a switch in interpretation mid-sentence. While allowing interpretation to change mid-sentence may be an attractive option for other reasons it is not an option that is available on Montague's account of meaning. The normal assumption is that in cases where two individuals have the same name the language contains two expressions which are pronounced the same, for example, *Sam*₁ and *Sam*₂. This would make the treatment of proper names somewhat like Montague's treatment of pronouns in that they have silent numerical subscripts attached to them. How many *Sam*_{*i*} should the language contain? One for each person named Sam, now, in the past and future and who could be named Sam in some non-actual world? If we follow the strategy with variables we would introduce countably many *Sam*_{*i*} so that we would always have enough. But with assignments to variables we can always assign individuals to more than one variable without this causing a problem. But the consequence of doing this with proper names would be to say that an individual can have many names that are pronounced the same. (Sam says, "My name is Sam", not "My names are Sam".) Similarly no two individuals would have the same name, although they would be able to have distinct names which are pronounced the same. This would mean that the interpretation of *have the same name* would have to mean "have names which are pronounced the same". This might cause difficulties distinguishing between a case where we have two people named Sam and a case where people really do have distinct names which are pronounced the same such as *Ann* and *Anne* (unless you want to count this as a case of spelling the same name differently).

¹This claim has been called into question by more recent research. See Maier (2009) for discussion.

In contrast the analysis of proper names we presented in Chapter 3 is sign-based. It allows several sign types to share the same phonology but be associated with different contents. Treating the language in terms of signs eliminates the need for arbitrary indexing of proper names. It also allows us to individuate names in a sensible way. One way to individuate names is by the phonologies occurring in proper name sign types. Thus if we have two proper name sign types with the same phonology but contents associated with different individuals, then we have two individuals with the same name. Note that this proposal would make *Ann* and *Anne* different spellings of the same name since they are both associated with the same phonological type. How we individuate names can be different in different contexts if we follow the kind of proposal for counting discussed by Cooper (2011). We could, for example, introduce a field into lexical sign types for an orthographical type and allow the individuation of names by either phonology or orthography or a combination of both depending on what is most useful to the purpose at hand.

Using signs in this way seems to give us a clear, if rather simple, advantage over Montague's formal language approach, even though we have so far essentially just transplanted Montague's analysis of proper names into our variant of a sign-based approach. However, there is a remaining question within sign-based approaches which is a kind of correlate to the need on Montague's approach to create many different names Sam_i . We are tempted to think of a "language" as being defined as a collection of sign types. Thus a person who knows English will know sign types which pair the phonological type "Sam" with various individuals who are called Sam. The problem with this is that different speakers of English will know different people named Sam and thus technically we would have to say that they speak different languages. This may well be a coherent technical notion of language. In the terminology of Chapter 3 we would say that the two agents indeed have different linguistic resources available to them. But there is also a resource which the two agents share, even if they do not have any overlap in the people named Sam that they are aware of. This is the knowledge that *Sam* is a proper name in English and can be used to name individuals. Arguably it is this knowledge which is constitutive of English, rather than the knowledge of who is actually called Sam, important though that might be for performing adequately in linguistic situations. In Chapter 3 we introduced sign type construction operations and in particular $\text{Lex}_{\text{PropName}}$ which maps a phonological type and an individual to an appropriate proper name sign type (see Appendix B). We called this a universal resource since it represents the general knowledge that utterances can be used to name individuals. In the English resources we defined there we named sign types such as $\text{Lex}_{\text{PropName}}(\text{"Sam"}, \text{sam})$, where we specify both the phonological type and the individual associated with it. But, given the power of functional abstraction, we can identify (6) as an English resource where the phonological type is specified but not the particular individual.

$$(6) \quad \lambda x:\text{Ind} . \text{Lex}_{\text{PropName}}(\text{"Sam"}, x)$$

Saying that an agent has this function available as an English resource could be argued to encode the fact that the agent has the knowledge that *Sam* is a proper name in English. An agent who has this resource has a recipe for constructing an appropriate sign type in their resources whenever

they meet somebody called Sam. Knowing that *Sam* is a proper name in English is not a matter of knowing who is called Sam but rather a matter of knowing what to do linguistically when you encounter somebody called Sam. Thus while we have so far just taken over Montague's original analysis of proper names we have given ourselves the opportunity to recast it in terms of a theory which enables agents to update their linguistic resources as they become aware of new facts about the world.

4.3 Proper names and communication

However, what we have done so far tells us little about the communicative processes associated with utterances of proper names. In Cooper (2013b) we pointed out that this kind of analysis does not give us any way of placing the requirement on the interlocutor's gameboard that there already be a person named Sam available in order to integrate the new information onto the gameboard. As Ginzburg (2012) points out, the successful use of a proper name to refer to an individual *a* requires that the name be publically known as a name for *a*. We will follow the analysis of Cooper (2013b) in parametrizing the content. A *parametric content* is a function which maps a context to a content. As such it relates to Montague's technical notion of *meaning* in his paper 'Universal Grammar' (Montague, 1970, 1974) where he regarded meaning as a function from possible worlds and contexts of use to denotations. This also corresponds to the notion of *character* in Kaplan (1978).

We will take a context to be a situation modelled as a record. A simple proposal for a parametric content for a proper name might be (7).

$$(7) \quad \lambda r: [x:Ind] . \lambda P:Ppty . P(r)$$

This would allow any record with an individual labelled 'x' to be mapped to a proper name content. Recall that the label 'x' is picked up by the notion of property that we defined in Chapter 3 as being of type $([x:Ind] \rightarrow RecType)$, an example being (8).

$$(8) \quad \lambda r: [x:Ind] . [e:run(r.x)]$$

Associating the phonological type "Sam" with (7) would essentially be a way of encapsulating in the interpretation of *Sam* what is expressed by (6) — namely, that potentially any individual can be called Sam. We want the parametric content of *Sam* to be more restrictive than this. It is going to be the tool that we use to help us identify an appropriate referent when we are confronted with an utterance of type "Sam". The obvious constraint that we should place is that the referent is indeed named Sam. Thus we can restrict (7) so that it is an appropriate parametric content for

Sam rather than something that appears to be a parametric content appropriate to proper names in general. The modification is given in (9).

$$(9) \quad \lambda r: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \cdot \\ \lambda P:P_{pty} . P(r)$$

This is closely related to treatments of proper names that were proposed earlier in situation semantics (Gawron and Peters, 1990; Cooper, 1991; Barwise and Cooper, 1993). A more recent close relation is Maier's (2009) proposal for the treatment of proper names in terms of layered discourse representation theory (LDRT). Maier points out in a useful overview of the history of semantic treatments of proper names that this view of proper names is a hybrid of the descriptivist and referential approaches: it uses a description like "named Sam" to provide a presuppositional restriction on the kind of referent which can be assigned to the proper name. (9) maps a context in which there is an individual named Sam to a proper name content based on that individual. Care has to be taken with the predicate 'named' on this kind of analysis. It is important that it not be too restrictive, for example, requiring the legal registering of the name. It may be sufficient that someone at some point has called the individual by the name. The exact conditions under which a situation may be of a type constructed with this predicate will vary depending on the needs associated with the conversation at hand. We may, for example, take a stricter view of what it means to have a certain name if we are talking in a court of law than if we are trying to attract somebody's attention to avoid an accident on a mountainside. This flexibility of meaning "in flux" has been discussed in Cooper and Kempson (2008); Cooper (2012b); Ludlow (2014); Ginzburg and Cooper (2014); Kracht and Klein (2014) among many other places and we will return to it several times in the following chapters.

An alternative to the use of parametric contents is to use parametric signs. This could be formulated as in (10) where $\text{Lex}_{\text{PropName}}$ is the function for associating lexical content with phonological types that was introduced in Chapter 3 and summarized in Appendix B.1.4.

$$(10) \quad \lambda r: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \cdot \\ \text{Lex}_{\text{PropName}}(\text{"Sam"}, r.x)$$

Intuitively, (10) says that given a situation in which there is an individual named by the phonological type "Sam" we can construct a sign type in which the phonological type "Sam" is associated with that individual. From the point of view of the formal semantics tradition (10) is a much more radical proposal than (9). The function (9) is a close relative of Montague's *meaning* and Kaplan's *character*. It is a function from contexts to contents, although our theory of what contexts and contents are differs from both Montague's and Kaplan's proposals. The function in

(10), however, is something that creates a kind of linguistic resource on the basis of a context. That is, given a context in which ‘sam’ is named by “Sam” we derive the information that linguistic signs can be used which associate “Sam” with ‘sam’. If we did not know this before we are extending the collection of linguistic resources we have available. We suspect that both parametric contents and parametric sign types could be of importance for a theory of linguistic interpretation and learning. For now, we will work with the less radical notion of parametric content.

Parametric contents as we have presented them so far are problematic for compositional semantics because the domain type of the function (representing the “presupposition”) which is the parametric content varies from case to case depending on what the intuitive presupposition of the phrase is. According to our rules it will always be some subtype of *RecType* (since we are thinking of contexts as records/situations) but it would not be possible to state a single type of parametric content for proper names or other syntactic categories. For this reason we will say that a parametric content is a pair (that is, a record with two fields) containing a type and a function whose domain type is that type. We can create such a parametric content by using a redefined version of ‘SemPropName’ which we introduced in Chapter 3, see Appendix B.1. Whereas the version from Chapter 3 took an individual as argument and created the content of a name of that individual, the new version will take a phonological type as argument and create a parametric content requiring an individual named by that phonological type. The new version is given in (11).

(11) $\text{SemPropName}(T)$, where T is a phonological type, is

$$\left[\begin{array}{lcl} \text{bg} & = & \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, T) \end{array} \right] \\ \text{fg} & = & \lambda r: \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, T) \end{array} \right] \cdot \\ & & \lambda P:P_{\text{pty}} . P(r) \end{array} \right]$$

Here the field labelled ‘bg’ (“background”) contains a record type and the field labelled ‘fg’ (“foreground”) is a function whose domain type is the background record type. From now on we will mean records of this kind by *parametric content*. The type of a parametric content of proper names is thus (12).

$$(12) \left[\begin{array}{lcl} \text{bg} & : & \text{RecType} \\ \text{fg} & : & (\text{bg} \rightarrow \text{Quant}) \end{array} \right]$$

That is, the foreground is a function from records of the background type (modelling contexts) to quantifiers. We will refer to this type as *PQuant* (“parametric quantifiers”). The universal resource $\text{Lex}_{\text{PropName}}$ for associating proper name content with phonological types, creating a

sign type for a proper name (see Appendix B.1), will now be redefined so that it only takes a phonological type as argument as in (13).

- (13) $\text{Lex}_{\text{PropName}}(T_{\text{Phon}})$, where T_{Phon} is a phonological type,
is defined as
 $\text{Lex}(T_{\text{Phon}}, NP) \wedge [\text{cnt}=\text{SemPropName}(T_{\text{Phon}}):PQuant]$

Note that the phonological type plays a dual role here. It figures once as determining the phonology of the sign and again as determining the presupposition associated with the parametric content.

There are two main questions that need to be answered about parametric contents. One concerns how the compositional semantics works and the other concerns the nature of contexts and how you compute with them. We will take the compositionality issue first. Let us assume that all signs provide us with a parametric content rather than a content. In those cases where there is no constraint on what the context must be we will use a trivial parametric content, that is, one that maps any context (modelled as a record) to the same content. Thus, for example, if we wish to represent a theory in which the intransitive verb *leave* does not place any restrictions on the context, we could represent its parametric content as (14a) which is of the type for parametric properties (*PPpty*) given in (14b).

- (14) a.
$$\left[\begin{array}{lcl} \text{bg} & = & \text{Rec} \\ \text{fg} & = & \lambda r_1:\text{Rec}.\lambda r_2:[x:\text{Ind}]. [e:\text{leave}(r_2.x)] \end{array} \right]$$

b.
$$\left[\begin{array}{lcl} \text{bg} & : & \text{RecType} \\ \text{fg} & : & (\text{bg} \rightarrow \text{Ppty}) \end{array} \right]$$

The foreground of this parametric property will map any context r_1 to the function $\lambda r_2:[x:\text{Ind}]. [e:\text{leave}(r_2.x)]$ which does not depend in any way on r_1 . Such a content could be introduced by a resource for lexical content construction ‘*SemIntransVerb*’ as characterized in (15), where T_{bg} , the “background” or “presupposition” type, is a record type and p is a predicate with arity $\langle \text{Ind} \rangle$.

- (15) $\text{SemIntransVerb}(T_{\text{bg}}, p)$ is
$$\left[\begin{array}{lcl} \text{bg} & = & T_{\text{bg}} \\ \text{fg} & = & \lambda r_1:T_{\text{bg}}.\lambda r_2:[x:\text{Ind}]. [e : p(r_2.x)] \end{array} \right]$$

Note that if (15) is the only way of constructing parametric content for lexical intransitive verbs, then although it is possible to place restrictions on the context by choosing a non-trivial record type (something other than *Rec*) for T_{bg} this will not have any effect on the property returned as

the content. As we are not here concerned with presuppositions introduced by lexical intransitive verbs we will leave open whether it is necessary to change this. ‘SemIntransVerb’ will be used by the universal resource ‘Lex_{IntransVerb}’ defined in (16), where T_{phon} is a phonological type and p is a predicate with arity $\langle \text{Ind} \rangle$.

$$(16) \text{ Lex}_{\text{IntransVerb}}(T_{\text{phon}}, T_{\text{bg}}, p) \\ \text{is defined as} \\ \text{Lex}(T_{\text{phon}}, N) \wedge [\text{cnt}=\text{SemIntransVerb}(T_{\text{bg}}, p):PP\text{pty}]$$

This means that the English resource corresponding to the lexical entry for *leave* can be defined as (17).

$$(17) \text{ Lex}_{\text{IntransVerb}}(\text{“leave”}, \text{Rec}, \text{leave})$$

A standard strategy for dealing with compositional semantics when using parametric contents is to use a version of what is known in combinatorial logic as the S-combinator. In its λ -calculus version this is (18).

$$(18) \lambda z. \alpha(z)(\beta(z))$$

Our version of the S-combinator including different type requirements on the context arising from the function and the argument will be (19).

$$(19) \text{ If } \alpha : \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow (T_1 \rightarrow T_2)) \end{array} \right] \text{ and } \beta : \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow T_1) \end{array} \right] \text{ then the combination of } \alpha \text{ and } \beta \\ \text{based on functional application, } \alpha @ \beta, \text{ is}$$

$$\left[\begin{array}{ll} \text{bg} & = \left[\begin{array}{l} \text{f:}\alpha.\text{bg} \\ \text{a:}\beta.\text{bg} \end{array} \right] \\ \text{fg} & = \lambda r: \left[\begin{array}{l} \text{f:}\alpha.\text{bg} \\ \text{a:}\beta.\text{bg} \end{array} \right] . \alpha.\text{fg}(r.\text{f})(\beta.\text{fg}(r.\text{a})) \end{array} \right]$$

Note that in the background for the result we have kept the backgrounds of α and β separated in their own fields labelled ‘f’ (“function”) and ‘a’ (“argument”).² This means that we avoid an

²While textually this statement of the combination will be correct, we need to take account of the fact that the abbreviatory notation for labels in argument positions to predicates now represent path-names in $\alpha.\text{bg}$ and $\beta.\text{bg}$ to which the labels ‘f’ and ‘a’ have been prefixed respectively. To be precise we could notate this as $[\alpha.\text{bg}]^f$ and $[\beta.\text{bg}]^a$.

unwanted clash of labels if $\alpha.bg$ and $\beta.bg$ should happen to share labels.³ We could use (19) to combine the contents (9) and (14). The foreground of the result is given in (20) where we can show by successive applications of β -reduction that (20a–d) all represent the same function.

$$\begin{aligned}
 (20) \text{ a. } & \lambda r_1: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] . \\
 & \quad (\lambda r_2: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] . \lambda P:Ppty . P(r_2))(r_1.f) \\
 & \quad (\lambda r_3:Rec . \lambda r_4: \left[\begin{array}{l} x:Ind \end{array} \right] . [e:leave(r_4.x)](r_1.a)) \\
 \text{b. } & \lambda r_1: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] . \\
 & \quad \lambda P:Ppty . P(r_1.f) \\
 & \quad (\lambda r_4: \left[\begin{array}{l} x:Ind \end{array} \right] . [e:leave(r_4.x)]) \\
 \text{c. } & \lambda r_1: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] . \\
 & \quad \lambda r_4: \left[\begin{array}{l} x:Ind \end{array} \right] . [e:leave(r_4.x)](r_1.f) \\
 \text{d. } & \lambda r_1: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] . \\
 & \quad [e:leave(r_1.f.x)]
 \end{aligned}$$

(20) represents the foreground of the parametric content of *Sam leaves*. Given a situation containing an individual, a , named by “Sam” it returns a type of situation in which a leaves. As usual this type can play the role of a “proposition”. It is, for example, “true” if there is a situation of the type and “false” if there is no situation of the type.

The background of the parametric content, that is the domain type of its foreground, is to be thought of as placing a constraint on the context. The idea is that you can only get to the non-parametric content if you have an appropriate situation available. The background of the parametric content is a type which represents a kind of *presupposition*. We shall treat presuppositions as constraints on the resources available to dialogue participants. In Chapter 2 we introduced the notion of a dialogue gameboard as a type of dialogue information state. The most obvious place to look for the referent of an utterance of a proper name is in the shared commitments represented

³This new method of combination for parametric contents means that we also have to adjust the sign combination operation `CntForwardApp` (“forward application of contents”) used in the definition of interpreted phrase structure rules. See Appendix B.1.4.2 for details.

on the gameboard representing what has been committed to in the dialogue so far. If an individual named Sam has already been introduced in the dialogue, then a subsequent utterance of *Sam* in that dialogue is most likely to refer to that individual unless there is an explicit indication to the contrary. The shared commitments on an agent's dialogue gameboard represent information that is particularly *salient* to the agent. The notion of salience in semantics was first introduced by Lewis (1979b) in connection with the analysis of definite descriptions. As Lewis says, "There are various ways for something to gain salience. Some have to do with the course of conversation, others do not." We wish to suggest that a way of gaining salience in a conversation is by figuring in the shared commitments on the gameboard. (Ginzburg, 2012, argues that being on shared commitments, or FACTS in his terminology, is not always sufficient to indicate salience.)

A reasonable strategy, then, is to look at the shared commitments on the dialogue gameboard first and then look elsewhere if that fails. We will first explore what we need to do to match the background type of a parametric content against the type which models the shared commitments of the dialogue and then we will discuss what needs to be done if there is not a successful match with the shared commitments. In Chapter 2 we treated the gameboard as a record type. In Chapter 2, example (74), for instance, the shared commitments were represented as the type (21).

$$(21) \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \textit{Rec} \\ \text{e:conductor(dudamel)} \end{array} \right] \\ \text{e:composer(beethoven)} \end{array} \right] \\ \text{e:pianist(uchida)} \end{array} \right]$$

Recall that with each successive updating of the shared commitments the record type representing the previous state of shared commitments was embedded under the label 'prev' ("previous"). This prevented label clash and also kept a record of the order in which information was introduced. As Lewis (1979b) observed, information introduced later in the dialogue tends to be more salient than information introduced earlier. Thus keeping track of the order also gives us one measure of relative salience.

In Chapter 2 we were using the Montague treatment of proper names that did not introduce the naming predicate. In this chapter we will work towards shared commitments where the naming associated with proper names is made explicit, as in (22).

$$(22) \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \text{Rec} \\ \text{bg:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{"Dudamel"}) \end{array} \right] \\ \text{fg:} [e:\text{conductor}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{"Beethoven"}) \end{array} \right] \\ \text{fg:} [e:\text{composer}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{"Uchida"}) \end{array} \right] \\ \text{fg:} [e:\text{pianist}(\uparrow\text{bg}.x)] \end{array} \right] \end{array} \right]$$

Here we are using the label ‘bg’ to represent background information in the manner suggested by Larsson (2010) and we see also that this labelling corresponds to our use of ‘bg’ and ‘fg’ in parametric contents. Note that in this version of the shared commitments we have lost the connection with the actual individuals ‘dudamel’, ‘beethoven’ and ‘uchida’. This can be seen as an advantage if we are representing the information state of an agent in the kind of situation described in Chapter 2. If we simply inform an agent with no previous knowledge of Dudamel that Dudamel is a conductor, then the information that this agent will get is that there is somebody named Dudamel who is a conductor. There will be no connection to a particular individual of whom the agent is aware. If this is not the case, we can reinstate the connection to the individuals by using manifest fields to anchor the information as in (23).

$$(23) \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \text{Rec} \\ \text{bg:} \left[\begin{array}{l} x=\text{dudamel:Ind} \\ e:\text{named}(x, \text{"Dudamel"}) \end{array} \right] \\ \text{fg:} [e:\text{conductor}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} x=\text{beethoven:Ind} \\ e:\text{named}(x, \text{"Beethoven"}) \end{array} \right] \\ \text{fg:} [e:\text{composer}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} x=\text{uchida:Ind} \\ e:\text{named}(x, \text{"Uchida"}) \end{array} \right] \\ \text{fg:} [e:\text{pianist}(\uparrow\text{bg}.x)] \end{array} \right] \end{array} \right]$$

The ‘bg’-fields in (22) can be thought of as corresponding to the internal anchors of Kamp (1990); Kamp *et al.* (2011). The use of manifest fields in (23) would then correspond to the association of what they call external anchors with those internal anchors.

The task we have before us is to try to match the domain type of the function in (20), that is, the type which is the background of the parametric content, repeated in (24), against the types of shared commitments in (22) or (23).

$$(24) \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:named(x, "Sam") \end{array} \right] \\ a:Rec \end{array} \right]$$

Intuitively, this attempt at matching should fail since there is no commitment to an individual named Sam in the shared commitments. Suppose now that we add to (22) as in (25).

$$(25) \left[\begin{array}{l} prev: \left[\begin{array}{l} prev: \left[\begin{array}{l} prev:Rec \\ bg: \left[\begin{array}{l} x:Ind \\ e:named(x, "Dudamel") \end{array} \right] \\ fg: [e:conductor(\uparrow bg.x)] \end{array} \right] \\ bg: \left[\begin{array}{l} x:Ind \\ e:named(x, "Beethoven") \end{array} \right] \\ fg: [e:composer(\uparrow bg.x)] \end{array} \right] \\ bg: \left[\begin{array}{l} x:Ind \\ e:named(x, "Uchida") \end{array} \right] \\ fg: [e:pianist(\uparrow bg.x)] \end{array} \right] \\ bg: \left[\begin{array}{l} x:Ind \\ e:named(x, "Sam") \end{array} \right] \\ fg: [e:singer(\uparrow bg.x)] \end{array} \right]$$

Intuitively, this should enable a match since this does commit to an individual named Sam. However, there is not a direct formal relationship between (24) and (25) corresponding to this intuition. We will use relabelling of record types (see Appendix A.14) in order to capture the relationship.

In order to match (24) against (25) we look for a relabelling, η , of (24) that would make (25) be a subtype of (24). Such a relabelling is given in (26a) which we will write as (26b). and the result of applying it to (24) is given in (26c).

(26) a. η is a function with domain $\{f.x, f.e, a\}$ such that

$$\eta(f.x) = bg.x$$

$$\eta(f.e) = bg.e$$

$$\eta(a) = prev^4$$

where $prev^4$ stands for $prev.prev.prev.prev$

b. $f.x \rightsquigarrow bg.x$

$$f.e \rightsquigarrow bg.e$$

$$a \rightsquigarrow prev^4$$

$$c. \left[\begin{array}{l} \text{bg} \\ \text{prev} \end{array} : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Sam"}) \end{array} \right] \right]$$

$$\left[\begin{array}{l} \text{prev} : \left[\text{prev} : \left[\text{prev} : \left[\text{prev} : \text{Rec} \right] \right] \right] \end{array} \right]$$

This means, then, that any situation which is of the type required by the shared commitments would, modulo the relabelling, be of the type which is the background of the parametric content under consideration, spelled out in (27).

$$(27) \left[\begin{array}{l} \text{bg} \\ \text{fg} \end{array} = \left[\begin{array}{l} \text{f:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{"Sam"}) \end{array} \right] \\ a:\text{Rec} \end{array} \right] \right]$$

$$\left[\begin{array}{l} \text{fg} \\ \text{fg} \end{array} = \lambda r: \left[\begin{array}{l} \text{f:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{"Sam"}) \end{array} \right] \\ a:\text{Rec} \end{array} \right] \right] . \left[e:\text{leave}(r.f.x) \right]$$

The background of the parametric content is being used as a presupposition which is being matched against the hearer's current information state.

Given that we have now found a match, how can we go about updating the shared commitments with the new information represented by the parametric content?

If we are updating (25) with the parametric content (27a) then the result should be (28) where (25) has been embedded under the label 'prev' and the new information provided by the parametric content has been added at the top level of the new type, suitably relabelled so as to pick out the individual named Sam which has been previously introduced.

$$(28) \quad \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \left[\begin{array}{l} \text{prev:} \text{Rec} \\ \text{bg:} \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Dudamel")} \end{array} \right] \\ \text{fg:} \left[\text{e:conductor}(\uparrow \text{bg.x}) \end{array} \right] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Beethoven")} \end{array} \right] \\ \text{fg:} \left[\text{e:composer}(\uparrow \text{bg.x}) \end{array} \right] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Uchida")} \end{array} \right] \\ \text{fg:} \left[\text{e:pianist}(\uparrow \text{bg.x}) \end{array} \right] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Sam")} \end{array} \right] \\ \text{fg:} \left[\text{e:singer}(\uparrow \text{bg.x}) \end{array} \right] \end{array} \right] \\ \text{bg:} \left[\begin{array}{l} \text{f:} \left[\begin{array}{l} \text{x} = \uparrow^2 \text{prev.bg.x:Ind} \\ \text{e} = \uparrow^2 \text{prev.bg.e:named(x, "Sam")} \end{array} \right] \\ \text{a} = \uparrow \text{prev}^5 \text{:Rec} \end{array} \right] \\ \text{fg:} \left[\text{e:leave}(\uparrow \text{bg.f.x}) \end{array} \right] \end{array} \right]$$

Note that this both achieves a link to a previous mention of Sam and simultaneously ensures that Sam is the most salient individual in shared commitments in virtue of the new mention.

We can achieve this update by the following method. Suppose that T_{comm} is the type representing shared commitments that we wish to update with a parametric content given in (29).

$$(29) \quad \left[\begin{array}{l} \text{bg} = T_{\text{bg}} \\ \text{fg} = f \end{array} \right]$$

where $f : (T_{\text{bg}} \rightarrow \text{RecType})$. We need to find a relabelling, η , of T_{bg} such that $T_{\text{comm}} \sqsubseteq [T_{\text{bg}}]_{\eta}$. Suppose that we have a record, r_{comm} , of type T_{comm} . We can use this together with η to anchor T_{bg} . The result of the anchoring is notated as $T_{\text{bg}} \parallel_{\eta} r_{\text{comm}}$. The operation $T \parallel_{\eta} r$ (defined explicitly in Appendix A.15) replaces fields in T , $[\ell:T']$, such that ℓ is in the domain of the relabelling, η , and for which η returns a path, π , in r such that $r.\pi : T'$ with a manifest field $[\ell=r.\pi:T']$. For example, suppose that T_{bg} is (30a), $r_{\text{comm}} : T_{\text{comm}}$ and η is (30b). Then $T_{\text{bg}} \parallel_{\eta} r_{\text{comm}}$ is (30c).

$$(30) \quad \begin{array}{l} \text{a.} \left[\begin{array}{l} \text{x} : \text{Ind} \\ \text{e} : \text{named(x, "Sam")} \end{array} \right] \\ \text{b. } \text{x} \rightsquigarrow \text{bg.x} \\ \quad \text{e} \rightsquigarrow \text{bg.e} \\ \text{c.} \left[\begin{array}{l} \text{x} = r_{\text{comm}}.\text{bg.x} : \text{Ind} \\ \text{e} = r_{\text{comm}}.\text{bg.e} : \text{named(x, "Sam")} \end{array} \right] \end{array}$$

In the type of the updated shared commitments the background will be the background of the parametric content anchored to the previous shared commitments and the foreground will be the result of applying the function which is the foreground of the parametric content to this background. The updated type of the shared commitments will thus be that given in (31).

$$(31) \quad \left[\begin{array}{ll} \text{prev} & : T_{\text{comm}} \\ \text{bg} & : T_{\text{bg}} \parallel_{\eta} \text{prev} \\ \text{fg} & : f(\text{bg}) \end{array} \right]$$

We can tie all this together in a single update function, given in (32), a preliminary version which we will revise slightly later in the light of other accommodation functions which we will introduce in Section 4.4.

(32) **AccGB**(η) – preliminary version

$$\begin{aligned} & \lambda T: \text{GameBoard} . \\ & \quad \lambda f: \left[\begin{array}{l} \text{bg}: \text{RecType} \\ \text{fg}: (\text{bg} \rightarrow \text{RecType}) \end{array} \right] . \\ & \quad \lambda r: T . \\ & \quad T \left[\bigwedge \right] \left[\text{shared}: \left[\text{commitments} = \left[\begin{array}{l} \text{prev}: r.\text{shared.commitments} \\ \text{bg}: f.\text{bg} \parallel_{\eta} \text{prev} \\ \text{fg}: f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \right] \right] \end{aligned}$$

This function takes a game-board, T , (recall that a gameboard is a type of an information state which in turn is a record) and a parametric content and returns a function that will map an information state of type T to a new type which is the result of an asymmetric merge of T with a type that will replace the type representing shared commitments according to T with a new type where the old shared commitments is labelled with ‘prev’ and new ‘bg’ and ‘fg’ fields are added as described above.

4.4 Proper names, salience and accommodation

What we have presented so far enables us to find a match for presuppositions introduced by a parametric content when such a match is present in shared commitments. Suppose there is more than one such match. In that case there will be a choice of relabellings η . In this case we may wish to choose the relabelling that corresponds to a match with the most salient match in terms of recency of introduction into the shared commitments. Technically, this means that we choose the relabelling which introduces labels with the least number of occurrences of ‘prev’. Note that the most recent match may be anchored to a match that was introduced earlier in the manner we have just described. There may be other factors than recency which contribute to salience, for example, the kinds of factors that are discussed in centering theory (Joshi and Weinstein, 1981;

Grosz *et al.*, 1983, 1995; Walker *et al.*, 1998; Poesio *et al.*, 2004). We will leave it to future work to give a more detailed account of saliency in the current framework.

What happens when there is no match for *Sam* in the shared commitments? Here we need some kind of accommodation in order to use the parametric content to update the gameboard. There are two kinds of accommodation we will consider. The first is where the agent knows of a person named *Sam* independently of the current conversation. That is, a match for *Sam* can be found in the agent's resources corresponding to long term memory. We will not attempt a detailed account of the structure of long term memory. We assume that it is complex and constantly in flux not only in terms of new information being added but also in terms of what is salient in the old information, depending on which part of the memory is being focussed on at any particular time. Here we will content ourselves with a simple model of long term memory as a record type of a similar kind to that we have proposed for shared commitments. This means that the techniques we need for matching will be the same as those discussed above. In reality the notion of salience with respect to long term memory will be a good deal more complicated than salience with respect to the shared commitments on the dialogue gameboard. You have to take into account not only recency but also likelihood based on other knowledge that it is this particular *Sam* that is being referred to. For example, if you believe that your interlocutor could not possibly know of the *Sam* in your memory who is otherwise the most likely candidate you should not choose that *Sam* as a match. Choosing an appropriate match involves a great deal of world knowledge and common sense. We will ignore these matters and concentrate our attention on what needs to be done if we find a suitable match. The idea is that if you have failed to find a match in shared commitments on the gameboard but you do find a match in long term memory, then you need to load the item from long term memory into the shared commitments on your gameboard. This is what will constitute accommodation in this case.

We will introduce the notion of a *total information state* (cf. Larsson, 2002) which includes a record type corresponding to long term memory, represented by the 'ltm'-field in (33) and a dialogue gameboard, represented by the 'gb'-field in (33). Up until now we have thought of the gameboard as a record type. Now, however, we want to be able to make links from the gameboard to long term memory and we will achieve this by making the gameboard be a dependent type which maps records (situations) of the type representing long term memory to the record type representing the gameboard. Thus a total information state will be of the type (33).

$$(33) \quad \left[\begin{array}{ll} \text{ltm} & : \text{RecType} \\ \text{gb} & : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right]$$

Here we use *GameBoard* as the type of types which are a subtype of *InfoState* (as defined in Appendix C.1.1), that is, a gameboard is a type of information states. Formally, this is expressed as in (34).

(34) $T : \text{GameBoard}$ iff $T \sqsubseteq \text{InfoState}$

An example of a type corresponding to long term memory is given in (35).

(35)
$$\left[\begin{array}{l} \text{id}_0: \text{Rec} \\ \text{id}_1: \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Dudamel")} \end{array} \right] \\ \text{id}_2: \left[\text{e:conductor}(\uparrow \text{id}_1.\text{x}) \right] \\ \text{id}_3: \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Beethoven")} \end{array} \right] \\ \text{id}_4: \left[\text{e:composer}(\uparrow \text{id}_3.\text{x}) \right] \\ \text{id}_5: \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Uchida")} \end{array} \right] \\ \text{id}_6: \left[\text{x:pianist}(\uparrow \text{id}_5.\text{x}) \right] \\ \text{id}_7: \left[\begin{array}{l} \text{x:Ind} \\ \text{e:named(x, "Sam")} \end{array} \right] \\ \text{id}_8: \left[\text{e:singer}(\uparrow \text{id}_7.\text{x}) \right] \end{array} \right]$$

(35) is one way of putting the information in shared commitments represented by (??) into a type corresponding to long term memory. We are assuming that in long term memory information is indexed by unique identifiers modelled here by the labels ‘ id_n ’ (of which we assume there is a countably infinite stock, one for each natural number, n). It is important that in long term memory paths are persistent under updating, that is, the old paths do not change when we add information to long term memory. This is in contrast to the kind of updating we proposed for the gameboard, adding the label ‘prev’ to the path for the old gameboard. This meant that all paths within the old gameboard were adjusted by an update. When we link from the gameboard to long term memory we want to make sure that the link uses a persistent path which will still be correct if the long term memory should get updated. When long term memory is updated we prefix the path to the new information with the identifier ‘ id_{i+1} ’, where i is the highest index on an ‘id’-label in the long term memory type we are updating. (This is the same technique we used for ‘e’-labels in our treatment of chart parsing in Chapter 3.) The way of achieving the link is illustrated schematically in (36) where we use M to represent the long term memory (35) and leave out all irrelevant details of the gameboard.

(36)

$$\left[\begin{array}{l} \text{ltm} = M : \text{RecType} \\ \text{gb} = \lambda r : \text{ltm} . \left[\begin{array}{l} \dots \\ \text{shared} : \left[\begin{array}{l} \text{commitments} = \left[\begin{array}{l} \dots \\ \text{bg} : \left[\begin{array}{l} x = r.\text{id}_7.x : \text{Ind} \\ e = r.\text{id}_7.e : \text{named}(x, \text{"Sam"}) \end{array} \right] : \text{RecType} \\ \text{fg} : [e : \text{leave}(\uparrow \text{bg}.x)] \end{array} \right] \end{array} \right] \end{array} \right] : \end{array} \right] \end{array} \right] : \end{array} \right] \quad (\text{ltm} \rightarrow \text{RecType})$$

The intuition expressed in (36) is as follows: given a situation, r , of the type represented by our long term memory, that is one in which a particular appropriate individual is labelled by ‘id₇’, the gameboard will be a type of information state where the background of the parametric content used to update the shared commitments is anchored to ‘id₇’. Two agents are aligned in their shared commitments to the extent that we can find an equivalence between the two types which represent their respective view of the shared commitments obtained by applying their respective functions labelled ‘gb’ to a situation of their respective memory types.

The link represented by the dependence on the long term memory type corresponds to what Kamp (1990); Kamp *et al.* (2011) call an internal anchor. We are representing here how individual roles in an agent’s view of shared commitments can be anchored in that agent’s long term memory. In a more complete treatment we could in addition make the gameboard depend on a type for the current visual scene and also types for other sensory input. Our use of dependent types and Kamp *et al.*’s use of internal anchors allow us to link different components of cognitive structure. Cognitive structure can also be linked to objects in the external world, giving rise to what Kamp *et al.* call external anchors. Our manifest fields can be used to correspond to their external anchors. Suppose, for example, that we have an individual ‘sam’ who is named Sam. We can use a manifest field to restrict the long term memory type (35) so that any record (“situation”) of that type has ‘sam’ in the ‘id₇.x’-field. This is represented in (37) where for convenience we have omitted all but the ‘id₇’-field in (35).

$$(37) \quad \left[\begin{array}{l} \dots \\ \text{id}_7 : \left[\begin{array}{l} x = \text{sam} : \text{Ind} \\ e : \text{named}(x, \text{"Sam"}) \end{array} \right] \\ \dots \end{array} \right]$$

If M in (36) is the type (37) then for any $r : M$, it will be the case that $r.\text{id}_7.x$ will be ‘sam’. Thus the shared commitment is that ‘sam’ leaves. Given that manifest fields can occur in any record type, this kind of external anchoring is not restricted to long term memory but could also be directly in the gameboard if that is desired.

Let us now consider how the update of a gameboard dependent on long term memory can be carried out when there is a match between the parametric content used for updating and an item in long term memory. Suppose that the current total information state, ι_{curr} , is of the type in (38)

$$(38) \quad \left[\begin{array}{ll} \text{ltm} & : \text{RecType} \\ \text{gb}=\lambda r:\text{ltm} . T_{\text{gb}}(r) & : (\text{ltm} \rightarrow \text{RecType}) \end{array} \right]$$

and that we wish to update this with the parametric content, f , given in (39) (where $T_{\text{bg}} \sqsubseteq [x:\text{Ind}]$).

$$(39) \quad \left[\begin{array}{ll} \text{bg} & = T_{\text{bg}} \\ \text{fg} & = \lambda r:T_{\text{bg}} . T_{\text{upd}}(r) \end{array} \right]$$

In order to find a match between $f.\text{bg}$, that is, T_{bg} and $\iota_{\text{curr}}.\text{ltm}$ (that is, to ascertain that the presupposition associated with the parametric content is met by the long term memory of the current total information state) we need to find a relabelling, η , of T_{bg} such that (40) holds.

$$(40) \quad \iota_{\text{curr}}.\text{ltm} \sqsubseteq [T_{\text{bg}}]_{\eta}$$

Then we can derive (41) as a type of the updated total information state.

$$(41) \quad \left[\begin{array}{l} \text{ltm}=\iota_{\text{curr}}.\text{ltm}:\text{RecType} \\ \text{gb}=\lambda r:\text{ltm} . (T_{\text{gb}}(r) \sqcap \left[\begin{array}{l} \text{shared:} \left[\begin{array}{l} \text{commitments}=\left[\begin{array}{l} \text{prev}:\iota_{\text{curr}}.\text{gb}.\text{shared}.\text{commitments} \\ \text{bg}:T_{\text{bg}} \parallel_{\eta} r \\ \text{fg}:f(\text{bg}) \end{array} \end{array} \right] : \text{RecType} \end{array} \right] \end{array} \right]) \end{array} \right] : (\text{ltm} \rightarrow \text{RecType}) \end{array} \right]$$

Here the notation $T_{\text{bg}} \parallel_{\eta} r$ represents the specification or anchoring of the type T_{bg} by the record r according to the relabelling η . That is, we replace fields in T_{bg} with manifest fields according to the matches we have in the ‘ltm’-field. Thus, for example, if T_{bg} is (24), repeated as (42a), r is a record representing long term memory of type (35), repeated as (42b) and η is the relabelling in (42c), then $T_{\text{bg}} \parallel_{\eta} r$ is (42d).

$$(42) \quad \text{a.} \quad \left[\begin{array}{l} \text{f:} \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, \text{“Sam”}) \end{array} \right] \\ \text{a:Rec} \end{array} \right]$$

$$\begin{array}{l}
\text{b. } \left[\begin{array}{l} \text{id}_0: \text{Rec} \\ \text{id}_1: \left[\begin{array}{l} x: \text{Ind} \\ e: \text{named}(x, \text{"Dudamel"}) \end{array} \right] \\ \text{id}_2: \left[e: \text{conductor}(\text{id}_1.x) \right] \\ \text{id}_3: \left[\begin{array}{l} x: \text{Ind} \\ e: \text{named}(x, \text{"Beethoven"}) \end{array} \right] \\ \text{id}_4: \left[e: \text{composer}(\text{id}_3.x) \right] \\ \text{id}_5: \left[\begin{array}{l} x: \text{Ind} \\ e: \text{named}(x, \text{"Uchida"}) \end{array} \right] \\ \text{id}_6: \left[x: \text{pianist}(\text{id}_5.x) \right] \\ \text{id}_7: \left[\begin{array}{l} x: \text{Ind} \\ e: \text{named}(x, \text{"Sam"}) \end{array} \right] \\ \text{id}_8: \left[e: \text{singer}(\text{id}_7.x) \right] \end{array} \right] \\
\text{c. } f.x \rightsquigarrow \text{id}_7.x \\
\quad f.e \rightsquigarrow \text{id}_7.e \\
\quad a \rightsquigarrow \text{id}_0 \\
\text{d. } \left[\begin{array}{l} f: \left[\begin{array}{l} x=r.\text{id}_7.x: \text{Ind} \\ e=r.\text{id}_7.e: \text{named}(x, \text{"Sam"}) \end{array} \right] \\ a=r.\text{id}_0: \text{Rec} \end{array} \right]
\end{array}$$

A precise and general definition of this notation is in Appendix A.15.

We can now put all this together as the update function in (43), which we call **AccLTM**(η) (“accommodate match with long term memory”).

$$\begin{array}{l}
(43) \quad \mathbf{AccLTM}(\eta) = \\
\quad \lambda r: \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] \\
\quad \lambda f: \left[\begin{array}{l} \text{bg} : \text{RecType} \\ \text{fg} : (\text{bg} \rightarrow \text{RecType}) \end{array} \right] \cdot \\
\quad \left[\begin{array}{l} \text{ltm} = r.\text{ltm}: \text{RecType} \\ \text{gb} = \lambda r_1: \text{ltm} . ((r.\text{gb})(r_1) \boxed{\wedge} \\ \quad \left[\begin{array}{l} \text{shared: } \left[\begin{array}{l} \text{commitments} = \left[\begin{array}{l} \text{prev}: (r.\text{gb})(\uparrow^3 \text{ltm}).\text{shared.commitments} \\ \text{bg}: f.\text{bg} \parallel_{\eta} r_1 \\ \text{fg}: f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \end{array} \right] \end{array} \right] : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right]
\end{array}$$

Here *GameBoard* is as defined in (34).

We have used accommodation from long term memory to represent the kind of accommodation where the agent has a resource which provides a match. In a more complete treatment we could use this technique for accommodation from other available resources such as the visual scene. We now turn our attention to accommodation where there is no appropriate match with other resources. This corresponds to the case where the hearer does not know any appropriate person named Sam but merely adds that there is a person named Sam to the shared dialogue commitments.

The first step in this update is to create a type from the parametric content under consideration so that we can merge it with $[\text{prev}:T]$, where T is the type representing the current shared commitments. Suppose we are considering the parametric content, ξ , given in (44a). Then the type we will create from ξ is defined as in (44b) which is identical with (44c).

$$\begin{aligned}
 (44) \quad a. \quad \xi &= \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] . [e:\text{leave}(r.f.x)] \end{array} \right] \\
 b. \quad \left[\begin{array}{l} \text{bg} : \xi.\text{bg} \\ \text{fg} : [e : \xi.\text{fg}(\text{bg})] \end{array} \right] \\
 c. \quad \left[\begin{array}{l} \text{bg}: \left[\begin{array}{l} f: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Sam"}) \end{array} \right] \\ a:Rec \end{array} \right] \\ \text{fg}: [e:\text{leave}(\text{bg}.f.x)] \end{array} \right]
 \end{aligned}$$

Suppose now that the current shared commitments are given by the type in (45).

$$(45) \quad \left[\begin{array}{l} \text{prev}: \left[\begin{array}{l} \text{prev}:Rec \\ \text{bg}: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Dudamel"}) \end{array} \right] \\ \text{fg}: [e:\text{conductor}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg}: \left[\begin{array}{l} x:Ind \\ \text{fg}: [e:\text{named}(x, \text{"Beethoven"})] \\ \text{fg}: [e:\text{composer}(\uparrow\text{bg}.x)] \end{array} \right] \\ \text{bg}: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Uchida"}) \end{array} \right] \\ \text{fg}: [e:\text{pianist}(\uparrow\text{bg}.x)] \end{array} \right]$$

Then the new shared commitments will be (46a) which is (46b).



We can now put this together as the update function in (47), which we call **AccNM** (“accommodate no match”). [???? Fix formatting]

(47) **AccNM** =

$$\lambda r: \left[\begin{array}{ll} \text{ltm} & : \text{RecType} \\ \text{gb} & : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] \\ \lambda f: \left[\begin{array}{ll} \text{bg} & : \text{RecType} \\ \text{fg} & : (\text{bg} \rightarrow \text{RecType}) \end{array} \right].$$

$$\left[\begin{array}{l} \text{ltm} = r.\text{ltm} : \text{RecType} \\ \text{gb} = \lambda r_1 : \text{ltm} . ((r.\text{gb})(r_1)) \boxed{\wedge} \\ \left[\text{shared} : \left[\text{commitments} = \left[\begin{array}{l} \text{prev} : (r.\text{gb})(\uparrow^3 \text{ltm}).\text{shared}.\text{commitments} \\ \text{bg} : f.\text{bg} \\ \text{fg} : f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \right] \right] \end{array} \right] : (\text{ltm} \rightarrow \text{GameBoard})$$

This is the same as **AccLTM** in (43) except that in the update for shared commitments there is no anchoring to long term memory.

We can now adjust the preliminary version of **AccGB** given in (32) which was the update function for cases where there is a match on the gameboard so that it is accommodated in the general format of update functions for total information states.

(48) **AccGB**(η) – final version

$$\lambda r : \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] . \\ \lambda f : \left[\begin{array}{l} \text{bg} : \text{RecType} \\ \text{fg} : (\text{bg} \rightarrow \text{RecType}) \end{array} \right] . \\ \left[\begin{array}{l} \text{ltm} = r.\text{ltm} : \text{RecType} \\ \text{gb} = \lambda r_1 : \text{ltm} . r.\text{gb}(r_1) \boxed{\wedge} \\ \left[\text{shared} : \left[\text{commitments} = \left[\begin{array}{l} \text{prev} : r.\text{gb}.\text{shared}.\text{commitments} \\ \text{bg} : f.\text{bg} \parallel_{\eta} \text{prev} \\ \text{fg} : f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \right] \right] \end{array} \right] : (\text{ltm} \rightarrow \text{RecType})$$

The three update functions for accommodation that we have defined are governed by the single licensing condition given in (49).

(49) If A is an agent, s_i is A 's current information state, f is a parametric content of type T_f such that

$$T_f \sqsubseteq \left[\begin{array}{l} \text{bg} : \text{RecType} \\ \text{fg} : (\text{bg} \rightarrow \text{RecType}) \end{array} \right]$$

and $s_i :_A T_i$ for some T_i such that

$$T_i \sqsubseteq \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : \left[\text{shared} : \left[\begin{array}{l} \text{commitments} : \text{RecType} \\ \text{latest-move} : [\text{cont} = f : T_f] \end{array} \right] \right] \end{array} \right]$$

then

if there is some η which is a relabelling of $f.bg$ such that

$$s_i.gb.shared.commitments \sqsubseteq [f.bg]_\eta$$

then $s_{i+1} :_A T_i \sqcap \mathbf{AccGB}(\eta)(s_i)(f)$ is licensed

else if there is some η which is a relabelling of $f.bg$ such that $s_i.ltm \sqsubseteq [f.bg]_\eta$

then $s_{i+1} :_A T_i \sqcap \mathbf{AccLTM}(\eta)(s_i)(f)$ is licensed

else $s_{i+1} :_A T_i \sqcap \mathbf{AccNM}(s_i)(f)$ is licensed

This account of accommodation for proper names where a new item is allowed to be created in memory when attempts at matching have failed is similar to a proposal by de Groote and Lebedeva (2010) to treat accommodation as error handling when a match has failed to be found. Our information states can be thought of as corresponding to their environment which they consider to be not simply a list of individuals but individuals with their properties, thus providing objects similar to those like the record types which can be found in our information states. One difference between the two proposals, apart from the obvious fact that our aim here has been to embed the theory in a more general theory of dialogue, is that de Groote and Lebedeva use a selection function to select the matches thus apparently assuming an algorithm which yields a unique result. We, on the other hand, talk in terms of matches being licensed and thereby allow for the possibility of non-deterministic selection. What we have in common, though, is that in order to account for the way accommodation is carried out we both add an additional layer to a type theory based semantics and talk in procedural terms of actions to be carried out: we with our licensing conditions for type acts and de Groote and Lebedeva with their error handling mechanism.

4.5 Paderewski

Kripke (1979) discusses the case of Peter who hears about a pianist called Paderewski. Later, in a different context, he learns of a Polish national leader and Prime Minister called Paderewski. In reality there was a single (remarkable) man called Paderewski who was both a famous concert pianist and a distinguished statesman. But Peter does not realize this and thinks that he has learned about two distinct people, both named Paderewski. Thus, in our terms, Peter's long term memory might be a subtype of (50) for some natural numbers i, j, k and l .

$$(50) \quad \left[\begin{array}{l} id_i : \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Paderewski"}) \end{array} \right] \\ id_j : \left[\begin{array}{l} e:pianist(\uparrow id_i.x) \end{array} \right] \\ id_k : \left[\begin{array}{l} x:Ind \\ e:named(x, \text{"Paderewski"}) \end{array} \right] \\ id_l : \left[\begin{array}{l} e:statesman(\uparrow id_k.x) \end{array} \right] \end{array} \right]$$

(50) technically allows for the two Paderewskis to be the same individual but if there is nothing in Peter's long term memory that requires them to be the same individual we will count that as corresponding to his view of them as distinct. If Peter were in this state and asked whether the pianist Paderewski and the statesman Paderewski were the same person Peter might reply, "Well, I wouldn't have thought so, but I suppose they could be the same person. I don't know." On being told that the two Paderewskis are in fact the same person he might update his long term memory by carrying out the merge in (51a), that is, his long term memory would now be (51b).

$$(51) \quad \begin{array}{l} \text{a.} \quad \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_j: \left[e:\text{pianist}(\uparrow\text{id}_i.x) \right] \\ \text{id}_k: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_l: \left[e:\text{statesman}(\uparrow\text{id}_k.x) \right] \end{array} \right] \hat{\wedge} \left[\begin{array}{l} \text{id}_i: \left[x:Ind \right] \\ \text{id}_k: \left[x=\uparrow\text{id}_i.x:Ind \right] \end{array} \right] \\ \\ \text{b.} \quad \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_j: \left[e:\text{pianist}(\uparrow\text{id}_i.x) \right] \\ \text{id}_k: \left[\begin{array}{l} x=\uparrow\text{id}_i.x:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_l: \left[e:\text{statesman}(\uparrow\text{id}_k.x) \right] \end{array} \right] \end{array}$$

Eventually, his long term memory may be restructured to the type in (52) which is set equivalent to that in (51), though not multiset equivalent to it since in any record of this type the individual named Paderewski will only occur once, not twice as in (51).

$$(52) \quad \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_j: \left[e:\text{pianist}(\uparrow\text{id}_i.x) \right] \\ \text{id}_l: \left[e:\text{statesman}(\uparrow\text{id}_i.x) \right] \end{array} \right]$$

We might think of the two types (51b) and (52) as representing two subtly different states of mind which Peter could be in. In (51b) he has two concepts of Paderewski, one concept associated with him being a pianist and perhaps other associated properties, such as practicing hard, wearing tails when he is performing, and so on and the other concept where he is a statesman, and perhaps associated with other properties such as being a dynamic national leader, a driver of hard political bargains or whatever. In (52) he has a single concept of Paderewski including all he knows about him. The first state is perhaps a natural one to be in after just learning that the two Paderewskis are in fact the same, before you have fully assimilated the identity. It is harder to discover contradictions between the two concepts here since it will only be the manifest field linking the two concepts which will reveal the contradiction. Suppose, for example, Peter's concept of

the statesman Paderewski has him always late for appointments and pressed for time whereas his concept of the pianist Paderewski has him never late for appointments and not pressed for time. There is no contradiction in the state when Peter believes there to be two Paderewskis. Checking for the inconsistency in the two concept state involves reasoning about the identity expressed by the manifest field. One could imagine a simple consistency checker that does not do this – logically inadequate, of course, but human perhaps. The single concept state could however involve a direct conflict between type and its negation which, one imagines, even the simplest of consistency checkers would find. Thus if Peter finds himself in such a state he might need to refine the properties that he was ascribing to the two Paderewskis in order to make the unified concept of the single Paderewski consistent, for example, by modifying the properties to be always late for political meetings and pressed for time in his political life but never late to a musical event and not pressed for time in concerts.

Note that the link that we have expressed between the two concepts in (51b) does not involve anything like an external anchor. An alternative offered us by the type theory to represent that the two Paderewskis are identical is (53), where we are using p to represent the individual Paderewski.

$$(53) \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x=p:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_j: \left[e:\text{pianist}(\text{id}_i.x) \right] \\ \text{id}_k: \left[\begin{array}{l} x=p:Ind \\ e:\text{named}(x, \text{"Paderewski"}) \end{array} \right] \\ \text{id}_l: \left[e:\text{statesman}(\text{id}_k.x) \right] \end{array} \right]$$

Here the link between Peter's two concepts goes through the world since both his Paderewski concepts are linked to the individual p . If an agent's long term memory is a subtype of (53), then Ind_p figures in the long term memory type (recall that the manifest field $[x=p:Ind]$ is a notation for $[x:Ind_p]$, where Ind_p is a type whose only witness is p (see Appendix A.7)). We take this to mean that the agent has a direct way of identifying Paderewski but that he has not in this case become conscious of the identity of the object involved in different perceptions of Paderewski.⁴ The situation could be that Peter observes Paderewski on the concert platform in tails and then sees him later in the parliament building. His observations are connected to the same individual although without him realizing that he has observed the same Paderewski twice. Thus the situation is similar to that described for *Hesperus* and *Phosphorus* in Frege (1892). In Frege's case the agent was visually aware of the planet Venus on different occasions, conceived of as the Evening Star (*Hesperus*) and the Morning Star (*Phosphorus*) without being aware that the same heavenly body was being observed in the morning as in the evening. The difference between Frege's example and that represented by (53) is that in Frege's case two different proper names were associated with the different observations of the same individual whereas here the

⁴One could choose to interpret such types differently in cognitive terms.

same proper name is being used for the same individual, though without awareness that the proper name is being associated with the same individual on both occasions.

Ludlow (2014) has discussed Kripke's Paderewski recently and argues that the reason that proper names can be used to refer to different individuals can be due to the fact that our lexicons are dynamic and that we use different microlanguages on different occasions. In this discussion he is building on previous work by Larson and Ludlow (1993) although in that work the emphasis is on interpreted logical forms (pairs of abstract syntactic representations and semantic values such as truth values for sentences) rather than on local microlanguages constructed for use in a particular situation as argued for on the basis of a number of different kinds of examples in Ludlow (2014). In general the idea of local microlanguages being constructed on the fly during the course of dialogues and for the purposes at hand is something for which I have a great deal of sympathy and have argued for in the past (Cooper and Ranta, 2008; Larsson and Cooper, 2009; Cooper, 2010, 2012b). And indeed Ludlow (2014) is right to argue that proper names provide support for this view of language. The argument is straightforward in the case of proper names and does not involve the kinds of subtleties of meaning variation which can lead some people to suspicion of this view in the case of other words. If somebody says to me at a party, "I'd like to introduce you to my friend Sam" and indeed I have never met Sam before, I can, as a competent speaker of English, immediately form an association between the phonological type "Sam" and the individual to whom I have been introduced. It is obviously not part of my competence as a speaker of English to know all of the individuals in the universe named Sam. Our competence lies rather in our ability to make the connection between the phonological type (a name) and an individual as the need arises. The competence involves a *dynamic* process of acquiring a linguistic coupling of a speech event type with another part of the world and not a *static* knowledge of all the available couplings. Once I have added this pairing, modelled in our terms as a sign type, to my resources, I have in a technical sense modified my language.⁵ An advantage of sign-based approaches of the kind we are proposing is that you do not have to resort to subscripts in some logical language in order to distinguish between pairings of the same phonological type with different individuals. This is a trap which Larson and Ludlow (1993) fall into when they claim that there are two (or more) names in such cases distinguished by subscripts in logical form. A disadvantage of this analysis is that no two individuals could have the same name in logical form and thus we would have to use something else to analyze sentences like (54).

- (54) My wife's sister, one of my graduate students and our neighbour all have the same name:
Karin

(54) describes a confusing situation which I have to contend with on a daily basis. If the logical form theory with subscripts were correct this sentence would be necessarily false and one might

⁵In my case the resource is quite likely to disappear again shortly afterwards. People vary in their ability to remember names.

have expected that the natural way to describe this situation would rather have been (55).

- (55) My wife's sister, one of my graduate students and our neighbour all have similar names in that they are pronounced "Karin"

(55), according to my intuitions, is not a natural way of describing the situation. This suggests to me that one would need something in addition to, or in place of, a logical form with subscripts to explain how speakers of natural languages individuate names.

One interpretation of Ludlow's proposal is that when a proper name is used to refer to different individuals, different microlanguages are used for the references to the different individuals. Thus when Elisabet says *Karin* and means her sister, she is using a slightly different language than when she says *Karin* and means our neighbour. While I am much in sympathy with the idea of different microlanguages in general it seems to me that such a proposal could not be quite right. Consider dialogues like (56), a kind of dialogue which is not infrequent in our house.

- (56) Elisabet: Karin called
 Robin: Karin?
 Elisabet: My sister

My utterance in (56) is an example of what is called a clarification request in the dialogue literature (Ginzburg and Cooper, 2004; Ginzburg, 2012, and much other literature). According to that literature one of the uses of a clarification request such as *Karin?* is to ask for further identification of the referent of the use of the proper name in the previous dialogue turn. It might initially seem tempting to regard such a request as being in effect a request for (partial) identification of the microlanguage Elisabet is talking. But if we take that route then we have to ask ourselves what language the clarification request itself is in. Assuming that we have three variants of microlanguages available, one where *Karin* refers to Elisabet's sister, one where it refers to our neighbour and one where it refers to my graduate student, then if the request is in any of those languages the answer to the question is selfevident and it is hard to see why I would ask it. And in particular if I was thinking of Karin, my graduate student, I might be justified in saying that Elisabet's answer was wrong. This, of course, is not at all what is going on. It seems that the clarification request is part of a microlanguage in which *Karin* can be used to refer to any of the three and I am interested in finding out which was meant here. This is the kind of option that might be offered by our sign-based approach where a single (micro)language can contain several different signs with the same phonology but with different contents. The exact treatment of this needs, of course, an account of questions and clarification questions in particular which we will not undertake here.

One can understand, however, why the idea of a single referent for a proper name in a single microlanguage might seem attractive. When Kripke (1979) introduces the puzzle about Peter and

Paderewski he is careful to point out the circumstances under which Peter came to the conclusion that there were two Paderewskis. Peter first learns the name Paderewski in connection with the famous pianist. Then: “Later, in a different circle, Peter learns of someone called ‘Paderewski’ who was a Polish nationalist leader and prime minister.” Kripke’s example would not have been at all as convincing if Peter had learned about Paderewski, the pianist and Paderewski, the statesman from the same person in the same conversation. Ludlow (2014) makes a similar point in criticising Kripke’s construction of the apparent contradiction that Peter believes, namely that Paderewski both is and is not a pianist. “The fallacy involves the conjunction of two sentences that have the appearance of contradicting each other. . . but they do not contradict because they come from different microlanguages.” (p. 148). The fact of the matter is that we do tend to use proper names to refer uniquely within the same dialogue, all other things being equal. Suppose we are involved in a conversation about pianists and have been, say, comparing the relative merits of Paderewski and Ashkenazy, and at some point I say (57)

(57) Paderewski was a leading statesman in Poland

You would naturally infer that I was talking about the same Paderewski, unless I explicitly point out that I intend to refer to a different person with the same name. It is, of course, possible to refer to two different people with the same name within the same dialogue and even within the same sentence, even though it may lead to confusion. The assumption is normally, though, that within the space of a dialogue a name will refer to a unique individual unless it is explicitly stated otherwise. One way of being explicit is to say something like (58)

(58) I know another person named Paderewski

If both dialogue participants are aware of the two people with the same name it is possible to use the names together in a construction which normally requires different intended referents as in (59).⁶

(59) Churchland and Churchland think that replacement of symbol manipulation
computer-like devices. . . with connectionist machines hold (*sic*) great promise

(Globus, 1995, p. 21)

Two people named John engaged in conversation with a third person can refer to each other with the name *John* when addressing the third person without risk of confusion as in (60)

⁶I am grateful to Anders Tolland and Stellan Petersson for calling my attention to the fact that the Churchlands are often referred to as “Churchland and Churchland”.

- (60) John E: I remember John as an inspiring professor when I was a student
 John P: Well, I remember John as an extremely bright student
 Third person: I didn't realize you'd known each other that long

When addressing a person you can always use their name as a vocative even if the message you wish to convey involves a person with the same name as in (61).

- (61) A: John, I'd like to introduce you to my good friend John
 B: Glad to meet you. Another John, eh?

It is conceivable that somebody would want to argue that all of these cases where the same name is used twice to refer to different people are examples of code-switching between microlanguages within the space of a dialogue or sentence. Since code-switching does take place even between different languages like English and Portuguese within single dialogues and sentences it is hard to say that such an analysis is impossible. However, given that a sign-based analysis of proper names does not require these examples to be cases of code-switching perhaps the onus is on the proponent of the code-switching analysis to motivate this more complex analysis.

Puzzles about proper names and reference such as the Paderewski puzzle and Frege's (1892) original puzzle about *Hesperus* and *Phosphorus* are standardly presented as puzzles about belief reports. Indeed the matters we have discussed in this section do give rise to puzzles in belief reports and we will return to this later. However, we would like to claim that the discussion here shows that the basis of these puzzles does not lie in the analysis of belief reports *per se* but in the nature of communication in dialogue and the resulting organization of memory. While these phenomena seem puzzling from a Fregean or Montagovian formal language perspective, from the point of view of a dialogic approach employing a sign-based analysis they seem to be a natural consequence of the way that communication takes place and knowledge gets stored.

4.6 Summary

In this chapter we have looked at the analysis of proper names. We started by showing how Montague's analysis of proper names could be recast in TTR and we showed that there was an advantage in the sign-based approach that we have adopted in accounting for the fact that different individuals can have the same name. Montague's original analysis did not say anything about the presupposition-like nature of proper names in that they seem to require interlocutors to be able to identify appropriate referents for the use of a proper name from among a number of potential referents which might be available. We showed how this could be treated by introducing parametric contents for proper names and we showed how accommodation phenomena could be accounted for including a simple-minded analysis of salience analyzed in terms of the information states of agents. Finally, we discussed Kripke's puzzle concerning Paderewski and

its possible relation to a theory of microlanguages as discussed recently by Ludlow. While in general we find the idea of microlanguages appealing we suggested that it plays a role in the analysis of proper names in a rather different way to that suggested by Ludlow.

Chapter 5

Common nouns, intransitive verbs, frames, the Partee puzzle and passengers

5.1 Montague’s treatment of common nouns and individual concepts

The treatment of common nouns in Chapter 3 is encapsulated in $\text{Lex}_{\text{CommonNoun}}$ in Appendix B.1.4.1. The idea is that for a common noun such as *dog* there should be a corresponding predicate ‘dog’ with arity $\langle \text{Ind} \rangle$ as well as a phonological type “dog”. Then an utterance event of the type “dog” will be associated with the content in (1a) which is of type (1b).

- (1) a. $\lambda r: [\text{x:Ind}] . [\text{e:dog}(r.\text{x})]$
b. $([\text{x:Ind}] \rightarrow \text{RecType})$

Montague (1973) introduces predicates corresponding to common nouns which in his type system are of the type $\langle \langle s, e \rangle, t \rangle$. The type $\langle s, e \rangle$ for Montague is the type of individual concepts. These are modelled as functions from world-time pairs (of type s) to individuals (of type e). The reason that Montague used this type rather than the simpler type $\langle e, t \rangle$, that is, the type of functions from individuals to truth-values, has to do with his treatment of the Partee puzzle concerning temperatures and prices which we will take up below. Much subsequent research has abandoned Montague’s analysis using individual concepts and used the simpler type $\langle e, t \rangle$. This alternative would correspond to (2) in our terms.

- (2) a. $\lambda x: \text{Ind} . [\text{e:dog}(x)]$
b. $(\text{Ind} \rightarrow \text{RecType})$

We will argue that (1) is preferable to (2) in that records which are arguments to such a function are frames and that, among other things, frames as arguments enable us to account for the Partee puzzle. We made this proposal in previous work (Cooper, 2010, 2012b). Here we will present a modification of that proposal which uses frames to introduce scales and measure functions and yields a more general treatment of the semantics of verbs like *rise* than we were able to give in the earlier treatment. In addition it gives us a way of treating nouns like *passenger* where, at least on some readings, we seem to be predicating of passenger events, rather than individual passengers. We will also relate our treatment to other recent work on the introduction of frame semantics into formal semantics.

5.2 The Partee puzzle

Perhaps the most recent discussion of the Partee puzzle is that of Löbner (in prep). As we will see, his proposal is closely related to our own. The puzzle is one that Barbara Partee raised while sitting in on an early presentation of the material that led to Montague (1973). In its simplest form it is that (3c) should follow from (3a,b) given some otherwise apparently harmless assumptions.

- (3) a. The temperature is rising
- b. The temperature is ninety
- c. Ninety is rising

Clearly, our intuitions are that (3c) does not follow from (3a,b). The assumptions that the error relies on are those given in (4).

- (4) a. *temperature* is a predicate of individuals
- b. *is* in (3b) represents identity between individuals

Montague's solution was to abandon (4a) and say that 'temperature' is a predicate not of individuals but of individual *concepts*, in his terms functions from world-time pairs to individuals, thus introducing intensionality into predication by common nouns. When we say (3a) we are predicating 'rise' not of an individual but of a function. When we say (3b) we are saying that the *value* of the function at the current world and time is identical with ninety. The technical machinery that Montague uses to achieve this involves his predilection for general treatments. He treats all common nouns as being predicates of individual concepts. But in the case of all nouns other than *price* or *temperature* in his fragment he requires that the individual concepts are rigid designators, that is, they are constant functions which return the same individual for every world-time pair. Similarly intransitive verbs will correspond to predicates of individual concepts but in the case of verbs other than *rise* and *change* (in his fragment) there will be a predicate of the value of the individual concept which holds just in case the verb predicate holds of the

individual concept. Finally *be* is treated as representing identity of the values of individual concepts and a given time and world and not identity of the individual concepts. Thus two distinct individual concepts can have identical values at a given world and time.

Given this machinery we can analyze the Partee puzzle represented in (3) as follows. When we say that the temperature is rising we are predicating ‘rise’ of an individual concept, a function from world-time pairs. Montague does not say what it might mean for such a function to rise. There is, however, something obvious that we could say, namely that if f is such that $\text{temperature}(f)$ at world w and time t , then $\text{rise}(f)$ is true at world w and time t just in case there is some time t' , $t' < t$ (“ t' is earlier than t ”), and some time t'' , $t < t''$, such that $f(w, t')$ is less than $f(w, t'')$. (We may assume that f returns a (real) number for any world and time.) When we say that the temperature f is ninety at world w and time t , what we mean is that $f(w, t) = 90$ (assuming that the interpretation of *ninety* is an individual concept g such that for any world, w , and time, t , $g(w, t) = 90$). From this it does not follow that *ninety* is rising, that is, $\text{rise}(g)$. After all, we have just said that *ninety* corresponds to a constant function which always returns the same value and rising functions have to return different values at different times.

We have now shown that Montague’s analysis prevents the offending inference from going through but we must also show that the inference does go through in “normal” cases according to his analysis. Consider (5).

- (5) a. The dog is barking
- b. The dog is Fido
- c. Fido is barking

Here we do want (5c) to follow as a conclusion from the premises (5a,b). When we say that the dog is barking we are predicating ‘bark’ of a constant function f since for an individual concept to fall under the predicate ‘dog’ it must be rigid, i.e. return the same object for each world and time. Furthermore, there is a predicate, call it ‘bark_{*}’, such that for any w and t , ‘bark_{*}’ holds of $f(w, t)$ just in case ‘bark’ holds of f . So in effect by predicating ‘bark’ of f at w and t , we are predicating ‘bark_{*}’ of $f(w, t)$. (Given Montague’s notion of proposition, $\text{bark}(f)$ and $\text{bark}_*(f(w, t))$ are the *same* proposition since they are true at exactly the same possible worlds and times.) When we say that the dog is Fido at w and t what we mean is that $f(w, t) = g(w, t)$ where g is the individual concept corresponding to *Fido*. According to Montague’s theory of proper names g too will be a constant function always returning the same individual, say, ‘fido’. Is Fido barking given these assumptions, that is, is $\text{bark}(g)$ true at w and t ? There are a couple of ways to make the argument. Since both f and g are constant functions if they have the same value at any world and time they will have the same value at all worlds and times, that is, given the classical set theoretic view of functions that Montague is using, f and g will in fact be the same function. Thus if we predicate anything of f it will also hold of g , since they are identical. The other argument involves the nature of the predicate ‘bark’. Since $\text{bark}(f)$ is equivalent to

$\text{bark}_*(f(w, t))$ and, given that $f(w, t) = g(w, t)$, $\text{bark}_*(f(w, t))$ is equivalent to $\text{bark}_*(g(w, t))$, which in turn is equivalent to $\text{bark}(g)$, then $\text{bark}(f)$ and $\text{bark}(g)$ are equivalent. Thus if $\text{bark}(f)$ is true, then so is $\text{bark}(g)$.

Despite the obvious ingenuity and formal correctness of this solution it fell into disuse. As Löbner (in prep) points out one objection is to the interpretation of (3) as an identity statement rather than the location of the temperature value on a scale. This point was made by Jackendoff (1979), a paper which has given rise to a trickle of remarks and replies in *Linguistic Inquiry* over a period of thirty years: Löbner (1981); Lasersohn (2005); Romero (2008). Part of Jackendoff's argument is that in addition to (6a) we can also say (6b), just as we can say (6c).

- (6) a. The temperature is ninety
- b. The temperature is at ninety
- c. The airplane is at 6000 feet

We do not, he argues, feel the temptation to conclude (7c) from (7a,b).

- (7) a. The airplane is at 6000 feet
- b. The airplane is rising
- c. 6000 feet is/are rising

So neither should we feel the temptation to draw the offending conclusion in the temperature puzzle since even though we say *the temperature is ninety* we mean *the temperature is at ninety*. Jackendoff does not point out, however, that there is an important difference between the temperature and the airplane case, namely that (8) does not mean the same as (7a), and to the extent that it means anything it means something absurd which involves an equality between an airplane and 6000 feet.

- (8) The airplane is 6000 feet

If Jackendoff were right that *is* can mean *is at* why would this be the case? Löbner (1981) has a stronger argument against Jackendoff. He points out that we cannot conclude (9c) from (9a,b)

- (9) a. The temperature of the air in my refrigerator is the same as the temperature of the air in your refrigerator
- b. The temperature of the air in my refrigerator is rising

- c. The temperature of the air in your refrigerator is rising

Lasersohn (2005) gives the example in (10) based on Löbner's example.

- (10) a. The temperature in Chicago is rising
 b. The temperature in Chicago is the very same as the temperature in St. Louis
 c. The temperature in St. Louis is rising

These examples are meant to show that there are similar cases to the original Partee puzzle where the construction seems clearly equative rather than locative. Note that we can mention identity explicitly as in (11).

- (11) The temperature in Chicago is identical with the temperature in St. Louis

Romero (2008) discusses examples with prices where it seems intuitive that there are two readings, one where the inference does not go through and one where it does.

- (12) a. The prices in supermarket *A* are (the very same as) the prices in supermarket *B*
 b. Most prices in supermarket *A* are rising
 c. Most prices in supermarket *B* are rising

On one reading (not the preferred one, I think) (12a) means that at the current time the prices just happen to be the same. In this case the inference does not go through. The other reading is that the prices in the two supermarkets are pegged to each other, perhaps because they are owned by the same chain even though they have different names. (Note that this is not quite the same as saying that the prices are *necessarily* the same which is the case that Romero discusses. This is a matter of business strategy, not logic. The supermarket owners *could* have chosen not to peg the prices to each other.) In this case the inference does go through.¹

Despite all this discussion there is an important intuition in Jackendoff's observation that the interpretation of *the temperature is ninety* involves the placement of the temperature on a scale. In a sense Montague was recognizing this by modelling temperature in terms of his individual

¹Actually, there is a further complication with these examples involving plural quantifiers, which Romero does not discuss. We also need an assumption that the two supermarkets have sufficiently similar stock. If most of the prices are rising in supermarket *A* and supermarket *B* only stocks those items whose prices are not rising in supermarket *A*, then even though the prices in the two supermarkets are the same (and pegged to each other), the prices in supermarket *B* are not rising.

concepts. He was giving us a function which returns for each world and time an individual (presumably a number) representing the temperature. Thus he could account for the fact that the temperature is different at different times. The problem is, though, that possible worlds (that is, total ways the universe could be) do not have a single temperature, even at a single point of time. The notion of individual concept he has is simply not fine-grained enough to deal with temperature. One can understand why Montague might not have wanted to pursue this matter further in PTQ. He wanted to include the treatment of temperature in his general treatment of intensions (functions from possible worlds and times to objects of various types) but in order to get temperature right he would have had to change this. One strategy would be to use possible situations (parts of possible worlds). Another strategy would have been to use an additional index, not just worlds and times but also locations. But if he had done this for temperature and maintained a general theory of intensions he would have had to make all intensions be functions defined on triples of worlds, times and locations and this would have raised issues about the relationship between intensionality and indexicality which he was probably wise to avoid at that point in the development. Nevertheless, it is an important issue which nags at some of the central assumptions of formal semantics as Montague was proposing it: namely, the use of possible worlds and evaluation with respect to a finite set of indices some of which are in the domain of intensions and some of which are contextual parameters.

Löbner's early work on this topic (Löbner, 1979, 1981) treated this problem by removing what he called *functional concepts* (*Funktionalbegriffe*) from the general notion of intension and allowing them to have different numbers and types of argument roles. These insights have led him in later work (Löbner, 2014, in prep) to adopt a frame semantic approach where the parameters that are relevant for interpretation can vary between different words and phrases and there is no fixed set of indices as there was in the original work on formal semantics. This is very much the same kind of proposal as in Cooper (2010, 2012b) although the historical precursors we had in mind were different. In my case, the precursors were early work on situation semantics such as Barwise and Perry (1983) and frame semantics of the kind suggested in Fillmore (1982, 1985) and taken as a foundation for FrameNet (Ruppenhofer *et al.*, 2006, <https://framenet.icsi.berkeley.edu>). In Löbner's case, the inspiration for frames comes from the psychological work of Barsalou (1992a,b, 1999).

5.3 Frames as records

Our leading idea in modelling frames is that they correspond to records and that the *roles* (or *frame elements* in the terminology of FrameNet) are represented by the record fields. Records are in turn what we use to model situations so frames and situations in our view turn out to be the same. Given that we are working in a type theory which makes a clear distinction between types and the objects which belong to those types it is a little unclear whether what we call frame should be a record or a record type. We need both and we will talk of frames (records) and frame types (record types). For example, when we look up the frame `Ambient_temperature` (<https://framenet2.icsi.berkeley.edu/fnReports/data/>

`frameIndex.xml?frame=Ambient_temperature`) in FrameNet we will take that to be an informal description of a frame type which can be instantiated by the kinds of situations which are described in the examples there. In our terms we can characterize a type corresponding to a very stripped down version of FrameNet’s *Ambient_temperature* which is sufficient for us to make the argument we wish to make. This is the type *AmbTempFrame* defined in (13).

$$(13) \quad \left[\begin{array}{ll} x & : \textit{Real} \\ \textit{loc} & : \textit{Loc} \\ e & : \textit{temp}(\textit{loc}, x) \end{array} \right]$$

This is different from the earlier proposal we made in Cooper (2012b) which is given in (14).

$$(14) \quad \left[\begin{array}{ll} x & : \textit{Ind} \\ \textit{e-time} & : \textit{Time} \\ \textit{e-location} & : \textit{Loc} \\ c_{\textit{temp_at_in}} & : \textit{temp_at_in}(\textit{e-time}, \textit{e-location}, x) \end{array} \right]$$

The new proposal in (13) differs from the old one in two ways. Firstly we have removed the field for time. This is because we now want to treat time in terms of strings of events rather than introducing time-points as such. This follows Fernando’s strategy (for example in Fernando, 2011) and relates to the discussion of the Russell-Wiener construction of time in Kamp (1979). Secondly we have made the type in the ‘x’-field (the field which will contain ‘ninety’ in our example) be *Real* (“real number”) rather than *Ind* (“individual”). As Lasersohn (2005) points out the issue was raised early in the literature as to whether numbers (or temperature measurements at any rate) should be treated as individuals in these examples or should be counted as belonging to a separate type (Bennett, 1974; Thomason, 1979). In our earlier work we assumed that temperatures were to be considered as individuals because we had no reason to do otherwise. In the current analysis, however, we want to build in a notion of scale which involves a mapping to real numbers and therefore we will model temperatures as real numbers. As we will see this will lead to a slight complication in the compositional semantics so there is still an open issue as to whether this is the right decision.

A scale is a function which maps frames (situations) to a real number. Thus a scale for ambient temperature will be of the type (15a) and the obvious function to choose of that type is the function in (15b) which maps any ambient temperature frame to the real number in its ‘x’-field.

$$(15) \quad \begin{array}{ll} \text{a. } (\textit{AmbTempFrame} \rightarrow \textit{Real}) \\ \text{b. } \lambda r:\textit{AmbTempFrame} . r.x \end{array}$$

Let us call (15b) ζ_{temp} . As a first approximation we can take an event of a temperature rise to be a string of two temperature frames, $r_1 \widehat{\ } r_2$, where $\zeta_{\text{temp}}(r_1) < \zeta_{\text{temp}}(r_2)$. Using a notation where T^n is the type of strings of length n each of whose members are of type T and where for a given string, s , $s[0]$ is the first member of s , $s[1]$ the second and so on, a first approximation to the type of temperature rises could be (16).

$$(16) \quad \left[\begin{array}{ll} \mathbf{e} & : \text{AmbTempFrame}^2 \\ \mathbf{c}_{\text{rise}} & : \zeta_{\text{temp}}(\mathbf{e}[0]) < \zeta_{\text{temp}}(\mathbf{e}[1]) \end{array} \right]$$

In the \mathbf{c}_{rise} -field of (16) we are using $<$ as an infix notation for a predicate ‘less-than’ with arity $\langle \text{Real}, \text{Real} \rangle$ which obeys the constraint in (17).

$$(17) \quad \text{less-than}(n, m) \text{ is non-empty (“true”) iff } n < m$$

A more general type for temperature rises is given by (18) where we abstract away from the particular temperature scale used by introducing a field for the scale into the record type. This, for example, allows for an event to be a temperature rise independent of whether it is measured on the Fahrenheit or Celsius scales.

$$(18) \quad \left[\begin{array}{ll} \text{scale} & : (\text{AmbTempFrame} \rightarrow \text{Real}) \\ \mathbf{e} & : \text{AmbTempFrame}^2 \\ \mathbf{c}_{\text{rise}} & : \text{scale}(\mathbf{e}[0]) < \text{scale}(\mathbf{e}[1]) \end{array} \right]$$

This type, though, is now too general to count as the type of temperature rising events. To be of this type, it is enough for there to be some scale on which the rise condition holds and the scale is allowed to be any arbitrary function from temperature frames to real numbers. Of course, it is possible to find some arbitrary function which will meet the rise condition even if the temperature is actually going down. For example, consider a function which returns the number on the Celsius scale but with the sign (plus or minus) reversed making temperatures above 0 to be below 0 and *vice versa*. There are two ways we can approach this problem. One is to make the type in the scale-field a subtype of $(\text{AmbTempFrame} \rightarrow \text{Real})$ which limits the scale to be one of a number of standardly accepted scales. This may be an obvious solution in the case of temperature where it is straightforward to identify the commonly used scales. However, scales are much more generally used in linguistic meaning and people create new scales depending on the situation at hand. This makes it difficult to specify the nature of the relevant scales in advance and we therefore prefer our second way of approaching this problem.

The second way is to parametrize the type of temperature rising events. By this we mean using a dependent type which maps a record providing a scale to a record type modelling the type of

temperature rising events according to that scale. The function in (19) is a dependent type which is related in an obvious way to the record type in (18).

$$(19) \quad \lambda r: [\text{scale}: (\text{AmbTempFrame} \rightarrow \text{Real})] . \\ \left[\begin{array}{ll} e & : \text{AmbTempFrame}^2 \\ c_{\text{rise}} & : r.\text{scale}(e[0]) < r.\text{scale}(e[1]) \end{array} \right]$$

According to (18) an event will be a temperature rise if there is some scale according to which the appropriate relation holds between the temperatures of the two stages of the event which we are comparing. According to (19) on the other hand, there is no absolute type of a temperature rise. We can only say whether an event is a temperature rise with respect to some scale or other. If we choose some non-standard scale like the one that reverses plus and minus temperatures as we suggested above then what we normally call a fall in temperature will in fact be a rise in temperature *according to that scale*. You are in principle allowed to choose whatever scale you like, though if you are using the type in a communicative situation you had better make clear to your interlocutor what scale you are using and perhaps also why you are using this scale as opposed to one of the standardly accepted ones. Like the parametric contents we introduced in Chapter 4, the dependent types introduce a presupposition-like component to communicative situations. We are assuming the existence of some scale in the context.

Why do we characterize the domain of the function in (19) in terms of records containing a scale rather than just scales as in (20)?

$$(20) \quad \lambda \sigma: (\text{AmbTempFrame} \rightarrow \text{Real}) . \\ \left[\begin{array}{ll} e & : \text{AmbTempFrame}^2 \\ c_{\text{rise}} & : \sigma(e[0]) < \sigma(e[1]) \end{array} \right]$$

The intuitive reason is that we want to think of the arguments to such functions as being contexts, that is situations (frames) modelled as records. The scale will normally be only one of many informational components which can be provided by the context and the use of a record type allows for there to be more components present. In practical terms of developing an analysis it is useful to use a record type to characterize the domain even if we have only isolated one parameter since if further analysis should show that additional parameters are relevant this will mean that we can add fields to the domain type thereby restricting the domain of the function rather than giving it a radically different type.

And indeed in this case we will now show that there is at least one more relevant parameter that needs to be taken account of before we have anything like a reasonable account of the type of temperature rise events. In (13) we specified that an ambient temperature frame relates a real number (“the temperature”) to a spatial location. And now we are saying that a temperature rise

is a string of two such frames where the temperature is higher in the second frame. But we have not said anything about how the locations in the two frames should be related. For example, suppose I have a string of two temperature frames where the location in the first is London and the location in the second is Marrakesh. Does that constitute a rise in temperature (assuming that the temperature in the second frame is higher than the one in the first)? Certainly not a temperature rise in London, nor in Marrakesh. If you want to talk about a temperature rise in a particular location then both frames have to have that location and we need a way of expressing that restriction. Of course, you can talk about temperature rises which take place as you move from one place to another and which therefore seem to involve distinct locations. However, it seems that even in these cases something has to be kept constant between the two frames. One might analyse it in terms of a constant path to which both locations have to belong or as a constant relative location such as the place where a particular person (or car, or airplane) is. You cannot just pick two arbitrary temperature frames without holding something constant which ties them together. We will deal here with the simple case where the location is kept constant.² We will say that the background information for judging an event as a temperature rise has to include not only a scale but also a location which is held constant in the two frames. This is expressed in (21).

$$(21) \quad \lambda r: \left[\begin{array}{l} \text{fix:} [\text{loc:} \text{Loc}] \\ \text{scale:} (\text{AmbTempFrame} \rightarrow \text{Real}) \end{array} \right] \cdot \left[\begin{array}{l} e \quad : \quad (\text{AmbTempFrame} \wedge [\text{loc} = r.\text{fix}.\text{loc:} \text{Loc}])^2 \\ c_{\text{rise}} \quad : \quad r.\text{scale}(e[0]) < r.\text{scale}(e[1]) \end{array} \right]$$

Here the ‘fix’-field in the context is required to be a record which provides a location. One reason for making the ‘fix’-field a record rather than simply a location is that we will soon see an example where more than one parameter needs to be fixed. It will also help us ultimately in characterizing a general type for a rising event (not just a rise in temperature) if we can refer to the type in the ‘fix’-field as *Rec* (“record”) rather than to list a disjunction of all the various types of the parameters that can be held constant in different cases.

The temperature rise event itself is now required to be a string of two frames which belong to a subtype of *AmbTempFrame*, namely where the ‘loc’-field has been made manifest and is specified to have the value specified for ‘loc’ in the ‘fix’-field. Here we are using the record in the ‘fix’-field of the argument to the function to partially specify the type *AmbTempFrame* by fixing values for some of its fields. One can think of the ‘fix’-record as playing the role of a partial assignment of values to fields in the type. To emphasize this important role and to facilitate making general statements without having to name the particular fields involved, we shall introduce an operation which maps a record type, *T*, and a record, *r* to the result of specifying *T* with *r*, which we will notate as *T* || *r*. (22) provides an abstract example of how it works.

²Although in astronomical terms, of course, even a location like London is a relative location, that is, where London is according to the rotation of the earth and its orbit around the sun. Thus the simple cases are not really different from the cases apparently involving paths.

$$(22) \quad \begin{bmatrix} \ell_1:T_1 \\ \ell_2:T_2 \\ \ell_3:T_3 \end{bmatrix} \parallel \begin{bmatrix} \ell_2=a \\ \ell_3=b \\ \ell_4=c \end{bmatrix} = \begin{bmatrix} \ell_1:T_1 \\ \ell_2=a:T_2 \\ \ell_3=b:T_3 \end{bmatrix}$$

provided that $a : T_2$ and $b : T_3$

In a case where for example $a : T_2$ but not $b : T_3$ we would have (23).

$$(23) \quad \begin{bmatrix} \ell_1:T_1 \\ \ell_2:T_2 \\ \ell_3:T_3 \end{bmatrix} \parallel \begin{bmatrix} \ell_2=a \\ \ell_3=b \\ \ell_4=c \end{bmatrix} = \begin{bmatrix} \ell_1:T_1 \\ \ell_2=a:T_2 \\ \ell_3:T_3 \end{bmatrix}$$

The result (23) would also have obtained if T_3 had not been a type but a pair consisting of a dependent type and a sequence of paths, that is, the kind of thing which in our standard abbreviation we represent as a predicate with a label as argument such as ‘walk(ℓ_1)’. A precise definition of this operation is given in Appendix A.15.

Using this notation we can now rewrite (21) as (24).

$$(24) \quad \lambda r: \begin{bmatrix} \text{fix}: [\text{loc}: \text{Loc}] \\ \text{scale}: (\text{AmbTempFrame} \rightarrow \text{Real}) \end{bmatrix} \cdot \begin{bmatrix} \text{e} & : & (\text{AmbTempFrame} \parallel r.\text{fix})^2 \\ \text{c}_{\text{rise}} & : & r.\text{scale}(\text{e}[0]) < r.\text{scale}(\text{e}[1]) \end{bmatrix}$$

This is still a very simple theory of what a temperature rise event may be but it will be sufficient for our current purposes. We move on now to price rise events. We will take (25) to be the type of price frames, *PriceFrame*.

$$(25) \quad \begin{bmatrix} \text{x} & : & \text{Real} \\ \text{loc} & : & \text{Loc} \\ \text{commodity} & : & \text{Ind} \\ \text{e} & : & \text{price}(\text{commodity}, \text{loc}, \text{x}) \end{bmatrix}$$

The fields represented here are based on a much stripped down version of the FrameNet frame `Commerce_scenario` where our ‘commdodity’-field corresponds to the frame element called ‘goods’ and the ‘x’-field corresponds to the frame element ‘money’. A price rise is a string of two price frames where the value in the ‘x’-field is higher in the second. Here, as in the case of a temperature rise, we need to keep the location constant. It does not make sense to say that a price rise has taken place if we compare a price in Marrakesh with a price in London, even though

the price in London may be higher. In the case of price we also need to keep the commodity constant, something that does not figure at all in ambient temperature. We cannot say that a price rise has taken place if we have the price of tomatoes in the first frame and the price of oranges in the second frame. Thus, following the model of (24), we can characterize the dependent type of price rises as (26).

$$(26) \quad \lambda r: \left[\begin{array}{l} \text{fix:} \left[\begin{array}{l} \text{loc:} Loc \\ \text{commodity:} Ind \end{array} \right] \\ \text{scale:} (PriceFrame \rightarrow Real) \end{array} \right] . \\ \left[\begin{array}{ll} e & : (PriceFrame \| r.\text{fix})^2 \\ c_{\text{rise}} & : r.\text{scale}(e[0]) < r.\text{scale}(e[1]) \end{array} \right]$$

Finally we consider a third kind of rising event discussed in Cooper (2012b) based on the example in (27).

- (27) As they get to deck, they see the Inquisitor, calling out to a Titan in the seas. The giant Titan rises through the waves, shrieking at the Inquisitor.

[http://en.wikipedia.org/wiki/Risen_\(video_game\)](http://en.wikipedia.org/wiki/Risen_(video_game))

accessed 4th February, 2010

Here what needs to be kept constant in the rising event is the Titan. What needs to change between the two frames in the event is the height of the location of the Titan. Thus in this example the location is *not* kept constant. In order to analyze this we can use location frames of the type *LocFrame* as given in (28).

$$(28) \quad \left[\begin{array}{ll} x & : Ind \\ \text{loc} & : Loc \\ e & : \text{at}(x, \text{loc}) \end{array} \right]$$

The dependent type for a rise in location event is (29).

$$(29) \quad \lambda r: \left[\begin{array}{l} \text{fix:} [x: Ind] \\ \text{scale:} (LocFrame \rightarrow Real) \end{array} \right] . \\ \left[\begin{array}{ll} e & : (LocFrame \| r.\text{fix})^2 \\ c_{\text{rise}} & : r.\text{scale}(e[0]) < r.\text{scale}(e[1]) \end{array} \right]$$

Here the obvious scale function does not simply return the value of a field in the location frame. What is needed is a scale based on the height of the location. One way to do this would be to characterize the type of locations, *Loc*, as the type of points in three-dimensional Euclidean space. That is, we consider *Loc* to be an abbreviation for (30).

$$(30) \quad \left[\begin{array}{ll} \text{x-coord} & : \text{Real} \\ \text{y-coord} & : \text{Real} \\ \text{z-coord} & : \text{Real} \end{array} \right]$$

Each of the fields in (30) corresponds to a coordinate in Euclidean space. A more adequate treatment would be to consider locations as regions in Euclidean space but we will not pursue that here. Treating *Loc* as (30) means that we can characterize the scale function as returning the height of the location in the location frame, as in (31).

$$(31) \quad \lambda r:\text{LocFrame} . r.\text{loc.z-coord}$$

If we wish to restrict the dependent type of rising events to vertical rises we can fix the *x* and *y*-coordinates of the location as in (32).

$$(32) \quad \lambda r: \left[\begin{array}{l} \text{fix:} \left[\begin{array}{l} \text{x:Ind} \\ \text{loc:} \left[\begin{array}{l} \text{x-coord:Real} \\ \text{y-coord:Real} \end{array} \right] \end{array} \right] \\ \text{scale:}(\text{LocFrame} \rightarrow \text{Real}) \end{array} \right] . \\ \left[\begin{array}{ll} \mathbf{e} & : (\text{LocFrame} \| r.\text{fix})^2 \\ \mathbf{c}_{\text{rise}} & : r.\text{scale}(\mathbf{e}[0]) < r.\text{scale}(\mathbf{e}[1]) \end{array} \right]$$

We have now characterized three kinds of rising events. In Cooper (2010, 2012b) we argued that there is in principle no limit to the different kinds of rising events which can be referred to in natural language and that new types are created on the fly as the need arises. The formulation in those works did not allow us to express what all these particular meanings have in common. We were only able to say that the various meanings seem to have some kind of family resemblance. Now that we have abstracted out scales and parameters to be fixed we have an opportunity to formulate something more general. There are two things that vary across the different dependent types that we have characterized for risings. One is the frame type being considered and the other is the type of the record which contains the parameters held constant in the rising event. If we abstract over both of these we have a characterization of rising events in general. This is given in (33).

$$(33) \lambda r: \left[\begin{array}{l} \text{frame_type: } RecType \\ \text{fix_type: } RecType \\ \text{fix: fix_type} \\ \text{scale: (frame_type} \rightarrow Real) \end{array} \right] \cdot \left[\begin{array}{ll} e & : (r.\text{frame_type} || r.\text{fix})^2 \\ c_{rise} & : r.\text{scale}(e[0]) < r.\text{scale}(e[1]) \end{array} \right]$$

(33) is so general (virtually everything of content has been parametrized) that it may be hard to see it as being used in the characterization of the meaning of *rise*. What seems important for characterizing the meanings of *rise* that a speaker has acquired is precisely the collection of frame types, and associated fix types and scales which an agent has developed through experience. (33) seems to be relevant to a kind of meta-meaning which specifies what kind of contents can be associated with the word *rise*. In this sense it seems related to the notion of *meaning potential*, a term which has its origins in the work of Halliday (1977) where meanings are spoken of informally as being “created by the social system” and characterized as “integrated systems of meaning potential” (p. 199). The notion is much discussed in more recent literature, for example, Linell (2009), where meaning potential is discussed in the following terms: “Lexical meaning potentials are (partly) open meaning resources, where actual meanings can only emerge in specific, situated interactions” (p. 330).

5.4 Using frames in a compositional semantics for the Partee puzzle

A central aspect of our analysis of the Partee puzzle is that the contents of common nouns are functions that take frames, that is records, as arguments. Nevertheless, we make a distinction between individual level predicates like ‘dog’ whose arity is $\langle Ind \rangle$ and frame level predicates like ‘temperature’ whose arity is $\langle Rec \rangle$. Leaving aside for now the need for parametric contents, the content for associated with an utterance event of type “dog” would be (1a) repeated here as (34a). This is contrasted with the content for an utterance of type “temperature” given in (34b).

$$(34) \text{ a. } \lambda r: [x: Ind] \cdot [e : \text{dog}(r.x)] \\ \text{ b. } \lambda r: [x: Rec] \cdot [e : \text{temperature}(r.x)]$$

We make an exactly similar distinction between individual level and frame level verb phrases. In (35) we present contents which can be associated with utterances of type “run” and “rise” respectively.

$$(35) \text{ a. } \lambda r: [x: Ind] \cdot [e : \text{run}(r.x)] \\ \text{ b. } \lambda r: [x: Rec] \cdot [e : \text{rise}(r.x)]$$

The types which we associate with the individual level and frame level properties in (34) and (35) are given in (36).

- (36) a. $([x:Ind] \rightarrow RecType)$
 b. $([x:Rec] \rightarrow RecType)$

While these types are distinct, they are nevertheless related in that they both have the same range type and the domain types of (36a) and (36b) are both record types requiring a field with the label ‘x’. Up until now we have used *Ppty* (“property”) to designate (36a). Now we might be more specific and designate it as *IndPpty* (“property of individuals”) and use *FramePpty* (“property of frames”) to designate (36b). In our previous treatment of the temperature puzzle both individual level and frame level properties were of the type (36a) because we treated numbers as individuals, that is, as being of type *Ind*. On this view *AmbTempFrame* can be defined as (37a) rather than our current proposal repeated in (37b).

- (37) a.
$$\left[\begin{array}{lll} x & : & Ind \\ loc & : & Loc \\ e & : & temp(loc, x) \end{array} \right]$$

 b.
$$\left[\begin{array}{lll} x & : & Real \\ loc & : & Loc \\ e & : & temp(loc, x) \end{array} \right]$$

Choosing (37a) rather than (37b) would mean that the distinction between individual level and frame level properties would not be one of the type of properties as such (since they would both be of type (36a)) but rather in the arity of the predicate used within the record type that they returned, that is, for example, $\langle Ind \rangle$ for ‘dog’ and $\langle Rec \rangle$ for ‘temperature’. This represents an appealing feature of using record types with subtyping, namely that fine-grained type distinctions can be introduced by predicates within record types which all belong to the same type *RecType*. For a compositional semantics this means that fine grained type distinctions associated with lexical items need not be reflected in the types of the contents of the phrases in which those lexical items occur. This is in significant contrast to the simple type theory used by Montague where there was not subtyping and any type distinction introduced in a constituent would be reflected as a type distinction in higher level phrases. To exploit this feature of the type system here we would have to treat (real) numbers as individuals. This would not necessarily mean abandoning the type *Real* but it would mean stipulating that *Real* is a subtype of *Ind*. For example, we could let *AmbTempFrame* be (37a) but require that the predicate ‘temp’ used in this type have arity $\langle Loc, Real \rangle$. This, together with the requirement $Real \sqsubseteq Ind$, would mean that (37a) and (37b) would be equivalent in the sense that they would have the same set of witnesses.

The alternative sketched above where numbers are treated as individuals has much to commend it. But nevertheless we have not chosen it here for a number of intuitive and practical reasons:

1. There is a fairly robust intuition that numbers are not, in fact, individuals.
2. The proposed solution involves stipulating a subtype relation between basic types. While this is not ruled out in TTR, we would like to keep it to a minimum and not use it where there is an alternative of analyzing the subtyping in terms of record types. Using record types you can see (and compute) whether one type is a subtype of another whereas subtyping between basic types is not explicit in the representation of the types.
3. There are types in TTR of which both types in (36) are subtypes and these types are candidates for the general type *Ppty*, i.e. properties in general as opposed to properties of particular types of objects.

The last point here requires some explanation. Given that TTR has join (disjunctive) types (Appendix A.8) we always have the option of forming the join of those types which we want to represent types of properties. Thus, given the two types of properties we have seen so far we can form the join type in (38).

$$(38) \ (([x:Ind] \rightarrow RecType) \vee ([x:Rec] \rightarrow RecType))$$

If there are more types of properties we wish to add to the general type of properties we can form a larger join type to include them. We can always form a join type based on any finite collection of types. Using join types in this way we can create a type which has all the witnesses of any finite collection of types. We cannot express a type corresponding to an infinite set of types. Also it does not make explicit any relationship between the various types in the collection, in this case that all the types in the collection are function types whose range type is *RecType* and whose domain type is a subtype of *Rec* with an ‘x’-field. In order to deal with this kind of case, we will use the same technique as we used for parametric contents in Chapter 4. We will treat properties as a pair consisting of a type (labelled with ‘bg’) and the function (labelled with ‘fg’) which we have up to now been calling a property. (39a) is an example of the new kind of property and (39b) is the new definition of the type, *Ppty*, of properties.³

$$(39) \text{ a. } \left[\begin{array}{ll} \text{bg} & = \text{Ind} \\ \text{fg} & = \lambda r: [x:Ind] . [e:\text{dog}(r.x)] \end{array} \right]$$

³For a similar kind of case, though a different approach to treating it, see the discussion in Ginzburg *et al.* (2014), p. 93, of the type of Austinian questions. We have also treated this kind of case in terms of a limited kind of polymorphism, for example, in Ginzburg and Cooper (2014).

$$\text{b. } \left[\begin{array}{ll} \text{bg} & : \text{Type} \\ \text{fg} & : ([x:\text{bg}] \rightarrow \text{RecType}) \end{array} \right]$$

Suppose that a situation e is of type $\text{temperature}(r)$. What does that tell us about r ? Given what we have seen so far we might expect that r is an ambient temperature frame, that is, $r : \text{AmbTempFrame}$. We might be tempted to express this as a constraint on assignments to types, that is, a constraint on possibilities in the sense of Appendix A.10. This might be expressed as in (40).

- (40) We restrict attention to those assignments to types (possibilities) such that for any situations e and r , if $e : \text{temperature}(r)$ then $r : \text{AmbTempFrame}$.

Intuitively (40) is similar to Montague's (1973) constraints on the interpretations of intensional logic to which we should restrict attention. These constraints are standardly referred to as *meaning postulates* in the literature although Montague himself did not give them this label. (40) is fine if we are only concerned with ambient temperature or wish the predicate 'temperature' to relate only to ambient temperature. In general, however, we must take account of the fact that there are other kinds of temperature such as the temperature of objects, substances and human bodies. If we choose to have separate predicates for all of these then (40) is a possible constraint to have. On the other hand, if we want to have a single predicate that covers all of these cases then (40) will have to be modified. One thing we might be tempted to do in this case is turn the implication around. Instead of saying "If it's a temperature then it's an ambient temperature frame" as in (40) we say "If it's an ambient temperature frame then it's a temperature". This might look like (41).

- (41) We restrict attention to those assignments to types (possibilities) such that for any situation r , if $r : \text{AmbTempFrame}$ then there is a situation e , such that $e : \text{temperature}(r)$.

Note that it is important here that we changed the quantification over e to existential quantification with scope over the consequent of the conditional. It would have been wrong to have universal quantification as in (42).

- (42) For any situations r and e , if $r : \text{AmbTempFrame}$ then $e : \text{temperature}(r)$

(42) would require that every situation would have to be of the type $\text{temperature}(r)$ if the antecedent holds and this would have the unintuitive consequence that every situation would be a proof object for the temperature in every available location. This would be particularly unwieldy if we wish to consider uncountably many locations as would be natural when considering geometric spaces.

One advantage of (41) is that it fits well into the kind of cognitive approach to semantics that we are trying to promote where we focus on how an agent with limited knowledge will use language rather than on a complete mathematical treatment of how a language considered as an abstract object relates to the world at large as Montague did. (41) expresses one way in which a situation can be considered to be a temperature situation. It leaves open whether there are other kinds of situations which can be considered as temperature situations. Consider an agent acquiring the concept of temperature, that is, what the temperature predicate can be applied to. Such an agent may first become aware of the relevance of ambient temperature as expressed by (41) and then subsequently add similar constraints on the same predicate for, say, the temperature of food, body temperature and so on. By adding constraints to its resources in this way, an agent can incrementally build up an increasingly rich appreciation of a concept like temperature. Another advantage of this is that an agent which has a collection of such constraints as resources can focus on some subset of those constraints or one particular constraint in a given context. Thus in Montague's example *The temperature is 90°* we know that the temperature that is being referred to is ambient temperature.

This suggests that the appropriate notion for these constraints is that of *topos* as discussed by Breitholtz (2014). According to Breitholtz a topos can be modelled as a function returning a type, that is, a dependent type similar to those we have used as update functions in this work. Thus we can replace the prose statement (41) with the function in (43a) which is associated with the licensing condition (43b).

- (43) a. $\lambda r: [x: AmbTempFrame] . [e : temperature(r.x)]$
 b. If $f : (T \rightarrow Type)$ is a topos and r is a record (situation or frame) then for any agent A , $r :_A T$ licenses $:_A f(r)$

(43) absorbs the prose statement in (41) into our theory of dependent types (which we assume can be implemented in memory) and our general theory of action which is supervenient on the type system. The licensing condition (43b) says that a topos will license an agent which judges a situation to be of the domain type of the topos to make a judgement that there is something of the type which the topos returns for that situation. Different agents will have different topoi in their resources and may choose to act or not on the basis of a particular judgement and topos. Different collections of the topoi available in an agent's resources may be activated in different circumstances. And, of course, the collection of resources is dynamic in the sense that an agent may be learning new topoi or adjusting old ones depending on input from the environment. Thus while topoi can be used to replace Montague's meaning postulates they represent a much more flexible tool which can be used to model the reasoning mechanisms available to an agent at a given time.

Another advantage of (43a) is that it is also the right kind of object to be used as a more specified content of *temperature*. It represents a restriction of the function in (34b) obtained by replacing

its domain type with a subtype. This is a natural restriction of (34b) given that we have (43a) as a resource. While it might seem intuitive that a particular utterance of the noun *temperature* might be restricted to ambient temperature, something that might seem puzzling for a formulation of compositional semantics is that this restriction is passed on to the verb as well although intuitively obvious. When we say *the temperature is rising* we are talking about an event which is a temperature rise, not a price rise or any other kind of rise. Somehow we have to coordinate the frame which is chosen in connection with the interpretation of temperature with the frame which is chosen in connection with the interpretation of *rise*. The solution to this that we wish to propose rests on the treatment of generalized quantifiers proposed in Cooper (2011, 2013a).

5.5 Definite descriptions as dynamic generalized quantifiers

In Chapter 3 we showed how to treat indefinite descriptions (consisting of an indefinite article and a common noun phrase) as generalized quantifiers. We will now do something similar for definite descriptions (consisting of a definite article and a common noun phrase). We will then show how to modify this static interpretation of generalized quantifiers so that it becomes a dynamic treatment as presented in Cooper (2011). We will see that the dynamic treatment accounts for how the frame associated with the noun is passed to the verb.

We introduce first a function ‘SemDefArt’ on the model of ‘SemIndefArt’ which was defined in Chapter 3, example (34). Given the new definition of properties as pairs of a type and a function (labelled ‘bg’ and ‘fg’, respectively) we have to specify that the arguments to the quantifier relation are the functions (i.e. ‘fg’). This is given in (44).

$$(44) \quad \lambda Q:Ppty . \quad \lambda P:Ppty . \left[\begin{array}{ll} \text{restr}=Q & : Ppty \\ \text{scope}=P & : Ppty \\ e & : \text{the}(\text{restr}, \text{scope}) \end{array} \right]$$

This is exactly the same as ‘SemIndefArt’ except that instead of the predicate ‘exist’ we have ‘the’. The constraint on ‘the’ which relates it to classical generalized quantifier theory is a refinement of the constraint given for ‘exist’ in Chapter 3, example (52). This is given in (45).

$$(45) \quad [\text{the}(P, Q)] \neq \emptyset \text{ iff } |\downarrow P| = 1 \text{ and } \downarrow P \cap \downarrow Q \neq \emptyset$$

(45) is like the constraint on ‘exist’ except that it adds the additional requirement that the property extension of the first argument to ‘the’ has cardinality exactly one. This replicates Montague’s Russellian treatment of the definite article. We could equivalently define this constraint as (46).

$$(46) \quad [\text{the}(P, Q)] \neq \emptyset \text{ iff } |\downarrow P| = 1 \text{ and } \downarrow P \subseteq \downarrow Q$$

Since we require that there be exactly one object which has P it does not make a difference whether we require that there is some object which has P which also has Q or that every object that has P also has Q . (46) is the way the constraint is stated in Cooper (2013b).

However, it is not quite right now that we have allowed $Ppty$ to be polymorphic. The problem has to do with the definition of $\llbracket \downarrow \cdot \rrbracket$, that is, property extension. The definition we gave in Chapter 3, example (43), is repeated in (47).⁴

$$(47) \quad \llbracket \downarrow P \rrbracket = \{a \mid \exists r[r : [x:Ind] \wedge r.x = a \wedge \llbracket P(r) \rrbracket \neq \emptyset]\}$$

This definition is based on the assumption that all properties are of type $([x:Ind] \rightarrow RecType)$. Now we need to modify it so that we will have a notion of property extension for our new definition of $Ppty$. This is done in (48).

$$(48) \quad \llbracket \downarrow P \rrbracket = \{a \mid \exists r[r : [x:P.bg] \wedge r.x = a \wedge \llbracket P.fg(r) \rrbracket \neq \emptyset]\}$$

We can meet the constraint in (46) by defining the witness condition for $the(P, Q)$ as in (49), using the definition of restricted properties introduced in Chapter 3 and Appendix B.1.

(49) If $P, Q:Ppty$ then,

$$s : the(P, Q) \text{ iff } |\llbracket \downarrow P \rrbracket s| = 1 \text{ and } \llbracket \downarrow P \rrbracket s \subseteq \llbracket \downarrow Q \rrbracket s$$

This gives the predicate ‘the’ the flavour of Russell’s ι -operator.⁵ (See Elbourne, 2012 for recent discussion of the use of this in the semantics of definite descriptions.)

It is well-known that the uniqueness condition in the Russellian treatment of definite descriptions is not quite right for natural language. (For a recent detailed discussion of the issues involved see Elbourne, 2012.) We can, for example, use the definite description *the dog* even though there are several dogs. It is not a simple matter of restricting ourselves to a particular situation that we are describing since we may be describing a situation with several dogs but still refer to some particular dog in the situation as *the dog*. Such examples are discussed in Cooper (1996) citing (50) from McCawley (1979).

⁴Recall that the notation $\llbracket T \rrbracket$ is defined by

$$\llbracket T \rrbracket = \{a \mid a : T\}$$

⁵An alternative is to maintain the intuition that witnesses for ptypes are situation-like (i.e. records) and let the witnesses be represented as $[x=r]$ and $[x=a]$ respectively.

(50) The dog had a fight with another dog yesterday

Our solution to this is to introduce resource situations (Barwise and Perry, 1983; Cooper, 1996). (A similar proposal is made by Elbourne, 2012.) We follow the analysis in Cooper (2013b) and exploit the fact that properties can be restricted to a particular situation by introducing a restricted field in the foreground as in (51).

$$(51) \left[\begin{array}{lcl} \text{bg} & = & \text{Ind} \\ \text{fg} & = & \lambda r: [\text{x:Ind}]. [\text{e} \in s: \text{dog}(r.x)] \end{array} \right]$$

For the restricted field notation $[\text{e} \in s: \text{dog}(r.x)]$ see Chapter 3, p. 89. (51) can be glossed as “the property of being a dog in s ”. We will abbreviate this as ‘dog’ $\upharpoonright s$ where we use ‘dog’ to abbreviate the property without the restricted field. These abbreviations are represented in (52).

$$(52) \begin{array}{l} \text{a. dog' abbreviates } \left[\begin{array}{lcl} \text{bg} & = & \text{Ind} \\ \text{fg} & = & \lambda r: [\text{x:Ind}]. [\text{e}: \text{dog}(r.x)] \end{array} \right] \\ \text{b. dog' } \upharpoonright s \text{ abbreviates } \left[\begin{array}{lcl} \text{bg} & = & \text{Ind} \\ \text{fg} & = & \lambda r: [\text{x:Ind}]. [\text{e} \in s: \text{dog}(r.x)] \end{array} \right] \end{array}$$

General definitions of these abbreviations are given in Appendix B.1.

In Cooper (2013b) we introduced a predicate ‘unique’ which takes two arguments, a property and a record (situation). That is, its arity is $\langle P\text{pty}, \text{Rec} \rangle$. We required that the constraint in (53) hold of ‘unique’.

$$(53) \text{ If } P:P\text{pty}, T \text{ is a type and } s:T, \\ \text{then } [\text{unique}(P, s)] \neq \emptyset \text{ iff } |[\downarrow P \upharpoonright s]| = 1$$

The constraint in (53) expresses that uniqueness holds between a property and a situation just in case the result of restricting the property to that situation is a property whose property extension is a singleton set.

Here we will do things slightly differently. We introduce a one-place predicate ‘unique’ whose arity is $\langle P\text{pty} \rangle$ and which is associated with the witness condition in (54).

$$(54) \text{ If } P:P\text{pty} \text{ and } s:\text{Rec}, \\ \text{then } s : \text{unique}(P) \text{ iff } |[\downarrow P \upharpoonright s]| = 1$$

(54) says that a situation, s , is of the type ‘unique(P)’ (where P is a property) just in case there is exactly one component of s which has the property P .

The reason that we need a uniqueness predicate of this kind has to do with the nature of our type theory. The typing mechanism allows us to say for example what is given in (55).

(55) $s : \text{dog}(a)$

One way to paraphrase (55) is “ a is a dog in s ”. It says that s is of type ‘dog(a)’ but does not rule out that s can be of other types as well including possibly ‘dog(b)’ where b is distinct from a . We do not have a way of saying that ‘dog(a)’ is the only type to which s belongs. This would correspond to Schubert’s (2000) notion of characterizing a situation, that is, in our terms, presenting an exhaustive list of types to which it belongs, which given that we have meet types (Appendix A.9), corresponds to providing a single type to which it belongs such that there is no other type to which it belongs. We have made this choice because it would be very hard if not impossible to guarantee that anything belongs to just one type in the kind of type system we have introduced. Consider, for example, join types. Given our definition of join types in Appendix A.8 if any object a is of some type T it will also be of type $(T \vee T')$ for any type T' . Introduction of this classical kind of disjunction into the system makes it difficult to define a useful notion of a type that completely characterizes an object or situation in the way that Schubert wants.⁶

Introducing the predicate ‘unique’ in the way that we have allows us to place a constraint on the types to which a situation belongs without having to give a complete characterization of all the types to which it belongs. Defining it as a predicate whose argument is a property means that its argument, the property, involves a type. A property is a function which returns a type. Technically, we call it a dependent type. In Chapter 6 we will suggest that allowing types or dependent types as arguments to predicates is a characteristic of evolutionary higher organisms (at least humans). It seems intuitive that the kind of uniqueness involved in the semantics of definite descriptions should belong to this higher kind of reasoning. We can imagine simple organisms (perhaps even as simple as an amoeba) which respond to situations of certain types in certain ways, for example, eating behaviour when confronted with a situation in which an item of food is present. However, it seems unintuitive that such a simple organism would be programmed to engage in eating behaviour when exactly one item of food is present and not otherwise.

⁶Schubert’s argument for needing the notion of characterization has to do with defining a causation relation between events. It seems to me that an analysis of causality must involve a type of the causing event. Thus in addition to a two-place cause relation between two events, “ e_1 caused e_2 ”, we need a three-place cause relation between two events and a type of the first event, “ e_1 caused e_2 in virtue of the fact that $e_1 : T$ ”. Thus, to take an example that Schubert discusses, *John’s singing in the shower caused Mary to wake up in virtue of the fact that it was a singing event* but not *John’s singing in the shower caused Mary to wake up in virtue of the fact that it was an event in the shower*. Allowing types to be arguments to predicates in the way that we do provides a different solution to the problem that Schubert presents.

We shall use uniqueness to create a presuppositional account of definite descriptions using the techniques for parametric contents which we developed in Chapter 4. The presupposition type (a version of that proposed in Cooper, 2013b adjusted to the new one-place predicate ‘unique’) is given in (56).

$$(56) \quad [e:\text{unique}(\text{dog}')]]$$

This is the type that, according to the techniques developed in Chapter 4, will need to be matched against an agent’s resources (gameboard or long term memory) or, if a match is not available, will need to be accommodated into the agent’s gameboard. It requires there to be a situation which has exactly one dog in it. Satisfying the uniqueness presupposition on this view is not so much a question of determining the way the world is (i.e. whether the dog is in some objective sense unique) as determining how the agent has carved up the world into situations.

(56) will, then, be the background of the parametric content of the noun-phrase *the dog*. Three different options for the foreground of this parametric content present themselves, as in (57).

$$(57) \quad \begin{array}{ll} \text{a. } \lambda r: [e:\text{unique}(\text{dog}')] . \\ \quad \lambda P:Ppty . [e : \text{the}(\text{dog}' \upharpoonright r.e, P)] \\ \text{b. } \lambda r: [e:\text{unique}(\text{dog}')] . \\ \quad \lambda P:Ppty . [e : \text{exist}(\text{dog}' \upharpoonright r.e, P)] \\ \text{c. } \lambda r: [e:\text{unique}(\text{dog}')] . \\ \quad \lambda P:Ppty . [e : \text{every}(\text{dog}' \upharpoonright r.e, P)] \end{array}$$

It does not make much difference which of these you choose for the analysis of singular definite descriptions. Since we are relating the generalized quantifier predicates to the classical set relations given in Barwise and Cooper (1981) and we are requiring uniqueness by the presupposition, it will not make a difference in terms of which objects are required to have which properties whether we choose (57a) (using the predicate constraint in (46)), (57b) (using the predicate constraint in (58), which repeats Chapter 3, example (52))

(58) If $P, Q:Ppty$ then

$$[\text{exist}(P, Q)] \neq \emptyset \text{ iff } [\downarrow P] \cap [\downarrow Q] \neq \emptyset$$

or (57c) using the constraint in (59).

(59) If $P, Q:Ppty$ then

$$[\text{every}(P, Q)] \neq \emptyset \text{ iff } [\downarrow P] \subseteq [\downarrow Q]$$

In Cooper (2013b) we chose the option corresponding to (57c), which offers some vague hope of being able to draw a parallel with plural definites. Note that choosing (57b) or (57c) eliminates the need for the predicate ‘the’.

What we have characterized so far is a static treatment of generalized quantifiers. Dynamic generalized quantifiers as presented in Cooper (2011) involving changing the constraint on the quantifier predicate so that the information represented by the first argument to the quantifier predicate is passed on as a restriction to the second argument of the predicate. What we mean by the information associated with a property P is the fixed point type $\mathcal{F}(P.\text{fg})$ as introduced in Chapter 4 and defined in Appendix A.13, p. 292. We shall use the fixed point type of the first argument to restrict the second argument. We define the restriction of a function by a type as in (60).

(60) If f is a function $\lambda v:T_1 . \phi$, then the *restriction of f by a type T_2* , $f|_{T_2}$, is $\lambda v:(T_1 \wedge T_2) . \phi$

We can extend this notation to properties as in (61).

(61) If $P:P\text{pty}$, then $P|_T$ is the property

$$\begin{bmatrix} \text{bg} & = & P.\text{bg} \wedge T \\ \text{fg} & = & P.\text{fg}|_T \end{bmatrix}$$

We can then define dynamic versions of the constraints on the quantifier predicates and their witness conditions as in (62).

(62) a. If $P, Q:P\text{pty}$ then

$$[\text{exist}(P, Q)] \neq \emptyset \text{ iff } [\downarrow P] \cap [\downarrow Q|_{\mathcal{F}(P.\text{fg})}] \neq \emptyset$$

b. If $P, Q:P\text{pty}$ then

$$[\text{every}(P, Q)] \neq \emptyset \text{ iff } [\downarrow P] \subseteq [\downarrow Q|_{\mathcal{F}(P.\text{fg})}]$$

c. If $P, Q:P\text{pty}$ then

$$s : \text{exist}(P, Q) \text{ iff } [\downarrow P \upharpoonright s] \cap [\downarrow Q|_{\mathcal{F}(P.\text{fg})} \upharpoonright s] \neq \emptyset$$

d. If $P, Q:P\text{pty}$ then

$$s : \text{every}(P, Q) \text{ iff } [\downarrow P \upharpoonright s] \subseteq [\downarrow Q|_{\mathcal{F}(P.\text{fg})} \upharpoonright s]$$

The original motivation for treating generalized quantifiers dynamically was to be able to treat the kind of “donkey-anaphora” binding that occurs in sentences like *every farmer who owns a donkey likes it*. Our version of dynamic generalized quantifiers essentially replicates the treatment in Chierchia (1995), though in our own terms. A similar analysis of generalized quantifiers,

exploiting contexts in type theory, is given in Fernando (2001). In order to see how our strategy here will facilitate the treatment of donkey anaphora we will have to wait until we have a treatment of anaphora in Chapter 7. The basic strategy is to exploit the conservativity of generalized quantifiers and treat *every farmer who owns a donkey likes it* as *every farmer who owns a donkey is a farmer who owns a donkey and likes it*. This is achieved by restricting the second argument of the quantifier predicate in the manner indicated in (62).

For present purposes the advantage of dynamicizing the generalized quantifiers is that if the first argument property is restricted to be a property of ambient temperature then that restriction will be passed on to the second argument. Let us look in detail at how this will happen. Consider the type in (63).

$$(63) \text{ every} \left(\begin{array}{l} \text{bg} = \text{AmbTempFrame} \\ \text{fg} = \lambda r: [x:\text{AmbTempFrame}] . [e \in s: \text{temperature}(r.x)] \end{array} \right), \\ \left(\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda r: [x:\text{Rec}] . [e: \text{rise}(r.x)] \end{array} \right)$$

This type might arise as a result of determining the content of *the temperature rises* using the parametric content for *the temperature* in (64) (based on (57)).

$$(64) \lambda r: \left[e: \text{unique} \left(\begin{array}{l} \text{bg} = \text{AmbTempFrame} \\ \text{fg} = \lambda r: [x:\text{AmbTempFrame}] . [e: \text{temperature}(r.x)] \end{array} \right) \right] . \\ \lambda P: \text{Ppty} . \left[e: \text{every} \left(\begin{array}{l} \text{bg} = \text{AmbTempFrame} \\ \text{fg} = \lambda r_1: [x:\text{AmbTempFrame}] . [e \in r.e: \text{temperature}(r_1.x)] \end{array} \right), P \right]$$

The result of applying \mathcal{F} to the foreground of the first argument of (63) in order to obtain a fixed point type is given in (65).

$$(65) \left[\begin{array}{ll} x & : \text{AmbTempFrame} \\ e \in s & : \text{temperature}(x) \end{array} \right]$$

The condition on ‘every’ in (62d) requires that we compare the first argument to ‘every’ with the result of restricting the second argument with (65). The foreground of this is given in (66a), which is identical with (66b) (by the definition of restriction) and (66c) (by the definition of merge) and to (66d) (by the definition of merge⁷ because *AmbTempFrame* is a subtype of *Rec*).

⁷For this step we need to take the version of merge in Appendix A.13 which contains the two additional clauses taking account of subtypes.

- (66) a. $\lambda r: [x:Rec] \cdot [e:rise(r.x)] \Big| \begin{matrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{matrix}$
- b. $\lambda r: [x:Rec] \wedge \begin{matrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:rise(r.x)]$
- c. $\lambda r: \begin{matrix} x:Rec \wedge AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:rise(r.x)]$
- d. $\lambda r: \begin{matrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:rise(r.x)]$

Thus intuitively by choosing to restrict the first argument property to ambient temperature frames we are also restricting the second argument property to ambient temperature frames.

This technique for dynamic quantifiers also has an important consequence if we try to combine frame level and individual level properties. Suppose for example that we are trying to compute the witness condition for *the temperature runs* where *runs* corresponds to the content given in (34a). Then we will have (67) as the foreground of the second argument property.

- (67) a. $\lambda r: [x:Ind] \cdot [e:run(r.x)] \Big| \begin{matrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{matrix}$
- b. $\lambda r: [x:Ind] \wedge \begin{matrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:run(r.x)]$
- c. $\lambda r: \begin{matrix} x:Ind \wedge AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:run(r.x)]$
- d. $\lambda r: \begin{matrix} x:Ind \wedge AmbTempFrame \\ e \in s:temperature(x) \end{matrix} \cdot [e:run(r.x)]$

Here since neither *Ind* nor *AmbTempFrame* are a subtype of the other the final step of merging represented in (67d) is the meet type (without the dot!) whose components are the two types which were merged. The property represented in (67) is thus necessarily empty (that is, its property extension is the empty set no matter what we assign to the basic types), if we have the assumption that individuals are non-records. This would be a way of requiring that the content of *runs* be coerced to something which could hold for temperature frames in order to prevent the sentence from being anomalous. Similarly, if we wish to find a content for *the dog rises* then we have to associate *rises* with an individual property or alternatively associate *dog* with a frame property.

What then should be the content of *is ninety*? An obvious modification to the treatment of *be* in Chapter 3 (see Appendix B.1.4.1), substituting the type *Real* for the type *Ind*, would lead to the property foreground in (68).

$$(68) \quad \lambda r: [x:Real] \cdot \left[\begin{array}{ll} x=r.x, 90 & : \text{Ind} \\ e & : \text{be}(x) \end{array} \right]$$

A property with this as foreground might be the content you need if you are treating a sentence like *2 times 45 is 90*. However, if we use this content with *the temperature* we will run into a similar problem as that represented in (68). This is spelled out in (69)

$$(69) \quad \begin{array}{l} \text{a. } \lambda r: [x:Real] \cdot \left[\begin{array}{ll} x=r.x, 90:Ind \\ e:\text{be}(x) \end{array} \right] \Big| \left[\begin{array}{l} x:AmbTempFrame \\ e \in s:\text{temperature}(x) \end{array} \right] \\ \text{b. } \lambda r: [x:Real] \wedge \left[\begin{array}{l} x:AmbTempFrame \\ e \in s:\text{temperature}(x) \end{array} \right] \cdot \left[\begin{array}{ll} x=r.x, 90:Ind \\ e:\text{be}(x) \end{array} \right] \\ \text{c. } \lambda r: \left[\begin{array}{l} x:Real \wedge AmbTempFrame \\ e \in s:\text{temperature}(x) \end{array} \right] \cdot \left[\begin{array}{ll} x=r.x, 90:Ind \\ e:\text{be}(x) \end{array} \right] \\ \text{d. } \lambda r: \left[\begin{array}{l} x:Real \wedge AmbTempFrame \\ e \in s:\text{temperature}(x) \end{array} \right] \cdot \left[\begin{array}{ll} x=r.x, 90:Ind \\ e:\text{be}(x) \end{array} \right] \end{array}$$

Assuming that real numbers are not records, we have the same problem as we had in (67) in that the property turns out to be necessarily empty. What we need instead is a property of frames (records) that will make reference to a scale, ζ , of the kind we defined for *AmbTempFrame* in (15), for example, a property with the foreground given in (70).

$$(70) \quad \lambda r: [x:Rec] \cdot \left[\begin{array}{ll} x=\zeta(r.x), 90:Ind \\ e:\text{be}(x) \end{array} \right]$$

If ζ is fixed to be the scale in (15) then (70) is identical with (71).

$$(71) \quad \lambda r: [x:Rec] \cdot \left[\begin{array}{ll} x=r.x.x, 90:Ind \\ e:\text{be}(x) \end{array} \right]$$

That is, what is checked for being identical with 90 is the ‘x’-field of the temperature frame which is in the ‘x’-field of the argument to the property. If we choose this property as the content for *is ninety* then the restriction of the property as second argument to the quantifier will give a property as result which is not necessarily empty. The foreground of this property is shown in (72).

$$(72) \quad \text{a. } \lambda r: [x:Rec] \cdot \left[\begin{array}{ll} x=\zeta(r.x), 90:Ind \\ e:\text{be}(x) \end{array} \right] \Big| \left[\begin{array}{l} x:AmbTempFrame \\ e \in s:\text{temperature}(x) \end{array} \right]$$

- b. $\lambda r: [x:Rec] \wedge \begin{bmatrix} x:AmbTempFrame \\ e=s:temperature(x) \end{bmatrix} \cdot \begin{bmatrix} x=\zeta(r.x), 90:Ind \\ e:be(x) \end{bmatrix}$
- c. $\lambda r: \begin{bmatrix} x:Rec \wedge AmbTempFrame \\ e \in s:temperature(x) \end{bmatrix} \cdot \begin{bmatrix} x=\zeta(r.x), 90:Ind \\ e:be(x) \end{bmatrix}$
- d. $\lambda r: \begin{bmatrix} x:AmbTempFrame \\ e \in s:temperature(x) \end{bmatrix} \cdot \begin{bmatrix} x=\zeta(r.x), 90:Ind \\ e:be(x) \end{bmatrix}$

Now, as in (66), (72d) is equivalent to (72c) in the sense that exactly the same objects will have the properties of which these functions are the foreground. This is because *AmbTempFrame* is a subtype of *Rec*. In the functions in (72) there are two parameters which will need to be determined by context in compositional semantics, that is, will need to be found by matching the domain type of a parametric content against an agent's resources. These are the resource situation, *s*, and the scale, ζ .

5.6 Individual vs. frame level nouns

We have made a distinction between individual level nouns like *dog* and frame level nouns like *temperature*, differentiating their contents as in (34) and motivating the distinction with the Partee puzzle. Now consider (73).

- (73) a. The dog is nine
 b. The dog is getting older/aging
 c. Nine is getting older/aging

We have the same intuitions about (73) as we do about the original temperature puzzle. We cannot conclude (73c) from (73a,b). Does this mean that *dog* is a frame level noun after all? Certainly, if we think of frames as being like entries in relational databases it would be natural to think of age (or information allowing us to compute age such as date of birth) as being a natural field in a dog-frame.⁸

Our strategy to deal with this will be to say that contents of individual level nouns can be coerced to frame level contents, whereas the contents of frame level nouns cannot be coerced “down” to individual level contents. Thus in addition to (34a), repeated as (74a) we have (74b).

- (74) a. $\lambda r: [x:Ind] \cdot [e : dog(r.x)]$

⁸Curiously, it does not seem to figure in FrameNet for *dog* (as of 2nd March, 2015). The noun *dog* is associated with the frame *Animals* which inherits from the frame *Biological entity*. But in neither of these frames is there a frame element corresponding to age or date of birth. There is a frame *Age* but this does not seem to be related to *Animals* or *Biological entity*.

$$\text{b. } \lambda r: [x:Rec] . [e : \text{dog_frame}(r.x)]$$

The predicate ‘dog_frame’ is related to the predicate ‘dog’ by the constraint in (75).

$$(75) \text{ dog_frame}(r) \text{ is non-empty implies } r : \begin{bmatrix} x:Ind \\ e:\text{dog}(x) \end{bmatrix}$$

There are several different kinds of dog frames with additional information about a dog which an agent may acquire or focus on. Here we will consider just frames which contain a field labelled ‘age’ as specified in (76).

$$(76) \text{ If } r : \begin{bmatrix} x:Ind \\ e:\text{dog}(x) \\ \text{age}:Real \\ c_{age}:\text{age_of}(x,\text{age}) \end{bmatrix} \text{ then dog_frame}(r) \text{ is non-empty}$$

An age scale, ζ_{age} , for individuals can then be defined as the function in (77).

$$(77) \zeta_{age} = \lambda r: \begin{bmatrix} x:Ind \\ \text{age}:Real \\ c_{age}:\text{age_of}(x,\text{age}) \end{bmatrix} . r.\text{age}$$

The content foreground for *is nine in the dog is nine* is then like (70) with ζ set to ζ_{age} and 9 replacing 90. Thus *be* followed by a numeral can be coerced to a content depending on some scale which is available as a resource.

We can think of the sentence *the dog is nine* as involving two coercions: one coercing the content of *dog* to a frame level property and the other coercing the content of *be* to a function which when applied to a number will return a frame level property depending on an available scale. Such coercions do not appear to be universally available in languages. For example, in German it is preferable to say *die Temperatur ist 35 Grad* “the temperature is 35 degrees” rather than *#die Temperatur ist 35* “the temperature is 35”. Similarly *der Hund ist neun Jahre alt* “the dog is nine years old” is preferred over *#der Hund ist neun* “the dog is nine”. We will return to the matter of coercion or creation of new contents in Section 5.8.

We now turn our attention to the formulation of the compositional semantics.

5.7 Defining a compositional semantics for the Partee puzzle

We will now make precise the resources that are needed in order to account for the data expressed in (78).

- (78) a. a/the dog runs
 b. the dog is nine
 c. the temperature is ninety
 d. the temperature rises

We start with the determiners. The background type (“presupposition”) introduced by *the* (56) depends on the common noun. This means that the contents of *the* and *dog* cannot be combined by the S-combinator strategy for the combination of parametric contents that was introduced in Chapter 4, defined in Appendix B.1.4.2. This combination method passes up the background requirements of the two daughters without modifying them as depicted graphically in (79).

(79)

$$\begin{array}{c}
 X \\
 \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f = T_{F_{\text{bg}}} \\ a = T_{A_{\text{bg}}} \end{array} \right] \\ fg = \lambda r: \left[\begin{array}{l} f = T_{F_{\text{bg}}} \\ a = T_{A_{\text{bg}}} \end{array} \right] . F_{fg}(r.f)(A_{fg}(r.a)) \end{array} \right] \\
 \swarrow \quad \searrow \\
 F \quad A \\
 \left[\begin{array}{l} \text{bg} = T_{F_{\text{bg}}} \\ fg = F_{fg} \end{array} \right] \quad \left[\begin{array}{l} \text{bg} = T_{A_{\text{bg}}} \\ fg = A_{fg} \end{array} \right]
 \end{array}$$

Instead what we need is (80) where \mathfrak{F} is a function from parametric contents to parametric contents and \mathcal{A} is a parametric content.

(80)

$$\begin{array}{c}
 X \\
 \mathfrak{F}(\mathcal{A}) \\
 \swarrow \quad \searrow \\
 F \quad A \\
 \mathfrak{F} \quad \mathcal{A}
 \end{array}$$

Here the function \mathfrak{F} has to do the S-combinator like work which was achieved by the combination method in (79). In (81) we show a schematic version of the function and on the node X we show the schematic result of applying the function to the argument.

(81)

$$\begin{array}{c}
 X \\
 \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f: T_{F_{\text{bg}}}(A_{\text{fg}}) \\ a: T_{A_{\text{bg}}} \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} f: T_{F_{\text{bg}}}(A_{\text{fg}}) \\ a: T_{A_{\text{bg}}} \end{array} \right] . F_{\text{fn}}(A_{\text{fg}}(r.a)) \end{array} \right] \\
 \swarrow \quad \searrow \\
 F \qquad A \\
 \lambda p: \left[\begin{array}{l} \text{bg}: \text{RecType} \\ f: T_{A_{\text{fg}}} \end{array} \right] . \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f: T_{F_{\text{bg}}}(p.\text{fg}) \\ a: p.\text{bg} \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} f: T_{F_{\text{bg}}}(p.\text{fg}) \\ a: p.\text{bg} \end{array} \right] . F_{\text{fn}}(p.\text{fg}(r.a)) \end{array} \right] \quad \left[\begin{array}{l} \text{bg} = T_{A_{\text{bg}}} \\ \text{fg} = A_{\text{fg}} \end{array} \right]
 \end{array}$$

We will call a function from parametric contents to parametric contents a *dependent parametric content*. That is, it depends on another parametric content to yield a parametric content. The content of definite articles, SemDefArt , will be defined as an instance of the schema under F in (81). The definition is given in (82a) whose type is (82b).

$$\begin{array}{l}
 (82) \text{ a. } \lambda Q: PPty . \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f: [e: \text{unique}(Q.\text{fg}(\uparrow a))] \\ a: Q.\text{bg} \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} f: [e: \text{unique}(Q.\text{fg}(\uparrow a))] \\ a: Q.\text{bg} \end{array} \right] . \lambda P: Ppty . \left[\begin{array}{l} \text{restr} = Q.\text{fg}(r.a): Ppty \\ \text{scope} = P: Ppty \\ e: \text{every}(\text{restr}, \text{scope}) \end{array} \right] \end{array} \right] \\
 \text{b. } (PPty \rightarrow PQuant)
 \end{array}$$

The lexical resource we need to include in the English lexicon is (83a) where $\text{Lex}_{\text{DefArt}}$ is a universal resource defined in (83b). (Lex is defined in Appendix B.1.4.1.)

(83) a. $\text{Lex}_{\text{DefArt}}(\text{“the”})$

b. $\text{Lex}_{\text{DefArt}}(T_{\text{phon}})$, where T_{phon} is a phonological type, is defined as
 $\text{Lex}(T_{\text{phon}}, \text{Det}) \wedge [\text{cnt} = \text{SemDefArt}: (PPty \rightarrow PQuant)]$

For the sake of consistency we shall adjust the definition of the SemIndefArt so that it too is a dependent parametric content of the same form as SemDefArt even though it does not introduce a presupposition that depends on the following noun. SemIndefArt is defined as (84).

$$(84) \quad \lambda Q:PPpty . \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} \text{f:Rec} \\ \text{a:Q.bg} \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} \text{f:Rec} \\ \text{a:Q.bg} \end{array} \right] . \lambda P:Ppty . \left[\begin{array}{l} \text{restr} = Q.\text{fg}(r.\text{a}):Ppty \\ \text{scope} = P:Ppty \\ \text{e:exist}(\text{restr}, \text{scope}) \end{array} \right] \end{array} \right]$$

The lexical resource we need to include in the English lexicon is (85a) where $\text{Lex}_{\text{IndefArt}}$ is a universal resource defined in (85b).

- (85) a. $\text{Lex}_{\text{IndefArt}}(\text{"a"})$
 b. $\text{Lex}_{\text{IndefArt}}(T_{\text{phon}})$, where T_{phon} is a phonological type, is defined as
 $\text{Lex}(T_{\text{phon}}, \text{Det}) \wedge [\text{cnt} = \text{SemIndefArt}: (PPpty \rightarrow PQuant)]$

We now move on to common nouns. We define $\text{SemCommonNoun}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type representing the background requirements, as (86).

$$(86) \quad \left[\begin{array}{l} \text{bg} = T_{\text{bg}} \\ \text{fg} = \lambda c:T_{\text{bg}} . \left[\begin{array}{l} \text{bg} = T_{\text{restr}} \\ \text{fg} = \lambda r: [x:T_{\text{restr}}] . [e:p(r.x)] \end{array} \right] \end{array} \right]$$

The lexical resources for common nouns we will include in the English lexicon are (87a) where $\text{Lex}_{\text{CommonNoun}}$ is a universal resource defined in (87b).

- (87) a. $\text{Lex}_{\text{CommonNoun}}(\text{"dog"}, \text{dog}, \text{Ind}, \text{Ind}, \text{Rec})$
 $\text{Lex}_{\text{CommonNoun}}(\text{"temperature"}, \text{temperature}, \text{Rec}, \text{Rec}, \text{Rec})$
 b. $\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where T_{phon} is a phonological type, p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type, is defined as
 $\text{Lex}(T_{\text{phon}}, N) \wedge [\text{cnt} = \text{SemCommonNoun}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}): PPpty]$

We can think of the common noun sign types in (87a) as unmodulated in something like the sense of modulation discussed by Recanati (2010) in that the restriction type yielding the type of the domain of the property is identical with the type that represents the arity of the predicate. We will see a way to modulate the content of the noun by choosing a subtype of the type of the predicate argument as the domain type of the property (that is, T_{restr} above).

To this we add an operation $\text{CommonNounIndToFrame}$ which is defined on individual level common noun sign types and “raises” them to frame level common noun sign types. This is defined in (88).

- (88) If T_{phon} is a phonological type, p is a predicate and T_{bg} is a record type (the “background type” or “presupposition”) then

$$\begin{aligned} &\text{CommonNounIndToFrame}(\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, \text{Ind}, \text{Ind}, T_{\text{bg}})) = \\ &\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p_frame, \text{Rec}, \text{Rec}, T_{\text{bg}}) \end{aligned}$$

This operation is a universal resource which may or may not be used by individual languages. Given the discussion in Section 5.6, we suggest that it is used productively in English but not in German, for example. This gives us a way of generating new lexical resources from already existing resources. Similarly, we can think of p_frame as being the result of applying a “raising” operation to the predicate p where the new predicate is associated with the general constraint expressed in (89).

- (89) If p is a predicate with arity $\langle \text{Ind} \rangle$, then for any e and r ,

$$e : p_frame(r) \text{ implies } r : \begin{bmatrix} x:\text{Ind} \\ e:p(x) \end{bmatrix}$$

Another way to generate new lexical resources from basic common noun sign types is to restrict the domain of the common noun by some type (perhaps related to a topos as suggested in Section 5.4). This is formulated in (90).

- (90) If T_{phon} is a phonological type, p is a predicate, T_{arg} is a type and that arity of p is $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$, T_{bg} is a record type and $T_{\text{mod}} \sqsubseteq T_{\text{restr}}$ then

$$\begin{aligned} &\text{RestrictCommonNoun}(\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}), T_{\text{mod}}) = \\ &\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{mod}}, T_{\text{bg}}) \end{aligned}$$

This will enable us, for example, to restrict the basic lexical entry for *temperature* (repeated in (91a)) to obtain the additional lexical resource (91b), which is needed to produce the noun-phrase content in (64).

- (91) a. $\text{Lex}_{\text{CommonNoun}}(\text{“temperature”}, \text{temperature}, \text{Rec}, \text{Rec}, \text{Rec})$
 b. $\text{Lex}_{\text{CommonNoun}}(\text{“temperature”}, \text{temperature}, \text{Rec}, \text{AmbTempFrame}, \text{Rec})$

We can combine restriction coercion with frame coercion. While frame coercion gives us a general frame level property of records we can restrict the frame to be of a certain type corresponding to a particular type of frame that we have as a resource. For example, suppose that we have resource which is a frame type for dog frames, *DogFrame* as introduced in (76) repeated in (92).

$$(92) \begin{bmatrix} x:Ind \\ e:dog(x) \\ age:Real \\ c_{age}:age_of(x,age) \end{bmatrix}$$

We can use *DogFrame* to restrict the result of coercing our frame level dog sign type, that is there can be a two-step coercion from the basic lexical entry in (87a) as represented in (93).

$$(93) \text{Lex}_{\text{CommonNoun}}(\text{"dog"}, dog, Ind, Ind, Rec) \rightsquigarrow \text{Lex}_{\text{CommonNoun}}(\text{"dog"}, dog_frame, Rec, Rec, Rec) \rightsquigarrow \text{Lex}_{\text{CommonNoun}}(\text{"dog"}, dog_frame, Rec, DogFrame, Rec)$$

We treat intransitive verbs in a parallel fashion to common nouns. Thus

$$\text{SemIntransVerb}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$$

where p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is the record type (86) as for common nouns. We define $\text{Lex}_{\text{IntransVerb}}$ similarly to $\text{Lex}_{\text{CommonNoun}}$ in (94).

$$(94) \text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}), \text{ where } T_{\text{phon}} \text{ is a phonological type, } p \text{ is a predicate with arity } \langle T_{\text{arg}} \rangle, T_{\text{restr}} \sqsubseteq T_{\text{arg}} \text{ and } T_{\text{bg}} \text{ is a record type, is defined as } \text{Lex}(T_{\text{phon}}, VP) \wedge [\text{cnt}=\text{SemIntransVerb}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}):PPty]$$

The basic lexical resources that we need for *runs* and *rises* are given in (95).

$$(95) \text{Lex}_{\text{IntransVerb}}(\text{"runs"}, run, Ind, Ind, Rec) \\ \text{Lex}_{\text{IntransVerb}}(\text{"rises"}, rise, Rec, Rec, Rec)$$

We need a treatment of *is* which will allow it to combine with numerals like *nine* and *ninety* to form a frame level predicate as indicated in (70). We start from a parametrized version of the definition of *SemBe* which we introduced in Chapter 3 (repeated in Appendix B.1.4.1). We give the parametrized version in (96).

$$(96) \lambda r:Rec . \\ \lambda Q:Quant .$$

$$\left[\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r_1: [\text{x:Ind}] . \\ \quad \mathcal{Q} \left(\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r_2: [\text{x:Ind}] . \left[\begin{array}{ll} \text{x}=r_1.\text{x}, r_2.\text{x} & : \text{Ind} \\ \text{e} & : \text{be}(\text{x}) \end{array} \right] \end{array} \right) \end{array} \right]$$

Here the context represented by the first argument to the function, r , does not contribute anything to the final content of be which represents straightforward equality. In this chapter we want to allow equality not only between individuals but also objects of other types as well as introducing a type for the background. We will thus give SemBe two arguments and represent (96) as $\text{SemBe}(\text{Ind}, \text{Rec})$. In general we will define $\text{SemBe}(T_{\text{arg}}, T_{\text{bg}})$ to be (97).

$$(97) \quad \lambda r: T_{\text{bg}} . \\ \lambda \mathcal{Q}: \text{Quant} . \\ \left[\begin{array}{l} \text{bg} = T_{\text{arg}} \\ \text{fg} = \lambda r_1: [\text{x}: T_{\text{arg}}] . \\ \quad \mathcal{Q} \left(\begin{array}{l} \text{bg} = T_{\text{arg}} \\ \text{fg} = \lambda r_2: [\text{x}: T_{\text{arg}}] . \left[\begin{array}{ll} \text{x}=r_1.\text{x}, r_2.\text{x} & : T_{\text{arg}} \\ \text{e} & : \text{be}(\text{x}) \end{array} \right] \end{array} \right) \end{array} \right]$$

We will say that this holds if T_{bg} does not require a scale, more precisely if T_{bg} is not a subtype of $[\text{sc}: (T_{\text{arg}} \rightarrow \text{Real})]$. If $T_{\text{bg}} \sqsubseteq [\text{sc}: (T_{\text{arg}} \rightarrow \text{Real})]$ then $\text{SemBe}(T_{\text{arg}}, T_{\text{bg}})$ is (98).

$$(98) \quad \lambda r: T_{\text{bg}} . \\ \lambda \mathcal{Q}: \text{Quant} . \\ \left[\begin{array}{l} \text{bg} = T_{\text{arg}} \\ \text{fg} = \lambda r_1: [\text{x}: T_{\text{arg}}] . \\ \quad \mathcal{Q} \left(\begin{array}{l} \text{bg} = [\text{x}: \text{Real}] \\ \text{fg} = \lambda r_2: [\text{x}: \text{Real}] . \left[\begin{array}{ll} \text{x}=r.\text{sc}(r_1.\text{x}), r_2.\text{x} & : \text{Real} \\ \text{e} & : \text{be}(\text{x}) \end{array} \right] \end{array} \right) \end{array} \right]$$

Using the scale ζ_{age} from (77), repeated as (99b), we can use the domain type of that function, which we refer to as *AgeFrame* (as specified in (99a)) as T_{arg} and construct a meaning for be $\text{SemBe}(\text{AgeFrame}, [\text{sc}: (\text{AgeFrame} \rightarrow \text{Real})])$, given in (99c).

$$(99) \quad \text{a. } \text{AgeFrame} = \left[\begin{array}{l} \text{x:Ind} \\ \text{age:Real} \\ \text{c}_{\text{age}}: \text{age_of}(\text{x}, \text{age}) \end{array} \right]$$

$$\text{b. } \zeta_{\text{age}} = \lambda r: \text{AgeFrame} . r.\text{age}$$

$$\begin{aligned}
\text{c. } & \lambda r: [\text{sc}:(\text{AgeFrame} \rightarrow \text{Real})] . \\
& \lambda Q: \text{Quant} . \\
& \left[\begin{array}{l} \text{bg} = \text{AgeFrame} \\ \text{fg} = \lambda r_1: [\text{x:AgeFrame}] . \\ \quad Q \left(\begin{array}{l} \text{bg} = [\text{x:Real}] \\ \text{fg} = \lambda r_2: [\text{x:Real}] . \left[\begin{array}{ll} \text{x}=r.\text{sc}(r_1.\text{x}), r_2.\text{x} & : \text{Real} \\ \text{e} & : \text{be}(\text{x}) \end{array} \right] \end{array} \right) \end{array} \right]
\end{aligned}$$

The idea is that (99c) will be created as a resource on the basis of the existence of (99b) which will in turn rely on the fact that (99a) is a resource. (99c) is available as a meaning of *be* in English but not for German.

To complete the picture we need to account for *nine* and *ninety*. We will treat these as logically proper names of real numbers. Thus we will not treat them as introducing presuppositions in the manner in which we suggested in Chapter 4 but rather in the Montague-like manner which we used for proper names in Chapter 3, except that we now adjust it to take account of parametric contents and the new definition of *Ppty* (property), repeated in (100).

$$(100) \left[\begin{array}{ll} \text{bg} & : \text{Type} \\ \text{fg} & : ([\text{x:bg}] \rightarrow \text{RecType}) \end{array} \right]$$

If T is a type we let $Ppty(T)$ represent the partial specification of *Ppty* presented in (101).

$$(101) \left[\begin{array}{ll} \text{bg}=T & : \text{Type} \\ \text{fg} & : ([\text{x:bg}] \rightarrow \text{RecType}) \end{array} \right]$$

In n is a (real) number, then $\text{SemNumeral}(n)$ (the content for a number expression such as *nine*) is as given in (102).

$$\begin{aligned}
(102) & \lambda r: \text{Rec} . \\
& \lambda P: \text{Ppty}(\text{Real}) . \\
& P.\text{fg}([\text{x}=n])
\end{aligned}$$

Then we can define $\text{Lex}_{\text{numeral}}$ as an operation which takes a phonological type T_{phon} and a (real) number n and returns the sign type (103).

$$\begin{aligned}
(103) \quad & \text{Lex}_{\text{numeral}}(T_{\text{phon}}, n) = \\
& \text{Lex}(T_{\text{phon}}, NP) \wedge [\text{cnt}=\text{SemNumeral}(n): \text{PQuant}]
\end{aligned}$$

The two sign types that we need as resources for our small fragment are given in (104).

- (104) a. $\text{Lex}_{\text{numeral}}(\text{"nine"}, 9)$
 b. $\text{Lex}_{\text{numeral}}(\text{"ninety"}, 90)$

Since we are now using two methods of semantic combination for contents, simple function application for *Det N* constructions and our variant of S-combination for other constructions we need to use the variant of CntForwardApp from Chapter 3 for the former and the variant from Chapter 4 for the latter. We will here call the Chapter 3 version CntForwardApp as before and rename the Chapter 4 version to CntSForwardApp . Details are given in Appendices B.1.4.2 and B.2.4.

5.8 Passengers and ships

Gupta (1980) points out examples such as (105).

- (105) a. National Airlines served at least two million passengers in 1975
 b. Every passenger is a person
 c. National Airlines served at least two million persons in 1975

His claim is that we cannot conclude (105c) from (105a,b). There is a reading of (105a) where what is being counted is not passengers as individual people but passenger events, events of people taking flights, where possibly the same people are involved in several flights. Gupta claims that it is the only reading that this sentence has. While it is certainly the preferred reading for this sentence (say, in the context of National Airlines' annual report or advertizing campaign), I think the sentence also has a reading where individuals are being counted. Consider (106).

- (106) National Airlines served at least two million passengers in 1975. Each one of them signed the petition.

While (106) could mean that a number of passengers signed the petition several times our knowledge that people normally only sign a given petition once makes a reading where there are two million distinct individuals involved more likely. Similarly, while (105c) seems to prefer the individual reading where there are two million distinct individuals it is not impossible to get an event reading here. Krifka (1990) makes a similar point. Gupta's analysis of such examples involves individual concepts and is therefore reminiscent of the functional concepts used by Löbner (1979, 1981) to analyze the Partee puzzle.

Carlson (1982) makes a similar point about Gupta's examples in that nouns which appear to normally point to individual related readings can in the right context get the event related readings. One of his examples is a traffic engineer's report as in (107).

- (107) Of the 1,000 cars using Elm St. over the past 49 hours, only 12 cars made noise in excess of EPA recommended limits.

It is easy to interpret this in terms of 1,000 and 12 car events rather than individual cars. Carlson's suggestion is to use his notion of *individual stage*, what he describes intuitively as "things-at-a-time". Krifka (1990) remarks that "Carlson's notion of a stage serves basically to reconstruct events". While this is not literally correct, the intuition is nevertheless right. Carlson was writing at a time when times and time intervals were used to attempt to capture phenomena that in more modern semantics would be analyzed in terms of events or situations. Thus Carlson's notion of stage is related to a frame-theoretic approach which associates an individual with an event.

Consider the noun *passenger*. It would be natural to assume that passengers are associated with journey events. FrameNet⁹ does not have an entry for *passenger*. The closest relevant frame appears to be TRAVEL which has frame elements for traveller, source, goal, path, direction, mode of transport, among others. The FrameNet lexical entry for *journey* is associated with this frame. Let us take the type *TravelFrame* to be the stripped down version of the travel frame type in (108a). Then we could take the type *PassengerFrame* to be (108b).

- (108) a.
$$\left[\begin{array}{ll} \text{traveller} & : \text{Ind} \\ \text{source} & : \text{Loc} \\ \text{goal} & : \text{Loc} \end{array} \right]$$
- b.
$$\left[\begin{array}{ll} x & : \text{Ind} \\ e & : \text{passenger}(x) \\ \text{journey} & : \text{TravelFrame} \\ c_{\text{travel}} & : \text{take_journey}(x, \text{journey}) \end{array} \right]$$

A natural constraint to place on the predicate 'take_journey' is that in (109).

- (109) If $a:\text{Ind}$ and $e:\text{TravelFrame}$, then the type $\text{take_journey}(a, e)$ is non-empty just in case $e.\text{traveller} = a$.

Let us suppose that the basic lexical entry for *passenger* is (110a). This will mean that its (parametric) content is (110b) (with vacuous dependence on the context).

⁹As of 13th May 2015.

(110) a. $\text{Lex}_{\text{CommonNoun}}(\text{"passenger"}, \text{passenger}, \text{Ind}, \text{Ind}, \text{Rec})$

$$\text{b. } \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda c:\text{Rec} . \left[\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r: [\text{x}:\text{Ind}] . \text{passenger}(r.x) \end{array} \right] \end{array} \right]$$

That is, its non-parametric content is a property of individuals. Given the coercion *CommonNounIndToFrame* we have defined we can coerce this lexical item to (111a) which means that its parametric content will be (111b).

(111) a. $\text{Lex}_{\text{CommonNoun}}(\text{"passenger"}, \text{passenger_frame}, \text{Rec}, \text{Rec}, \text{Rec})$

$$\text{b. } \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda c:\text{Rec} . \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda r: [\text{x}:\text{Rec}] . \text{passenger_frame}(r.x) \end{array} \right] \end{array} \right]$$

This means that the non-parametric content is a property of frames. An agent who has the frame type *PassengerFrame* available as a resource can use it to restrict the domain of the property using the coercion *RestrictCommonNouns*. This produces (112a) which means that its parametric content will be (112b).

(112) a. $\text{Lex}_{\text{CommonNoun}}(\text{"passenger"}, \text{passenger_frame}, \text{Rec}, \text{PassengerFrame}, \text{Rec})$

$$\text{b. } \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda c:\text{Rec} . \left[\begin{array}{l} \text{bg} = \text{PassengerFrame} \\ \text{fg} = \lambda r: [\text{x}:\text{PassengerFrame}] . [\text{e}:\text{passenger_frame}(r.x)] \end{array} \right] \end{array} \right]$$

This means that the non-parametric content will now be a property of passenger frames of type *PassengerFrame*. This introduces not only a passenger but also a journey, an event in which in which the passenger is the traveller.

It seems that we have now done something which Krifka (1990) explicitly warned us against. At the end of his discussion of Carlson's analysis he comes to the conclusion that it is wrong to look for an explanation of event-related readings of these sentences in terms of a noun ambiguity. One of Krifka's examples is (113) (which gives the title to his paper).

(113) Four thousand ships passed through the lock

This can either mean that four thousand distinct ships passed through the lock or that there were four thousand ship-passing-through-the-lock events a number of which involved the same

ships. The problem he sees is that if we treat *ship* as being ambiguous between denoting individual ships or ship stages in Carlson's sense then there will be too many stages which pass through the lock. For example, suppose that a particular ship passes through the lock twice. This gives us two stages of the ship which pass through the lock. But then, Krifka claims, there will be a third stage, the sum of the first two, which also passes through the lock. It is not clear to me that this is an insuperable problem for the stage analysis. We need to count stages that pass through the lock exactly once. Let us see how the frame analysis fares.

We will start with a singular example in order to avoid the additional problems offered by the plural. Consider (114).

(114) Every passenger gets a hot meal

Suppose that an airline has this as part of its advertizing campaign. Smith, a frequent traveller, takes a flight with the airline and as expected gets a hot meal. A few weeks later she takes another flight with the same airline and does not get a hot meal. She sues the airline for false advertizing. At the hearing, her lawyer argues, citing Gupta (1980), that the advertizing campaign claims that every passenger gets a hot meal on every flight they take. The lawyer for the airline company argues, citing Krifka (1990), that the sentence in question is ambiguous between an individual and an event reading, that the airline had intended the individual reading and thus the requirements of the advertizing campaign had been met by the meal that Smith was served on the first flight. Smith's lawyer then calls an expert witness, a linguist who quickly crowdsources a survey of native speakers' interpretations of the sentence in the context of the campaign and discovers that there is an overwhelming preference for the meal-on-every-flight reading. (The small percentage of respondents who preferred the individual reading over the event reading gave their occupation as professional logician.) Smith wins the case and receives an additional hot meal.

What is important for us at the moment is the fact that there is an event reading of this sentence. We will return to the matter of preferred readings below. We will treat the content of *every* on the model of the content of the indefinite article, except that the quantifier relation will be 'every' instead of 'exist'. Thus we will define SemUniversal on the model of SemIndefArt in Appendix B.1.4.1.¹⁰ This is given in (115).

$$(115) \lambda Q:PPpty. \left[\begin{array}{l} bg = \left[\begin{array}{l} f:Rec \\ a:Q.bg \end{array} \right] \\ fg = \lambda r: \left[\begin{array}{l} f:Rec \\ a:Q.bg \end{array} \right] . \lambda P:Ppty. \left[\begin{array}{l} restr = Q.fg(r.a):Ppty \\ scope = P:Ppty \\ e:every(restr, scope) \end{array} \right] \end{array} \right]$$

¹⁰We leave to one side the issue of whether *every* should introduce a background constraint that there are at least three objects which have the property associated with the noun.

If we use the content associated with *passenger* in (112) the non-parametric content associated with *every passenger* will be (116).

$$(116) \lambda P:Ppty . \left[\begin{array}{l} \text{restr} = \left[\begin{array}{l} \text{bg} = \text{PassengerFrame} \\ \text{fg} = \lambda r: [x:\text{PassengerFrame}] . \text{passenger_frame}(r.x) \end{array} \right] : Ppty \\ \text{scope} = P:Ppty \\ \text{e:every}(\text{restr}, \text{scope}) \end{array} \right]$$

In order to simplify matters let us treat *gets a hot meal* as if it were an intransitive verb corresponding to a single predicate ‘get_a_hot_meal’. This is a predicate whose arity is $\langle \text{Ind} \rangle$. It is individuals, not frames (situations), that get hot meals. Thus the non-parametric content of *gets a hot meal* will be (117).

$$(117) \left[\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r: [x:\text{Ind}] . [e:\text{get_a_hot_meal}(r.x)] \end{array} \right]$$

While (117) is the right type of argument for (116) since it is a property it will lead us eventually into problems because there is nothing which is both a passenger frame and an individual for the reasons discussed in Section 5.5. What we need is a coercion which will obtain a frame level intransitive verb to match the frame level noun. This would be a coercion *IntransVerbIndToFrame* exactly parallel to *CommonNounIndToFrame* defined in (88). Thus *IntransVerbIndToFrame* is defined as in (118).

$$(118) \text{ If } T_{\text{phon}} \text{ is a phonological type, } p \text{ is a predicate and } T_{\text{bg}} \text{ is a record type (the “background type” or “presupposition”) then} \\ \text{IntransVerbIndToFrame}(\text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, p, \text{Ind}, \text{Ind}, T_{\text{bg}})) = \text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, \\ p_frame, \text{Rec}, \text{Rec}, T_{\text{bg}})$$

Thus the new non-parametric content derived for *get_a_hot_meal* will be (119).

$$(119) \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda r: [x:\text{Rec}] . [e:\text{get_a_hot_meal_frame}(r.x)] \end{array} \right]$$

Recall that if p is a predicate of individuals then p_frame is a predicate of frames that contain an individual of which p holds (as required in (89)). This means that an argument, r , to ‘get_a_hot_meal_frame’ which makes the type ‘get_a_hot_meal_frame(r)’ non-empty will be of type (120).

$$(120) \begin{bmatrix} x & : & Ind \\ e & : & get_a_hot_meal(x) \end{bmatrix}$$

Thus intuitively the ‘every’ relation holding between the two frame-level coerced individual properties corresponding to *passenger* and *get_a_hot_meal* will mean “every frame (situation) containing an individual in the ‘x’-field who is a passenger taking a journey will be a frame where the individual in the ‘x’-field gets a hot meal”. Or, more formally, (121).

$$(121) \text{ every } r \text{ of type } \begin{bmatrix} x & : & Ind \\ e & : & passenger(x) \\ journey & : & TravelFrame \\ c_{travel} & : & take_journey(x, journey) \end{bmatrix} \text{ is of type } \begin{bmatrix} x & : & Ind \\ e & : & get_a_hot_meal(x) \end{bmatrix}$$

This means that every frame of type *PassengerFrame* will be of type (122a), that is (122b) which is identical with (122c).

$$(122) \text{ a. } PassengerFrame \wedge \begin{bmatrix} x:Ind \\ e:get_a_hot_meal(x) \end{bmatrix}$$

$$\text{b. } \begin{bmatrix} x & : & Ind \\ e & : & passenger(x) \\ journey & : & TravelFrame \\ c_{travel} & : & take_journey(x, journey) \end{bmatrix} \wedge \begin{bmatrix} x & : & Ind \\ e & : & get_a_hot_meal(x) \end{bmatrix}$$

$$\text{c. } \begin{bmatrix} x & : & Ind \\ e & : & passenger(x) \wedge get_a_hot_meal(x) \\ journey & : & TravelFrame \\ c_{travel} & : & take_journey(x, journey) \end{bmatrix}$$

Thus even though we have coerced to a frame-level reading it is still the passengers (i.e. individuals) in the frames who are getting the hot meal not the situation which is the frame.

Things go less well with cardinality quantifiers, however. Consider *2000 passengers get a hot meal* which corresponds to (123).

$$(123) \text{ 2000 } r \text{ of type } \left[\begin{array}{ll} x & : \text{Ind} \\ e & : \text{passenger}(x) \\ \text{journey} & : \text{TravelFrame} \\ c_{\text{travel}} & : \text{take_journey}(x, \text{journey}) \end{array} \right] \text{ are of type } \left[\begin{array}{ll} x & : \text{Ind} \\ e & : \text{get_a_hot_meal}(x) \end{array} \right]$$

The problem is not exactly the same as the problem which Krifka foresaw with the summing of stages although it is intuitively related. It has to do with the way we have set up subtyping with record types. Given a record of a type we can always add a new field to the record and obtain a distinct record of the same type. Trivially the field we add could contain an object already occurring in a field in the original record. As we are assuming that the set of labels is countably infinite if there is one record of a given type there will be infinitely many records of the same type. We illustrate this with an abstract example in (124).

$$(124) \text{ a. } \left[\begin{array}{ll} \ell_1 & : T_1 \\ \ell_2 & : T_2 \end{array} \right]$$

$$\text{ b. } \left[\begin{array}{ll} \ell_1 & : a \\ \ell_2 & : b \end{array} \right]$$

$$\text{ c. } \left[\begin{array}{ll} \ell_1 & : a \\ \ell_2 & : b \\ \ell_3 & : a \end{array} \right]$$

$$\text{ d. } \left[\begin{array}{ll} \ell_1 & : a \\ \ell_2 & : b \\ \ell_3 & : a \\ \ell_4 & : a \end{array} \right]$$

$$\text{ e. } \dots$$

If (124b) is of type (124a) (i.e. $a : T_1$ and $b : T_2$), then so are (124c) and (124d) and so on as we successively “grow” the record without changing the fields that make the records a witness for the type and without necessarily adding anything new in the new fields. If records model events (situations) then this corresponds to the intuition that given any event there will always be a larger event of which it is a part. For example, if I wash my hands that is part of an event in which I wash my hands and stand at the washbasin. This is in turn part of an event in which I wash my hands, stand at the washbasin and breathe and so on. We want this to be true but still there is the robust intuition that we are only talking about one event of washing my hands here which is part of infinitely many larger events.

Fortunately, this problem is easy to fix. Recall that records are sets of fields (Appendix A.12). As a first approximation we can say that a record, r_1 , is a proper part of a record, r_2 , $r_1 < r_2$, just

in case r_1 is a proper subset of r_2 . This definition is not quite sufficient, however, since records can contain records and we wish the proper part of relation to be recursive. Consider (125). We would like to say that (125a) is a proper part of (125b) even though (125a) is not a proper subset of (125b).

$$(125) \text{ a. } \left[\begin{array}{l} \ell_1 = \left[\begin{array}{l} \ell_1 = a \\ \ell_2 = b \end{array} \right] \\ \ell_3 = d \end{array} \right]$$

$$\text{b. } \left[\begin{array}{l} \ell_1 = \left[\begin{array}{l} \ell_1 = a \\ \ell_2 = b \\ \ell_3 = c \end{array} \right] \\ \ell_3 = d \end{array} \right]$$

We achieve this by defining the proper part of relation in terms of the flattening operation on records (represented by φ , see Appendix A.12). Thus if we take the flattenings of the records in (125) as in (126) we see that (126a) is a proper subset of (126b).

$$(126) \text{ a. } \left[\begin{array}{l} \ell_1.\ell_1 = a \\ \ell_1.\ell_2 = b \\ \ell_3 = d \end{array} \right]$$

$$\text{b. } \left[\begin{array}{l} \ell_1.\ell_1 = a \\ \ell_1.\ell_2 = b \\ \ell_1.\ell_3 = c \\ \ell_3 = d \end{array} \right]$$

We define the proper part of relation in (127).

- (127) a. If r_1 and r_2 are records then r_1 is a proper part of r_2 , $r_1 < r_2$, just in case $\varphi(r_1) \subset \varphi(r_2)$.
- b. If o_1 and o_2 are objects of some type and at least one of them is not of type *Rec*, then o_1 is not a proper part of o_2 , $o_1 \not< o_2$

This notion yields a notion of minimal object of a given type which is related to Schubert's 2000 notion of characterization discussed in Section 5.5. It is different from Schubert's notion in that we do not say that there are no other types to which the situation belongs but rather that no proper part of the situation is of the type. In this way it is related to the notion of minimal situation discussed by Kratzer (2014) and elsewhere in earlier work. It is also, of course, related to mereological approaches that have been used, for example, in approaches to the analysis of

the plural as in Krifka (1990) and much other literature. It is this we will exploit in our analysis of the plural cardinality quantifiers.

We characterize a notion of plurality types as in (128).

- (128) a. If T is a type, then $\{ T \}$ is also a type (the type of pluralities such that every element of the plurality is of type T)
- b. $A : \{ T \}$ iff
1. $A : \{ T \}$
 2. if $a \in A$ then for any b such that $a < b$, $b \notin A$

The definition in (128) requires that a plurality is a set of objects of the relevant type but that it does not contain two objects where one is a proper part of the other. It might seem natural to require that a plurality contains at least two objects. The choice not to place this requirement on a plurality makes this analysis number neutral in the sense of Zweig (2008, 2009). Zweig (2008) contains a useful overview of some of the variants of analyses of the plural that have been proposed in the literature, including the distinction between set-based and sum-based analyses. In the type theory we have proposed we have sets already available and a kind of mereology based on the structure of records, as illustrated in (127), and we have used a combination of these in our characterization of plurality. Whether this proposal would survive an in-depth investigation of the plural in this framework is an open question. In particular the work on mass terms by Sutton and Filip (????) suggests that we will in any case need an additional sum-structure.

We propose here a treatment of basic plural quantification cases involving cardinality quantification that will allow us to say something about the content of *2000 passengers get a hot meal*. We first distinguish between singular and plural properties. The definition of $Ppty$ given in (39b) and repeated in (129a) becomes our definition of the type of singular properties, $SgPpty$.

- (129) a. $SgPpty \equiv \left[\begin{array}{ll} bg & : \text{Type} \\ fg & : ([x:bg] \rightarrow RecType) \end{array} \right]$
- b. $PlPpty \equiv \left[\begin{array}{ll} bg & : \text{Type} \\ fg & : ([x:\{bg\}] \rightarrow RecType) \end{array} \right]$
- c. $Ppty \equiv SgPpty \vee PlPpty$

The type of plural properties, $PlPpty$, given in (129b), requires the foreground of the property to be a function which has as its domain records whose ‘x’-field contains a plurality of objects of the background type. We then redefine $Ppty$ in (129c) as being the type of objects which are either singular or plural properties.

Note that according to (129) there is no constraint on the kinds of types which can be the backgrounds of either singular or plural properties. Thus a singular property could have a plurality type as its background. This could be used for nouns like *committee* which seems to represent a property of a plurality of people. Note that we can create a plurality type of any type including plurality types so we have an infinite hierarchy of plurality types. This can be seen in examples like *league* (of baseball teams) where each element in the league, i.e. a team, is itself a plurality of baseball players. These kinds of examples were noted in the earliest literature on treating the plural in Montague semantics as problematic for Montague's approach since they appeared to involve an infinite hierarchy of types which would have to correspond to an infinite hierarchy of syntactic categories (Bennett, 1974). Montague's type system lacked the option of creating a single type corresponding to the infinite hierarchy in the way we are doing here. Plural properties too can have plurality types as their backgrounds and this could correspond to plural nouns like *committees* and *leagues*. In terms of the domain of the foreground function of properties there is largely an overlap between singular and plural properties. Singular properties allow for both pluralities and non-pluralities whereas plural properties allow only for pluralities. It may seem strange from a formal point of view not to make a non-overlapping division between the two, but this does seem to correspond to the way the plural works in natural languages.¹¹

We consider cardinality quantifiers such as *two* and *two thousand* to correspond to predicates whose arity is $\langle PlPty, PlPty \rangle$. If ν is a natural number (that is, an object of type *Nat*), let ν_{pred} be such a predicate corresponding to ν . We also introduce a predicate 'card' ("cardinality") with arities $\langle \{T\}, Card \rangle$ for any type, T , where *Card* is the type of cardinal numbers (the natural numbers together with the transfinite cardinals, $\aleph_0, \aleph_1, \dots$). This predicate obeys the constraint in (130).

$$(130) \quad [\text{card}(X, \nu)] \neq \emptyset \text{ iff } |X| = \nu$$

The cardinality predicates obey the constraint in (131).

$$(131) \quad [\nu_{pred}(P, Q)] \neq \emptyset \text{ iff}$$

$$[\mathcal{F}(Q.fg |_{\mathcal{F}(P.fg)}) \wedge \left[\begin{array}{l} x: \{P.bg\} \\ c: \text{card}(x, \nu) \end{array} \right]] \neq \emptyset$$

Let us take this definition through our example *2000 passengers get a hot meal*. The relevant type corresponding to the content of this sentence is given in (132).

¹¹It is not currently clear how *pluralia tantum* such as *scissors* and *trousers* fit into this story.

$$(132) \text{ } 2000_{\text{pred}} \left(\begin{array}{l} \text{bg} = \text{PassengerFrame} \\ \text{fg} = \lambda r: [x: \{\text{PassengerFrame}\}] \cdot [e: \text{passenger_frame_pl}(r.x)] \end{array} \right), \\ \left(\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda r: [x: \{\text{Rec}\}] \cdot [e: \text{get_a_hot_meal_frame_pl}(r.x)] \end{array} \right)$$

The first argument to ‘ 2000_{pred} ’ is a pluralized version of the property in the restriction field in (116). The second argument is a pluralized version of the property in (119). (132) makes use of a pluralization operation, ‘ _pl ’ on predicates which can be introduced as in (133).

(133) If p is a predicate with arity $\langle T \rangle$, then $p\text{_pl}$ is a predicate with arity $\langle \{ T \} \rangle$

The cases we are considering here are distributive plurals, that is, the constraint in (134) holds.

$$(134) \text{ } [p\text{_pl}(A)] \neq \emptyset \text{ iff } a \in A \text{ implies } [p(a)] \neq \emptyset$$

Let us instantiate (131) bit by bit with (132). We first compute the fixed point type of the foreground of the first argument. That is, (135a) which is identical to (135b).

$$(135) \text{ a. } \mathcal{F}(\lambda r: [x: \{\text{PassengerFrame}\}] \cdot [e: \text{passenger_frame_pl}(r.x)]) \\ \text{b. } \left[\begin{array}{ll} x & : \{\text{PassengerFrame}\} \\ e & : \text{passenger_frame_pl}(x) \end{array} \right]$$

Then we compute the result of restricting the second argument to the quantifier predicate by (135b). This is given in (136a) which is identical to (136b).

$$(136) \text{ a. } \lambda r: [x: \{\text{Rec}\}] \cdot [e: \text{get_a_hot_meal_frame_pl}(r.x)] \left[\begin{array}{l} x: \{\text{PassengerFrame}\} \\ e: \text{passenger_frame_pl}(x) \end{array} \right] \\ \text{b. } \lambda r: \left[\begin{array}{l} x: \{\text{PassengerFrame}\} \\ e: \text{passenger_frame_pl}(x) \end{array} \right] \cdot [e: \text{get_a_hot_meal_frame_pl}(r.x)]$$

We then compute the fixed point type of (136b), given in (137a) which is identical with (137b).

$$(137) \text{ a. } \mathcal{F}(\lambda r: \left[\begin{array}{l} x: \{\text{PassengerFrame}\} \\ e: \text{passenger_frame_pl}(x) \end{array} \right] \cdot [e: \text{get_a_hot_meal_frame_pl}(r.x)])$$

$$\text{b. } \left[\begin{array}{l} x: \{ \text{PassengerFrame} \} \\ e: \text{passenger_frame_pl}(x) \wedge \text{get_a_hot_meal_frame_pl}(x) \end{array} \right]$$

The final step specified in (131) involves merging (137b) with (138a), that is, (138b) which is identical with (138c).

$$\begin{aligned} (138) \text{ a. } & \left[\begin{array}{l} x : \{ \text{PassengerFrame} \} \\ c : \text{card}(x, 2000) \end{array} \right] \\ \text{b. } & \left[\begin{array}{l} x: \{ \text{PassengerFrame} \} \\ e: \text{passenger_frame_pl}(x) \wedge \text{get_a_hot_meal_frame_pl}(x) \end{array} \right] \wedge \left[\begin{array}{l} x: \{ \text{PassengerFrame} \} \\ c: \text{card}(x, 2000) \end{array} \right] \\ \text{c. } & \left[\begin{array}{l} x: \{ \text{PassengerFrame} \} \\ e: \text{passenger_frame_pl}(x) \wedge \text{get_a_hot_meal_frame_pl}(x) \\ c: \text{card}(x, 2000) \end{array} \right] \end{aligned}$$

It is thus the type (138c) which is required to be non-empty by the content of an utterance of *2000 passengers get a hot meal*. This means that it is required that there is a plurality of passenger frames where the passenger gets a hot meal and this plurality has the cardinality 2000, or slightly more colloquially, there are 2000 separate events of a passenger getting a hot meal.

Requiring that there is a plurality with cardinality 2000 is to say that there are at least 2000 objects meeting whatever conditions are invoked. It does not rule out the possibility of there being a larger plurality of objects meeting the same conditions. If we want to express *exactly* ν we can use the condition in (139).

$$(139) \quad [\text{exactly_}\nu_{\text{pred}}(P, Q)] \neq \emptyset \text{ iff}$$

1. $[\nu_{\text{pred}}(P, Q)] \neq \emptyset$
2. $[\text{at_most_}\nu_{\text{pred}}(P, Q)] \neq \emptyset$

where ‘at_most_ ν ’ obeys the constraint in (140).

$$(140) \quad [\text{at_most_}\nu_{\text{pred}}(P, Q)] \neq \emptyset \text{ iff}$$

$$[\mathcal{F}(Q.\text{fg} \mid \mathcal{F}(P.\text{fg})) \wedge \left[\begin{array}{l} x: \{ P.\text{bg} \} \\ n: \text{Nat} \\ c_n: n > \nu \\ c: \text{card}(x, n) \end{array} \right]] = \emptyset$$

5.9 Conclusion

In this chapter we have proposed an analysis of frames as records which model situations (including events) and we have suggested that frame types (record types) are important in both the analysis of the Partee puzzle concerning rising temperatures and prices and in the analysis of quantification which involves counting events rather than individuals like passengers or ships passing through a lock.

Our original inspiration for frames comes from the work of Fillmore (1982, 1985) and work on FrameNet (<https://framenet.icsi.berkeley.edu>). An important aspect of our approach to frames is that we treat them as first class objects. That is, they can be arguments to predicates and can be quantified over. While this is important, it is not surprising once we decide that frames are in fact situations (here modelled by records) or situation types (here modelled by record types). The distinction between frames and frame types is not made in the literature deriving from Fillmore’s work but it seems to be an important distinction to draw if we wish to apply the notion of frame to the kind of examples we have discussed in this chapter.

The proposal that we have made for solving the Partee puzzle is closely related to the work of Löbner (2014, in prep) whose inspiration is from the work of Barsalou (1992b,a, 1999) rather than Fillmore. Barsalou’s approach embedded in a theory of cognition based on perception and a conception of cognition as dynamic, that is, a system in a constant state of flux (Prinz and Barsalou, 2014), seems much in agreement with what we are proposing in this book. Barsalou’s (1999) characterization of basic frame properties constituting a frame as: “(1) predicates, (2) attribute-value bindings, (3) constraints, and (4) recursion” seem to have a strong family resemblance with our record types. Our proposal for incorporating frames into natural language semantics is, however, different from Löbner’s in that he sees the introduction of a psychological approach based on frames as a reason to abandon a formal semantic approach whereas we see type theory as a way of combining the insights we have gained from model theoretic semantics with a psychologically oriented approach.

Our approach to frames has much in common with that of Kallmeyer and Osswald (2013) who use feature structures to characterize their semantic domain. We have purposely used record types in a way that makes them correspond both to feature structures and discourse representation structures which allows us to relate our approach to more traditional model theoretic semantics at the same time as being able to merge record types corresponding to unification in feature-based systems. However, our record types are included in a richer system of types including function types facilitates a treatment of quantification and binding which is not available in a system which treats feature structures as a semantic domain.¹²

¹²It is possible to code up a notation for quantification in feature structures but that is not the same as giving a semantics for it.

Chapter 6

Modality and intensionality without possible worlds

6.1 Possible worlds, modality and intensionality

Montague (1973) uses possible worlds to analyze both modality (represented in his fragment by the adverbs *possibly* and *necessarily*) and a variety of intensional constructions in addition to the temperature and price examples discussed in Chapter 5: intensional transitive verbs such as *seek*, intensional adverbs such as *voluntarily*, verbs of propositional attitudes such as *believe* and *assert* and verbs taking infinitival complements such as *try (to)* and *wish (to)*.

A short introduction to the use of possible worlds in modal logic and philosophical conceptions of possible worlds is given by Menzel (2015). As he points out at the beginning of this article possible worlds are considered to be totalities (or at least a limit) which include the situations which we are aware of around us.

The notion of possible world is intuitively appealing. We talk of living in the best (or worst) of all possible worlds. But equally we talk of the best (or worst) possibility. When we talk in such terms we normally have a small finite number of possibilities in mind which we are contrasting. This has led some authors to use the term “possible world” to refer not to a total universe but to a small set of facts that might obtain in some version of the world. This appears to be standard usage in probability theory (e.g. Halpern, 2003). It is important not to confuse this notion with the notion of possible world as a totality which is used in semantics, inherited from modal logic. This point is made by Cooper *et al.* (2014a) and Lappin (2015).

Problems have been raised for the notion possible world. These have to do with how you individuate and count them and how many possible worlds there must be. Rescher (1999) takes up these problems from a philosophical perspective. He argues that it is impossible to individuate possible worlds and therefore impossible to count them. Lappin (2015) takes up the representation

problem for possible worlds. If you cannot represent possible worlds then you cannot individuate them. The central problem for possible worlds as they are talked about in the semantics literature seems to be that the intuitive way to distinguish one possible world from another is to find a proposition that is true in the first world but false in the second. This would be fine except that we now have the corresponding problem for propositions. Unfortunately the intuitive way of distinguishing between one proposition and another (if you are a possible worlds theorist) is to find a possible world in which the first proposition is true and the other is false. This, of course, is circular and will not give us an individuation of either possible worlds or propositions. The standard version of possible worlds semantics as proposed by Montague does not, of course, fall into this obvious trap. Worlds are not represented in terms of sets of propositions which are true in them. Rather we just define an interpretation to include a set of possible worlds and leave aside the question of how they have been individuated. In a sense it is fine from a technical point of view to have an arbitrary set whose membership we cannot represent as a central component of our semantic theory. But it leaves us with the suspicion that we are left with an abstract theory which we do not really know how to connect to any empirical observations of the world. If you take a mathematical view of the semantic enterprise as Montague did, this may be acceptable. But if you are interested in semantics as an aspect of human cognitive ability it can appear problematic. Traditional possible world semantics is a theory based on an assumed set of possible worlds. But it is not a theory of the possible worlds as such, beyond the claim that there is a set of them.

Despite this, there is an intuition about the set of possible worlds which possible world theorists hold onto: that they represent all the logical possibilities. This, at least, gives us a way of considering the required cardinality of the set of possible worlds. The issue of the cardinality of the set of possible worlds and its relationship to a psychological theory of language is something that is already taken up by Partee (1977). Here she refers to Lewis's (1973) argument that there must be at least \beth_2 (the cardinality of the power set of the power set of natural numbers) possible worlds. The argument¹ goes like this: suppose we have a family that goes on for ever. That is, there would be \aleph_0 members of the family. Now consider that in a logically possible world (though possibly not in biologically possible worlds) any subset of these family members might have blue eyes (none of them, all of them and all the possibilities in between). This gives us a set of possible worlds whose cardinality is the same as the power set of the natural numbers 2^{\aleph_0} or \beth_1 , that is, the cardinality of the set of real numbers. Now consider the logical possibility that each of those possible worlds is biologically plausible. Again, logically speaking, any subset of those worlds could be biologically plausible. This will yield a set of possible worlds of cardinality $2^{2^{\aleph_0}}$ or \beth_2 . In principle one could create sets of possible worlds of any of the infinitely many infinite cardinalities although as Lewis claims \beth_2 is probably sufficient for normal purposes.

Another argument for the uncountability of the set of possible worlds comes from usual assumptions about space and time. We normally assume that the set of moments of time has the same cardinality as the set of points on the real line, that is, that time is continuous. Similarly we

¹which I first heard from Barbara Partee but for which I cannot find a published reference

also assume that space is continuous. Now for any possible world where an object is at a certain location at a certain time there is another logically possible world where that object is located at a different location or occupies its location in the first world at a different time. For each such world there are uncountably many different logically possible worlds in which the object is located elsewhere.

How do we manage to reason about such large numbers of possibilities? The answer we want to propose here is that we reason in terms of types. A single type has a set of witnesses and there are no constraints on the cardinality of the set of witnesses. Types which have infinitely many witnesses are not more complex than types which have a small finite number of witnesses. Reasoning with a type involves manipulating the structural object which is the type itself not the set of its witnesses. Thus, for example, reasoning with a record type may be more complex than reasoning with a basic type that has no components. But still a record type is always a finite structure and so we are not entering into the complexity of manipulating uncountable sets, even though the record type may be thought of as a “representation” for its set of witnesses which may indeed be an uncountable set. It is here that our approach connects with proof theoretic approaches. In proof theory we manipulate expressions in a language which may represent sets of objects. Our types are not expressions in a language but they are objects in our type theoretic universe which could be thought of as “representing” the set of their witnesses. This approach also makes it possible to have a learning theory where agents can be acquainted with a type without being acquainted with the complete set of its witnesses. Knowing a type whose witnesses are dogs does not mean that you are acquainted with the set of all dogs, but rather that you know a dog when you see one, that is, you have a reliable dog *classifier*. An important aspect of human cognitive processing is that it involves reasoning with the types themselves, treating them as first class citizens which can be arguments to predicates. This is what gives rise to modality and intensionality. Possibly this higher level reasoning is unique, or at least, most fully developed in humans.

We think of types like record types as being types of situations. If we want to keep to the idea of possible worlds as total universes it is straightforward to convert a type of situations, T , to a type of worlds, T^W , as long as we have a way of defining worlds as maximal situations. We could say that a world, w , is of type T^W just in case some part, s of w is of type T . Actually, we do not need to do this because of the way we have set up subtyping. If T is a record type and $s : T$, then if $s < s'$, that is s is a proper part of s' in the sense defined in Chapter 5, then $s' : T$. If we had a way of defining maximal situations, that is, situations s such that there is no s' such that $s < s'$, we could take these to be our worlds. The problem is, though that it is not clear that it is desirable, or even possible, to characterize a notion of maximal situation in this sense. Certainly, there is no notion of maximal record so our choice of modelling situations as records suggests that there is no notion of maximal situation. Our axioms say that given any record it is always possible to add a new field to it.²

²This fact is parallel to Proposition 2 in Barwise (1989), Ch. 8: *Every situation, s , is a proper part of some other situation, s' .*

6.2 Modality without possible worlds

Montague (1973) introduces *necessarily* and *possibly* as sentence adverbs, that is, they combine with a sentence to produce another sentence. If α is a sentence, then *necessarily* α is true in a possible world, w , just in case α is true in every possible world and *possibly* α is true in a possible world, w , just in case there is some possible world in which α is true.³

In Chapter 1, Section 1.6 and Appendix A.10 we introduce modal type systems which are families of type systems, which we call *possibilities*, differing in their assignments of witnesses to basic types and ptypes. The important difference between possible worlds and possibilities is that for possibilities the parameters along which they can vary are fixed by the available types introduced in the type system, a well-defined notion, and one which varies depending on the particular type system. Thus we have a way of characterizing the dimensions along which the possibilities associated with a given type system vary and thus we have a way of representing the possibilities whereas we do not have such a way of characterizing possible worlds. We introduced modal notions relating to such modal type systems: essentially a type is necessary if it has a non-empty set of witnesses in every possibility and a type is possible if there is some possibility in which it has a non-empty set of witnesses. (For precise definitions see Appendix A.10.) Corresponding to the operators in modal logic we can introduce type constructors ‘ \Box ’ and ‘ \Diamond ’ as in (1).

- (1) If T is a type, then $\Box T$ and $\Diamond T$ are types

These types should obey the constraints in (2).

- (2) a. $\Box T$ is non-empty iff T is necessary (non-empty in all possibilities)
 b. $\Diamond T$ is non-empty iff T is possible (non-empty in some possibility)

In order to see how we can meet these constraints we have to first note that in a modal type system we cannot talk of an object a being of a type T *tout court* as we have done so far. a may be of type T in some possibilities but not others. This means that we have to relativize being of a type to possibilities, p which are members of a modal type system, \mathcal{P} . Instead of writing $a : T$, we will write $a :_{p,\mathcal{P}} T$ (“ a is of type T in possibility p within modal type system \mathcal{P} ”).⁴ We also correspondingly relativize our notation for the set of witnesses of a type as in (3).

³This simple treatment of modality corresponds to the modal logic system S5 where there is no restriction on accessibility between possible worlds (Hughes and Cresswell, 1968, 1996).

⁴Note that in Appendix A we have throughout the formal development of TTR always relativized the of-type relation to the type system being considered, and in the case of modal type systems in addition to the possibility (identified by the model associated with the possibility).

$$(3) \quad [T]_{p,\mathcal{P}} = \{a \mid a :_{p,\mathcal{P}} T\}$$

We introduce two basic types of types, *Nec* and *Poss*, the types of necessary and possible propositions respectively. The witness conditions for these types are given in (4).⁵

$$(4) \quad \begin{array}{ll} \text{a. } T :_{p,\mathcal{P}} \text{Nec} & \text{iff for all } p' \in \mathcal{P}, [T]_{p',\mathcal{P}} \neq \emptyset \\ \text{b. } T :_{p,\mathcal{P}} \text{Poss} & \text{iff for some } p' \in \mathcal{P}, [T]_{p',\mathcal{P}} \neq \emptyset \end{array}$$

Now consider that the inclusion of singleton types in our system (Appendix A.7) allows for the types Nec_T and Poss_T for any type, T . These types have a single witness, T , if T is necessary or possible respectively and otherwise have no witnesses. These types thus meet the constraints on $\Box T$ and $\Diamond T$ given in (2). We propose therefore to make the identifications given in (5).

$$(5) \quad \begin{array}{ll} \text{a. } \Box T = \text{Nec}_T \\ \text{b. } \Diamond T = \text{Poss}_T \end{array}$$

Note that we could also have defined ptypes corresponding to (5) by introducing predicates of ‘nec’ and ‘poss’ with arity $\langle \text{Type} \rangle$ which obey the constraints in (6).

$$(6) \quad \begin{array}{ll} \text{a. } [\text{nec}(T)] \neq \emptyset & \text{iff } T : \text{Nec} \\ \text{b. } [\text{poss}(T)] \neq \emptyset & \text{iff } T : \text{Poss} \end{array}$$

The option of using ptypes will become important below where we wish to add additional arguments to the predicate.

How many possibilities are there in a modal type system? The answer to this question is that there can be as many as you choose for the given type system, ranging from a small finite number of possibilities to a higher order infinity. The definition of a modal type system given in Appendix A.10 only requires that there be a family of possibilities. Thus this definition includes the kind of restricted sets of “possible worlds” differing along a small finite set of parameters which probability theorists talk of and indeed also linguistic semanticists talk of informally when they are in pedagogical explanatory mode (see, for example, Dowty *et al.*, 1981 and a lot of recent literature on inquisitive semantics such as Groenendijk and Roelofsen, 2012).

⁵Note that it is important for these types that we have introduced stratified types (Appendix A.11) since *Nec* and *Poss* can themselves be necessary and possible types. For example, instead of $\text{Nec} :_{p,\mathcal{P}} \text{Nec}$ we have $\text{Nec}^n :_{p,\mathcal{P}} \text{Nec}^{n+1}$ to avoid the danger of running into a version of Russell’s paradox. As usual we will suppress discussion of stratification in the text in order to simplify the presentation.

It is important in a modal type system that the identity criteria for the possibilities are determined by the types provided by the system. Two possibilities are distinct only if they differ in the witnesses associated with some basic type or ptype. It is not possible to make distinctions for which you do not have appropriate types available. Thus the range of possibilities is limited by the types which are available to classify objects.

This is not to say that we have eliminated all potential decidability problems from modal type systems. Of course, if the types that we use to construct the system are not decidable it may not be possible to decide on identity for possibilities. Even if all the types are guaranteed to be decidable, given an infinite set of possibilities there cannot be any general guarantee that we can decide whether an arbitrary type is necessary or possible or not since we cannot visit every possibility in a finite amount of time. We can only be sure if we have some general argument about the possibilities which does not involve inspecting each possibility individually. But having a way of distinguishing between possibilities which may in the limit be undecidable is better than not having a way of distinguishing between possibilities, other than that they are distinct members of a set.

The work on modality in natural language which has followed after Montague's original work all points to a more restricted kind of modality which involves arguing from some basic assumptions to a conclusion rather than considering all logical possibilities. This view of modality in natural language has been put forward by Kratzer in a body of work beginning with Kratzer (1977). This and other papers by Kratzer on modality are collected in revised and commented form in Kratzer (2012) and there is much other literature which builds on Kratzer's ideas. An excellent introduction to Kratzer's work is given in Chapter 3 of Portner (2009). The essential idea is that modals like *must* (corresponding to necessity) and *can* (corresponding to possibility) must be interpreted relative to a "conversational background" which in Kratzer (1981) (Chapter 2 of Kratzer, 2012) is split into two components, a *modal base* and an *ordering source*. The modal base is a set of propositions⁶ which characterize the assumptions from which we are arguing. The ordering source is a set of propositions⁷ which determine an ideal which we are trying to get as close to as possible. It is called an ordering source because Kratzer, following Lewis (1981), thinks of it as inducing a partial ordering on possible worlds, in terms of their closeness to the ideal. Kratzer's insight is that necessity and possibility in natural language should be defined relative to a modal base and an ordering source. In simple terms, a proposition, p , is necessary with respect to a modal base, b , and an ordering source (ideal), i , just in case p follows from the conjunction of b and i . A proposition, p , is possible with respect to b and i just in case p is consistent with the conjunction of b and i .

We shall construe Kratzer's propositions as types and we shall take modal bases and ideals to be types as well. To recreate a Kratzerian semantics for necessity and possibility we let the predicates 'nec' and 'poss' have arity $\langle \text{Type}, \text{Type}, \text{Type} \rangle$ and require that they obey the constraints in

⁶Actually, a function which determines a set of propositions for each possible world.

⁷Again relativized to possible worlds.

(7).

- (7) a. $[\text{nec}(T, B, I)]_{p, \mathcal{P}} \neq \emptyset$ iff for any $p' \in \mathcal{P}$, if both $[B]_{p', \mathcal{P}} \neq \emptyset$ and $[I]_{p', \mathcal{P}} \neq \emptyset$ then $[T]_{p', \mathcal{P}} \neq \emptyset$
 b. $[\text{poss}(T, B, I)] \neq \emptyset$ iff for some $p' \in \mathcal{P}$, $[B]_{p', \mathcal{P}} \neq \emptyset$, $[I]_{p', \mathcal{P}} \neq \emptyset$ and $[T]_{p', \mathcal{P}} \neq \emptyset$

Building on a basic example from Portner, 2009, p. 49, suppose that T is *Mary-eat-her-broccoli*, B is *Mary-has-broccoli-on-her-plate* and I is *Mary-eats-everything-on-her-plate*. Then according to the definitions in (7) $\text{nec}(T, B, I)$ is non-empty just in case for any of the possibilities we are considering if both B and I are non-empty then T is non-empty, that is if there's a situation where Mary has broccoli on her plate and there's a situation where Mary eats everything on her plate then there's a situation in which Mary eats her broccoli. Similarly, $\text{poss}(T, B, I)$ is non-empty just in case there is some possibility that we are considering where there's a situation in which Mary has broccoli on her plate, a situation in which Mary eats everything on her plate and a situation in which Mary eats her broccoli.

A more restrictive notion of necessity than is given in (7a) would be in terms of subtyping as in (8).

- (8) $[\text{nec}(T, B, I)]_{p, \mathcal{P}} \neq \emptyset$ iff $B \not\perp I$ and $(B \wedge I) \sqsubseteq_{\mathcal{P}} T$

(Here we are using $\not\perp$ for “does not preclude”, that is, it is possible for something to be of both types.)

(8) requires that anything of type $B \wedge I$ will also be of type T (no matter what gets assigned to the basic types and ptypes) whereas (7) only requires that if there is something of type B and there is something of type I then there will also be something of type T , though not necessarily the same thing. The subtyping variant is interesting because if you have a way of (at least approximately) computing whether one type is a subtype of another simply by looking at the types, then you will not have to look at the different possibilities. Similarly, for possibility we may have a way of computing (at least approximately) that a type is instantiable simply by looking at the type and doing a consistency check without having to inspect the possibilities. This points towards a more proof theoretic oriented approach to modality. Part of the important insight of Kratzer's approach to modality is that it involves arguments which can be constructed from the modal base and the ideal.

(8) also seems to fit better with the particular broccoli example we are discussing. $\text{nec}(T, B, I)$ will be non-empty just in case Mary having broccoli on her plate does not preclude her eating everything on her plate and in all of the possibilities under consideration any situation in which

she has broccoli on her plate and eats everything on her plate is also a situation in which she eats her broccoli.

When Kratzer talks of the conversational background consisting of the base and the ideal she often talks about rules that might be encoded there (bodies of laws or regulations in the case of deontic modality). This idea of rules being involved actually fits better with the broccoli example. It is not so much that we are considering possibilities where Mary eats everything on her plate, but rather that we are considering possibilities where there is a rule that Mary eats whatever is on her plate.

It is important for Kratzer that such rules not be logical laws in the sense that they always hold true. For example, a law that cars not park on double yellow lines does not entail that cars do not park on double yellow lines – this is only something that holds true in deontically ideal worlds. This suggests that there could be a role for what Breitholtz (2014) calls *topoi*. A *topos* in her terms is a dependent type, that is, a function which maps an object of some type to a type. Given a situation of the domain type of the *topos*, the *topos* will return a new type. The standard licensing condition associated with a *topos* is similar to the licensing condition we have for, for example, sign combination functions in Chapter 3 (see also Appendix B.1.4.2). This is given in (9).

(9) If $\tau : (T \rightarrow \text{Type})$ is a *topos* available to agent A , then for any $s, s :_A T$ licenses $:_A \tau(s)$

That is, if an agent, A , judges a situation s to be of the domain type of the *topos*, then A is licensed to judge that there is something of type $\tau(s)$.

We will use the same trick that we used for the polymorphism of properties in Chapter 5 in characterizing the type *Topos*. That is, we will define *Topos* to be the type in (10).

(10)
$$\left[\begin{array}{ll} \text{bg} & : \text{Type} \\ \text{fg} & : (\text{bg} \rightarrow \text{Type}) \end{array} \right]$$

This means that we can reformulate the licensing condition in (9) as (11).

(11) If τ is a *topos* available to agent A , then for any $s, s :_A \tau.\text{bg}$, licenses $:_A \tau.\text{fg}(s)$

We will say that *topoi* associated with this condition are *epistemic*. The condition has to do with increasing our knowledge on the basis of a previous judgement. If we judge something, s , to be of the type which is the background of the *topos* then we can judge that there is something of the type resulting from applying the foreground of the *topos* to s .

Topoi can also be *deontic*, that is, they are associated with a condition which involves an obligation to carry out a certain act (create something of a given type). This condition is as in (12).

(12) If τ is a topos available to agent A , then for any s , $s :_A \tau.\text{bg}$ obliges $:_A \tau.\text{fg}(s)!$

That is, if an agent, A , judges a situation, s , to be of the background type of the topos, then A is obliged to create (contribute to the creation of) something which is of the type resulting from applying the foreground of the topos to s .

Topoi can be associated with either of these conditions and they can be associated with both which means that they can be used either epistemically or deontically.

We now replace the third “ideal” type argument to the predicates ‘nec’ and ‘poss’ with a topos argument, giving them the arity $\langle \text{Type}, \text{Type}, \text{Topos} \rangle$. If we need to recreate the option provided by the type rather than the topos we can use a topos whose background type is the type Rec . That is, it does not place any constraints on the situations in its domain and thus will return a type for any situation. If such a function is a constant function, that is, it returns the same type for any situation, then this will give us the same effect as we obtained when the argument was a type rather than a topos.

We define the witness conditions in (13) for the new versions of ‘nec’ and ‘poss’.

(13) a. If T and B are types and τ is a topos, then

$s : \text{nec}(T, B, \tau)$ iff

$s : B$,

$B \sqsubseteq \tau.\text{bg}$ and

$\tau.\text{fg}(s) \sqsubseteq T$

b. If T and B are types and τ is a topos, then

$s : \text{poss}(T, B, \tau)$ iff

$s : B$,

$B \sqsubseteq \tau.\text{bg}$ and

$\tau.\text{fg}(s) \not\sqsubseteq T$

In informal terms, (13) says that a situation, s , witnesses that a type, T , is necessary with respect to a background type, B , and a topos, τ , just in case s is of the type B , τ is defined on situations of type B and the type resulting from the application of τ to s is such that any situation of that type will be of type T . It says that T is possible under the same conditions except that the third condition is changed to requiring that the type resulting from the application of τ to s does not preclude T , i.e. that it is possible for a situation to be of both types.

Let us see how this might play out in our basic example (taken from Portner, 2009, p. 49). Consider (14).

(14) Mary should eat her broccoli

Portner points out that this sentence can receive a bouletic (having to do with desires) interpretation if “we are talking about the fact that Mary loves broccoli” while “if we are trying to enforce the idea that children should eat everything on their plates, it naturally receives a deontic interpretation”. Suppose that b is the broccoli on Mary’s plate. For simplicity we will assume $b : Ind$. Let m be Mary and p her plate. Then the type, B , of the base situation could be (15).

$$(15) \quad \left[\begin{array}{ll} x=b & : \quad Ind \\ c_1 & : \quad \text{broccoli}(x) \\ y=m & : \quad Ind \\ c_2 & : \quad \text{child}(y) \\ z=p & : \quad Ind \\ c_3 & : \quad \text{plate}(z) \\ e_1 & : \quad \text{have}(y,z) \\ e_2 & : \quad \text{on}(x,z) \end{array} \right]$$

Let us in addition assume that broccoli is food, that is, (16) holds.

(16) For any a , $\text{broccoli}(a) \sqsubseteq \text{food}(a)$

Now let us introduce two topoi, τ_1 and τ_2 . We represent their foregrounds in (17a and b) respectively.

$$(17) \quad \text{a. } \lambda r: \left[\begin{array}{l} x:Ind \\ c_1:\text{food}(x) \\ y:Ind \\ c_2:\text{child}(y) \\ z:Ind \\ c_3:\text{plate}(z) \\ e_1:\text{have}(y,z) \\ e_2:\text{on}(x,z) \end{array} \right] \cdot \left[e : \text{eat}(r.y, r.x) \right]$$

$$\text{b. } \lambda r: \left[\begin{array}{l} x:Ind \\ c_1:food(x) \\ y:Ind \\ c_2:child(y) \\ e:love(y,x) \end{array} \right] . \left[e : eat(r.y, r.x) \right]$$

(17a) associates the type of situation where a child has food on her plate with the type of situation where the child eats that food. This topos is naturally associated with a deontic condition, that is, a child is obliged to create a situation of the type returned by the topos, to eat the food on her plate. (17b) associates the type of situation where there is food which the child loves with the type of situation where the child eats that food. This topos is naturally associated with what we might call a *bouletic* condition, that is, we can use the topos to reason that the child has a desire to create a situation of the type returned by the topos, that is, the child wants to eat the food. This involves a kind of condition which we have not talked about yet which associates types with mental states rather than actions. We will discuss this more in Section 6.3.

The type corresponding to *Mary should eat her brocolli* based on these resources could be either of the types in (18), where T_{broc} is (15) and τ_1 and τ_2 are (17a and b) respectively.

$$\begin{aligned} (18) \quad \text{a. } & nec([e:eat(m, b)], T_{broc}, \tau_1) \\ \quad \text{b. } & nec([e:eat(m, b)], T_{broc}, \tau_2) \end{aligned}$$

We can now check the witness conditions in (13). Any s which is of the type (18a) has to fulfil the conditions in (19).

$$\begin{aligned} (19) \quad \text{a. } & s : T_{broc} \\ \quad \text{b. } & T_{broc} \sqsubseteq \left[\begin{array}{l} x:Ind \\ c_1:food(x) \\ y:Ind \\ c_2:child(y) \\ z:Ind \\ e_1:have(y,z) \\ e_2:on(x,z) \end{array} \right] \\ \quad \text{c. } & \tau_1(s) \sqsubseteq [e:eat(m, b)] \end{aligned}$$

Assuming that s meets (19a), we can check that (19b) holds by noting that anything of the first type will also be of the second type. (In this case, the two types are identical except for (i) ‘brocolli’ in the first type corresponds to ‘food’ in the second, but we know from (16) that brocolli is food (ii) the manifest fields in the first type correspond to non-manifest fields in the

second, but we know from the definition of singleton types represented by manifest fields that they are subtypes of the corresponding non-singleton type.) We can see that (19c) will hold given our characterization of τ_1 in (17) since $\tau_1(s)$ will be (20a) and given that $s : T_{\text{broc}}$, $s.y$ will be m and $s.x$ will be b . Thus $\tau_1(s)$ is identical with (20b).

- (20) a. $\left[e : \text{eat}(s.y, s.x) \right]$
 b. $\left[e : \text{eat}(m, b) \right]$

Thus (19) is checking that the type $\left[e : \text{eat}(m, b) \right]$ is a subtype of itself and, of course, any type is a subtype of itself.

We can make a similar argument for (18b).

This is an inferential view of modality in the sense that the topoi, which correspond to patterns of inference, have taken over the work of the accessibility relations between possible worlds which Kratzer uses. Note that while it might appear from our formulation of the witness conditions for ‘nec’ and ‘poss’ that we have a definition of modal predicates which does not use the previous notion of modality that we had in terms of possibilities defined in varying the assignments to basic types and ptypes, this is in fact not the case since our definitions of subtyping and preclusion rely on this kind of modality. Thus these definitions have both an inferential flavour (in that they use topoi which are similar to rules of inference) and also a Kripke model flavour in that they use sets of possibilities.

While the use of topoi here gives us something corresponding to accessibility relations in Kratzer’s treatment of modality in Kratzer (1977) (Kratzer, 2012, Chapter 1), it does not yet give us anything corresponding to the notion of ordering source introduced in Kratzer (1981) (Kratzer, 2012, Chapter 2) to deal with the different degrees of modality expressed in examples like

- (21) a. Mary absolutely must eat her broccoli
 b. Mary must eat her broccoli
 c. Mary ought to eat her broccoli
 d. Mary should eat her broccoli

While it is not obvious that there is a fixed order of strength in (21) it is nevertheless the case that speakers of English will perceive differences of strength in the modalities having to do with how necessary it is for Mary to eat her broccoli. For that we need the notion of preference structure as it is discussed in Condoravdi and Cooper (????).

Let us take a look at how these ideas can be exploited in a compositional semantics. In order to do a compositional semantics for modal verbs we need to distinguish between tensed and non-tensed verbs. Our strategy for the structure of sentences with modal verbs is represented by the informal tree in (22).



That is, we treat the modal *should* as combining with a non-tensed verb phrase to form a tensed verb-phrase.

For simplicity of discussion let us consider the intransitive verb *eat* rather than the complex verb phrase *eat her broccoli*. The version of the operation ‘SemIntransVerb’ defined in Chapter 5 (given in Appendix B.1.4.1) yields (23) when applied to the predicate ‘eat’ with arity $\langle Ind \rangle$ and no restrictions introduced on the domain of the content or background conditions on the context. That is, (23) is $\text{SemIntransVerb}(\text{eat}, Ind, Ind, Rec)$.

$$(23) \left[\begin{array}{l} \text{bg} = Rec \\ \text{fg} = \lambda c:Rec . \left[\begin{array}{l} \text{bg} = Ind \\ \text{fg} = \lambda r:[x:Ind] . [e : \text{eat}(r.x)] \end{array} \right] \end{array} \right]$$

This parametric content for an utterance of *eat* requires that a sentence such as *Mary eats* has a content which is the event type (24) (assuming *m* is Mary).

$$(24) [e : \text{eat}(m)]$$

This type does not require any relationship between the eating event and the utterance event. We therefore conclude that this corresponds best to a tenseless expression. It says nothing about when an event of this type needs to occur. Note that this is something that is natural in a system based on types whereas in a semantics based on the kind of tense operators we find in tense logic it is not so straightforward to represent that content of a non-tensed utterance. How could we

modify the type in (24) to represent the relationship of the eating event to some particular speech event, s ? In a simple-minded tense system there are basically three possibilities. The eating event is either required to be simultaneous with s , prior to s or after s . We model this by creating event types for events which have two components, the speech event, s and the eating event. The types are given in (25a–c) corresponding to *Mary eats*, *Mary ate* and *Mary will eat*, respectively.

- (25) a. $\left[\begin{array}{ll} \text{s-event}=s & : \text{SEvent} \\ e & : \text{eat}(m) \end{array} \right]$
 b. $\left[e : \text{eat}(m) \right] \cap \left[\text{s-event}=s : \text{SEvent} \right]$
 c. $\left[\text{s-event}=s : \text{SEvent} \right] \cap \left[e : \text{eat}(m) \right]$

Of course, we would expect an actual tense and aspect system for a natural language to involve more complex types than this, for example, allowing partial overlap between the eating event and the speech event. Our aim here is not to develop a realistic account of tense but rather to show how we can distinguish between tensed and tenseless contents in the kind of system we are proposing. The types in (25) can be derived from (24) by tense operators which take a speech event and a type as arguments and return a new type. These operators are defined in (26).

(26) If $s : \text{SEvent}$ and T is a type, then

1. $\text{pres}(s)(T) = T \wedge \left[\text{s-event}=s : \text{SEvent} \right]$
2. $\text{past}(s)(T) = T \cap \left[\text{s-event}=s : \text{SEvent} \right]$
3. $\text{fut}(s)(T) = \left[\text{s-event}=s : \text{SEvent} \right] \cap T$

In a more complete treatment of tense we might want to generalize these operators so that they can relate types to other kinds of events in addition to speech events in order to be able to deal with embedded tenses and phenomena like the historic present (as in *So I was in the pub and this man **comes** up to me ...*).

In addition to non-tensed contents for verbs, as illustrated by the result of applying ‘SemIntransVerb’ given in (23), we will also have tensed contents for verbs. Thus in addition to ‘SemIntransVerb’ we will also have ‘SemIntransVerb $_{\alpha}$ ’ where α is one of ‘pres’, ‘past’ and ‘fut’. These functions will return a function from speech events to parametric contents as given by the example in (27).

(27) $\text{SemIntransVerb}_{\alpha}(\text{eat}, \text{Ind}, \text{Ind}, \text{Rec}) =$

$$\lambda s : \text{SEvent} . \left[\begin{array}{ll} \text{bg} & = \text{Rec} \\ \text{fg} & = \lambda c : \text{Rec} . \left[\begin{array}{ll} \text{bg} & = \text{Ind} \\ \text{fg} & = \lambda r : [\text{x} : \text{Ind}] . \alpha(s)([e : \text{eat}(r.x)]) \end{array} \right] \end{array} \right]$$

This indicates that the contents of tensed expressions depend on speech event in a way that non-tensed expressions do not.

We now turn our attention to how information about tense plays a role in sign types. Recall that in Chapter 3 we defined *Sign* as a recursive type whose witness condition is as in (28). (See also Appendix B.1.)

$$(28) \quad \sigma : \textit{Sign} \text{ iff } \sigma : \left[\begin{array}{ll} \text{s-event} & : \textit{SEvent} \\ \text{syn} & : \textit{Syn} \\ \text{cnt} & : \textit{Cnt} \end{array} \right]$$

Here the type *Syn* (for “syntax”) was defined as in (29).

$$(29) \quad \left[\begin{array}{ll} \text{cat} & : \textit{Cat} \\ \text{daughters} & : \textit{Sign}^* \end{array} \right]$$

Now we are going to add a further field to this type to indicate whether a sign is tensed or non-tensed. The new definition of *Syn* is given in (30).

$$(30) \quad \left[\begin{array}{ll} \text{cat} & : \textit{Cat} \\ \text{tns} & : \textit{Bool} \\ \text{daughters} & : \textit{Sign}^* \end{array} \right]$$

The definitions of the category sign types in Chapter 3 (see Appendix B.1) for *S*, *V* and *VP* can remain the same, since these categories are underspecified for tense; they can be either tensed or non-tensed. We will use (31a) to represent the type (31b) and (31c) to represent the type (31d) and we will do similarly for *VP* and *S*.

$$(31) \quad \begin{array}{ll} \text{a. } V_{[+tns]} & \\ \text{b. } \textit{Sign} \wedge \left[\text{syn} : \left[\begin{array}{ll} \text{cat=v} & : \textit{Cat} \\ \text{tns=1} & : \textit{Bool} \end{array} \right] \right] & \\ \text{c. } V_{[-tns]} & \\ \text{d. } \textit{Sign} \wedge \left[\text{syn} : \left[\begin{array}{ll} \text{cat=v} & : \textit{Cat} \\ \text{tns=0} & : \textit{Bool} \end{array} \right] \right] & \end{array}$$

We will assume that the categories *NP*, *Det* and *N* are universally untensed⁸ and therefore take *NP* to be the type (32) and similarly for *Det* and *N*.

$$(32) \text{ Sign } \wedge \left[\text{syn} : \left[\begin{array}{ll} \text{cat=np} & : \text{Cat} \\ \text{tns=0} & : \text{Bool} \end{array} \right] \right]$$

We now define tensed versions of the universal resource for lexical sign type construction, $\text{Lex}_{\text{IntransVerb}}$ as defined in Chapter 5 (also in Appendix B.1.4.1). Letting α stand for ‘past’, ‘pres’ or ‘fut’ we characterize $\text{Lex}_{\text{IntransVerb}_\alpha}$ as in (33).

- (33) $\text{Lex}_{\text{IntransVerb}_\alpha}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where T_{phon} is a phonological type, p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type is defined as

$$\text{Lex}(T_{\text{phon}}, VP) \wedge \left[\begin{array}{l} \text{s-event:SEvent} \\ \text{syn:}[\text{tns=1:Bool}] \\ \text{cnt=SemIntransVerb}_\alpha(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})(\text{s-event}):PPpty \end{array} \right]$$

We will use ‘ $\text{Lex}_{\text{IntransVerb}}$ ’ (without the α) to construct sign types for non-finite verbs characterized by (34).

- (34) $\text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where T_{phon} is a phonological type, p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type is defined as

$$\text{Lex}(T_{\text{phon}}, VP) \wedge \left[\begin{array}{l} \text{syn:}[\text{tns=0:Bool}] \\ \text{cnt=SemIntransVerb}_\alpha(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}):PPpty \end{array} \right]$$

We now turn our attention to the modal verbs. The parametric content of a modal verb (such as *should*) is a function which requires a background with a modal base (a type) and a topos. Given such a background this function returns a function from properties (such as *eat*) to properties (such as *should eat*). We define ‘ $\text{SemModalVerb}_{\text{nec}}$ ’ and ‘ $\text{SemModalVerb}_{\text{poss}}$ ’ as (35a and b) respectively.

⁸This is something of an open question. See Tonhauser (2007) for discussion.

$$\begin{array}{l}
(35) \text{ a. } \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} \text{base} : \text{Type} \\ \text{topos} : \text{Topos} \end{array} \right] \\ \text{fg} = \lambda c: \left[\begin{array}{l} \text{base:Type} \\ \text{topos:Topos} \end{array} \right] . \\ \left[\begin{array}{l} \text{bg} = \text{Ppty} \\ \text{fg} = \lambda P:\text{Ppty} . \\ \left[\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r:[x:\text{Ind}] . \\ [e : \text{nec}(P.\text{fg}(r), c.\text{base}, c.\text{topos})] \end{array} \right] \end{array} \right] \end{array} \right] \\
\text{b. } \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} \text{base} : \text{Type} \\ \text{topos} : \text{Topos} \end{array} \right] \\ \text{fg} = \lambda c: \left[\begin{array}{l} \text{base:Type} \\ \text{topos:Topos} \end{array} \right] . \\ \left[\begin{array}{l} \text{bg} = \text{Ppty} \\ \text{fg} = \lambda P:\text{Ppty} . \\ \left[\begin{array}{l} \text{bg} = \text{Ind} \\ \text{fg} = \lambda r:[x:\text{Ind}] . \\ [e : \text{poss}(P.\text{fg}(r), c.\text{base}, c.\text{topos})] \end{array} \right] \end{array} \right] \end{array} \right]
\end{array}$$

The type, *Modal*, of modal parametric contents, that is, a type of objects like those in (35), is given in (36).

$$(36) \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} \text{base:Type} \\ \text{topos:Topos} \end{array} \right] : \text{Type} \\ \text{fg} : (\text{bg} \rightarrow (\text{Ppty} \rightarrow \text{Ppty})) \end{array} \right]$$

We will introduce a syntactic category for modal verbs, ‘vm’. Thus we will now characterize the type, *Cat* as in (37).

$$(37) \text{ s, np, det, n, v, vp, vm} : \text{Cat}$$

We will use the symbol $V_{[+M]}$ to represent the type (38).

$$(38) \text{ Sign} \wedge \left[\text{syn} : \left[\text{cat=vm} : \text{Cat} \right] \right]$$

We can now characterize a universal resource, $\text{Lex}_{\text{ModalV}}$, for creating lexical sign types for modal verbs. This is done in (39).

- (39) If T_{phon} is a phonological type and p is either ‘nec’ or ‘poss’, then $\text{Lex}_{\text{ModalV}}(T_{\text{phon}}, p)$ is defined as

$$\text{Lex}(T_{\text{phon}}, \underset{[+M]}{V}) \wedge \left[\begin{array}{l} \text{s-event: } SEvent \\ \text{syn: } [tns=1: Bool] \\ \text{cnt= } \left[\begin{array}{l} \text{bg= } \left[\begin{array}{l} \text{base: } Type \\ \text{topos: } Topos \end{array} \right] \\ \text{fg= } \lambda c: \left[\begin{array}{l} \text{base: } Type \\ \text{topos: } Topos \end{array} \right] . \text{pres(s-event)(SemModalVerb}_p.\text{fg}(c)) \end{array} \right] : Modal \end{array} \right]$$

The lexical resources for English can now tell us that *should* is a modal verb of necessity, as in (40).

- (40) $\text{Lex}_{\text{ModalV}}(\text{“should”}, \text{nec})$

Finally, English resources will need to include the tense and modal sensitive phrase structure rules in (41) (using the abbreviatory conventions of Appendix B.2.4).

- (41) a. $\underset{[+tns]}{S} \longrightarrow \underset{[+tns]}{NP} \underset{[+tns]}{VP} \mid \underset{[+tns]}{NP'} @ \underset{[+tns]}{VP'}$
 b. $\underset{[+tns]}{VP} \longrightarrow \underset{[+M]}{V} \underset{[-tns]}{VP} \mid \underset{[+M]}{V'} @ \underset{[-tns]}{VP'}$

6.3 Intensionality without possible worlds

In Section 6.1 we discussed problems that have to do with individuating and counting possible worlds. Here we discuss well-known problems that arise when you consider propositions to be the sets of possible worlds⁹ which make them true. The central problem is that the sets of possible worlds provide a too coarse-grained analysis of propositions. There are intuitively distinct propositions which are true in the same sets of possible worlds. Standard examples of this are mathematical propositions. Mathematical propositions are not contingent, that is, they are either true in every possible world or false in every possible world. The view of propositions as sets of possible worlds has the consequence that there are only two mathematical propositions: the necessarily true proposition and the necessarily false proposition. It seems unintuitive to reduce a rich field of continuing investigation where new “propositions” are still being discovered and proved or disproved to a field where just two propositions are being discussed. Clearly, mathematics involves a different intuitive notion of proposition that is not modelled by a set of possible worlds. One might be tempted to think that this is a problem about mathematics rather than natural language and that for normal every day dialogue we can ignore this problem.

⁹Or, if we are concerned with tensed propositions, sets of pairs of possible worlds and times.

Perhaps we just do not normally talk about necessary propositions or at least what we think of as being necessarily true is in fact relativized in the way that we discussed above in relation to Kratzer's semantics for modality. This is a dangerous route to pursue, not least perhaps because, although many of us do not spend a lot of our time talking about mathematical propositions, we are nevertheless able to express mathematical propositions in natural language and to ignore them would be to rule out something that is part of linguistic activity. There are many of us who are not mathematicians who can nevertheless understand that there is a difference in the content of the two examples in (42).

- (42) a. Andrew Wiles proved that two plus two equals four
 b. Andrew Wiles proved that Fermat's last theorem is true

If the correct notion of proposition for natural language was that propositions are sets of possible worlds then we should have difficulty in distinguishing the content of these two sentences.

There are non-mathematical candidates for propositions that would be true in all possible worlds. King (2014) points to examples like (43).

- (43) a. Bachelors are unmarried
 b. Brothers are male siblings

These are examples of what are sometimes called analytic sentences, true in virtue of their meaning. Despite the considerable difficulties with the notion of analyticity (see Rey, 2015, for discussion), it is nevertheless hard to think of a possible world where one of these sentences is true and the other is false. Yet they seem to correspond to different propositions. It does not seem attractive to say that all analytic sentences express the one and only analytic proposition (which in addition is identical with the true mathematical proposition).

There are also examples of sentences, such as those in (44), which we can argue that they express different propositions although they are true in the same possible worlds.

- (44) a. Kim sold *Syntactic Structures* to Sam
 b. Sam bought *Syntactic Structures* from Kim

An early reference to the equivalence relationship between *buy* and *sell* in the linguistic literature is Fillmore (1970) where it is stated:

There are no situations that can in themselves be distinguished as buying situations or selling situations; but the choice of one or another of these verbs seems to make it possible to speak of a buying/selling transaction from one of the participant's point of view.

In our terms we would want to say that the ptypes $\text{buy}(a,b,c)$ and $\text{sell}(c,b,a)$ are distinct types which have the same witnesses. In terms of propositions as sets of possible worlds we would be committed to claim that these sentences express the same proposition.

The problem is not just a matter of what we intuitively consider to be distinct propositions. It has consequences for the truth of sentences with sentential complements after verbs like *believe* and *know*, the verbs of propositional attitude. If we analyze these verbs in terms of relations between individuals and propositions and we treat propositions as sets of possible worlds then for some individual, a , if a believes/knows p and p is logically equivalent to q (that is, is true in the same possible worlds which in turn means that p and q are the same proposition) then a believes/knows q . This has the unfortunate consequence that once you know one logical truth you know them all. So, for example, somebody who knows that the sum of 2 and 2 is 4 also knows any other mathematical truth (since they are all the same proposition), as well as any analytic truth and any logically valid truths. The problem extends beyond propositions that are true in all, or no, possible worlds. For any two propositions that are true in the same possible worlds (that is, are logically equivalent) if you know or believe one of them then you also know or believe the other. It interacts with the idea (originally advanced by Kripke, 1972) that proper names should be rigid designators, that is, that they should have the same denotation in every possible world. One of the puzzles goes back to discussion by Frege (1892). In the ancient world people believed that the morning star and the evening star were distinct heavenly bodies, whereas they are in fact both the planet Venus. The “morning star” had the name *Phosphorus* and the “evening star” had the name *Hesperus*. If both these names refer to the same planet Venus in all possible worlds then *Phosphorus rose in the morning* expresses the same proposition as *Hesperus rose in the morning*, that is, the two sentences are true in the same possible worlds, though they are not true in all possible worlds. Yet it seems reasonable to say that the Ancients believed that Phosphorus rose in the morning but that they did not believe that Hesperus rose in the morning. Frege's original puzzle, which is also problematic for the view that propositions are sets of possible worlds concerned the difference between *The Ancients believed that Hesperus is Hesperus* (true if they believed in the law of self identity which they presumably did) and *The Ancients believed that Hesperus is Phosphorus* (false, since it was an astronomical discovery that both Hesperus and Phosphorus denote the planet Venus). Yet both *Hesperus is Hesperus* and *Hesperus is Phosphorus* represent the same proposition, the one that is true in all possible worlds. As we noted in Chapter 4 this problem is related to Kripke's Paderewski puzzle which we discussed there and we will build on our analysis of proper names in that chapter in our analysis of the attitudes in this chapter.

The example of the equivalence of *buy* and *sell* may initially seem like an argument for the

straightforward possible worlds approach when we consider propositional attitudes like *believe* and *know*. It seems impossible that any rational agent who believes or knows one of (44) would not know or believe the other. However, there are other attitude predicates where it does seem feasible to make the distinction. The sentences in (45) do not seem to be contradictory.

- (45) a. Chris was happy that Kim bought *Syntactic Structures* from Sam
 b. Chris was not happy that Sam sold *Syntactic Structures* to Kim

There are other non-attitude predicates which also make the distinction. For example, in Sweden it is illegal to buy sex but not illegal to sell sex which has important consequences for who gets punished in a situation where sex is bought and sold. Thus the sentences in (46) are consistent when considering Swedish law.

- (46) a. It was illegal that Kim bought sex from Sam
 b. It was not illegal that Sam sold sex to Kim

These problems have been well known since the early days of formal semantics. There is an excellent overview of the discussion up to the end of 1970's in Dowty *et al.* (1981), 170ff. Partee (1979) provides an important account of relevant issues. For a modern update of Partee's view see Partee (2014). For some modern philosophical views of propositions which go in somewhat similar directions to the proposals here, linking propositions to perception and action, see King *et al.* (2014).

Our basic strategy here is to replace the notion of propositions as sets of possible worlds with the notion of propositions as types, which goes back to work in intuitionistic logic (see discussion by Ranta, 1994, for a relation of this idea to linguistic semantics, and Wadler, 2015, for an overview of the history of the idea from the perspective of logic and computer science). There is a more sophisticated view of propositions in TTR which was advanced by Ginzburg (2012) and used, for example, in Cooper *et al.* (2015). This is that we should regard propositions as pairs of a situation and a type (that is, a record with two fields). This is the notion of Austinian proposition which goes back to Barwise and Perry (1983) who coined the term because of the proposal in Austin (1961) that propositions should incorporate the part of the world which they are true (or false) of. Both of these notions of proposition exploit the intensionality of types, the fact that you can have two distinct types with the same set of witnesses. A type used as a proposition is true just in case there is something of the type. This makes types as propositions parallel to what was called a Russellian proposition in Barwise (1989), Chap. 11. An Austinian proposition is true just in case the situation in the proposition is of the type of the proposition. An Austinian proposition is a way of reifying a judgement, that is, it gives us an object in our type theoretic universe which corresponds to the act of judging a particular situation to be of a type (a record of such a judgement). This means that if a Russellian proposition is true then there is an Austinian

proposition containing the same type which is true. If an Austinian proposition is true then the corresponding Russellian proposition is true. If a Russellian proposition is false then any Austinian proposition containing the same type is also false. However, if an Austinian proposition is false, then we cannot conclude from this either the truth or falsity of the corresponding Russellian proposition. We know that the particular situation in the Austinian proposition is not of the type in the Austinian proposition but this tells us nothing about whether there is some other situation of the type.

Neither “proposition” nor “Russellian proposition” are technical terms in TTR. This is because we can judge any type to be non-empty (“true”) or empty (“false”) and thus any type can be used as a proposition. In practice, however, we will take record types (intuitively, types of situations) to be what corresponds to the intuitive notion of propositions that can be expressed in natural language. The simplest theory of verbs of propositional attitude like *believe* and *know* on this kind of view would be that they correspond to predicates which express relations between individuals and record types, that is, there are predicates ‘believe’ and ‘know’ with arity $\langle \text{Ind}, \text{RecType} \rangle$. This means that we will have a ptype like (47) where a is an individual and T is a record type.

(47) $\text{believe}(a, T)$

What does it mean for this type to be non-empty? We will say that it involves finding a match, in the sense introduced in Chapter 4, for T in a ’s long term memory. In the terms introduced in Chapter 4 this means that if r is a ’s total information state, then a ’s long term memory will be $r.\text{ltm}$, which is a record type, a type representing how the world would be if a ’s long term memory were true. Thus we are matching the type T , a record type which is the second argument of ‘believe’, against another record type corresponding to a ’s long term memory. Note that according to the proposal for matching in Chapter 4 this involves finding a relabelling for T . The match obtains if there is a relabelling, η , of T , such that $r.\text{ltm} \sqsubseteq [T]_\eta$, where $[T]_\eta$ is the result of relabelling T by η (see Appendix A.14). Let us introduce an abbreviatory notation for this as in (48).

(48) $T_1 \sqsubseteq_{\rightsquigarrow} T_2$ just in case there is some relabelling, η , of T_2 such that $T_1 \sqsubseteq [T_2]_\eta$.

Our preliminary witness conditions for $\text{believe}(a, T)$ are given in (49). (We will modify this below.)

(49) $e : \text{believe}(a, T)$ iff
 $e : \text{ltm}(a, T')$
 and $T' \sqsubseteq_{\rightsquigarrow} T$

The fact that relabelling is involved in the matching process is important for the analysis of belief because it means that (50) holds.

- (50) If $\text{believe}(a, T)$ is non-empty, then for any relabelling, η , of T , $\text{believe}(a, [T]_\eta)$ is non-empty

This means that the choice of particular labels in a record type is not relevant when we compute whether an agent stands in the belief (or other attitude) relation to a record type. Note also that, given the way we have defined relabelling in Appendix A.14 via relabellings of a flattened representation of the type, that record types which are structured differently, as in (51a,b), will also count as relabellings of each other, in this example in virtue of the relabelling (51c).

$$\begin{aligned}
 (51) \quad & \text{a. } \left[\begin{array}{c} \ell_1 : \left[\begin{array}{c} \ell_2 : T_1 \\ \ell_3 : T_2 \end{array} \right] \\ \ell_4 : T_3 \end{array} \right] \\
 & \text{b. } \left[\begin{array}{c} \ell_1 : T_1 \\ \ell_2 : \left[\begin{array}{c} \ell_3 : T_2 \\ \ell_4 : T_3 \end{array} \right] \end{array} \right] \\
 & \text{c. } \ell_1.\ell_2 \rightsquigarrow \ell_1 \\
 & \quad \ell_1.\ell_3 \rightsquigarrow \ell_2.\ell_3 \\
 & \quad \ell_4 \rightsquigarrow \ell_2.\ell_4
 \end{aligned}$$

Thus any agent who stands in the belief-relation to (51a) will also stand in the belief-relation to (51b) and *vice versa*. The intuition here is that two agents will have the same beliefs even though they structure the information differently in their separate long term memories.

This can be contrasted with proposals for structured meanings in the possible worlds literature, starting with Lewis (1972), who based his idea on the notion of intensional isomorphism from Carnap (1956), and developed by Cresswell (1985). The idea here is that you alleviate the coarse-grainedness of the possible worlds analysis of propositions by keeping around the functions and arguments that are used to compute the set of possible worlds corresponding to a sentence during its derivation. (A computer scientist could usefully compare this notion of structured meaning to *lazy evaluation*, discussed in relation to computational semantics by van Eijck and Unger, 2010.) The structured meaning is then a semantic derivation structure which is used to calculate synonymy and as the second argument of predicates like *believe*. One problem with this approach is that sentences with radically different structure which nevertheless intuitively express the same proposition may correspond to different structured meanings. One possible example is the active and passive sentences in (52).

- (52) a. Kim sold the book to Sam

- b. Sam was sold the book by Kim

It is hard to think of a way in which a competent native speaker of English could believe one of these but not the other. Such examples depend very much on the way in which you analyze them and how you set up the relation between syntax and semantics. For example, if you believe that compositional semantics is not defined directly on English syntax but on a logical form derived from English syntax and you are careful to relate both sentences to the same logical form, then, of course, both sentences could be related to the same structured meaning. Another kind of example which is possibly more difficult to handle with such machinery is cases of speakers of different languages with radically different structure who nevertheless intuitively share the same belief.

This kind of theory when viewed from the perspective of the theory presented in this book presents a rather odd view of the phenomena. It first proposes a theory of propositions which is obviously too coarse-grained to model the propositional attitudes. It then tries to fix this by using the derivational structure involved in reading these propositions off the syntax of the natural language. When this turns out to be too fine-grained a wholly new representation, logical form, is introduced to fix this new problem. The status of logical form is in our terms mysterious. It is neither based on the utterance situation nor on the situation types used to construct the content associated with the utterance situation. It is an additional language introduced in order to fix problems involved in interpreting utterances directly, a language which mediates between the utterance and the content. If logical form is more amenable to semantic interpretation than natural language one might raise the question why we do not speak in logical forms rather than the way we do. It is hard to imagine what the realistic status of this intermediate language should be either in terms of the utterance situation, the type of situation associated with the content or neurological events associated with perceiving or conceiving either of these.

A second problem for the structured meaning approach is that it tells us nothing about cases where no syntactic structure is involved, for example, proper names which have the same referent like *Hesperus* and *Phosphorus* or synonymous words like *groundhog* and *woodchuck*. This is pointed out by Dowty *et al.* (1981).

The fact that matching is involved in the logic of belief rules out two important ways (relating to labelling and the internal structure of record types) in which record types could be too fine-grained to give an analysis of intuitive propositions. In general it seems preferable to start from objects that are too fine-grained since we can then set about finding ways of collapsing distinctions rather than starting out with something (like sets of possible worlds) which are not fine-grained enough and try to add things to it to make the finer distinctions.

Another advantage of this strategy is that it offers possibilities for varying the fineness of the grain for different cases. Thus while we can understand that (46) can be consistent, it is much harder to think that both of (53a,b) could be true.

- (53) a. Chris believes that Kim bought sex from Sam
 b. Chris does not believe that Sam sold sex to Kim

The best we can do to make sense of (53) as a pair of consistent sentences is that Chris is either irrational in her beliefs or does not have sufficient understanding of the language, or that somehow the equivalence between *buy* and *sell* has been suspended. This seems very different from (46).

In the case of *believe* we have suggested that the type represented by the complement has to be matched against the long-term memory of the believer in order for the sentence to be true. The kind of matching introduced in Chapter 4 involves not only relabelling but also subtyping. Suppose that Chris's long term memory is modelled by the type (54a) and that the content of an utterance of *Sam sold sex to Kim* is the type (54b).

$$(54) \text{ a. } \left[\begin{array}{c} \vdots \\ \text{id}_i : [e : \text{buy}(\text{kim}, \text{sex}, \text{sam})] \\ \vdots \end{array} \right]$$

b. $[e : \text{sell}(\text{sam}, \text{sex}, \text{kim})]$

Is there a match for (54b) in (54a)? The answer is “yes”. The relevant relabelling is (55a) and the result of applying that relabelling to (54b) is (55b).

- (55) a. $e \rightsquigarrow \text{id}_i.e$
 b. $[\text{id}_i : [e : \text{sell}(\text{sam}, \text{sex}, \text{kim})]]$

We can see that (54a) is a subtype of (55b) in virtue of the fact in (56) — any event of buying is also an event of selling.

$$(56) \text{ buy}(\text{kim}, \text{sex}, \text{sam}) \sqsubseteq \text{sell}(\text{sam}, \text{sex}, \text{kim})$$

In this way we can obtain the correct level of granularity for *believe*. Consider now (57a) where we have the verb *say* instead of *believe* and a situation where the actual utterance that Chris made was (57b).

- (57) a. Chris said that Sam sold sex to Kim
 b. Kim bought sex from Sam

Is (57a) true in this case? It seems that we answer this question differently depending on how close the match between the reported speech and the original utterance has to be for the purposes at hand. Ginzburg and Cooper (2014) treat direct quotation in terms of a similarity metric on types which is associated with the context. In different contexts we require different similarity metrics. In some contexts (58) might be considered close enough given that what Chris had said originally was (57b).

(58) Chris said, “Sam sold sex to Kim.”

This might be especially be true if Chris’s original utterance was in a language other than English. Here I would like to say that indirect speech cases like (57) also involve a similarity metric given by the context and that similarity metrics associated with indirect speech in general can be looser than those associated with direct speech where we often look not only at the content of the original utterance but also its exact form of words. So according to some similarity metrics (57a) will be true and for others it will be false. It will be true intuitively if the content of its complement is close enough for current purposes to the content of Chris’s original utterance.

We can assimilate our treatment of belief to this general treatment involving similarity metrics by defining a similarity metric that says that the type representing an agent’s long term memory is similar to the type which is the content of the belief complement if the complement content matches the long term memory type in the way we have described. We will argue below that there is an advantage to making this assimilation since the criteria we use for whether an agent has a certain belief seem to vary according to the purposes we have at hand in the current context.

One of the distinctions that it seems to be possible to make in similarity metrics involves different kinds of subtyping. We have defined subtyping so that for two types, T_1 and T_2 , T_1 is a subtype of T_2 just in case for any a , if $a : T_1$ then $a : T_2$ and that this holds no matter what assignment is made to basic types and ptypes (Appendix A.2). Now consider the two examples of subtyping in (59).

- (59) a. $\left[\begin{array}{cc} \ell_1 & : & T_1 \\ \ell_2 & : & T_2 \end{array} \right] \sqsubseteq \left[\ell_1 \quad : \quad T_1 \right]$
 b. $\text{sell}(a, b, c) \sqsubseteq \text{buy}(c, b, a)$

(59a) holds because of the general characterization of our type theory. It is, if you like, “hard-wired” into the type theoretic system. There is no way that you could construct a type system of the kind TTR characterizes which does not require (59a). (59b), on the other hand, holds only in virtue of a “postulate” that we have added to the general system relating to the particular predicates ‘buy’ and ‘sell’. Just as Montague (1973) introduced what have come to be known as meaning postulates in his system as “restrict[ing] attention to those interpretations of intensional

logic in which the following formulas are true”, a postulate concerning the equivalence of selling and buying events in TTR means that we are restricting attention to possibilities (assignments to basic types and ptypes) in which the equivalence holds. According to the general definitions of TTR (not including such postulates) it is possible to construct a system where the equivalence does not hold. We will refer to (59a) as an instance of *structural subtyping* and (59b) as an instance of *postulated subtyping*. It appears that natural languages can distinguish between these different kinds of subtyping in the kind of matching that is required by predicates which take types as arguments. In the case of $\text{believe}(a, T)$ we say that this is instantiated (non-empty) just in case a 's long term memory is characterized by a type which, modulo relabelling, is a subtype (either structural or postulated) of T . On the other hand, if we think of a set of laws as characterizing, among other things, a set of forbidden types of situations, then $\text{illegal}(T)$ would be instantiated just in case T is, modulo relabelling, a structural subtype of one of the forbidden types.

The distinction between structural and postulated subtyping also gives us a clue on how to deal with groundhogs and woodchucks. Structural subtyping is hardwired into the system. Any cognitive system which implements types will also have structural subtyping, assuming TTR is the right type theory for cognitive systems. Any such system will also have the capability to include postulated subtyping. But exactly which postulates the system has is a matter of learning. Different agents will acquire different postulates depending on their experience. While it is hard to imagine a competent speaker of English not knowing the equivalence between buying and selling it is very easy to suppose that a competent speaker does not know the equivalence between woodchucks and groundhogs. Indeed it would be natural for speakers to assume that the words *woodchuck* and *groundhog* are associated with distinct types and an agent would need some kind of evidence to establish an equivalence between the types. It would be possible for an agent who has not acquired the postulates that establish the equivalence to believe that a woodchuck is in the garden but not to believe that a groundhog is in the garden. However, an agent who has acquired the equivalence would have to believe or disbelieve both. Thus the claim in (60) seems contradictory.

- (60) Kim knows that woodchucks are the same as groundhogs and believes that a woodchuck is in the garden but does not believe that a groundhog is in the garden

The only way we can make sense of Kim believing something about a woodchuck but not about a groundhog is that Kim is unaware that woodchucks and groundhogs are the same animal. Thus getting the semantics of these attitude reports right is not simply a matter of having a finegrained enough semantics to distinguish between *woodchuck* and *groundhog* but also in linking this fine-grainedness to a lack of knowledge about equivalence on the agent's part.

Suppose that Kim believes a woodchuck is in the garden and does not have the postulated equivalence between woodchuck and groundhog. It would seem from what we have said above that it does not follow that Kim believes that a groundhog is in the garden, and indeed there is a sense in

which this is right, if we are taking account of subtyping according to Kim's postulates. Suppose, however, that I do know that woodchucks and groundhogs are the same animal. It seems that I can truthfully report that Kim believes that a groundhog is in the garden, using my knowledge that woodchucks and groundhogs are the same, even though Kim would not herself necessarily assent to a claim: "There's a groundhog in the garden". There is a systematic ambiguity in reports of this kind as to whether the match with Kim's long term memory is computed using the postulates available in Kim's resources or the postulates available in the reporter's resources. Most of the time we do not notice this distinction because it only arises in the case where there is this particular discrepancy between the resources available to the two agents. But it is important to note that in this case there is no one answer to the question *Does Kim believe that a groundhog is in the garden?*. In one sense she does not, and in another sense she does. On the reading where the reporter uses her own postulates it seems that there is a relationship with quotation in translation. Suppose that Kim is a monolingual speaker of German and has a belief which would be reported in German as "Ein Waldmurmeltier ist im Garten". The way in which this belief should be reported in English has to depend entirely on the reporter's resources concerning the correspondences between the contents of *Waldmurmeltier*, *groundhog* and *woodchuck*.

There is a similar systematic ambiguity to that we saw with reporting beliefs about woodchucks and groundhogs in our reporting of ancient beliefs about Hesperus and Phosphorus. Did the ancients believe that Venus rose in the morning? In one sense they did not, since they did not know that the heavenly body which they called Hesperus was in fact Venus. In another sense they did, since the heavenly body which they called Hesperus is in fact (according to the reporter's resources) Venus. The change in long term memory of an ancient who learns that Hesperus and Phosphorus are identical is parallel to that discussed in relation to example (50) and subsequent examples in Chapter 4 except that two proper names are involved rather than one. The type of the ancients' long term memory in their state of ignorance could be a subtype of (61) for some natural numbers i, j, k and l .

$$(61) \left[\begin{array}{l} id_i: \left[\begin{array}{l} x:Ind \\ e:named(x, "Hesperus") \end{array} \right] \\ id_j: [e:rise_in_the_evening(\uparrow id_i.x)] \\ id_k: \left[\begin{array}{l} x:Ind \\ e:named(x, "Phosphorus") \end{array} \right] \\ id_l: [e:rise_in_the_morning(\uparrow id_k.x)] \end{array} \right]$$

Upon the ancients' learning that Hesperus and Phosphorus are the same object (61) would be updated to (62a) which is identical with (62b).

$$\begin{aligned}
 (62) \quad & \text{a.} \quad \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_j: \left[\begin{array}{l} e:\text{rise_in_the_evening}(\uparrow\text{id}_i.x) \end{array} \right] \\ \text{id}_k: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \\ \text{id}_l: \left[\begin{array}{l} e:\text{rise_in_the_morning}(\uparrow\text{id}_k.x) \end{array} \right] \end{array} \right] \hat{\cdot} \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \end{array} \right] \\ \text{id}_k: \left[\begin{array}{l} x=\uparrow\text{id}_i.x:Ind \end{array} \right] \end{array} \right] \\
 & \text{b.} \quad \left[\begin{array}{l} \text{id}_i: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_j: \left[\begin{array}{l} e:\text{rise_in_the_evening}(\uparrow\text{id}_i.x) \end{array} \right] \\ \text{id}_k: \left[\begin{array}{l} x=\uparrow\text{id}_i.x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \\ \text{id}_l: \left[\begin{array}{l} e:\text{rise_in_the_morning}(\uparrow\text{id}_k.x) \end{array} \right] \end{array} \right]
 \end{aligned}$$

Note that (62) could be construed as corresponding to a state of mind where an ancient would still refer to Venus as “Hesperus” in connection with evening rising events and “Phosphorus” in connection with morning rising events even though she was aware that they were the same object. The structure of the memory associates the different name with certain types of events. This seems intuitively correct.

Recall that *SemPropName* is defined in Chapter 4, example (11), also Appendix B.1.4.1, as (63).

$$(63) \quad \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, T) \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, T) \end{array} \right] \cdot \lambda P:Pty . P(r) \end{array} \right]$$

According to the account we gave in Chapter 4 (see example (49)), the type in the ‘bg’-field has to be matched against the gameboard or failing that against the long-term memory or failing that added to the gameboard before the new information can be integrated. The assumption in that discussion was that the relevant long-term memory was that of the agent integrating the utterance. Now we are raising the issue of whose long-term memory is the relevant one to check. There are three long-term memories which can be relevant in a belief report: that of the agent integrating the utterance of the report (that is, the same long-term memory as we were considering in Chapter 4), the long-term memory of the reporter and the long-term memory of the subject of the report (the “believer”). Obviously it is the information state of the agent integrating the report that we are primarily concerned with as it is this integration process which we are trying to explain. This agent does not, of course, have direct access to the long-term memories of either the reporter or the subject of the report. (The integrator’s brain is not wired to either the reporter’s or the subject’s brain.) However, the integrator can form views of the nature of their long-term memories using as evidence, among other things, utterances made by

them or utterances made by others about them. Such information about the long-term memories of the reporter and subject can be incorporated in the integrator's long-term memory. That is, among the beliefs we have encoded in long-term memory we have beliefs concerning what others believe. Consider the type characterizing long-term memory in (64).

$$(64) \quad \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_2: \left[e:\text{rise_in_the_evening}(\uparrow\text{id}_1.x) \right] \\ \text{id}_3: \left[e:\text{rise_in_the_morning}(\uparrow\text{id}_1.x) \right] \\ \text{id}_4: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Homer"}) \end{array} \right] \\ \text{id}_5: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_2: \left[e:\text{rise_in_the_evening}(\uparrow\text{id}_1.x) \right] \\ \text{id}_3: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \\ \text{id}_4: \left[e:\text{rise_in_the_morning}(\uparrow\text{id}_3.x) \right] \end{array} \right] :RecType \\ \text{id}_2: \left[e:\text{believe}(\uparrow^2\text{id}_4.x, \uparrow\text{id}_1) \right] \\ \text{id}_3 = \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x=\uparrow^2\text{id}_1.x:Ind \\ e:\text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_3: \left[\begin{array}{l} x=\uparrow^2\text{id}_1.x:Ind \\ e:\text{named}(x, \text{"Venus"}) \end{array} \right] \end{array} \right] :RecType \\ \text{id}_4: \text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] \end{array} \right]$$

Here the type in the 'id₅.id₃'-field in (64) is a *point of view* on the type in the 'id₅.id₁'-field. A point of view on a type, T , is a type which has labels which overlap with those of T and represents an alternative take on the fields with corresponding labels. In (64), we introduce a predicate 'pov' with arity $\langle RecType, RecType \rangle$ such that $\text{pov}(T_1, T_2)$ will have a witness just in case T_1 is a point of view on T_2 . Here what is represented is that both what Homer calls Hesperus and what Homer calls Phosphorus is what the agent whose long term memory is represented in (64) would call Venus. We can obtain a complete alternative version of the original type by taking the asymmetric merge (see Appendix A.13) of the original type with the point of view. Thus in this case we can obtain (65a) which is identical with (65b).

$$(65) \text{ a. } \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_2: \left[e:\text{rise_in_the_evening}(\uparrow\text{id}_1.x) \right] \\ \text{id}_3: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \\ \text{id}_4: \left[e:\text{rise_in_the_morning}(\uparrow\text{id}_3.x) \right] \end{array} \right] \boxed{\wedge}$$

$$\begin{array}{c}
 \left[\begin{array}{c} \text{id}_1: \left[\begin{array}{c} x = \uparrow^2 \text{id}_1.x : \text{Ind} \\ e : \text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_3: \left[\begin{array}{c} x = \uparrow^2 \text{id}_1.x : \text{Ind} \\ e : \text{named}(x, \text{"Venus"}) \end{array} \right] \end{array} \right] \\
 \text{b. } \left[\begin{array}{c} \text{id}_1: \left[\begin{array}{c} x = \uparrow^2 \text{id}_1.x : \text{Ind} \\ e : \text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_2: \left[e : \text{rise_in_the_evening}(\uparrow \text{id}_1.x) \right] \\ \text{id}_3: \left[\begin{array}{c} x = \uparrow^2 \text{id}_1.x : \text{Ind} \\ e : \text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_4: \left[e : \text{rise_in_the_morning}(\uparrow \text{id}_3.x) \right] \end{array} \right]
 \end{array}$$

In order to account for belief when a point of view is involved we need to revise the witness conditions for ‘believe’ which were given in (49). The revision is given in (66) and involves replacing the original biconditional with a conditional and adding an additional conditional to cover the case for a point of view:

$$\begin{array}{l}
 (66) \quad e : \text{believe}(a, T) \text{ if} \\
 \quad \quad e : \text{ltm}(a, T') \\
 \quad \quad \text{and } T' \sqsubseteq_{\sim} T \\
 \quad e : \text{believe}(a, T) \text{ if} \\
 \quad \quad e : \text{believe}(a, T_1) \\
 \quad \quad e : \text{pov}(T_2, T_1) \\
 \quad \quad \text{and } T_1 \sqsubseteq_{\sim} T_2 \sqsubseteq_{\sim} T
 \end{array}$$

Suppose, contrary to fact, that Homer encountered Pythagoras, who believed that the morning star and the evening star were identical and Homer, while perfectly aware of Pythagoras’ belief, maintained a distinction between the two objects and that Pythagoras used the name “Venus”¹⁰ to refer to both Hesperus and Phosphorus. A type of Homer’s long-term memory could be (67).

¹⁰Or more correctly from the historical point of view: “Aphrodite”.

$$(67) \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_2: \left[e:\text{rise_in_the_evening}(\uparrow\text{id}_1.x) \right] \\ \text{id}_3: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \\ \text{id}_4: \left[e:\text{rise_in_the_morning}(\uparrow\text{id}_3.x) \right] \\ \text{id}_5: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Pythagoras"}) \end{array} \right] \\ \text{id}_6: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Venus"}) \end{array} \right] \\ \text{id}_2: \left[e:\text{rise_in_the_morning}(\uparrow\text{id}_1.x) \right] \\ \text{id}_3: \left[e:\text{rise_in_the_evening}(\uparrow\text{id}_1.x) \right] \end{array} \right] \\ \text{id}_2: \left[e:\text{believe}(\uparrow^2\text{id}_5.x, \uparrow\text{id}_1) \right] \\ \text{id}_3 = \left[\text{id}_1: \left[x = \uparrow^3\text{id}_1.x, \uparrow^3\text{id}_3.x:Ind \right] \right] :RecType \\ \text{id}_4: \text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] :RecType \end{array} \right]$$

Note that here we have generalized the convenient notation for manifest fields. The manifest field in the type (under ‘id₆.id₃’) $[x = \uparrow^3\text{id}_1.x, \uparrow^3\text{id}_3.x:Ind]$ requires that the value in the ‘x’-field is identical with the value of the two values in the fields at the top level on the paths ‘id₁.x’ and ‘id₃.x’. We use the notation $[\ell = a, b, \dots : T]$ to represent $[\ell: T_a \wedge T_b \wedge \dots]$.

A more complex point of view in place of the type given under ‘id₆.id₃’ is (68).

$$(68) \left[\begin{array}{l} \text{id}_1: \left[x = \uparrow\text{id}_4.x, \uparrow\text{id}_5.x:Ind \right] \\ \text{id}_4: \left[\begin{array}{l} x = \uparrow^3\text{id}_1.x:Ind \\ e:\text{named}(x, \text{"Hesperus"}) \end{array} \right] \\ \text{id}_5: \left[\begin{array}{l} x = \uparrow^3\text{id}_3.x:Ind \\ e:\text{named}(x, \text{"Phosphorus"}) \end{array} \right] \end{array} \right]$$

With (68) substituted for the record type in ‘id₆.id₃’ in (67), Homer might now truthfully report (69).

- (69) Pythagoras believes that Hesperus rises in the evening and Phosphorus rises in the morning (though, of course, he believes they are both something called Venus)

In case (69) does not seem convincing let us consider a more modern story (given in (70)) where there actually are two individuals who are mistakenly considered to be one individual.

- (70) Tom and Bill Smith are identical twins who are both employed as teachers at the same school (a source of endless confusion for staff and students alike). Sam is a new girl

spending her first day at the school. Early in the morning Tom, for whom Sam has the name “Mr Smith”, tells her class, “There will be Geometry at 11”. Sam thinks he said ‘There will be Geography at 11’. Later in the morning Bill addresses her class and Sam thinks he is the same “Mr Smith” she saw earlier. Bill says, “There will be French at 11:30.” Sam, who had been too nervous to eat much breakfast and is already feeling quite hungry, thinks he said ‘There will be lunch at 11:30’. Later, in the staff room, Matti, the head teacher explains to some of her colleagues that one of the new girls was in tears in her office complaining that the Geography lesson was about strangely shaped countries which were difficult to understand and there was no lunch when she went to the dining hall. Matti says, “She thought Tom said ‘Geography at 11’ and Bill said ‘Lunch at 11:30’. And to add to the confusion, she thought they were the same person. Poor wee thing, she’s had a difficult day.”

In (70) it seems natural that Matti should use her names for the two teachers rather than Sam’s name “Mr Smith” for both of them when reporting Sam’s beliefs.

In summary, our approach to intensional constructions in natural language has two main components. Firstly, we use (hyper)intensional types rather than sets of possible worlds or situations as the objects of intensional predicates (like ‘believe’). Secondly, we characterize the truth-conditions of these constructions in terms of matching these types against other types (such as types characterizing the long-term memories of the believer, the reporter or the hearer of the report or, in the case of *illegal*, types characterizing a particular canon of law). This opens up possibilities for varying interpretations depending on both which types we match against and what kind of match is required. This makes the interpretation of intensional constructions much more context dependent than is normally assumed and interactive in the sense that we are often comparing (our view of) resources available to different agents.

In the rest of this section we will look at how these ideas can be applied to other intensional constructions that Montague (1973) treated: intensional transitive verbs, verbs taking infinitival complements and intensional adverbs.

Our basic strategy for recasting Montague’s analysis of transitive verbs in terms of TTR is to treat them in terms of a predicate whose arguments are an individual (of type *Ind*) and a quantifier (of type *Quant*). Here, in order to present the basic idea, we will first present a simplified treatment which does not take account of the parametric contents that we introduced in Chapter 5. We will return later to implementing this as an addition to the kind of grammar presented in that chapter. Two ptypes corresponding to *Sandy finds a unicorn* and *Sandy seeks a unicorn* would be as in (71).

- (71) a. $\text{find}(\text{sandy}, \lambda P:P\text{pty} . \text{exist}(\text{unicorn}_*, P))$
 b. $\text{seek}(\text{sandy}, \lambda P:P\text{pty} . \text{exist}(\text{unicorn}_*, P))$

Here we are using ‘unicorn_{*}’ to represent the foreground of non-parametric property of individuals: $\lambda r:[x:Ind] . [e:unicorn(r.x)]$. The predicate ‘find’, corresponding to an extensional transitive verb is related to another predicate ‘find[†]’ whose two arguments are both of type *Ind*. The relationship between the two predicates is expressed in (72).

$$(72) \quad e : \text{find}(a, Q) \text{ iff } e : Q(\lambda r:[x:Ind] . \text{find}^\dagger(a, r.x))$$

This corresponds more or less exactly to Montague’s meaning postulate in Montague (1973) for extensional transitive verbs (meaning postulate 4). In our version, however, the postulate is situation specific. It says that any situation which is of the type ‘find(*a*, *Q*)’ is itself of the type obtained by “quantifying in” *Q* over a type constructed from the predicate ‘find[†]’ as specified in (72) and furthermore, as it is a biconditional any situation of the second type with the quantifier “exported” will also be of the first type with the quantifier as the second argument to ‘find’.

In addition to this meaning postulate for extensional verbs, Montague also had a specific meaning postulate relating *seek* to *try to find* (his meaning postulate 9 in Montague, 1973). We will treat this rather differently in terms of what it means for a search to be successful. The intuitive idea is that a search is successful if you find what you are looking for. Our postulate is presented in (73).

$$(73) \quad \text{if } e : \text{successful}(e', \text{seek}(a, Q)), \text{ then } e' : \text{find}(a, Q)$$

This says that if any situation, *e*, is of a type which predicates success of a situation, *e'*, with respect to the type of situation where *a* seeks *Q* (that is, *e'* is a successful seeking of *Q* by *a*) then *e'* must be of the type where *a* finds *Q*. Note that this is a conditional but not a biconditional. It could be that *a* finds *Q* without looking for it. Finding something does not always represent a successful search.

Doing things this way gives us a rather different perspective on the relationship between extensional and intensional verbs than Montague’s analysis. Montague treats both intensional and extensional verbs as relations between individuals and quantifier intensions. Extensional verbs are extensional in virtue of the fact that they are associated with a meaning postulate which says that there is an equivalence between the relation holding between some individual and the quantifier intension and the quantifier having wide scope over a formula which involves the corresponding relation between individuals. This can seem unintuitive in that the extensional case, which intuitively seems simpler than the intensional case, involves an extra meaning postulate which is not available for the intensional case. Thus the intensional case is taken as the basic case and the extensional case involves an additional inference. On the kind of analysis we are proposing both extensional and intensional verbs involve inferences whereby the quantifier is given wide scope. In the case of an extensional verb the inference is more direct and in some

sense simpler. As exemplified by (72), it involves an equivalence and is an inference concerning the same situation which is of the type with the quantifier in argument position. The predicate involved is intuitively a version of the same predicate which takes two individuals as arguments rather than an individual and a quantifier. The inference involved with intensional verbs (as illustrated by (73)) is, however, more complex. Firstly, it does not involve an inference directly from the intensional verb holding between an individual and a quantifier, but rather concerning what would be a successful outcome of a situation of that type. Secondly, it involves a conditional inference rather than an equivalence, though perhaps it is unclear whether that has anything to do with intuitive complexity. And thirdly, the conclusion does not involve the original intensional predicate but a distinct extensional predicate from which one can draw a conclusion with the quantifier exported. It seems then that in an intuitive sense something more complex, or at least special, is going on in the case of the intensional predicates.

In the PTQ fragment (Montague, 1973) *worship* was treated as an intensional verb. Early on in the literature on formal semantics this was generally regarded as a mistake (Bennett, 1974) ????. It is indeed the case that a sentence like (74) does not entail the existence of a god.

(74) Kim worshipped a god

But on the other hand it seems that there has to be a specific (though possibly non-existent) god which Kim worshipped. This is different to the case with a true intensional verb like *seek*, *look for* or *need* where there is no requirement of specificity. The verb *worship* requires specificity but not existence of the object in this case. There are verbs which require existence but not specificity. An example of this is *book* as used in connection with booking a table at a restaurant. Compare the examples in (75), for example.

- (75) a. # Kim booked a table but there were no tables
 b. Kim was looking for / needed a table but there were no tables

(75a) seems inconsistent. If Kim booked a table then there must have been at least one table. (75b) on the other hand seems fine. Note, however, that you can book a table without booking a specific table. It may be that when you get to the restaurant there are several available tables and you get to choose which one you want to sit at. Booking a table can just mean that there will be a table available for you at the agreed time at the restaurant, not that any specific table has been reserved for you, although that, of course, is possible too and you may have specified which table you want (the one in the corner by the window). The verb *book* has thus the hallmarks of an intensional verb when it comes to specificity but nevertheless requires existence.

Let us deal first with *worship* and religious beliefs. Suppose that somebody says (76).

(76) Kim worships Zeus

This does not commit the speaker to the existence of Zeus, but it does commit the speaker to a system of religious belief in which there is a god Zeus and also commits her to the claim that Kim subscribes to such a belief system. Thus (77a) seems perfectly consistent whereas (77b) seems either inconsistent or at best to force a rather special meaning for *worship*.

- (77) a. Kim worships Zeus, but I don't believe in Zeus
 b. # Kim worships Zeus, but she doesn't believe in Zeus

If Kim worships Zeus, then she not only has to believe in Zeus, but she also has to believe that she worships Zeus. (78a) sounds strange or at least forces us into an unusual meaning for *worship*. In contrast (78b) with a standard extensional verb seems perfectly consistent.

- (78) a. # Kim worships Zeus, but she doesn't believe that she worships Zeus
 b. Kim found Harry, but she doesn't believe that she found Harry (because he was wearing a disguise)

This is another aspect of the content of *worship* which seems to suggest intensionality (or intentionality, with a 't'): it describes a conscious aspect of somebody's mental state. The verb *find*, on the other hand, can describe an external fact about an agent of which the agent itself is not aware.

We suppose that in our characterization of an agent's mental state by a type we can isolate a type that characterizes the agent's religious beliefs. We might regard this as a part of long term memory or as a separate component at the same level as long term memory in the type characterizing the agent's total information state, visual field and so on. The type corresponding to religious beliefs represents the way the world would be if the agent's religious beliefs were true. In (79) we give a type that could correspond to Kim's religious beliefs.

$$(79) \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{god}(\uparrow \text{id}_1.x) \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\text{kim}, \uparrow \text{id}_1.x) \end{array} \right] \end{array} \right]$$

Notice that the worship predicate we use is 'worship[†]' representing a relation between individuals. In this way *worship* is like the extensional verb *find* in that it is related by postulate to a †-variant of the same predicate, getting us the specificity. However, in the case of *worship*

the type constructed with the \uparrow -predicate is embedded within another type representing religious belief which gives the verb an intensional quality.¹¹

(76) commits the speaker not to the existence of Zeus but the type characterizing Kim's religious beliefs containing a match (in the sense we have discussed above) for the type (80).

$$(80) \left[\begin{array}{lcl} x & : & Ind \\ e & : & \text{named}(x, \text{"Zeus"}) \end{array} \right]$$

In the sense that the object of *worship* has to be matched against a mental state just like the complement of *believe* rather than checked against the world *worship* is behaving like an intensional verb. It is also like an intensional verb in that we do not have a postulate like the one we have for *find* in (72). Where it is different from the classical cases of intensional verbs is that we do not have postulates like the one for *seek* in (73) but rather require a specific match for *Zeus* in the religious beliefs of the subject. (81) is a type corresponding to the speaker's long term memory which would fulfil the commitments of (76).

$$(81) \left[\begin{array}{lcl} id_1 : & \left[\begin{array}{lcl} x & : & Ind \\ e & : & \text{named}(x, \text{"Kim"}) \end{array} \right] \\ id_2 = & \left[\begin{array}{lcl} id_1 & : & \left[\begin{array}{lcl} x & : & Ind \\ e & : & \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ id_2 & : & \left[\begin{array}{lcl} e & : & \text{god}(\uparrow id_1.x) \end{array} \right] \\ id_3 & : & \left[\begin{array}{lcl} e & : & \text{worship}^\dagger(\uparrow^2 id_1.x, \uparrow id_1.x) \end{array} \right] \end{array} \right] \\ id_3 : & \left[\begin{array}{lcl} e & : & \text{rbelieve}(\uparrow id_1.x, \uparrow id_2) \end{array} \right] \end{array} \right] : RecType$$

In the ' id_3 '-field here we are using the predicate 'rbelieve'. Informally, $\text{rbelieve}(a, T)$ is non-empty just in case a has T as a religious belief, that is the type identified as representing a 's religious beliefs in the type corresponding to her total information state is a subtype of T .

Armed with this view of the characterization of religious belief in information states we formulate a postulate for 'worship' in (82) which shows the intensionality and specificity requirements.

$$(82) \quad e : \text{worship}(a, Q) \text{ iff for some } T$$

1. $e : \text{rbelieve}(a, T)$
2. $T \sqsubseteq_{\sim} Q(\lambda r : [x : Ind] . \text{worship}^\dagger(a, r.x))$

¹¹Notice also that we are using 'kim' for the owner of the belief state, ignoring here *de se* issues.

Clause (1) of (82) represents the intensionality (or perhaps more appropriately *intentionality*) requirement in that it requires us to use a type which characterizes the religious belief of the first argument of ‘worship’. Clause (2) represents the specificity requirement in that it requires the religious belief type to be a subtype of a type (possibly relabelled) obtained by “exporting” the quantifier Q .

Things are, however, a little more complicated than this. Jupiter is the Roman name for the god which in Greek is called Zeus. Suppose that Kim is Greek oriented and does not know the name Jupiter. Suppose further that the speaker is Roman oriented and reports (83).

(83) Kim worships Jupiter

It seems that the speaker of (83) can be said to have spoken the truth even though Kim does not know the name Jupiter and the speaker does not believe that Jupiter exists or have any kind of religious belief in Jupiter. It is true because we know that Jupiter and Zeus are the same god in the Roman-Greek pantheon. How can this be when the speaker is not committed to the existence of Zeus/Jupiter?

First consider that the speaker might be committed to a type characterizing part of the Roman pantheon, for example, (84).

$$(84) \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{“Jupiter”}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right]$$

This can be incorporated into (81) as in (85).

$$(85) \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{“Kim”}) \end{array} \right] \\ \text{id}_2 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{“Zeus”}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{god}(\uparrow \text{id}_1.x) \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\uparrow^2 \text{id}_1.x, \uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{rbelieve}(\uparrow \text{id}_1.x, \uparrow \text{id}_2) \end{array} \right] \\ \text{id}_4 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{“Jupiter”}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \\ \text{id}_5 : \left[\begin{array}{l} e : \text{roman_pantheon}(\uparrow \text{id}_4) \end{array} \right] \end{array} \right] \begin{array}{l} : \text{RecType} \\ : \text{RecType} \end{array}$$

In (85) no connection is expressed between the types in the ‘id₂’-field and the ‘id₄’-field. Note that they are, however, aligned in their labelling. The individual named Zeus and the individual named Jupiter are required to have the same labelling in witnesses for the two types and similarly for the situations that show they have those names respectively as well as the situations that show the individual is a god and chief god respectively. We shall take this alignment of labelling of the two types to be a significant aspect of the type of information state that (85) represents. We shall say that the type in the ‘id₄’-field is a point of view on the type in the ‘id₂’-field. It represents the agent’s alternative perspective on certain aspects of what she believes to be Kim’s religious beliefs. We introduce an additional field to (85) in order to represent this connection between the two types and we reorganize the point of view and Kim’s religious beliefs into a single record type since it will be important to be able to refer to a single situation providing this information in computing the semantics of *worship*. This is given in (86).

$$(86) \quad \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Kim"}) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{god}(\uparrow \text{id}_1.x) \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\uparrow^3 \text{id}_1.x, \uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] : \text{RecType} \\ \text{id}_3 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Jupiter"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] : \text{RecType} \\ \text{id}_4: \left[\begin{array}{l} e : \text{roman_pantheon}(\uparrow \text{id}_3) \end{array} \right] \\ \text{id}_5: \left[\begin{array}{l} e : \text{pov}(\uparrow \text{id}_3, \uparrow \text{id}_1) \end{array} \right] \end{array} \right]$$

The idea of a point of view is that it represents an alternative take on certain aspects of another type. As before we can obtain a complete alternative version of the original type by taking the asymmetric merge of the original type with the point of view. Thus in the case of the relevant types in (86) we can obtain (87a) which is identical with (87b).

$$(87) \text{ a. } \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{god}(\uparrow \text{id}_1.x) \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\text{kim}, \uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \boxed{\wedge} \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Jupiter"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right]$$

$$b. \left[\begin{array}{l} id_1 : \left[\begin{array}{l} x : Ind \\ e : \text{named}(x, \text{"Jupiter"}) \end{array} \right] \\ id_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow id_1.x) \end{array} \right] \\ id_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\text{kim}, \uparrow id_1.x) \end{array} \right] \end{array} \right]$$

We will call (87) a *complete point of view*. We will now allow two alternative witness conditions for ‘worship’. The first, as before, checks that the type corresponding to the religious beliefs of the subject is a subtype of the type resulting from exporting the quantifier. The second checks that a complete point of view fulfils this condition. The two witness conditions are given in (88).

- (88) a. $e : \text{worship}(a, Q)$ if for some T
1. $e : \text{rbelieve}(a, T)$
 2. $T \sqsubseteq_{\rightsquigarrow} Q(\lambda r : [x : Ind] . \text{worship}^\dagger(a, r.x))$
- b. $e : \text{worship}(a, Q)$ if for some T_1, T_2
1. $e : \text{rbelieve}(a, T_1)$
 2. $e : \text{pov}(T_2, T_1)$
 3. $T_1 \sqcap T_2 \sqsubseteq_{\rightsquigarrow} Q(\lambda r : [x : Ind] . \text{worship}^\dagger(a, r.x))$

We will allow the background conditions of proper names, for example, the type in (89) to match a complete resource and indeed (89) matches (87).

$$(89) \left[\begin{array}{l} x : Ind \\ e : \text{named}(x, \text{"Jupiter"}) \end{array} \right]$$

[How to get the matching apparatus from Ch. 4 to achieve a match for the proper name in this case.]

Now consider the case with the indefinite article. Suppose we have a situation of the type (81), where Kim worships Zeus, who is, according to Kim’s religious beliefs, a god. In order to show that Kim worships a god we may show, following (88a), that this type is a subtype of (90)¹².

$$(90) \text{ exist}(\text{god}_*, \lambda r : [x : Ind] . \text{worship}^\dagger(\text{kim}, r.x))$$

In virtue of the witness conditions associated with ‘exist’ introduced in Chapter 3, example (61), the subtype relation holds between (81) and (90). That is, any situation of type (81) will be of type (90) since in such a situation there will be a individual who is a god whom Kim worships.

¹²Where ‘god_{*}’ abbreviates $\lambda r : [x : Ind] . [e : \text{god}(r.x)]$.

Now let us consider a case which shows that we can use the witness condition (88b). In (91) it is presumably not the case that it is part of Kim's religious beliefs that the god she worships is a false god.

(91) Kim worships a false god

Rather the phrase *false god* can be used to represent a point of view on the part of the speaker. In (92) we add such a point of view to (86).

$$(92) \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Kim"}) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{god}(\uparrow \text{id}_1.x) \end{array} \right] \\ \text{id}_3 : \left[\begin{array}{l} e : \text{worship}^\dagger(\uparrow^3 \text{id}_1.x, \uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} e : \text{rbelieve}(\uparrow^2 \text{id}_1.x, \uparrow \text{id}_1) \end{array} \right] \\ \text{id}_3 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Jupiter"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{chief_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \\ \text{id}_4: \left[\begin{array}{l} e : \text{roman_pantheon}(\uparrow \text{id}_3) \end{array} \right] \\ \text{id}_5: \left[\begin{array}{l} e : \text{pov}(\uparrow \text{id}_3, \uparrow \text{id}_1) \end{array} \right] \\ \text{id}_6 = \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{named}(x, \text{"Zeus"}) \end{array} \right] \\ \text{id}_2 : \left[\begin{array}{l} e : \text{false_god}(\uparrow \text{id}_1.x) \end{array} \right] \end{array} \right] \\ \text{id}_7: \left[\begin{array}{l} e : \text{pov}(\uparrow \text{id}_6, \uparrow \text{id}_1) \end{array} \right] \end{array} \right] : \text{RecType} \end{array} \right]$$

One thing to note about (92) is that an agent can have more than one distinct point of view on the same type, here shown by the 'id₂.id₃'-field and 'id₂.id₆'-field which both introduce points of view on the type in the 'id₂'-field. In searching for a match for the content of (91) we may use the asymmetric merge of the 'id₂.id₁' type with either the 'id₂.id₃' type or the 'id₂.id₆' type. That is, there will be a relabelling of the content of (91), such that any situation of type (92) will also be a situation of the type which is the relabelled content. To make this concrete consider a putative (non-parametric) content for (91) represented in (93a) and the relabelling given in (93b), yielding the relabelled type in (93c).

$$(93) \text{ a. } \left[\begin{array}{l} x : \text{Ind} \\ c : \text{named}(x, \text{"Kim"}) \\ e : \text{worship}(x, \lambda P:P\text{pty} . [\text{e:exist}(\text{false_god}', P)]) \end{array} \right]$$

b. $x \rightsquigarrow \text{id}_1.x$

$$\begin{array}{l}
c \rightsquigarrow id_1.e \\
e \rightsquigarrow id_2 \\
c. \left[\begin{array}{l} id_1 : \left[\begin{array}{l} x : Ind \\ e : \text{named}(x, \text{"Kim"}) \end{array} \right] \\ id_2 : \text{worship}(x, \lambda P:Ppty . [e:\text{exist}(\text{false_god}', P)]) \end{array} \right]
\end{array}$$

We can see that any situation of type (92) would also be of type (93c) given the witness conditions associated with ‘worship’ and ‘exist’ that we have discussed.

Let us now consider sentences such as *Kim booked a table* where there is no inference to a specific table but nevertheless an inference to the existence of a table. This can be achieved by introducing the postulate in (94).

(94) If $a:Ind$ and Q is a monotone increasing quantifier, then

$\text{book}(a, Q)$

is non-empty implies

$Q(\lambda r: [x:Ind] . [e:\text{be}(r.x)])$

is non-empty

Intuively, (94) requires that if there is some situation in which a books a table, then there is some table which is a constituent of some situation (following the witness conditions for ‘be’ presented in Chapter 3). Notice that we do not get such an inference if the quantifier is monotone decreasing. Thus *Kim booked no table* (to the extent that this is an acceptable sentence of English) does not imply that there are no tables and neither does it imply that there are tables. We will discuss monotonicity in Chapter 7.

We now turn our attention to a phenomenon first discussed in the semantics literature by Fodor (1970), p. 226ff. She points out that (95) can mean that Charley does not want a specific coat even if Charley would not describe what he wants as a “coat like Bill’s”.

(95) Charley wants to buy a coat like Bill’s

The sentence could be true on what Montague would call a *de dicto* reading even though Charley does not know Bill, or a least does not know what kind of coat he has. This kind of example, seems straightforwardly treatable with the notion of point of view that we have developed. Consider the type (96), representing the long term memory of some agent.

$$(96) \quad \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Charley"}) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{coat}(x) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} e:\text{trenchcoat}(\uparrow\text{id}_1.x) \end{array} \right] \\ \text{id}_3: \left[\begin{array}{l} e:\text{has_big_pockets}(\uparrow\text{id}_1.x) \end{array} \right] \\ \text{id}_4: \left[\begin{array}{l} e:\text{buy}^\dagger(\uparrow^3\text{id}_1.x, \uparrow\text{id}_1.x) \end{array} \right] \end{array} \right] :RecType \\ \text{id}_2:\text{want}^\dagger(\uparrow^2\text{id}_1.x, \text{id}_1) \\ \text{id}_3 = \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{coat}(x) \end{array} \right] \\ \text{id}_5: \left[\begin{array}{l} e:\text{coat_like_Bill's}(\uparrow\text{id}_1.x) \end{array} \right] \end{array} \right] :RecType \\ \text{id}_4:\text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] \end{array} \right]$$

Here we use the predicate ‘buy[†]’ which is the buy-relation between individuals and the predicate ‘want[†]’ which has arity $\langle Ind, RecType \rangle$. It is related to two other want-predicates: ‘want_P’ with arity $\langle Ind, Ppty \rangle$ and ‘want_Q’ with arity $\langle Ind, Quant \rangle$ used to treat examples like, respectively, *want to buy a coat* and *want a coat*. These two predicates are related to ‘want[†]’ as shown in (97).

(97) a. If $a : Ind$ and $P : Ppty$, then

$$e : \text{want}_P(a, P) \text{ iff } e : \text{want}^\dagger(a, P(\uparrow x=a))$$

b. If $a : Ind$ and $Q : Quant$, then

$$e : \text{want}_Q(a, Q) \text{ iff } e : \text{want}^\dagger(a, Q(\lambda r: [x:Ind] . \text{have}(a, r.x)))$$

In order to characterize witness conditions for ‘want[†]’ we need to assume that an agent’s total information state include a record type representing the agent’s desires, parallel to long term memory and religious beliefs. Intuively the type for desires is a type which represents the way the world would be if the agent’s desires were fulfilled. Thus a total information state would belong to a subtype of (98).

$$(98) \quad \left[\begin{array}{ll} \text{ltm} & : RecType \\ \text{rbel} & : RecType \\ \text{des} & : RecType \end{array} \right]$$

We will introduce predicates ‘desire’ and ‘info_state’ which hold between an individual and a record type (that is, with arity $\langle Ind, RecType \rangle$). The witness conditions for ‘desire’ are given in (99).

- (99) $e : \text{desire}(a, T)$ iff
 $e : \text{info_state}(a, T')$
 and $r : T'$ implies $r.\text{des} = T$

If $e : \text{info_state}(a, T)$ then a 's information state is of type T .

Now we can characterize witness conditions for 'want[†]' as given in (100).

- (100) $e : \text{want}^\dagger(a, T)$ if
 $e : \text{desire}(a, T')$
 and $T' \sqsubseteq_{\sim} T$
 $e : \text{want}^\dagger(a, T)$ if
 $e : \text{want}^\dagger(a, T_1)$
 $e : \text{pov}(T_2, T_1)$
 and $T_1 \sqcap T_2 \sqsubseteq_{\sim} T$

Our analysis of Fodor's example is similar to Schwager (2009) in that the analysis using a point of view and asymmetric merge involves replacing part of the original content of the attitude with something from the perspective of the speaker. Schwager¹³ introduces what she calls the *replacement principle* quoted in (101).

- (101) For the sake of reporting an attitude, a property that is involved in the content of the attitude that is to be reported (the **reported property**) can be replaced by a different property (the **reported property**) as long as the reported property is a subset of the reporting property at all relevant worlds.

(Schwager, 2009, p. 409)

This addresses the important question of what replacements we can make. As Schwager points out we cannot make arbitrary replacements and use them to report an attitude. For example, we cannot report Charley's desire to buy a trenchcoat with big pockets by (102).

- (102) Charley wants to buy a unicorn

For us, this question concerns what must hold if the 'pov'-relation holds between two types. We could imitate something like Schwager's replace principle by requiring (103).

¹³currently named Magdalena Kaufmann

(103) If $\text{pov}(T_1, T_2)$ is witnessed then $T_2 \sqsubseteq T_2 \sqcap T_1$

This says that if the world is of the original attitude type then it is also of the type resulting from asymmetrically merging the attitude type with the point of view. This successfully rules out replacing a trenchcoat with big pockets with a unicorn. If the world is such that Charley has a trenchcoat with big pockets it does not follow that he has a unicorn.

Unfortunately, this constraint does not seem to hold for all of the examples for which we have suggested using points of view. For example, if the original attitude type involves two heavenly bodies, Hesperus and Phosphorus, which rise respectively in the evening and the morning and the point of view requires Hesperus and Phosphorus to correspond to one heavenly body, Venus which rises in the morning and the evening, then it is not the case that any situation which contains the two heavenly bodies would be one in which there is only one heavenly body (and not *vice versa* either). There is no subtyping relation here, just a disagreement about how many heavenly bodies are involved and what they are called.

Similarly, consider the example where, according to the original attitude, Kim worships a god called Zeus and according to the point of view Kim worships a false god called Zeus. It is not obvious that a situation in which the first holds is also a situation in which the second holds or *vice versa*, although one might argue that this depends on whether false gods are gods or not. I would tend to think that sometimes we make the inference from *false god* to *god* and sometimes we do not and that this is part of the general nature of semantic flux in language. However, an inference from *god* to *false god* seems unlikely. Again the point of view seems to reflect a disagreement about the status of Zeus rather than a subtyping relation.

A potential conflict in judgement between the reporter and the attitude bearer also arises in Fodor examples like those in (104).

- (104) a. Charley wants to buy something nonexistent
 b. Charley wants to buy an uncool coat

If Charley's attitude involves a trenchcoat with big pockets, it may in fact be the case that there are no such trench coats but that does not mean that the type *Trenchcoat with big pockets* is a subtype of *Nonexistent*. Recall that for T_1 to be a subtype of T_2 it has to be the case that no matter what we assign to basic types and ptypes (that is, no matter which possibility we consider), something of type T_1 would also be of type T_2 . But there are presumably possibilities in which there are trenchcoats with big pockets. What (104a) seems to commit the speaker to is that there are no trenchcoats with big pockets in the actual possibility we are considering, not that trenchcoats with big pockets are impossible. In actual fact the speaker seems to be committed to a falsehood here because trenchcoats in general have big pockets. We cannot say the same about (104b). However, cool we may think trenchcoats are the speaker is entitled to her opinion. The word

uncool is a predicate of personal taste and the concept of *faultless disagreement* (Kölbel, 2004) becomes relevant. For a suggestion of how predicates of personal taste could be treated using TTR and some references to other literature see Cooper (2015). What this points to is that the speaker is replacing a judgement by the attitude bearer with a judgement of her own. This could be expressed by placing the constraint on judgements in (105).

(105) For any agents A and B , types T_1 and T_2 and object (situation) s ,
 if $:_A \text{pov}(T_1, T_2)$ and $s :_B T_2$, then $s :_A T_2 \boxed{\wedge} T_1$

(105) says that if A judges $\text{pov}(T_1, T_2)$ to be witnessed (that is, A judges that T_1 is a point of view of T_2) and B judges s to be of type T_2 , then A judges s to be of the type resulting from asymmetrically merging T_2 with T_1 . Note that this is a constraint concerning the judgements that two agents make rather than the way things actually are in the world. It belongs to the realm of our theory of action based on type theory rather than to the type theory itself. This means that in principle there is nothing preventing a speaker reporting Charley's desire to buy a trenchcoat with large pockets as a desire to buy a unicorn, provided that the speaker is willing to commit to a claim that she would judge a trenchcoat with large pockets to be a unicorn, something that we would find unexpected given the normal meanings associated with these words.

The proposal here also has aspects in common with the proposal by Pross (ms) which presents a semantics in terms of DRT which takes account of how to represent the attitudes of an agent and analyzes the attitude report in terms of this. Our approach to representing the attitudes in TTR has a good deal in common with the DRT approach as set forth in Kamp (1990); Kamp *et al.* (2011).

Where our proposal differs from previous proposals in the literature is that we are not concerned with trying to identify objects in possible worlds in order to get the Fodorean reading. Rather we are concerned with how different agents might judge situations of certain types. Whether there are situations of the relevant types is not a question which is of relevance to the analysis. Let us see how this relates to examples which have been discussed in the literature. In order to do this we will first say in a little more detail what the content of an utterance of *Charlie wants to buy a coat like Bill's* will be. Let us first look at the (non-parametric) content of *Charlie bought a coat like Bill's* (ignoring issues of tense), given in (106).

(106) $[e : \text{buy}(\text{charlie}, \lambda P:P\text{pty} . [e : \text{exist}(\text{coat_like_Bill's}', P)])]$

Here we use 'coat_like_Bill's'' to represent the property in (107), which we do not analyze further for present purposes.

(107) $\lambda x:[x:Ind] . [e : \text{coat_like_Bill's}(x)]$

Since ‘buy’ is an extensional predicate it will obey the constraint in (108) which relates ‘buy’ to the corresponding predicate with two individual arguments, ‘buy[†]’.

$$(108) \ e : \text{buy}(a, Q) \text{ iff } e : Q(\lambda r : [x : \text{Ind}] . [e : \text{buy}^\dagger(a, r.x)])$$

This means that (106) is equivalent to (109), in the sense that any record of the former type will be of the latter type and *vice versa*.

$$(109) \ [e : \text{exist}(\text{coat_like_Bill's'}, \lambda r : [x : \text{Ind}] . [e : \text{buy}^\dagger(\text{charlie}, r.x)])]$$

Suppose e is of the type in (110).

$$(110) \ [e : \begin{bmatrix} x : \text{Ind} \\ c : \text{coat_like_Bill's}(x) \\ e : \text{buy}^\dagger(\text{charlie}, x) \end{bmatrix}]$$

Then, by the witness condition associated with ‘exist’ introduced in Chapter 3, example (61), e will also be of the type (109). Here we may assume that if the speaker who is asserting this content is speaking truthfully then she judges the coat the Charlie bought to be a coat like Bill’s (by Grice’s maxim of quality). One may, of course, disagree since the degree of similarity between two objects may be a matter of personal taste.

Now let us consider the non-parametric content of *Charlie wants to buy a coat like Bill’s* given in (111a) which, by (97a), is equivalent to (111b).

$$(111) \ \text{a. } [e : \text{want}_P(\text{charlie}, \lambda r : [x : \text{Ind}] . [e : \text{buy}(r.x, \lambda P : \text{Ppty} . [e : \text{exist}(\text{coat_like_Bill's'}, P)])])] \\ \text{b. } [e : \text{want}^\dagger(\text{charlie}, [e : \text{buy}(\text{charlie}, \lambda P : \text{Ppty} . [e : \text{exist}(\text{coat_like_Bill's'}, P)])])]$$

Note that ‘want[†]’, in virtue of (100), introduces the possibility of a point of view.

We will now examine whether the kind of content expressed in (111a) can be matched to a number of scenarios for Fodorean readings which have been discussed in the literature. Here we follow the recent survey of the literature presented by Pross (ms) in Section 1.2 of his paper. Consider the scenario in (112).

- (112) Suppose a store sells some jackets that all look like Malte’s and that Adrian does not know anything about Malte. Assume further that Adrian wants one of those jackets and any of them is an option.

(Romoli and Sudo (2009))

In (113) we exhibit a type which corresponds to how the speaker might represent this scenario in memory.

$$(113) \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Adrian"}) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1 = \text{jacket_in_the_shop}':Ppty \\ \text{id}_2: [e:\text{buy}(\uparrow^3 \text{id}_1.x, \lambda P:Ppty . [e:\text{exist}(\uparrow \text{id}_1, P)])] \end{array} \right] :RecType \\ \text{id}_2: \text{id}_2:\text{want}^\dagger(\uparrow^2 \text{id}_1.x, \text{id}_1) \\ \text{id}_3 = [\text{id}_1 = \text{jacket_like_Malte's}':Ppty] :RecType \\ \text{id}_4: \text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] \end{array} \right]$$

(113) requires that Adrian's desire is to buy something with the property of being a jacket in the shop. The alternative point of view is that the property of being a jacket in the shop can be replaced with the property of being a jacket like Malte's. If the world matches this type then the sentence *Adrian wants to buy a jacket like Malte's* is true.

The next scenario Pross considers is (114).

- (114) Suppose a store offers some jackets that all look like Malte's and that Adrian does not know anything about Malte. Assume that some of the jackets are on sale while others are not and that Adrian is aware of this. Assume further that Adrian wants one of the jackets on sale and any of them is an option.

The point of this is that it emphasizes that the property of jackets involved in Adrian's desire need not be coextensive with the property in the point of view. That is, it is still the case that all the jackets in the shop are like Malte's but Adrian has his sights set on a subset of them (those on sale) although he has not chosen any particular jacket among those. The type we exhibit for this scenario (in (115)) is almost exactly the same as the previous one.

$$(115) \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{"Adrian"}) \end{array} \right] \\ \text{id}_2: \left[\begin{array}{l} \text{id}_1 = \left[\begin{array}{l} \text{id}_1 = \text{jacket_on_sale_in_the_shop}':Ppty \\ \text{id}_2: [e:\text{buy}(\uparrow^3 \text{id}_1.x, \lambda P:Ppty . [e:\text{exist}(\uparrow \text{id}_1, P)])] \end{array} \right] :RecType \\ \text{id}_2: \text{id}_2:\text{want}^\dagger(\uparrow^2 \text{id}_1.x, \text{id}_1) \\ \text{id}_3 = [\text{id}_1 = \text{jacket_like_Malte's}':Ppty] :RecType \\ \text{id}_4: \text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] \end{array} \right]$$

The only difference between this and the previous type is that we have replaced the property ‘jacket_in_the_shop’ with ‘jacket_on_sale_in_the_shop’. Note that the constraint on ‘pov’ that we introduced in (105) did not require coextension in any way, only that if an agent were to judge a situation as of type T_1 then the agent with the point of view would judge the situation to be of the point of view type. This allows for the possibility that there could be additional situations of the point of view type which the first agent would not judge to be of type T_1 . This view of things was in fact already important for the first scenario, since, while Adrian is focussed on the jackets in the shop, the speaker presumably could consider that there are other jackets in addition to those in the shop which are like Malte’s. That is, the sentence would not be falsified by discovering a jacket like Malte’s which is not in the shop.

Pross’s third scenario is (116) which he offers as problematic for the proposal in von Fintel and Heim (2011) because there are no actual jackets like Malte’s as would be required by their analysis.

- (116) Suppose Adrian has seen a picture of a certain green Burberry jacket in a catalogue and wants to buy one. Unbeknownst to Adrian, Malte happens to own exactly such a green Burberry jacket. Unbeknownst to Adrian, the type of jacket in the picture which Adrian has seen is sold out and no further jackets of this type have been produced yet: there are no actual jackets like Malte’s.

This scenario could correspond to the type in (117).

$$(117) \left[\begin{array}{l} \text{id}_1: \left[\begin{array}{l} x:Ind \\ e:\text{named}(x, \text{“Adrian”}) \end{array} \right] \\ \text{id}_1 = \left[\begin{array}{l} \text{id}_1 = \text{jacket_like_the_one_in_the_catalogue’:}Ppty \\ \text{id}_2: [e:\text{buy}(\uparrow^3 \text{id}_1.x, \lambda P:Ppty . [e:\text{exist}(\uparrow \text{id}_1, P)])] \end{array} \right]:RecType \\ \text{id}_2: \left[\begin{array}{l} \text{id}_2:\text{want}^\dagger(\uparrow^2 \text{id}_1.x, \text{id}_1) \\ \text{id}_3 = [\text{id}_1 = \text{jacket_like_Malte’s’:}Ppty]:RecType \\ \text{id}_4:\text{pov}(\text{id}_3, \text{id}_1) \end{array} \right] \end{array} \right]$$

The only difference between this and the previous type is that we have replaced the property ‘jacket_on_sale_in_the_shop’ with the property ‘jacket_like_the_one_in_the_catalogue’. Note that the constraint on ‘pov’ that we introduced in (105) does not require that anything have either the property ‘jacket_like_the_one_in_the_catalogue’ or ‘jacket_like_Malte’s’, only that if an agent were to judge a situation as of the type involving the first property then the agent with the point of view would judge the situation to be of the type involving the second property. Whether there actually are such jackets is an independent question.

Pross (ms) introduces a further scenario for the Fodorian reading in Section 3.5 of his paper which we reproduce in (118).

- (118) Adrian has seen a jacket which has three stripes on its sleeves and wants to buy such a jacket. However, he has read that Adidas uses child labour in the production of its jackets, so the additional condition for his purchase is that the jacket is not from Adidas. If Adrian does not know that Adidas is the brand with the three stripes, he has a desire that he would paraphrase as “I want to buy a jacket from the brand with the three stripes but not from Adidas.” Fritz hears Adrian’s utterance and as he has seen Malte’s jacket which has three stripes and as he also knows about the problem with child labour and Adidas he believes that Malte would never buy a jacket which is made using child labour. Fritz also doesn’t know that Adidas is the brand with the three stripes. He reports Adrian’s desire as “Adrian wants to buy a jacket like Malte’s”.

This mixes in the problem of contradictory beliefs with that of Fodorean readings. The type corresponding to Fritz’s information state could be represented by the type in (119).

$$(119) \left[\begin{array}{l} id_1: \left[\begin{array}{l} x:Ind \\ e:named(x, \text{“Adrian”}) \end{array} \right] \\ id_2: \left[\begin{array}{l} id_1 = \left[\begin{array}{l} id_1=jacket_with_three_stripes_on_its_sleeves':Ppty \\ id_2=not_Adidas':Ppty \end{array} \right] \\ id_3: [e:buy(\uparrow^3 id_1.x, \lambda P:Ppty . [e:exist(\uparrow id_1 \wedge \uparrow id_2, P)])] \\ id_2:want^\dagger(\uparrow^2 id_1.x, id_1) \\ id_3=[id_1=jacket_like_Malte's':Ppty]:RecType \\ id_4:pov(id_3, id_1) \end{array} \right] :RecType \end{array} \right]$$

Adrian’s desire is perfectly rational given that he does not know that a jacket with three stripes on its sleeves is made by Adidas. Note that we can also truthfully report his desire even when we know about Adidas and the three stripes, as in (120).

- (120) Adrian wants to buy a jacket like Malte’s but not from Adidas. He doesn’t realize that having three stripes on the sleeve means that the jacket is from Adidas.

It seems that none of these successive complications of the scenario, increasing, so to speak, the degree of intensionality involved, provide a problem when you combine an theory of intensional types with the notion of point of view as we have described it.

6.4 Compositional semantics

6.5 Conclusion

In this chapter we first pointed out some conceptual and technical problems involving possible worlds as they are standardly used in semantics. We have considered the two main areas where possible worlds have been used: modality and intensionality involving the attitudes. We have suggested that both benefit from an analysis in terms of intensional types instead of possible worlds.

@ @

[How does pov interact with monotonicity? ok *Charley doesn't want to buy a coat like mine*]

[What to do with predicative cases? *Charley thinks it is a coat like mine, Charley thinks he bought a coat like mine, Charley thinks it's like mine, Charley doesn't think it's like mine, Charley denies it's like mine*]

[Liz' references on buy and sell]

[matching a parametric content for a proper name. How to get the case where the name is in the reporter's resources?]

[Jonathan's comments on ch 1 - reply?, believe vs. legal, matching vs relabelling, believe vs know (Austinian propositions)]

@ @

Chapter 7

Quantification, anaphora and underspecification

7.1 The interpretation of unbound pronouns

We will consider how to recreate a simple interpretation of pronouns ranging over individuals, first treating them in a similar way to free variables in logic and then showing how they can be bound by quantifiers. The central idea is to use records as *pronominal contexts* which correspond to partial assignments to variables in logical treatments. Consider first a simple sentence with a deictic pronoun as in (1).

(1) he left

In our initial pass we will ignore matters of gender to make things simpler. The content of (1) is a type which depends on a context (a situation) which provides a value of the pronoun *he*. Thus it will have a parametric content which is a function from a context assigning a value to the pronoun to a type.¹ We will use the variable ‘*s*’ to represent such contexts. Thus a parametric content for *he* could be that given in (2a), for *left* (ignoring tense) as (2b) and their combination, using a variant of S-combination which we will discuss as we progress, is represented in (2c).

- (2) a. $\lambda s:[x_0:Ind] . \lambda P:Ppty . P([x=s.x_0])$
b. $\lambda s:Rec . \lambda r:[x:Ind] . [e : leave(r.x)]$
c. $\lambda s:[x_0:Ind] . [e : leave(s.x_0)]$

We follow Montague’s strategy in allowing the content of *he* to be (3) for any natural number *i*.

¹We will consider later how to combine types of context which assign values to pronouns with other context types which we have introduced with parametric contents.

$$(3) \quad \lambda s: [x_i:Ind] . \lambda P:Ppty . P([x=s.x_i])$$

Thus considering the boy and the dog, a content for (4a) will be (4b).

$$(4) \quad \begin{array}{ll} \text{a. he hugged it} \\ \text{b. } \lambda s: \begin{bmatrix} x_0:Ind \\ x_1:Ind \end{bmatrix} . [e : \text{hug}(s.x_0, s.x_1)] \end{array}$$

An advantage of using record types to characterize pronominal contexts rather than variable assignments is that we can add further information represented by the pronoun such as gender. Thus a simple treatment of gender for (4) might be given by making the content be (5).

$$(5) \quad \lambda s: \begin{bmatrix} x_0:Ind \\ c_0:\text{male}(x_0) \\ x_1:Ind \\ c_1:\text{neuter}(x_1) \end{bmatrix} . [e : \text{hug}(s.x_0, s.x_1)]$$

As we will see below, there are some complications with this simple idea when it comes to the interpretation of pronouns which are bound by quantifiers. Even for deictic pronouns there are problems determining which predicates should be used in a semantic treatment of gender. Even for a language like English which apparently has semantic gender (as opposed to grammatical gender like German or French), neuter can be used for objects which do not have gender (like tables) and for animals other than humans which do have gender and which can be referred to with masculine and feminine pronouns.

7.2 Quantifier scope and underspecification

Given the kind of interpretation rules we have so far we can obtain a reading for (6a) which corresponds to the kind of parametric content using pronominal contexts in (6b), using the abbreviations ‘boy’ and ‘dog’ as introduced in Chapter 5.

$$(6) \quad \begin{array}{ll} \text{a. a boy hugged every dog} \\ \text{b. } \lambda s:Rec . \\ \quad [e:\text{exist}(\text{boy}', \lambda r_1: [x:Ind] . [e:\text{every}(\text{dog}', \lambda r_2: [x:Ind] . [e:\text{hug}(r_1.x, r_2.x)])])] \end{array}$$

This, of course, represents the reading where there is a boy such that he hugs every dog. Notice that here we have vacuous abstraction over pronominal contexts and the constraint on them is that they are a record of some kind without requiring any particular fields.

In order to obtain the reading where *every dog* has wide scope, we follow Montague and pretty much everybody else in basing our treatment of quantifier scope on the treatment of free pronouns, though without the contribution of any gender information. Let us imagine just for a moment that such a pronoun existed in English and is written as *it*^{*} with the kind of pronoun interpretations given in (3). Then the content for (7a) could be (7b).

- (7) a. a boy hugged *it*^{*}
 b. $\lambda s: [x_0: Ind] . [e: exist(boy', \lambda r: [x: Ind] . [e: hug(r.x, s.x_0)])]$

Let us further imagine, contrary to fact, that English represented the fact that a noun phrase has wide scope over a sentence by placing it at the beginning of a sentence as in (8a) and giving it an interpretation where the interpretation of *it*^{*} gets bound as in (8b).

- (8) a. every dog, a boy hugged *it*^{*}
 b. $\lambda s: Rec . [e: every(dog', \lambda r_1: [x_0: Ind] . [e: exist(boy', \lambda r_2: [x: Ind] . [e: hug(r_2.x, r_1.x_0)])])]$

The imaginary English expression (8a) corresponds quite closely to the kind of representation for wide scope readings that are used in various theories of logical form. A major difference is that in logical form there is an index corresponding to the label '*x*₀' that we use in the interpretation which shows that *every dog* binds *it*^{*}. This might be represented something like in (9).

- (9) every dog_{*x*₀}, a boy hugged *it*_{*x*₀}^{*}

The imaginary sentence (8a) also corresponds closely to Montague's (1973) treatment of scope phenomena. Montague would also index the pronoun and use a quantification rule with the same index which would replace the pronoun with the noun-phrase being quantified in.

Neither of these options are open to us since our syntax is defined in terms of types of utterance situations and signs which related utterance situations to contents. Our realistic strategy does not allow for the use of additional imaginary structures. For this reason we will adapt the kind of storage technique used in Cooper (1983). In the original version of storage we moved from assigning a single content to a syntactic structure to assigning a set of contents in order to allow for the ambiguous interpretation of a single syntactic structure. In our sign-based approach using types the corresponding move is not such a major change and the result yields a theory involving underspecified content rather than a set of contents.

To see this consider the type *Sign* introduced in Chapter 3. Any object of type *Sign* will be of the type in (10).

$$(10) \quad \begin{bmatrix} \text{s-event} & : & SEvent \\ \text{syn} & : & Syn \\ \text{cont} & : & Cont \end{bmatrix}$$

The type in (10) is completely underspecified. Any sign will be of this type. We could specify it with respect to content by making the ‘cont’-field be a manifest field as in (11b), where c is as given in (11a).

$$(11) \quad \begin{array}{l} \text{a. } c = [e:\text{exist}(\text{boy}', \lambda r_1:[x:Ind] . [e:\text{every}(\text{dog}', \lambda r_2:[x:Ind] . [e:\text{hug}(r_1.x, r_2.x)]))] \\ \text{b. } \begin{bmatrix} \text{s-event} & : & SEvent \\ \text{syn} & : & Syn \\ \text{cont}=c & : & Cont \end{bmatrix} \end{array}$$

Now recall that the manifest field $[cont=c:Cont]$ is just a convenient way of writing $[cont:Cont_c]$ where $Cont_c$ is a singleton type whose only witness is c . It is in this sense that the content has been specified to be c . Suppose now that we do not have enough information to fully specify the content, that is, tie it down to be one particular content, but we know that it has to be one of either c or c' (as characterized in (12a)). This could be represented by using a join type of two singleton types, $Cont_c \vee Cont_{c'}$, as in (12b).

$$(12) \quad \begin{array}{l} \text{a. } c' = [e:\text{every}(\text{dog}', \lambda r_1:[x:Ind] . [e:\text{exist}(\text{boy}', \lambda r_2:[x:Ind] . [e:\text{hug}(r_2.x, r_1.x)]))] \\ \text{b. } \begin{bmatrix} \text{s-event} & : & SEvent \\ \text{syn} & : & Syn \\ \text{cont} & : & Cont_c \vee Cont_{c'} \end{bmatrix} \end{array}$$

(12) is the type of signs whose contents are either c or c' . This is, then, a single type, which corresponds to an “underspecified content”. Of course, the set of witnesses of the join type, $\{c, c'\}$, corresponds to the set of contents that could be generated by a storage algorithm. Our strategy, then, is to devise a way of computing such types on the basis of compositional interpretation. Notice that saying that the content is of type $Cont_c \vee Cont_{c'}$, that is, that it is either c or c' is distinct from saying that the content is the disjunction of c and c' , that is, if c and c' are types, the type $c \vee c'$. The witnesses for $c \vee c'$ will be situations s such that either $s : c$ or $s : c'$ whereas the witness for $Cont_c \vee Cont_{c'}$ will be the contents c and c' . This will make an important difference in compositional interpretation. For example, if a sentence S has a content of type $Cont_c \vee Cont_{c'}$ we would want to say that the sentence *Sam believes S* has a content of type $Cont_{\text{believe}(\text{sam}, c) \vee Cont_{\text{believe}(\text{sam}, c')}}$ rather than that it has the content ‘believe(sam, $c \vee c'$)’.

In order to make the relationship between sets of contents and join types smoother we will use generalized join types (Appendix A.8). Suppose we have computed a set of possible contents,

$\{c_1, \dots, c_n\}$, for a phrase, then the type of the content for that phrase is $\bigvee \{Cont_{c_1}, \dots, Cont_{c_n}\}$. If \mathbb{C} represents the set of contents we have computed, we can use the more convenient notation $\bigvee_{c \in \mathbb{C}} Cont_c$ to mean $\bigvee \{Cont_c \mid c \in \mathbb{C}\}$.

We will take content now to be divided into two fields labelled ‘unplugged’ and ‘plugged’. This terminology goes back to Bos (1996). The ‘unplugged’-field contains a set of unplugged interpretations consisting of a set of parametric quantifiers (possibly empty) and a core interpretation consisting of a parametric content. Thus we define a type *UInterp* of unplugged interpretations as in (13).

$$(13) \quad UInterp = \left[\begin{array}{ll} \text{quants} & : \{PQuant\} \\ \text{core} & : PCont \end{array} \right]$$

The ‘cont’-field in a sign will now be required to be of the type in (14).

$$(14) \quad \left[\begin{array}{ll} \text{unplugged} & : \{UInterp\} \\ \text{plugged} & : \bigvee_{c \in \text{unplugged}, c.\text{quants}=\emptyset} Cont_{c.\text{core}} \end{array} \right]$$

That is, the ‘unplugged’-field in the content will contain a set of unplugged interpretations (consisting of a quantifier store and a core interpretation) and the ‘plugged’-field will contain one of the cores of those unplugged interpretations whose quantifier store is empty.

Let us consider what this will look like for the example *every boy wants a dog* using *every’*, *boy’* etc. as abbreviations for non-parametric contents for these words. We show the ‘cont’-field of the sign type associated with constituents of this sentence.

(15) a. a dog

$$\text{b.} \quad \left[\begin{array}{l} \text{unplugged} = \left\{ \begin{array}{l} \text{quants} = \emptyset \\ \text{core} = \left[\begin{array}{l} \text{bg} = Rec \\ \text{fg} = \lambda s:Rec . a'(dog') \end{array} \right] \end{array} \right\}, \\ \left[\begin{array}{l} \text{quants} = \left\{ \begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \lambda s:[x_0:Ind] . a'(dog') \end{array} \right\} \\ \text{core} = \left[\begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \lambda s:[x_0:Ind] . \lambda P:Ppty . P(s.x_0) \end{array} \right] \end{array} \right\} : \{UInterp\} \\ \text{plugged} : Cont \left[\begin{array}{l} \text{bg} = Rec \\ \text{fg} = \lambda s:Rec . a'(dog') \end{array} \right] \end{array} \right]$$

(16) a. want a dog

$$\begin{array}{l}
 \text{b.} \left[\begin{array}{l}
 \text{unplugged} = \left\{ \begin{array}{l} \text{quants} = \emptyset \\ \text{core} = \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda s: \text{Rec} . \text{want}'(a'(\text{dog}')) \end{array} \right] \end{array} \right\}, \\
 \left[\begin{array}{l} \text{quants} = \left\{ \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \lambda s: [x_0: \text{Ind}] . a'(\text{dog}')) \end{array} \right] \right\} \\ \text{core} = \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \lambda s: [x_0: \text{Ind}] . \text{want}'(\lambda P: \text{Ppty} . P(s.x_0)) \end{array} \right] \end{array} \right\} : \{U\text{Interp}\} \\
 \text{plugged: Cont} \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda s: \text{Rec} . \text{want}'(a'(\text{dog}')) \end{array} \right]
 \end{array} \right]
 \end{array}$$

(17) a. every boy

$$\begin{array}{l}
 \text{b.} \left[\begin{array}{l}
 \text{unplugged} = \left\{ \begin{array}{l} \text{quants} = \emptyset \\ \text{core} = \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda s: \text{Rec} . \text{every}'(\text{boy}') \end{array} \right] \end{array} \right\}, \\
 \left[\begin{array}{l} \text{quants} = \left\{ \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \lambda s: [x_0: \text{Ind}] . \text{every}'(\text{boy}') \end{array} \right] \right\} \\ \text{core} = \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \lambda s: [x_0: \text{Ind}] . \lambda P: \text{Ppty} . P(s.x_0) \end{array} \right] \end{array} \right\} : \{U\text{Interp}\} \\
 \text{plugged: Cont} \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda s: \text{Rec} . \text{every}'(\text{boy}') \end{array} \right]
 \end{array} \right]
 \end{array}$$

(18) a. every boy wants a dog

$$\begin{array}{l}
 \text{b.} \left[\begin{array}{l}
 \text{unplugged} = \mathbb{C}_{\text{unplugged}} : \{U\text{Interp}\} \\
 \text{plugged: } \bigvee_{c \in \mathbb{C}_{\text{plugged}}} \text{Cont}_c
 \end{array} \right]
 \end{array}$$

where $\mathbb{C}_{\text{unplugged}}$ is the set whose members are listed in (19) and $\mathbb{C}_{\text{plugged}}$ is the set whose members are listed in (20).

$$\begin{array}{l}
 \text{(19) a.} \left[\begin{array}{l} \text{quants} = \emptyset \\ \text{core} = \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \text{fg} = \lambda s: \text{Rec} . \text{every}'(\text{boy}')(\text{want}'(a'(\text{dog}')))) \end{array} \right] \end{array} \right] \\
 \text{b.} \left[\begin{array}{l} \text{quants} = \left\{ \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \lambda s: [x_0: \text{Ind}] . \text{every}'(\text{boy}') \end{array} \right] \right\} \\ \text{core} = \left[\begin{array}{l} \text{bg} = [x_0: \text{Ind}] \\ \text{fg} = \left[\begin{array}{l} \lambda s: [x_0: \text{Ind}] . \\ \lambda P: \text{Ppty} . \\ P(s.x_0)(\text{want}'(a'(\text{dog}')))) \end{array} \right] \end{array} \right] \end{array} \right]
 \end{array}$$

$$\begin{array}{l}
\text{c.} \quad \left[\begin{array}{l} \text{quants} = \left\{ \begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \lambda s: [x_0:Ind] . \text{every}'(\text{boy}') \end{array} \right\}, \\ \left[\begin{array}{l} \text{bg} = [x_1:Ind] \\ \text{fg} = \lambda s: [x_1:Ind] . a'(\text{dog}') \end{array} \right] \end{array} \right] \\
\text{core} = \left[\begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \begin{array}{l} \lambda s: [x_0:Ind] . \\ \lambda P:Ppty . \\ P(s.x_0)(\text{want}'(\lambda P:Ppty . P(s.x_1))) \end{array} \end{array} \right] \\
\text{d.} \quad \left[\begin{array}{l} \text{quants} = \left\{ \begin{array}{l} \text{bg} = [x_1:Ind] \\ \text{fg} = \lambda s: [x_1:Ind] . a'(\text{dog}') \end{array} \right\} \\ \left[\begin{array}{l} \text{bg} = [x_1:Ind] \\ \text{fg} = \begin{array}{l} \lambda s: [x_1:Ind] . \\ \text{every}'(\text{boy}') \\ (\lambda r: [x_0:Ind] . \\ \lambda P:Ppty . P(r.x_0) \\ (\text{want}' \\ (\lambda P:Ppty . P(s.x_1)))) \end{array} \end{array} \right] \end{array} \right] \\
\text{e.} \quad \left[\begin{array}{l} \text{quants} = \left\{ \begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \lambda s: [x_0:Ind] . \text{every}'(\text{boy}') \end{array} \right\} \\ \left[\begin{array}{l} \text{bg} = [x_0:Ind] \\ \text{fg} = \begin{array}{l} \lambda s: [x_0:Ind] . \\ a'(\text{dog}') \\ (\lambda r: [x_1:Ind] . \\ \lambda P:Ppty . \\ P(r.x_1) \\ (\text{want}'(\lambda P:Ppty . P(s.x_0)))) \end{array} \end{array} \right] \end{array} \right] \\
\text{f.} \quad \left[\begin{array}{l} \text{quants} = \{ \} \\ \left[\begin{array}{l} \text{bg} = \text{Rec} \\ \lambda s: \text{Rec} \\ a'(\text{dog}') \\ (\lambda r_1: [x_1:Ind] . \\ \text{every}'(\text{boy}') \\ (\lambda r_2: [x_0:Ind] . \\ \lambda P:Ppty . P(r_1.x_0) \\ \text{wants}'(\lambda P:Ppty . P(r_1.x_1)))) \end{array} \right] \end{array} \right]
\end{array}$$

$$\begin{array}{l}
\text{g.} \quad \left[\begin{array}{l} \text{quants}=\{\} \\ \text{core} = \left[\begin{array}{l} \text{bg}=\text{Rec} \\ \lambda s:\text{Rec} \\ \text{every}'(\text{boy}') \\ (\lambda r_1:[x_0:\text{Ind}] . \\ \text{fg} = \begin{array}{l} a'(\text{dog}') \\ (\lambda r_2:[x_1:\text{Ind}] . \\ \lambda P:\text{Ppty} . P(r_1.x_0) \\ \text{wants}'(\lambda P:\text{Ppty} . P(r_1.x_1)))) \end{array} \end{array} \right] \end{array} \right]
\end{array}$$

(20) a. $\left[\begin{array}{l} \text{bg}=\text{Rec} \\ \text{fg}=\lambda s:\text{Rec} . \text{every}'(\text{boy}')(\text{want}'(a'(\text{dog}')))) \end{array} \right]$

b. $\left[\begin{array}{l} \text{bg}=\text{Rec} \\ \lambda s:\text{Rec} \\ a'(\text{dog}') \\ (\lambda r_1:[x_1:\text{Ind}] . \\ \text{fg} = \begin{array}{l} \text{every}'(\text{boy}') \\ (\lambda r_2:[x_0:\text{Ind}] . \\ \lambda P:\text{Ppty} . P(r_1.x_0) \\ \text{wants}'(\lambda P:\text{Ppty} . P(r_1.x_1)))) \end{array} \end{array} \right]$

c. $\left[\begin{array}{l} \text{bg}=\text{Rec} \\ \lambda s:\text{Rec} \\ \text{every}'(\text{boy}') \\ (\lambda r_1:[x_0:\text{Ind}] . \\ \text{fg} = \begin{array}{l} a'(\text{dog}') \\ (\lambda r_2:[x_1:\text{Ind}] . \\ \lambda P:\text{Ppty} . P(r_1.x_0) \\ \text{wants}'(\lambda P:\text{Ppty} . P(r_1.x_1)))) \end{array} \end{array} \right]$

We now take a look at what is needed in order achieve the analysis we have discussed above. We will express storage and retrieval as constraints on linguistically allowable sign types. As a preliminary we will define a bookkeeping function, \max_x , which enables us to determine the maximum i such that ' x_i ' is used as a label in a set of records, R . This will enable us to store quantifiers using labels ' x_0 ', ' x_1 ' and so on in order as in the examples discussed above. ' \max_x ' is defined as in (21) where R is an arbitrary set of records.

$$\begin{aligned}
(21) \quad \max_x(R) = & \\
& \max(\{i \mid \exists r \in R \exists v \langle x_i, v \rangle \in r.\text{bg}\}), \text{ if defined} \\
& \text{otherwise: } -1
\end{aligned}$$

That is, $\max_x(R)$ is the maximum i such that ' x_i ' is a label in the 'bg'-field of one of the records in R . If there is no such label $\max_x(R)$ returns -1. Thus in adding to a quantifier store, R , (as in the 'quants'-fields in the examples above) we will always add a new quantifier indexed by $\max_x(R)+1$. For brevity we will denote this $\text{incr}_x(R)$.

It will also be useful to define an operation that adds a field $[\ell=v]$ to a record r if there is no field with the label ℓ or replaces the ℓ -field in r with $[\ell=v]$. We will represent the new record thus derived as $r \oplus [\ell=v]$. This is defined in (22).

$$(22) \quad r \oplus [\ell=v] \text{ is} \\ r \cup \{\langle \ell, v \rangle\}, \text{ if there is no } v' \text{ such that } \langle \ell, v' \rangle \in \mathfrak{s} \\ (r - \{\langle \ell, r.l \rangle\}) \cup \{\langle \ell, v \rangle\}, \text{ otherwise}$$

Similarly we define $T \ominus [\ell:v]$ for record types, T , in (23).

$$(23) \quad T \ominus [\ell:v] \text{ is} \\ T - \{\langle \ell, v \rangle\}$$

Note that this version of subtraction for record types will only result in a well-formed record type if there are not fields in T which depend on ℓ . In order to rectify this we need to make $T \ominus [\ell:v]$ remove in addition all fields which depend on ℓ . A field, f , is dependent on ℓ just in case it is of the form $\langle \ell', \langle \mathcal{F}, \pi \rangle \rangle$ where π is a sequence of paths containing ℓ . Therefore we can refine the definition in (23) as that in (24).

$$(24) \quad T \ominus [\ell:v] \text{ is} \\ T - (\{\langle \ell, v \rangle\} \cup \{f \in T \mid f \text{ dependent on } \ell\})$$

This is technically not enough, since there may still be fields left which are dependent on the dependent fields we have removed. Therefore we need to make the definition recursive as in (25).

$$(25) \quad T \ominus [\ell:v] \text{ is} \\ (\dots (T - \{\langle \ell, v \rangle\} \ominus f_0) \dots) \ominus f_n \text{ where } \{f_0, \dots, f_n\} \text{ is } \{f \in T \mid f \text{ dependent on } \ell\}$$

We can now define storage as in (26), where σ is an arbitrary sign.

$$(26) \quad \textbf{Storage}$$

If $\alpha \in \sigma.\text{cont.unplugged}$, $\alpha.\text{core} : PQuant$ and
 $\alpha.\text{core.fg} \neq \lambda s: [x_i:Ind] . \lambda P:Ppty . P(s.x_i)$ for any i ,
 then

$$\left[\begin{array}{lcl} \text{quants} & = & \alpha.\text{quants} \cup \left\{ \left[\begin{array}{lcl} \text{bg} & = & [x_{\text{incr}_x(\alpha.\text{quants})}:Ind] \\ \text{fg} & = & \lambda s: [x_{\text{incr}_x(\alpha.\text{quants})}:Ind] . \alpha.\text{core.fg}(s) \end{array} \right] \right\} \\ \text{core} & = & \left[\begin{array}{lcl} \text{bg} & = & [x_{\text{incr}_x(\alpha.\text{quants})}:Ind] \\ \text{fg} & = & \lambda s: [x_{\text{incr}_x(\alpha.\text{quants})}:Ind] . \lambda P:Ppty . P(s.x_{\text{incr}_x(\alpha.\text{quants})}) \end{array} \right] \end{array} \right] \\ \in \sigma.\text{cont.unplugged}$$

If σ is a sign, $\sigma.\text{cont.unplugged}$ is a set of unplugged interpretations. (26) says that if any of those is a noun-phrase interpretation, that is, its core is of type $PQuant$, and is not the result of applying storage, that is, the foreground of its core is not an interpretation depending on a context, s , as introduced by storage, then $\sigma.\text{cont.unplugged}$ also contains an unplugged interpretation where the core has been stored.

We now characterize retrieval (in a version corresponding to quantification with scope over sentences, as opposed to verb phrases or common nouns) in (27), where again σ is an arbitrary sign.

(27) **Retrieval (S)**

If $\alpha \in \sigma.\text{cont.unplugged}$, for some i , $\lambda s: [x_i:Ind] . Q \in \alpha.\text{quants}$ and $\alpha.\text{core} : (T \rightarrow Type)$, where $T \sqsubseteq [x_i:Ind]$, then

$$\left[\begin{array}{lcl} \text{quants} & = & \alpha.\text{quants} - \{ \lambda s: [x_i:Ind] . Q \} \\ \text{core} & = & \lambda s: T - \{ [x_i:Ind] \} . Q(\lambda r: [x_i:Ind] . \alpha.\text{core}(s \oplus [x_i=r.x_i])) \end{array} \right] \\ \in \sigma.\text{cont.unplugged}$$

(27) says that if any of the unplugged interpretations in a sign contains a quantifier indexed by i , that is, it is a parametric quantifier whose domain type is $[x_i:Ind]$, and the core of this unplugged interpretation is a function from contexts to a type, that is, it corresponds to a (declarative) sentence interpretation, where the context is required to provide a value for ' x_i ', that is, the domain type is a subtype of $[x_i:Ind]$,² then the quantifier is removed from the store and given scope over the sentence interpretation in the core binding ' x_i '. The dependence on ' x_i ' in the context is removed from the core interpretation.

Finally, we show one way in which stored quantifiers can be “percolated” to higher constituents by showing in (28) how to construct an unplugged content from two constituent unplugged contents when the basic semantic composition involves function application.

²Note that this requirement will prevent the kind of vacuous binding that Keller storage (Keller, 1988) avoided.

(28) **Application**

If $\alpha \in \sigma_1.\text{cont.unplugged}$ and $\beta \in \sigma_2.\text{cont.unplugged}$ and η is a relabelling of $\beta.\text{quants}$ such that $\eta(x_n) = x_{\max_x(\alpha.\text{quants})+n+1}$, then

$$\alpha @ @ \beta = \begin{bmatrix} \text{quants} & = & \alpha.\text{quants} \cup [\beta.\text{quants}]_\eta \\ \text{core} & = & \alpha.\text{core} @ [\beta.\text{core}]_\eta \end{bmatrix}$$

(28) says that the application of an unplugged content, α , to another unplugged content, β , involves first relabelling the quantifier indices in β so that they increment the quantifier indices in α in the way that is illustrated in the examples discussed above. For example, if both α and β contain the quantifier index ‘ x_0 ’ and this is the maximum in α , then ‘ x_0 ’ will be relabelled to ‘ x_1 ’ in β . Now the result of application is an unplugged content whose ‘quants’ are the union of α ’s ‘quants’ and the relabelled ‘quants’ of β . The core is the result of applying α ’s core to β ’s core under the same relabelling.

7.3 Long distance dependencies

Let us consider how to derive the content of *who Sam hugged* as a relative clause. We use **hug_V** to represent the foreground of the content of *hug* as a transitive verb, characterized as in (29).

$$(29) \quad \mathbf{hug}_V = \lambda s: \text{Rec} . \lambda r_2: [x: \text{Quant}] . \lambda r_1: [x: \text{Ind}] . [e: \text{hug}(r_1.x, r_2.x)]$$

Our theory of syntax in terms of event types means that what have been called “gaps” or “traces” in other theories of syntax would somehow have to correspond to non-events in which nothing happens and do not have any temporal extent. This strongly suggests that they do not exist. Instead we will take a strategy similar to that pursued, for example, in combinatory categorial grammar (see, for example, Steedman, 2012) where *hug* can also be interpreted as a verb phrase whose content is “looking for” a *wh*-phrase content. We will represent the foreground of this content as **hug_{VP}** and characterize it in terms of **hug_V** as in (30a) which is identical with (30b).

$$(30) \quad \begin{aligned} \text{a. } \mathbf{hug}_{VP} &= \lambda s: [\text{wh}_0: \text{Ind}] . \mathbf{hug}_V(s)(\lambda P: \text{Ppty} . P([x=s.\text{wh}_0])) \\ \text{b. } \lambda s: [\text{wh}_0: \text{Ind}] & . \lambda r_1: [x: \text{Ind}] . [e: \text{hug}(r_1.x, \lambda P: \text{Ppty} . P([x=s.\text{wh}_0]))] \end{aligned}$$

Let us now consider the content of *Sam hugged*. In order to avoid the additional complexity of combining different kinds of context which we will undertake in Section 7.6 we will go back to the treatment of proper names in Chapter 3 rather than the treatment in Chapter 4 where proper names introduce a constraint on the context. Thus we will consider the foreground of the parametric content of *Sam*, which we will represent as **Sam**, to be (31) where ‘sam’ is a particular individual named Sam.

$$(31) \text{ Sam} = \lambda s:Rec . \lambda P:Ppty . P([x=sam])$$

The foreground of the content of *Sam hugged*, which we will represent as **Sam**[^]**hugged**, is (32a) which is identical with (32b).

$$(32) \text{ a. } \text{Sam}^{\wedge} \text{hugged} = \lambda s:[wh_0:Ind] . \text{Sam}(s)(\text{hug}_{VP}(s))$$

$$\text{ b. } \lambda s:[wh_0:Ind] . [e:\text{hug}(sam, \lambda P:Ppty . P([x=s.wh_0]))]$$

The foreground of the content of *who*, which we will represent as **who**, is given in (33).

$$(33) \text{ who} = \lambda s:[wh:Ind] . \lambda P:Ppty . P([x=s.wh])$$

In order to be able to characterize the kind of binding involved in combining *who* with *Sam hugged* we will define an operation, \oplus , which adds a field to a record or, if the label of the field being added is already present in the record, will replace that field with the new field. This operation is characterized in (34).

$$(34) \text{ If } r \text{ is a record, } \ell \text{ is a label and } v \text{ is some TTR object, then}$$

$$r \oplus [\ell = v] \text{ is}$$

$$r \cup \{\langle \ell, v \rangle\}, \text{ if there is no } v' \text{ such that } \langle \ell, v' \rangle \in r$$

$$(r - \{\langle \ell, r.\ell \rangle\}) \cup \{\langle \ell, v \rangle\}, \text{ otherwise}$$

We are using records here to do the work of partial variable assignments in logic and the \oplus -operation carries out the kind of modification which is performed on variable assignments when defining binding in the logical treatment. This now gives us what we need to characterize the foreground of the content of the relative clause *who Sam hugged*, which we will represent as **who**[^]**Sam**[^]**hugged**, as in (35a) which is identical with (35b).

$$(35) \text{ a. } \text{who}^{\wedge} \text{Sam}^{\wedge} \text{hugged} =$$

$$\lambda s:Rec .$$

$$\lambda r_1:[x:Ind] .$$

$$\text{who}(s \oplus [wh = r_1.x]) (\lambda r_2:[x:Ind] . \text{Sam}^{\wedge} \text{hugged}(s \oplus [wh_0 = r_2.x]))$$

$$\text{ b. } \lambda s:Rec . \lambda r_1:[x:Ind] . [e:\text{hug}(sam, \lambda P:Ppty . P([x=r_1.x]))]$$

7.4 Conservativity and dynamic generalized quantifiers

In Chapter 5, example (62), we introduced the notion of dynamic generalized quantifier and pointed out that one of the original motivations for them was what is known as donkey anaphora which we will discuss in Section 7.5. Here we will point out a connection between dynamic quantifiers and conservativity of quantifiers. The informal way to state conservativity for quantifiers is as in (36a) and an example is given in (36b).

- (36) a. $Q A B$ is true just in case $Q A A \& B$ is true
 b. *every farmer likes a donkey* is true just in case *every farmer is a farmer and likes a donkey* (or more naturally, *every farmer is a farmer who likes a donkey*) is true

Most, if not all, natural language quantifiers have this property.³

Now consider the definition of dynamic generalized quantifiers given in Chapter 5, Section 5.5. There in example (62) we gave dynamic versions of generalized quantifiers. We can give a general characterization of the witness conditions for dynamic generalized quantifiers as in (37), where q is a predicate corresponding to a quantifier and q^* is the relation between sets corresponding to q in classical generalized quantifier theory.

- (37) $e : q(P, Q)$ iff q^* holds between $[\downarrow P \uparrow e]$ and $[\downarrow Q|_{\mathcal{F}(P.\text{fg})} \uparrow e]$

According to the definitions given in Chapter 5, $[\downarrow P \uparrow e]$ is the set of individuals in e which have property P and $Q|_{\mathcal{F}(P.\text{fg})}$ is Q with its domain restricted by P . We spell this out for the type corresponding to *every dog runs* in (38). Here, as before, we only represent the foregrounds of the properties which are arguments to the predicate ‘every’ in order to make the presentation clearer.

7.5 Anaphora

[donkey anaphora]

7.6 Structured contexts, binding principles and projection

7.7 Compositional semantics and incremental processing

³For discussion, see Peters and Westerståhl (2006), p. 138f.

Appendix A

Type theory with records

Unless otherwise stated this is the version of TTR presented in Cooper (2012b).

A.1 Underlying set theory

In previous statements of this system such as Cooper (2012b) we tacitly assumed a standard underlying set theory such as ZF (Zermelo-Fraenkel) with urelements (as formulated for example in Suppes, 1960). This is what we take to be the common or garden working set theory which is familiar from the core literature on formal semantics deriving from Montague’s original work (Montague, 1974). When we introduced complex objects and types other than records and record types we were not explicit about exactly which structured set-theoretic object they represented. The reason for this was that, except in the case of records and record types, it did not seem important exactly how you code structured objects in the set theory and a detailed exposition would seem to provide another level of complication over and above an already complicated story.

In this version we will take advantage of the freedom provided by an appendix and spell out a set theoretic coding for all of our structured objects. We will use *labelled sets* to model our structured objects which is what we use for records and record types as well. We will assume that our set theory comes equipped with a set of *urelements* (objects which are not sets but which can be members of sets) which is partitioned into two countable subsets or *urelements proper* (intuitively “real atomic objects”) and *labels* (intuitively “objects that are used to label real objects, either atomic or sets”). A *labelled set* is a set of ordered pairs whose first member is a label and whose second element is either an urelement proper or a set (possibly a labelled set), such that no more than one ordered pair can contain any particular label as its first member. This means that a labelled set is the traditional set theoretic construction of an extensional function from a set of labels onto some set. Suppose that we have a set

$$\{a, b, c, d\}$$

and that $\ell_0, \ell_1, \ell_2, \ell_3$ are labels. Then examples of labelled sets would be

$$\{\langle \ell_0, a \rangle, \langle \ell_1, b \rangle, \langle \ell_2, c \rangle, \langle \ell_3, d \rangle\}$$

and

$$\{\langle \ell_0, \{\langle \ell_0, a \rangle, \langle \ell_1, b \rangle\} \rangle, \langle \ell_2, c \rangle, \langle \ell_3, d \rangle\}$$

Labelled sets where we identify particular distinguished labels will always give us enough structure to model the structured objects that we need and define operations on them as required by the type theory.

A.2 Basic types

A *system of basic types* is a pair:

$$\mathbf{TYPE}_B = \langle \mathbf{Type}, A \rangle$$

where:

1. **Type** is a non-empty set
2. A is a function whose domain is **Type**
3. for any $T \in \mathbf{Type}$, $A(T)$ is a set disjoint from **Type**
4. for any $T \in \mathbf{Type}$, $a :_{\mathbf{TYPE}_B} T$ iff $a \in A(T)$

A *modal system of basic types*¹ is a family of pairs:

$$\mathbf{TYPE}_{MB} = \langle \mathbf{Type}, A \rangle_{A \in \mathcal{A}}$$

where:

1. \mathcal{A} is a set of functions with domain **Type**
2. for each $A \in \mathcal{A}$, $\langle \mathbf{Type}, A \rangle$ is a system of basic types

¹This definition was not present in Cooper (2012b).

This enables us to define some simple modal notions:

If $\mathbf{TYPE}_{MB} = \langle \mathbf{Type}, A \rangle_{A \in \mathcal{A}}$ is a modal system of basic types, we shall use the notation \mathbf{TYPE}_{MB_A} (where $A \in \mathcal{A}$) to refer to that system of basic types in \mathbf{TYPE}_{MB} whose type assignment is A . Then:

1. for any $T_1, T_2 \in \mathbf{Type}$, T_1 is (necessarily) equivalent to T_2 in \mathbf{TYPE}_{MB} , $T_1 \approx_{\mathbf{TYPE}_{MB}} T_2$, iff for all $A \in \mathcal{A}$, $\{a \mid a :_{\mathbf{TYPE}_{MB_A}} T_1\} = \{a \mid a :_{\mathbf{TYPE}_{MB_A}} T_2\}$
2. for any $T_1, T_2 \in \mathbf{Type}$, T_1 is a subtype of T_2 in \mathbf{TYPE}_{MB} , $T_1 \sqsubseteq_{\mathbf{TYPE}_{MB}} T_2$, iff for all $A \in \mathcal{A}$, $\{a \mid a :_{\mathbf{TYPE}_{MB_A}} T_1\} \subseteq \{a \mid a :_{\mathbf{TYPE}_{MB_A}} T_2\}$
3. for any $T \in \mathbf{Type}$, T is necessary in \mathbf{TYPE}_{MB} iff for all $A \in \mathcal{A}$, $\{a \mid a :_{\mathbf{TYPE}_{MB_A}} T\} \neq \emptyset$
4. for any $T \in \mathbf{Type}$, T is possible in \mathbf{TYPE}_{MB} iff for some $A \in \mathcal{A}$, $\{a \mid a :_{\mathbf{TYPE}_{MB_A}} T\} \neq \emptyset$

A.3 Complex types

A.3.1 Predicates

We start by introducing the notion of a predicate signature.

A *predicate signature* is a triple

$$\langle \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle$$

where:

1. **Pred** is a set (of predicates)
2. **ArgIndices** is a set (of indices for predicate arguments, normally types)
3. **Arity** is a function with domain **Pred** and range included in the set of finite sequences of members of **ArgIndices**.

A *polymorphic predicate signature* is a triple

$$\langle \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle$$

where:

1. **Pred** is a set (of predicates)
2. **ArgIndices** is a set (of indices for predicate arguments, normally types)
3. *Arity* is a function with domain **Pred** and range included in the powerset of the set of finite sequences of members of **ArgIndices**.

A.3.2 Systems of complex types

A *system of complex types* is a quadruple:

$$\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$$

where:

1. $\langle \mathbf{BType}, A \rangle$ is a system of basic types
2. $\mathbf{BType} \subseteq \mathbf{Type}$
3. for any $T \in \mathbf{Type}$, if $a :_{\langle \mathbf{BType}, A \rangle} T$ then $a :_{\mathbf{TYPE}_C} T$
4. $\langle \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle$ is a (polymorphic) predicate signature
- 5.² $P(a_1, \dots, a_n) \in \mathbf{PType}$ iff $P \in \mathbf{Pred}, T_1 \in \mathbf{Type}, \dots, T_n \in \mathbf{Type}, \mathbf{Arity}(P) = \langle T_1, \dots, T_n \rangle$
 $((\langle T_1, \dots, T_n \rangle \in \mathbf{Arity}(P)) \text{ and } a_1 :_{\mathbf{TYPE}_C} T_1, \dots, a_n :_{\mathbf{TYPE}_C} T_n)$
6. $\mathbf{PType} \subseteq \mathbf{Type}$
7. for any $T \in \mathbf{PType}$, $F(T)$ is a set disjoint from \mathbf{Type}
8. for any $T \in \mathbf{PType}$, $a :_{\mathbf{TYPE}_C} T$ iff $a \in F(T)$

We call the pair $\langle A, F \rangle$ in a complex system of types the *model* because of its similarity to first order models in providing values for the basic types and the ptypes constructed from predicates and arguments. It is this pair which connects the system of types to the non-type theoretical world of objects and situations.

In Cooper (2012b) we did not define exactly what object is represented by $P(a_1, \dots, a_n)$. Here we will specify it to be the labelled set

²This clause has been modified since Cooper (2012b) where it was a conditional rather than a biconditional.

$$\{\langle \text{pred}, P \rangle, \langle \text{arg}_1, a_1 \rangle, \dots, \langle \text{arg}_n, a_n \rangle\}$$

where ‘pred’, ‘arg_i’ are reserved labels (not used except as required here).

What are the objects which belong to these types? The intuition is that, for example,

$$e : \text{run}(a)$$

means that e is an event or situation where the individual a is running. There are two competing intuitions about what e could be. One is that it is a “part of the world”, a non-set (urelement). That is, from the perspective of set theory and type theory it is an unstructured atom. The other intuition we have is that it is a structured object which contains a as a component and in which a running activity is going on which involves smaller events such as picking feet up off the ground, spending certain time in each step cycle with neither foot touching the ground and so on. We want to allow for both of these intuitions. That is, a witness for a ptype can be a non-set corresponding to our notion of an event of a certain type. Or it can be the kind of labelled set which we call a record. That is e does not only belong to the type ‘run(a)’ but also a record type which characterizes in more detail the structure of the event. We will argue in the text that both intuitions are important and that observers of the world shift between type theories where certain ptypes are regarded as types of non-sets and type theories where those ptypes are types of records.

A.4 Function types

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has function types if

1. for any $T_1, T_2 \in \mathbf{Type}$, $(T_1 \rightarrow T_2) \in \mathbf{Type}$
2. for any $T_1, T_2 \in \mathbf{Type}$, $f :_{\mathbf{TYPE}_C} (T_1 \rightarrow T_2)$ iff f is a function whose domain is $\{a \mid a :_{\mathbf{TYPE}_C} T_1\}$ and whose range is included in $\{a \mid a :_{\mathbf{TYPE}_C} T_2\}$

In Cooper (2012b) we did not specify exactly what object is represented by a function type $(T_1 \rightarrow T_2)$. Here we specify it to be the labelled set

$$\{\langle \text{dmn}, T_1 \rangle, \langle \text{rng}, T_2 \rangle\}$$

where ‘dmn’ (“domain”) and ‘rng’ (“range”) are reserved labels.

We also introduce a limited kind of polymorphism in function types which we did not have in Cooper (2012b).

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has polymorphic function types if

1. for any $T_1, T_2 \in \mathbf{Type}$, $\bigvee_{T \sqsubseteq T_1} (T \rightarrow T_2) \in \mathbf{Type}$
2. for any $T_1, T_2 \in \mathbf{Type}$, $f :_{\mathbf{TYPE}_C} \bigvee_{T \sqsubseteq T_1} (T \rightarrow T_2)$ iff there is some type T' such that $f : (T' \rightarrow T_2)$ and $T' \sqsubseteq T_1$

We specify the type $\bigvee_{T \sqsubseteq T_1} (T \rightarrow T_2)$ to be the labelled set

$$\{\langle \text{polydmn}, T_1 \rangle, \langle \text{rng}, T_2 \rangle\}$$

where ‘polydmn’ (“polymorphic domain”) and ‘rng’ (“range”) are reserved labels (‘rng’ being the same reserved label that was used for non-polymorphic function types).

In Cooper (2012b) we also left it open exactly what kind of object a function is and assumed there was some theory of functions which would allow us to characterize them in terms of their domain and range. In a classical set theoretic setting where functions are modelled extensionally as sets of ordered pairs, little more needs to be said. Ideally, we want a notion of function that is more like a program or a procedure. That is, functions can be intensional in the sense that two distinct functions can correspond to the same set of ordered pairs. However, it seems that for the purposes at hand the standard extensional notion of function as a set of ordered pairs is sufficient and a lot more straightforward to handle than a more intensional notion given the set-theoretic basis on which we are building. There appears to be sufficient intensionality introduced by our notion of type and the of-type relation. For this reason, we will model functions here as sets of ordered pairs in the classical set-theoretic way. Ultimately, we suspect that a more computational and intensional notion of function should be substituted, but at this point it is unclear what consequences this might have for the rest of the system.

This choice of modelling functions as sets of ordered pairs means that $f :_{\mathbf{TYPE}_C} (T_1 \rightarrow T_2)$ iff $f \subseteq \{a \mid a :_{\mathbf{TYPE}_C} T_1\} \times \{a \mid a :_{\mathbf{TYPE}_C} T_2\}$ such that if $b \in \{a \mid a :_{\mathbf{TYPE}_C} T_1\}$ then there is exactly one c , such that $\langle b, c \rangle \in f$. We shall say that in this case the result of applying the function f to b , in symbols, $f(b)$, is c .

We introduce a notation for functions which is borrowed from the λ -calculus as used by Montague (1973). Let $O[v]$ be the notation for some object of our type theory which uses the variable v and let T be a type. Then the function

$$\lambda v : T . O[v]$$

is to be the function

$$\{\langle v, O[v] \rangle \mid v : T\}$$

(Here we suppress the subscript \mathbf{TYPE}_C on the ‘.’.) For example, the function

$$\lambda v:Ind . \text{run}(v)$$

is the set of ordered pairs

$$\{\langle v, \text{run}(v) \rangle \mid v : Ind\}$$

Recall that ‘ $\text{run}(v)$ ’ is itself a representation for the labelled set

$$\{\langle \text{pred}, \text{run} \rangle, \langle \text{arg}_1, v \rangle\}$$

Note that if f is the function $\lambda v:Ind . \text{run}(v)$ and $a:Ind$ then $f(a)$ (the result of applying f to a) is ‘ $\text{run}(a)$ ’. Our definition of function-argument application guarantees what is called β -equivalence in the λ -calculus. When we discuss record types as arguments to functions we will need to introduce one slight complication to our notion of function application. We will introduce that complication when we discuss record types.

A.5 List types

List types were not included in Cooper (2012b).

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has list types if

1. for any $T \in \mathbf{Type}$, $[T] \in \mathbf{Type}$
2. for any $T \in \mathbf{Type}$,
 - a) $a \mid L : \mathbf{TYPE}_C [T]$ iff $a : \mathbf{TYPE}_C T$ and $L : \mathbf{TYPE}_C [T]$
 - b) $[] : \mathbf{TYPE}_C [T]$

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has non-empty list types if

1. for any $T \in \mathbf{Type}$, $_{ne}[T] \in \mathbf{Type}$

2. for any $T \in \mathbf{Type}$,

- a) $a \mid L : \mathbf{TYPE}_C \text{ }_{ne}[T]$ iff $a : \mathbf{TYPE}_C T$ and $L : \mathbf{TYPE}_C \text{ }_{ne}[T]$
- b) $[a] : \mathbf{TYPE}_C \text{ }_{ne}[T]$ iff $a : T$

If $a \mid L : \mathbf{TYPE}_C \text{ }_{ne}[T]$ for some system of complex types \mathbf{TYPE}_C and type T , then we use $\text{fst}(a \mid L)$ to refer to a and $\text{rst}(a \mid L)$ to refer to L .

In contrast to Cooper (2012b) we here make it explicit that $[T]$ represents $\{\langle \text{lst}, T \rangle\}$ and $\text{ }_{ne}[T]$ represents $\{\langle \text{nelst}, T \rangle\}$ where ‘lst’ and ‘nelst’ are reserved labels.

Lists are a common data structure used in computer science but they are not normally defined in basic set theory, although it is straightforward to define them in terms of sets. In Cooper (2012b) we did not specify an encoding of lists in terms of sets. Here we will use an encoding with labelled sets using the reserved labels ‘fst’ and ‘rst’ for the first member of the list and the remainder (“rest”) of the list respectively. We let the empty list, $[]$, be the empty set, \emptyset .³ If L is a list then $a \mid L$ is to be the labelled set $\{\langle \text{fst}, a \rangle, \langle \text{rst}, L \rangle\}$.

A.6 Set types

Set types were not included in Cooper (2012b).

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has set types if

- 1. for any $T \in \mathbf{Type}$, $\{T\} \in \mathbf{Type}$
- 2. for any $T \in \mathbf{Type}$, $A : \mathbf{TYPE}_C \{T\}$ iff A is a set and for all $a \in A$, $a : \mathbf{TYPE}_C T$

We let $\{T\}$ represent the labelled set $\{\langle \text{set}, T \rangle\}$ where ‘set’ is a reserved label.

We also introduce a special kind of set type known as a plurality type. The idea here is that a plurality is a set that does not contain any two objects such that one is a proper part of the other. The notion of proper part is characterized by:

- 1. If r_1 and r_2 are records then r_1 is a proper part of r_2 , $r_1 < r_2$, just in case $\varphi(r_1) \subset \varphi(r_2)$.

³If it is important to distinguish the empty list from the empty set we could use an additional reserved label, e.g. ‘lst’, and have the empty list be the labelled set $\{\langle \text{lst}, \emptyset \rangle\}$.

2. If o_1 and o_2 are objects of some type and at least one of them is not of type *Rec*, then o_1 is *not* a proper part of o_2 , $o_1 \not\prec o_2$

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Ariety} \rangle, \langle A, F \rangle \rangle$ has *plurality types* if

1. for any $T \in \mathbf{Type}$, $\{ T \} \in \mathbf{Type}$
2. for any $T \in \mathbf{Type}$, $A :_{\mathbf{TYPE}_C} \{ T \}$ iff
 - a) $A :_{\mathbf{TYPE}_C} \{ T \}$
 - b) if $a \in A$ then for any b such that $a < b$, $b \notin A$

We let $\{ T \}$ represent the labelled set $\{ \langle \text{plurality}, T \rangle \}$ where ‘plurality’ is a reserved label.

A.7 Singleton types

Singleton types were not included in the formal definition in Cooper (2012b).

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Ariety} \rangle, \langle A, F \rangle \rangle$ has *singleton types* if

1. for any $T \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_C} T$, $T_a \in \mathbf{Type}$
2. for any $T \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_C} T$, $b :_{\mathbf{TYPE}_C} T_a$ iff $a = b$

These clauses are those presented in Cooper (2012b). A more general version of these clauses seems useful for the uses we wish to make of singleton types, for example, the restriction of properties discussed in Appendix B.1. The more general version allows singleton types to be created using an object of any type but will guarantee that the type is empty if the object is not of the type being restricted:

1. for any $T, T' \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_C} T'$, $T_a \in \mathbf{Type}$
2. for any $T, T' \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_C} T'$, $b :_{\mathbf{TYPE}_C} T_a$ iff $b :_{\mathbf{TYPE}_C} T$ and $a = b$

As we now allow singleton types that are empty (because the object used to restrict them is not of the required type) it may seem that the name “singleton type” is a misnomer. The cases of empty types are those where we have failed to define a singleton type.

Note that these definitions allow the formation of singleton types from singleton types. We sometimes refer to these as *multiple* singleton types and notate them as $T_{a,b,\dots}$ rather than the typographically unfortunate $T_{ab,\dots}$. Following the definition above an object c will be of type $T_{a,b}$ just in case $c : T$ and $a = b = c$.

We let T_a represent the labelled set $\{\langle \text{singleton}, T, a \rangle\}$ where ‘singleton’ is a reserved label.

A.8 Join types

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has *join types* if

1. for any $T_1, T_2 \in \mathbf{Type}$, $(T_1 \vee T_2) \in \mathbf{Type}$
2. for any $T_1, T_2 \in \mathbf{Type}$, $a :_{\mathbf{TYPE}_C} (T_1 \vee T_2)$ iff $a :_{\mathbf{TYPE}_C} T_1$ or $a :_{\mathbf{TYPE}_C} T_2$

Here, but not in Cooper (2012b), we specify that $(T_1 \vee T_2)$ represents the labelled set $\{\langle \text{disj}_1, T_1 \rangle, \langle \text{disj}_2, T_2 \rangle\}$ where ‘disj₁’ and ‘disj₂’ are reserved labels (“disjunct”).

We add generalized join types which were not present in Cooper (2012b). A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has *generalized join types* if

1. for any finite set of types, \mathbb{T} , such that $\mathbb{T} \subseteq \mathbf{Type}$, $\bigvee \mathbb{T} \in \mathbf{Type}$
2. for any finite $\mathbb{T} \subseteq \mathbf{Type}$, $a :_{\mathbf{TYPE}_C} \bigvee \mathbb{T}$ iff $a :_{\mathbf{TYPE}_C} T$ for some $T \in \mathbb{T}$

We specify that $\bigvee \mathbb{T}$ represents the labelled set $\{\langle \text{disj}, \mathbb{T} \rangle\}$ where ‘disj’ is a reserved label (“disjunction”).

A.9 Meet types

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has *meet types* if

1. for any $T_1, T_2 \in \mathbf{Type}$, $(T_1 \wedge T_2) \in \mathbf{Type}$
2. for any $T_1, T_2 \in \mathbf{Type}$, $a :_{\mathbf{TYPE}_C} (T_1 \wedge T_2)$ iff $a :_{\mathbf{TYPE}_C} T_1$ and $a :_{\mathbf{TYPE}_C} T_2$

Here, but not in Cooper (2012b), we specify that $(T_1 \wedge T_2)$ represents the labelled set $\{\langle \text{conj}_1, T_1 \rangle, \langle \text{conj}_2, T_2 \rangle\}$ where ‘conj₁’ and ‘conj₂’ are reserved labels (“conjunct”).

We add generalized meet types which were not present in Cooper (2012b). A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has *generalized meet types* if

1. for any finite set of types, \mathbb{T} , such that $\mathbb{T} \subseteq \mathbf{Type}$, $\bigwedge \mathbb{T} \in \mathbf{Type}$
2. for any finite $\mathbb{T} \subseteq \mathbf{Type}$, $a :_{\mathbf{TYPE}_C} \bigwedge \mathbb{T}$ iff $a :_{\mathbf{TYPE}_C} T$ for all $T \in \mathbb{T}$

We specify that $\bigwedge \mathbb{T}$ represents the labelled set $\{\langle \text{conj}, \mathbb{T} \rangle\}$ where ‘conj’ is a reserved label (“conjunction”).

A.10 Models and modal systems of types

A modal system of complex types provides a collection of models, \mathcal{M} , so that we can talk about properties of the whole collection of type assignments provided by the various models $M \in \mathcal{M}$.

A *modal system of complex types based on \mathcal{M}* is a family of quadruples⁴:

$$\mathbf{TYPE}_{MC} = \langle \mathbf{Type}_M, \mathbf{BType}, \langle \mathbf{PType}_M, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, M \rangle_{M \in \mathcal{M}}$$

where for each $M \in \mathcal{M}$, $\langle \mathbf{Type}_M, \mathbf{BType}, \langle \mathbf{PType}_M, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, M \rangle$ is a system of complex types.

This enables us to define modal notions:

If $\mathbf{TYPE}_{MC} = \langle \mathbf{Type}_M, \mathbf{BType}, \langle \mathbf{PType}_M, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, M \rangle_{M \in \mathcal{M}}$ is a modal system of complex types based on \mathcal{M} , we shall use the notation \mathbf{TYPE}_{MC_M} (where $M \in \mathcal{M}$) to refer to that system of complex types in \mathbf{TYPE}_{MC} whose model is M . Let $\mathbf{Type}_{MC_{restr}}$ be $\bigcap_{M \in \mathcal{M}} \mathbf{Type}_M$, the “restrictive” set of types which occur in all possibilities, and $\mathbf{Type}_{MC_{incl}}$ be $\bigcup_{M \in \mathcal{M}} \mathbf{Type}_M$, the “inclusive” set of types which occur in at least one possibility. Then we can define modal notions either restrictively or inclusively (indicated by the subscripts r and i respectively):

⁴This definition has been modified since Cooper (2012b) to make \mathbf{PType} and \mathbf{Type} be relativized to the model M .

restrictive modal notions

1. for any $T_1, T_2 \in \mathbf{Type}_{MC_{restr}}$, T_1 is (necessarily) equivalent_r to T_2 in \mathbf{Type}_{MC} , $T_1 \approx_{\mathbf{Type}_{MC}} T_2$, iff for all $M \in \mathcal{M}$, $\{a \mid a : \mathbf{Type}_{MC_M} T_1\} = \{a \mid a : \mathbf{Type}_{MC_M} T_2\}$
2. for any $T_1, T_2 \in \mathbf{Type}_{MC_{restr}}$, T_1 is a subtype_r of T_2 in \mathbf{Type}_{MC} , $T_1 \sqsubseteq_{\mathbf{Type}_{MC}} T_2$, iff for all $M \in \mathcal{M}$, $\{a \mid a : \mathbf{Type}_{MC_M} T_1\} \subseteq \{a \mid a : \mathbf{Type}_{MC_M} T_2\}$
3. for any $T \in \mathbf{Type}_{MC_{restr}}$, T is necessary_r in \mathbf{Type}_{MC} iff for all $M \in \mathcal{M}$, $\{a \mid a : \mathbf{Type}_{MC_M} T\} \neq \emptyset$
4. for any $T \in \mathbf{Type}_{MC_{restr}}$, T is possible_r in \mathbf{Type}_{MC} iff for some $M \in \mathcal{M}$, $\{a \mid a : \mathbf{Type}_{MC_M} T\} \neq \emptyset$

inclusive modal notions

1. for any $T_1, T_2 \in \mathbf{Type}_{MC_{incl}}$, T_1 is (necessarily) equivalent_i to T_2 in \mathbf{Type}_{MC} , $T_1 \approx_{\mathbf{Type}_{MC}} T_2$, iff for all $M \in \mathcal{M}$, if T_1 and T_2 are members of \mathbf{Type}_M , then $\{a \mid a : \mathbf{Type}_{MC_M} T_1\} = \{a \mid a : \mathbf{Type}_{MC_M} T_2\}$
2. for any $T_1, T_2 \in \mathbf{Type}_{MC_{incl}}$, T_1 is a subtype_i of T_2 in \mathbf{Type}_{MC} , $T_1 \sqsubseteq_{\mathbf{Type}_{MC}} T_2$, iff for all $M \in \mathcal{M}$, if T_1 and T_2 are members of \mathbf{Type}_M , then $\{a \mid a : \mathbf{Type}_{MC_M} T_1\} \subseteq \{a \mid a : \mathbf{Type}_{MC_M} T_2\}$
3. for any $T \in \mathbf{Type}_{MC_{incl}}$, T is necessary_i in \mathbf{Type}_{MC} iff for all $M \in \mathcal{M}$, if $T \in \mathbf{Type}_M$, then $\{a \mid a : \mathbf{Type}_{MC_M} T\} \neq \emptyset$
4. for any $T \in \mathbf{Type}_{MC_{incl}}$, T is possible_i in \mathbf{Type}_{MC} iff for some $M \in \mathcal{M}$, if $T \in \mathbf{Type}_M$, then $\{a \mid a : \mathbf{Type}_{MC_M} T\} \neq \emptyset$

It is easy to see that if any of the restrictive definitions holds for given types in a particular system then the corresponding inclusive definition will also hold for those types in that system.

A.11 The type *Type* and stratification

An *intensional system of complex types* is a family of quadruples indexed by the natural numbers:

$$\mathbf{Type}_{IC} = \langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle_{n \in \mathcal{Nat}}$$

where (using \mathbf{TYPE}_{IC_n} to refer to the quadruple indexed by n):

1. for each n , $\langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle$ is a system of complex types
2. for each n , $\mathbf{Type}^n \subseteq \mathbf{Type}^{n+1}$ and $\mathbf{PType}^n \subseteq \mathbf{PType}^{n+1}$
3. for each n , if $T \in \mathbf{PType}^n$ and $p \in F^n(T)$ then $p \in F^{n+1}(T)$
4. for each $n > 0$, $Type^n \in \mathbf{Type}^n$
5. for each $n > 0$, $T :_{\mathbf{TYPE}_{IC_n}} Type^n$ iff $T \in \mathbf{Type}^{n-1}$

Here, but not in Cooper (2012b), we make explicit that $Type$ is a distinguished urelement and that $Type^n$ represents the labelled set $\{\langle \text{ord}, n \rangle, \langle \text{typ}, Type \rangle\}$ where ‘ord’ and ‘typ’ are reserved labels (“order”, “type”).

An intensional system of complex types \mathbf{TYPE}_{IC} ,

$$\mathbf{TYPE}_{IC} = \langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle_{n \in \mathbf{Nat}}$$

has dependent function types if

1. for any $n > 0$, $T \in \mathbf{Type}^n$ and $\mathcal{F} :_{\mathbf{TYPE}_{IC_n}} (T \rightarrow Type^n)$, $((a : T) \rightarrow \mathcal{F}(a)) \in \mathbf{Type}^n$
2. for each $n > 0$, $f :_{\mathbf{TYPE}_{IC_n}} ((a : T) \rightarrow \mathcal{F}(a))$ iff f is a function whose domain is $\{a \mid a :_{\mathbf{TYPE}_{IC_n}} T\}$ and such that for any a in the domain of f , $f(a) :_{\mathbf{TYPE}_{IC_n}} \mathcal{F}(a)$.

We might say that on this view dependent function types are “semi-intensional” in that they depend on there being a type of types for their definition but they do not introduce types as arguments to predicates and do not involve the definition of orders of types in terms of the types of the next lower order.

Here, in contrast to Cooper (2012b), we make explicit that $((a : T) \rightarrow \mathcal{F}(a))$ represents the labelled set $\{\langle \text{dmn}, T \rangle, \langle \text{deprng}, \mathcal{F} \rangle\}$ where ‘dmn’ as before for function types is a reserved label corresponding to “domain” and ‘deprng’ is a reserved label corresponding to “dependent range”.

Putting the definition of a modal type system and an intensional type system together we obtain:⁵

⁵This explicit definition was not present in Cooper (2012b).

An *intensional modal system of complex types based on \mathfrak{M}* is a family, indexed by the natural numbers, of families of quadruples indexed by members of \mathfrak{M} :

$$\mathbf{TYPE}_{IMC} = \langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \mathcal{M}_n \rangle_{\mathcal{M} \in \mathfrak{M}, n \in \mathbb{N}}$$

where:

1. for each n , $\langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \mathcal{M}_n \rangle_{\mathcal{M} \in \mathfrak{M}}$ is a modal system of complex types based on $\{\mathcal{M}_n \mid \mathcal{M} \in \mathfrak{M}\}$
2. for each $\mathcal{M} \in \mathfrak{M}$, $\langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \mathcal{M}_n \rangle_{n \in \mathbb{N}}$ is an intensional system of complex types

A.12 Records and Record types

In this section we will define what it means for a system of complex types to have record types. The objects of record types, that is, records, are themselves structured mathematical objects of a particular kind and we will start by characterizing them.


A.12.1 Records

A *record* is a finite set of ordered pairs (called *fields*) which is the graph of a function. If r is a record and $\langle \ell, v \rangle$ is a field in r we call ℓ a *label* and v a *value* in r and we use $r.\ell$ to denote v . $r.\ell$ is called a *path* in r . This means that a record is a labelled set as introduced in Appendix A.1.

We will use a tabular format to represent records. A record $\{\langle \ell_1, v_1 \rangle, \dots, \langle \ell_n, v_n \rangle\}$ is displayed as

$$\left[\begin{array}{cc} \ell_1 & = & v_1 \\ \dots & & \\ \ell_n & = & v_n \end{array} \right]$$

An alternative notation is as a tree (or directed graph) whose root is unlabelled:



We will sometime say that the occurrences ℓ_1, \dots, ℓ_n are *sisters* in the record.

A value may itself be a record and paths may extend into embedded records. A record which contains records as values is called a *complex record* and otherwise a record is *simple*. Values which are not records are called *leaves*. Consider a record r

$$\left[\begin{array}{l} f \\ g \end{array} = \left[\begin{array}{l} f \\ h \end{array} = \left[\begin{array}{l} ff \\ gg \\ g \\ h \end{array} = \begin{array}{l} a \\ b \\ c \\ a \\ d \end{array} \right] \right] \right]$$

Among the paths in r are $r.f$, $r.g.h$ and $r.f.f.ff$ which denote, respectively,


$$\left[\begin{array}{l} f \\ g \end{array} = \left[\begin{array}{l} ff \\ gg \end{array} = \begin{array}{l} a \\ b \end{array} \right] \right]$$

$$\left[\begin{array}{l} g \\ h \end{array} = \begin{array}{l} a \\ d \end{array} \right]$$

and a . We will make a distinction between *absolute paths*, such as those we have already mentioned, which consist of a record followed by a series of labels connected by dots and *relative paths* which are just a series of labels connected by dots, e.g. $g.h$. Relative paths are useful when we wish to refer to similar paths in different records. We will use *path* to refer to either absolute or relative paths when it is clear from the context which is meant. The set of leaves of r , also known as its *extension* (those objects other than labels which it contains), is $\{a, b, c, d\}$. The bag (or multiset) of leaves of r , also known as its *multiset extension*, is $\{a, a, b, c, d\}$. A record may be regarded as a way of labelling and structuring its extension.

An object, a , is a *component* of a record, r , in symbols, $a \in r$, just in case there is some path, π , in r such that $r.\pi = a$. Thus the record, r , above has the following components: $r.f$, $r.f.f$, $r.f.f.ff$, $r.f.f.gg$, $r.f.g$, $r.g$, $r.g.h$, $r.g.h.g$ and $r.g.h.h$. An object, a , is *present* in a record, r , in symbols, $a \in r$, just in case either $a = r$ or $a \in r$.

We can, of course, also think of complex records as trees, for example:



We can thus use tree terminology to describe relationships between occurrences of labels. For example, we can say that the topmost occurrence of f immediately dominates another occurrence of f and an occurrence of g .

Two records are *(multiset) extensionally equivalent* if they have the same (multiset) extension. Two important, though trivial, facts about records are:

Flattening. For any record r , there is a multiset extensionally equivalent simple record. We can define an operation of flattening on records which will always produce an equivalent simple record. In the case of our example, the result of flattening is

$$\left[\begin{array}{lcl} f.f.ff & = & a \\ f.f.gg & = & b \\ f.g & = & c \\ g.h.g & = & a \\ g.h.h & = & d \end{array} \right]$$

assuming the flattening operation uses paths from the original record in a rather obvious way to create unique labels for the new record.

Relabelling. For any record r , if $\pi_1.\ell.\pi_2$ is a path π in r , and $\pi_1.\ell'.\pi_2'$ is *not* a path in r (for any π_2'), then substituting ℓ' for the occurrence of ℓ in π results in a record which is multiset equivalent to r . We could, for example, substitute k for the second occurrence of g in the path $g.h.g$ in our example record.

$$\left[\begin{array}{lcl} f & = & \left[\begin{array}{lcl} f & = & \left[\begin{array}{lcl} ff & = & a \\ gg & = & b \end{array} \right] \\ g & = & c \end{array} \right] \\ g & = & \left[\begin{array}{lcl} h & = & \left[\begin{array}{lcl} k & = & a \\ h & = & d \end{array} \right] \end{array} \right] \end{array} \right]$$

We will first look in more detail at flattening. Flattened records have *complex labels* (or *paths*) of the form $\ell_1.\ell_2.\dots.\ell_n$. If L is a set of labels, the set of *complex labels*, L_π , based on L is the minimal set such that:

1. if $\ell \in L$, then $\ell \in L_\pi$
2. if $\pi \in L_\pi$ and $\ell \in L$, then $\pi.\ell \in L_\pi$

Recall that a record, r , is a set of fields, f , where each f , is an ordered pair, $\langle \ell, v \rangle$ where ℓ is a (complex) label and v is the value of the field which may itself be a record or some other object. If r is a record, then $\varphi(r)$, the *flattening* of r , is

$$\bigcup_{f \in r} \text{FlattenField}(f)$$

where $\text{FlattenField}(\langle \ell, v \rangle)$ is

$$\varphi(\{\langle \ell.\ell_1, v_1 \rangle, \langle \ell.\ell_2, v_2 \rangle, \dots, \langle \ell.\ell_n, v_n \rangle\})$$

if v is the record $\{\langle \ell_1, v_1 \rangle, \langle \ell_2, v_2 \rangle, \dots, \langle \ell_n, v_n \rangle\}$ and $\langle \ell, v \rangle$ otherwise.

We can show that if r is a record then $\varphi(r)$ is a record. It follows trivially from the requirement that records are graphs of functions from labels to values. If r is a record $\{\langle \ell_1, v_1 \rangle, \dots, \langle \ell_n, v_n \rangle\}$, then the *set of labels* of r , $\text{labels}(r)$, is $\{\ell_1, \dots, \ell_n\}$. For example, if r is the record

$$\left[\begin{array}{lcl} x & = & \left[\begin{array}{lcl} z & = & a \\ w & = & b \end{array} \right] \\ y & = & c \end{array} \right]$$

then $\text{labels}(r)$ is $\{x, y\}$. That is, ‘labels’ picks out the “highest level” of labels in a complex record. We use $\text{tpaths}(r)$ to represent the *set of total paths* of r , $\text{labels}(\varphi(r))$.

If L is a set of labels, then the set of *complex labels* (or *paths*), L_+ , based on L is the minimal set such that

1. if $\ell \in L$, then $\ell \in L_+$
2. if $\pi \in L_+$ and $\ell \in L$, then $\pi.\ell \in L_+$

If r is a record such that $\text{labels}(r) \subseteq L_+$, then the *inverse flattening* (or *expansion*) of r , $\varphi^-(r)$, is the minimal set such that

1. if $\langle \ell, v \rangle \in r$ and $\ell \in L$, then $\langle \ell, v \rangle \in \varphi^-(r)$
2. if $\ell \in L$, $\pi \in L_+$ and $\langle \ell.\pi, v \rangle \in r$, then $\langle \ell, \varphi^-(\{\langle \pi', v' \rangle \mid \langle \ell.\pi', v' \rangle \in r\}) \rangle \in \varphi^-(r)$

Claim: If r is a record, then $\varphi^-(\varphi(r)) = r$.

The *set of total paths of r* , $\text{tpaths}(r)$, is $\text{labels}(\varphi(r))$.

If r is a record, then the *set of paths of r* , $\text{paths}(r)$, is the minimal set such that

1. If $\pi \in \text{tpaths}(r)$, then $\pi \in \text{paths}(r)$
2. If ℓ is a label (in L), π is a non-empty path based on L and $\pi.\ell \in \text{paths}(r)$, then $\pi \in \text{paths}(r)$

If Π is a set of paths (in L_+), then the *set of longest paths in Π* , $\text{longest}(\Pi)$ is such that $\pi \in \text{longest}(\Pi)$ iff $\pi \in \Pi$ and there is no π' such that $\pi.\pi' \in \Pi$.

Claim: $\text{longest}(\text{paths}(r)) = \text{tpaths}(r)$

If R is a set of records, then the *set of common paths of R* , $\text{cpaths}(R)$, is

$$\bigcap_{r \in R} \text{paths}(r),$$

the *set of longest common paths of R* , $\text{lcpaths}(R)$, is

$$\text{longest}\left(\bigcap_{r \in R} \text{paths}(r)\right)$$

and the *set of common total paths of R* , $\text{ctpaths}(R)$, is

$$\bigcap_{r \in R} \text{tpaths}(r)$$

Claim: $\text{lcpaths}(R) = \text{ctpaths}(R)$

Consider the records

$$r_1 = \left[\begin{array}{l} x = \\ y = \end{array} \left[\begin{array}{l} x = a \\ y = b \\ x = c \\ y = d \end{array} \right] \right]$$

$$r_2 = \left[\begin{array}{l} x = \\ y = \end{array} \left[\begin{array}{l} z = e \\ w = f \\ x = g \\ y = h \end{array} \right] \right]$$

The following hold true:

$$\begin{aligned} \text{paths}(r_1) &= \{x, x.x, x.y, y, y.x, y.y\} \\ \text{paths}(r_2) &= \{x, x.z, x.w, y, y.x, y.y\} \\ \text{lcpaths}(\{r_1, r_2\}) &= \{x, y.x, y.y\} \end{aligned}$$

The notion of longest common paths for a set of records will be important when we discuss paths in record types below. [Is this still true?]

We now turn our attention to relabelling. An L -relabelling is a one-one function whose domain and range are included in L_+ . We will standardly represent a relabelling, η , whose domain is $\{\pi_1, \dots, \pi_n\}$ and which is defined by $\eta(\pi_1) = \pi'_1, \dots, \eta(\pi_n) = \pi'_n$ as:

$$\begin{array}{l} \pi_1 \rightsquigarrow \pi'_1 \\ \vdots \\ \pi_n \rightsquigarrow \pi'_n \end{array}$$

If r is a record $\{\langle \ell_1, v_1 \rangle, \dots, \langle \ell_n, v_n \rangle\}$ and η is a relabelling, the *application of η to r* , $\eta * r$ is $\{\langle \ell'_1, v_1 \rangle, \dots, \langle \ell'_n, v_n \rangle\}$ where for each i such that $1 \leq i \leq n$, ℓ'_i is $\eta(\ell_i)$ if ℓ_i is in the domain of η and ℓ_i otherwise. Suppose η is the relabelling $x \rightsquigarrow xx, z \rightsquigarrow zz$ and r is as above. Then $\eta * r$ is

$$\left[\begin{array}{l} xx = \\ y = c \end{array} \left[\begin{array}{l} z = a \\ w = b \end{array} \right] \right]$$

That is, the label ‘x’ is changed, but nothing else.

We say that a relabelling, η , is *limited* to a record, r , just in case $\text{dom}(\eta) \subseteq \text{labels}(r)$.

We are, however, mainly interested in replacing total paths in a record which is why we have defined relabellings on the set of paths, L_+ rather than the set of labels, L . The *result of relabelling a record, r , with a relabelling, η* , in symbols, $[r]_\eta$, is $\varphi^-(\eta * \varphi(r))$. We say that η is a *proper relabelling* for r just in case $[r]_\eta$ is a record.

Suppose r is the record

$$\left[\begin{array}{l} x = \\ y = \end{array} \left[\begin{array}{l} x = a \\ y = b \\ x = c \end{array} \right] \right]$$

and η is the relabelling

$$\begin{array}{l} x.x \rightsquigarrow y.z \\ x.y \rightsquigarrow x \end{array}$$

$\varphi(r)$ is the flattened record

$$\left[\begin{array}{l} x.x = a \\ x.y = b \\ y.x = c \end{array} \right]$$

$\eta * \varphi(r)$ is

$$\left[\begin{array}{l} y.z = a \\ x = b \\ y.x = c \end{array} \right]$$

Then $[r]_\eta$ is the expansion of this record, that is, the record

$$\left[\begin{array}{l} x = b \\ y = \end{array} \left[\begin{array}{l} x = c \\ z = a \end{array} \right] \right]$$

and η is a proper relabelling for r .

Now let us consider a case where the relabelling is not proper. Suppose we start with the same record, r , as above but take a new relabelling, η' :

$$x.x \rightsquigarrow x$$

Now $\eta' * \varphi(r)$ is the record

$$\left[\begin{array}{ll} x & = a \\ x.y & = b \\ y.x & = c \end{array} \right]$$

However, the expansion $\varphi^-(\eta' * \varphi(r))$ is the set of ordered pairs:

$$\{\langle x, a \rangle, \langle x, \{\langle y, b \rangle\} \rangle, \langle y, \{\langle x, c \rangle\} \rangle\}$$

which is not a record since it is not the graph of a function: it has two ordered pairs whose first member is 'x'. We can read the problem already from the flattened record since it has a label, 'x', which is a proper initial segment of another of its labels, 'x.y'. Therefore η' is not a proper relabelling for r . Note that we cannot read the improperness off η' alone. η' is a proper relabelling for the record r' :

$$\left[\begin{array}{ll} x & = \left[\begin{array}{ll} x & = a \end{array} \right] \\ y & = \left[\begin{array}{ll} x & = c \end{array} \right] \end{array} \right]$$

$[r']_{\eta'}$ is the record:

$$\left[\begin{array}{ll} x & = a \\ y & = \left[\begin{array}{ll} x & = c \end{array} \right] \end{array} \right]$$

We use η^- to represent the inverse of the labelling η . That is, for any $\pi, \pi', \pi \rightsquigarrow \pi'$ is in η just in case $\pi' \rightsquigarrow \pi$ is in η^- .

Claim: If η is a proper relabelling for r limited to $\varphi(r)$, then η^- is a proper relabelling for $[r]_\eta$ and $[[r]_\eta]_{\eta^-} = r$.

The requirement that η is limited to $\varphi(r)$ is important here since it prevents cases where η^- would change labelling not changed by η . Consider an example where η is

$$\begin{aligned} x &\rightsquigarrow z \\ x.x &\rightsquigarrow y \end{aligned}$$

and r is

$$\left[\begin{array}{l} x = \\ z = \end{array} \left[\begin{array}{l} x = a \\ b \end{array} \right] \right]$$

Note that in this case η is *not* limited to $\varphi(r)$ whose labels are ‘x.x’ and ‘z’ but do not include ‘x’. η is proper for r and $[r]_\eta$ is

$$\left[\begin{array}{l} y = a \\ z = b \end{array} \right]$$

Now consider η^- which is

$$\begin{aligned} z &\rightsquigarrow x \\ y &\rightsquigarrow x.x \end{aligned}$$

The relabelling of $[r]_\eta$ by η^- will be the set of ordered pairs

$$\{\langle x, b \rangle, \langle x, \{\langle x, a \rangle\} \rangle\}$$

which is not a record (since ‘x’ occurs twice as a label). Hence in this case η^- is not a proper relabelling for $[r]_\eta$ and $[[r]_\eta]_{\eta^-}$ is not identical with r . The reason for this is that η^- relabels a field that was not relabelled by η . Hence it is doing more than simply “undoing” the relabelling performed by η . This state of affairs will be prevented by requiring that η is limited to r , for example by removing ‘ $x \rightsquigarrow z$ ’ from η . Only paths changed by η will be changed back by η^- and our claim will hold true.

A.12.2 Record types

A (proper) record *type* is a labelled set where the objects labelled are types or, in some cases, certain kinds of mathematical objects which can be used to construct types.

A record r is *well-typed* with respect to a system of types **TYPE** with set of types **Type** and a set of labels L iff for each field $\langle \ell, a \rangle \in r$, $\ell \in L$ and either $a :_{\mathbf{TYPE}} T$ for some $T \in \mathbf{Type}$ or a is itself a record which is well-typed with respect to **TYPE** and L .

A system of complex types $\mathbf{TYPE}_C = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has record types based on $\langle L, \mathbf{RType} \rangle$, where L is a countably infinite set (of labels) and $\mathbf{RType} \subseteq \mathbf{Type}$, where \mathbf{RType} is defined by:

1. $Rec \in \mathbf{RType}$
2. $r :_{\mathbf{TYPE}_C} Rec$ iff r is a well-typed record with respect to \mathbf{TYPE}_C and L .
3. $ERec \in \mathbf{RType}$
4. $r :_{\mathbf{TYPE}_C} ERec$ iff $r = \emptyset$
5. if $\ell \in L$ and $T \in \mathbf{Type}$, then $\{\langle \ell, T \rangle\} \in \mathbf{RType}$.
6. $r :_{\mathbf{TYPE}_C} \{\langle \ell, T \rangle\}$ iff $r :_{\mathbf{TYPE}_C} Rec$, $\langle \ell, a \rangle \in r$ and $a :_{\mathbf{TYPE}_C} T$.
7. if $R \in \mathbf{RType} - \{Rec, ERec\}$, $\ell \in L$, ℓ does not occur as a label in R (i.e. there is no field $\langle \ell', T' \rangle$ in R such that $\ell' = \ell$) and $T \in \mathbf{Type}$, then $R \cup \{\langle \ell, T \rangle\} \in \mathbf{RType}$.
8. $r :_{\mathbf{TYPE}_C} R \cup \{\langle \ell, T \rangle\}$ iff $r :_{\mathbf{TYPE}_C} R$, $\langle \ell, a \rangle \in r$ and $a :_{\mathbf{TYPE}_C} T$.

We say that T is a *proper* record type if it is a non-empty set of fields.⁶

This gives us non-dependent record types in a system of complex types. We can extend this to intensional systems of complex types (with stratification).

An *intensional system of complex types* $\mathbf{TYPE}_{IC} = \langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle_{n \in \mathbf{Nat}}$ has record types based on $\langle L, \mathbf{RType}^n \rangle_{n \in \mathbf{Nat}}$ if for each n , $\langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle$ has record types based on $\langle L, \mathbf{RType}^n \rangle$ and

1. for each n , $\mathbf{RType}^n \subseteq \mathbf{RType}^{n+1}$
2. for each $n > 0$, $RecType^n \in \mathbf{RType}^n$

⁶This terminology was not introduced in Cooper (2012b).

3. for each $n > 0$, $T : \mathbf{TYPE}_{IC_n} \text{RecType}^n$ iff $T \in \mathbf{RType}^{n-1}$

Here, but not in Cooper (2012b), we make explicit that *RecType* is treated in a similar manner to *Type*, that is, it is a distinguished urelement and RecType^n represents the labelled set $\{\langle \text{ord}, n \rangle, \langle \text{typ}, \text{RecType} \rangle\}$ where ‘ord’ and ‘typ’ are reserved labels (“order”, “type”).

Intensional type systems may in addition contain *dependent* record types.

An *intensional system of complex types* $\mathbf{TYPE}_{IC} = \langle \mathbf{Type}^n, \mathbf{BType}, \langle \mathbf{PType}^n, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F^n \rangle \rangle_{n \in \text{Nat}}$ has *dependent record types based on* $\langle L, \mathbf{RType}^n \rangle_{n \in \text{Nat}}$, if it has record types based on $\langle L, \mathbf{RType}^n \rangle_{n \in \text{Nat}}$ and for each $n > 0$

1. if R is a member of \mathbf{RType}^n , $\ell \in L$ not occurring as a label in R , $T_1, \dots, T_m \in \mathbf{Type}^n$, $R.\pi_1, \dots, R.\pi_m$ are paths in R and \mathcal{F} is a function of type $(T_1 \rightarrow \dots \rightarrow (T_m \rightarrow \text{Type}^n) \dots)$, then $R \cup \{\langle \ell, \langle \mathcal{F}, \langle \pi_1, \dots, \pi_m \rangle \rangle \rangle\} \in \mathbf{RType}^n$.
2. $r : \mathbf{TYPE}_{IC_n} R \cup \{\langle \ell, \langle \mathcal{F}, \langle \pi_1, \dots, \pi_m \rangle \rangle \rangle\}$ iff $r : \mathbf{TYPE}_{IC_n} R$, $\langle \ell, a \rangle$ is a field in r , $r.\pi_1 : \mathbf{TYPE}_{IC_n} T_1, \dots, r.\pi_m : \mathbf{TYPE}_{IC_n} T_m$ and $a : \mathbf{TYPE}_{IC_n} \mathcal{F}(r.\pi_1), \dots, (r.\pi_m)$.

We represent a record type $\{\langle \ell_1, T_1 \rangle, \dots, \langle \ell_n, T_n \rangle\}$ graphically as

$$\left[\begin{array}{ll} \ell_1 & : \quad T_1 \\ \dots & \\ \ell_n & : \quad T_n \end{array} \right]$$

In the case of a manifest field, that is, one containing a singleton type, as in $\langle \ell, T_a \rangle$, we display this

$$\left[\begin{array}{ll} \ell=a & : \quad T \end{array} \right]$$

In the case of a multiple singleton type (a singleton type formed from a singleton type) as in $\langle \ell, T_{a,b,\dots} \rangle$, we display

$$\left[\begin{array}{ll} \ell=a,b,\dots & : \quad T \end{array} \right]$$

In the case of dependent record types we sometimes use a convenient notation representing e.g.

$$\langle \lambda u \lambda v \text{ love}(u, v), \langle \pi_1, \pi_2 \rangle \rangle$$

as

$$\text{love}(\pi_1, \pi_2)$$

Our systems now allow both function types and dependent record types and allow dependent record types to be arguments to functions. We have to be careful when considering what the result of applying a function to a dependent record type should be. Consider the following simple example:

$$\lambda v_0 : \text{RecType} . [c_0 : v_0]$$

What should be the result of applying this function to the record type

$$\left[\begin{array}{ll} x & : \text{Ind} \\ c_1 & : \langle \lambda v_1 : \text{Ind} . \text{dog}(v_1), \langle x \rangle \rangle \end{array} \right]$$

Given normal assumptions about function application the result would be

$$\left[c_0 : \left[\begin{array}{ll} x & : \text{Ind} \\ c_1 & : \langle \lambda v_1 : \text{Ind} . \text{dog}(v_1), \langle x \rangle \rangle \end{array} \right] \right] \text{ (incorrect!)}$$

but this would be incorrect. In fact it is not a well-formed record type since x is not a path in it. Instead the result should be

$$\left[c_0 : \left[\begin{array}{ll} x & : \text{Ind} \\ c_1 & : \langle \lambda v_1 : \text{Ind} . \text{dog}(v_1), \langle c_0.x \rangle \rangle \end{array} \right] \right]$$

where the path from the top of the record type is specified. However, in the abbreviatory notation we write just ‘ x ’ when the label is used as an argument and interpret this as the path from the top of the record type to the field labelled ‘ x ’ in the local record type. Thus we will write

$$\left[\begin{array}{ll} x & : \text{Ind} \\ c_1 & : \text{dog}(x) \end{array} \right]$$

(where the ‘x’ in ‘dog(x)’ signifies the path consisting of just the single label ‘x’) and

$$\left[\begin{array}{c} c_0 \end{array} : \left[\begin{array}{c} x : Ind \\ c_1 : dog(x) \end{array} \right] \right]$$

(where the ‘x’ in ‘dog(x)’ signifies the path from the top of the record type down to ‘x’ in the local record type, that is, ‘c₀.x’).⁷

Note that this adjustment of paths is only required when a record type is being substituted into a position that lies on a path within a resulting record type. It will not, for example, apply in a case where a record type is to be substituted for an argument to a predicate such as when applying the function

$$\lambda v_0 : RecType . [c_0 : appear(v_0)]$$

to

$$\left[\begin{array}{c} x : Ind \\ c_1 : \langle \lambda v : Ind . dog(v), \langle x \rangle \rangle \\ c_2 : \langle \lambda v : Ind . approach(v), \langle x \rangle \rangle \end{array} \right]$$

where the position of v_0 is in an “intensional context”, that is, as the argument to a predicate and there is no path to this position in the record type resulting from applying the function. Here the result of the application is

$$\left[\begin{array}{c} c_0 : appear(\left[\begin{array}{c} x : Ind \\ c_1 : \langle \lambda v : Ind . dog(v), \langle x \rangle \rangle \\ c_2 : \langle \lambda v : Ind . approach(v), \langle x \rangle \rangle \end{array} \right]) \end{array} \right]$$

with no adjustment necessary to the paths representing the dependencies.⁸ (Note that ‘c₀.x’ is not a path in this record type.)

Suppose that we wish to represent a type which requires that there is some dog such that it appears to be approaching (that is a *de re* reading). In the abbreviatory notation we might be tempted to write

⁷This convention of representing the path from the top of the record type to the “local” field by the final label on the path is new since Cooper (2012b).

⁸This record corresponds to the interpretation of *it appears that a dog is approaching*.

$$\left[\begin{array}{lcl} x & : & Ind \\ c_1 & : & \text{dog}(x) \\ c_0 & : & \text{appear}([c_2 : \text{approach}(x)]) \end{array} \right] \text{ (incorrect!)}$$

corresponding to

$$\left[\begin{array}{lcl} x & : & Ind \\ c_1 & : & \langle \lambda v:Ind . \text{dog}(v), \langle x \rangle \rangle \\ c_0 & : & \text{appear}([c_2 : \langle \lambda v:Ind . \text{approach}(v), \langle x \rangle \rangle]) \end{array} \right] \text{ (incorrect!)}$$

This is, however, incorrect since it refers to a path ‘x’ in the type which is the argument to ‘appear’ which does not exist. Instead we need to refer to the path ‘x’ in the record type containing the field labelled ‘c₀’:

$$\left[\begin{array}{lcl} x & : & Ind \\ c_1 & : & \langle \lambda v:Ind . \text{dog}(v), \langle x \rangle \rangle \\ c_0 & : & \langle \lambda v:Ind . \text{appear}([c_2 : \text{approach}(v)]), \langle x \rangle \rangle \end{array} \right]$$

In the abbreviatory notation we will use ‘ \uparrow ’ to indicate that the path referred to is in the “next higher” record type⁹:

$$\left[\begin{array}{lcl} x & : & Ind \\ c_1 & : & \text{dog}(x) \\ c_0 & : & \text{appear}([c_2 : \text{approach}(\uparrow x)]) \end{array} \right]$$

These matters arise as a result of our choice of using paths to represent dependencies in record types (rather than, for example, introducing additional unique identifiers to keep track of the positions within a record type as has been suggested by Thierry Coquand). It seems like a matter of implementation rather than a matter of substance and it is straightforward to define a path-aware notion of substitution which can be used in the definition of what it means to apply a TTR function to an argument. If f is a function represented by $\lambda v : T . \varphi$ and α is the representation of an object of type T , then the result of applying f to α , $f(\alpha)$, is represented by $\text{Subst}(\alpha, v, \varphi, \emptyset)$, that is, the result of substituting α for v in φ with respect to the empty path where for arbitrary α, v, φ, π , $\text{Subst}(\alpha, v, \varphi, \pi)$ is defined as

1. $\text{extend-paths}(\alpha, \pi)$, if φ is v

⁹This notation is new since Cooper (2012b).

2. φ , if φ is of the form $\lambda v : T . \zeta$, for some T and ζ (i.e. don't do any substitution if v is bound within φ)
3. $\lambda u : T . \text{Subst}(\alpha, v, \zeta, \pi)$, if φ is of the form $\lambda u : T . \zeta$ and u is not v .
4. $\left[\begin{array}{ll} \ell_1 & : \text{Subst}(\alpha, v, T_1, \pi.\ell_1) \\ \dots & \\ \ell_n & : \text{Subst}(\alpha, v, T_n, \pi.\ell_n) \end{array} \right]$, if φ is $\left[\begin{array}{ll} \ell_1 & : T_1 \\ \dots & \\ \ell_n & : T_n \end{array} \right]$
5. $P(\text{Subst}(\alpha, v, \beta_1, \pi), \dots, \text{Subst}(\alpha, v, \beta_n, \pi))$, if α is $P(\beta_1, \dots, \beta_n)$ for some predicate P
6. φ otherwise

$\text{extend-paths}(\alpha, \pi)$ is

1. $\langle f, \langle \pi.\pi_1, \dots, \pi.\pi_n \rangle \rangle$, if α is $\langle f, \langle \pi_1, \dots, \pi_n \rangle \rangle$
2. $\left[\begin{array}{ll} \ell_1 & : \text{extend-paths}(T_1, \pi) \\ \dots & \\ \ell_n & : \text{extend-paths}(T_n, \pi) \end{array} \right]$ if α is $\left[\begin{array}{ll} \ell_1 & : T_1 \\ \dots & \\ \ell_n & : T_n \end{array} \right]$
3. $P(\text{extend-paths}(\beta_1, \pi), \dots, \text{extend-paths}(\beta_n, \pi))$, if α is $P(\beta_1, \dots, \beta_n)$ for some predicate P
4. α , otherwise

A.13 Merges of record types

If T_1 and T_2 are record types then there will always be a record type (not a meet) T_3 which is necessarily equivalent to $T_1 \wedge T_2$. Let us consider some examples:

$$[f:T_1] \wedge [g:T_2] \approx \left[\begin{array}{l} f:T_1 \\ g:T_2 \end{array} \right]$$

$$[f:T_1] \wedge [f:T_2] \approx [f:T_1 \wedge T_2]$$

We define a function μ which maps meets of record types to an equivalent record type, record types to equivalent types where meets in their values have been simplified by μ and any other types to themselves:

1. If for some $T_1, T_2, T = T_1 \wedge T_2$ then $\mu(T) = \mu'(\mu(T_1) \wedge \mu(T_2))$.
2. If T is a record type then $\mu(T)$ is T' such that for any $\ell, v, \langle \ell, \mu(v) \rangle \in T'$ iff $\langle \ell, v \rangle \in T$.
3. Otherwise $\mu(T) = T$.

$\mu'(T_1 \wedge T_2)$ is defined by:

1. if T_1 and T_2 are record types, then $\mu'(T_1 \wedge T_2) = T_3$ such that
 - a) for any ℓ, v_1, v_2 , if $\langle \ell, v_1 \rangle \in T_1$ and $\langle \ell, v_2 \rangle \in T_2$, then
 - i. if v_1 and v_2 are $\langle \lambda u_1 : T'_1 \dots \lambda u_i : T'_i(\phi), \langle \pi_1 \dots \pi_i \rangle \rangle$ and $\langle \lambda u'_1 : T''_1 \dots \lambda u'_k : T''_k(\psi), \langle \pi'_1 \dots \pi'_k \rangle \rangle$ respectively, then $\langle \lambda u_1 : T'_1 \dots \lambda u_i : T'_i, \lambda u'_1 : T''_1 \dots \lambda u'_k : T''_k(\mu(\phi \wedge \psi)), \langle \pi_1 \dots \pi_i, \pi'_1 \dots \pi'_k \rangle \rangle \in T_3$
 - ii. if v_1 is $\langle \lambda u_1 : T'_1 \dots \lambda u_i : T'_i(\phi), \langle \pi_1 \dots \pi_i \rangle \rangle$ and v_2 is a type (i.e. not of the form $\langle f, \Pi \rangle$ for some function f and sequence of paths Π), then $\langle \lambda u_1 : T'_1 \dots \lambda u_i : T'_i(\mu(\phi \wedge v_2)), \langle \pi_1 \dots \pi_i \rangle \rangle \in T_3$
 - iii. if v_2 is $\langle \lambda u'_1 : T''_1 \dots \lambda u'_k : T''_k(\psi), \langle \pi'_1 \dots \pi'_k \rangle \rangle$ and v_1 is a type, then $\langle \lambda u'_1 : T''_1 \dots \lambda u'_k : T''_k(\mu(v_1 \wedge \psi)), \langle \pi'_1 \dots \pi'_k \rangle \rangle \in T_3$
 - iv. otherwise $\langle \ell, \mu(v_1 \wedge v_2) \rangle \in T_3$
 - b) for any ℓ, v_1 , if $\langle \ell, v_1 \rangle \in T_1$ and there is no v_2 such that $\langle \ell, v_2 \rangle \in T_2$, then $\langle \ell, v_1 \rangle \in T_3$
 - c) for any ℓ, v_2 , if $\langle \ell, v_2 \rangle \in T_2$ and there is no v_1 such that $\langle \ell, v_1 \rangle \in T_1$, then $\langle \ell, v_2 \rangle \in T_3$
2. if T_1 is *Rec* and T_2 is a record type, then $\mu'(T_1 \wedge T_2) = T_2$
3. if T_1 is a record type and T_2 is *Rec*, then $\mu'(T_1 \wedge T_2) = T_1$
4. If T_1 is $[T'_1]$ ($\{T'_1\}, \{\{T'_1\}\}$) and T_2 is $[T'_2]$ ($\{T'_2\}, \{\{T'_2\}\}$), then $\mu'(T_1 \wedge T_2) = [\mu(T'_1 \wedge T'_2)]$ ($\{\mu(T'_1 \wedge T'_2)\}, \{\{\mu(T'_1 \wedge T'_2)\}\}$)
5. Otherwise $\mu'(T_1 \wedge T_2) = T_1 \wedge T_2$

$T_1 \wedge T_2$ is used to represent $\mu(T_1 \wedge T_2)$. We call $T_1 \wedge T_2$ the *merge* of T_1 and T_2 .

The following two clauses could be added at the beginning of the definition of μ (after providing a characterization of the subtype relation, \sqsubseteq).

1. if for some $T_1, T_2, T = T_1 \wedge T_2$ and $T_1 \sqsubseteq T_2$ then $\mu(T) = T_1$
2. if for some $T_1, T_2, T = T_1 \wedge T_2$ and $T_2 \sqsubseteq T_1$ then $\mu(T) = T_2$

The current first clause would then hold in case neither of the conditions of these two clauses are met. The definition without these additional clauses only accounts for simplification of meets which have to do with merges of record types whereas the definition with the additional clauses would in addition have the effect, for example, that $\mu(T \wedge T_a) = T_a$ and $\mu(T_1 \wedge (T_1 \vee T_2)) = T_1$ (provided that we have an appropriate definition of \sqsubseteq) whereas the current definition without the additional clauses means that μ leaves these types unchanged.

We define also a notion of *asymmetric merge* of T_1 and T_2 which is defined by a function exactly like μ except that clause 5 of the definition of μ' is replaced by

5'. Otherwise $\mu'(T_1 \wedge T_2) = T_2$

We use $T_1 \boxed{\wedge} T_2$ to represent the asymmetric merge of T_1 and T_2 .

These definitions do not in general avoid the formation of ill-formed record types since they allow record types to be replaced with non-record types within a record type thus potentially removing paths that might be included in dependent type fields elsewhere in the resulting type. However, if merging is restricted to either two record types or two non-record types this problem should not occur since all paths from both types will be preserved. In the case of asymmetric merges we can allow the replacement of non-record types by record types without risk. Note that our definition of dependent record types in A.12 allows for dependencies to fields that have conflicting types. Such record types will be well-formed though will not have any witnesses.

Merging functions which return types $\lambda r : T_1 . T_2(r) \boxed{\wedge} \lambda r : T_3 . T_4(r)$ denotes the function $\lambda r : T_1 \boxed{\wedge} T_3 . T_2(r) \boxed{\wedge} T_4(r)$.

Constructing fixed point types for functions which return types If, for some type T_1 , $f : (T_1 \rightarrow \text{Type})$ then $\mathcal{F}(f)$ is a *fixed point type* for f , that is $a : \mathcal{F}(f)$ implies $a : f(a)$. \mathcal{F} is defined by

$$\mathcal{F}(\lambda r : T_1 . T_2(r)) = T_1 \boxed{\wedge} T'$$

where T' is like $T_2(r)$ except that any path $r.\pi$ is replaced by π .

Strictly speaking this definition is not quite correct since T' may not be a type because there may be a path occurring as an argument to a predicate which is not introduced in T' . A more correct, though less perspicuous, definition would be

$$\mathcal{F}(\lambda r : T_1 . T_2(r)) = [(\lambda r : T_1 . T_1 \boxed{\wedge} T_2(r))(r^*)]^{-r^*}$$

where r^* is a record of type T_1 such that there is no path occurring as an argument to a predicate in T_1 or $T_2(r)$ of the form $r^*.\pi$ for any π and $[T]^{-r}$ is the result of replacing any path of the form $r.\pi$, for any π , occurring as an argument to a predicate in T , with π .

This definition, however, fails to capture that \mathcal{F} must be defined as a partial function which is undefined on functions meeting a certain condition. This is taken account of in the following final definition:

\mathcal{F} is a partial function on functions such that if, for some type T_1 , $f : (T_1 \rightarrow \text{Type})$, $f = \lambda r : T_1 . T_2(r)$, g is the function

$$\lambda r : T_1 . T_1 \wedge T_2(r)$$

and for any $r_1, r_2 : T_1$

$$[g(r_1)]^{-r_1} = [g(r_2)]^{-r_2}$$

then if $r^* : T_1$,

$$\mathcal{F}(f) = [g(r^*)]^{-r^*}$$

For any function f not covered by the above, $\mathcal{F}(f)$ is undefined.

Let us take some concrete examples based on the discussion in Chapter 5, Section 5.5. Suppose that f is the function

$$\lambda r : [x:Ind] . [e : \text{dog}(r.x)]$$

Then g will be

$$\lambda r : [x:Ind] . [x:Ind] \wedge [e:\text{dog}(r.x)]$$

That is,

$$\lambda r : [x:Ind] . \left[\begin{array}{l} x : Ind \\ e : \text{dog}(r.x) \end{array} \right]$$

If we now compute $[g(r^*)]^{-r^*}$ for some $r^* : [x:Ind]$, that is we apply g to r^* and then remove r^* from all path-names that begin with it we will obtain

$$\left[\begin{array}{l} x : Ind \\ e : \text{dog}(x) \end{array} \right]$$

We would have obtained the same result no matter which record we chose to be r^* , since r only occurs in g at the head of path names. Now consider f to be a variant of what we propose to be the content of *temperature* in Chapter 5 (in fact, similar to a variant that we have proposed in previous work):

$$\lambda r:Rec . [e : temperature(r)]$$

Now g will be

$$\lambda r:Rec . Rec \wedge [e:temperature(r)]$$

That is,

$$\lambda r:Rec . [e : temperature(r)]$$

which happens to be identical with f . If we apply this to a record r^* we will obtain

$$[e : temperature(r^*)]$$

The result of removing all occurrences of r^* at the head of a path will be identical since r^* occurs here not as the head of a path but as an argument to a predicate. Consequently if we choose a different record for r^* the result will be a different type. Thus \mathcal{F} is not defined on this function. This is intuitively correct since defining fixed points for this function would involve a kind of non-well foundedness which we have not allowed in TTR.

The moral of this tale is that if you wish to define a dependent type (that is, a function returning a type), $\lambda r: T_1 . T_2(r)$, for which you will be able to compute a fixed point type, make sure that T_2 only depends on r in that r may be the head of path names in $T_2(r)$. Normally, you will also want to ensure that T_1 and T_2 do not share any labels in order to avoid unwanted clashes when T_1 and T_2 are merged.

A.14 Relabelling of record types

A.14.1 Some considerations

We extend the notion relabelling of records discussed in Appendix A.12 to record types.¹⁰ For non-dependent record types both flattening and relabelling can be regarded as exactly parallel to that of records since non-dependent record types are sets of ordered pairs of labels and types. They are just like records except that the second members of the pairs are types. However, in the case of dependent record types, flattening and relabelling introduces some complications which need to be taken account of. We shall start with an informal discussion of the complications.

¹⁰This was not made explicit in Cooper (2012b).

Labels in dependent fields First of all we need to note that labels not only occur as the first members of the ordered pairs that constitute the record type but also in the paths which occur in dependent fields. This introduces complications if we attempt to define an operation of flattening for record types. Also in relabelling we need to make sure that the relabelling applies to the paths in the dependent fields as well as the labels of the record type.

Consider the record type which in abbreviatory notation we write as:

$$\left[\begin{array}{c} x \\ \vdots \end{array} \left[\begin{array}{c} x : Ind \\ e : \text{dog}(x) \end{array} \right] \right]$$

It would be incorrect to represent the flattening of this type as:

$$\left[\begin{array}{c} x.x : Ind \\ x.e : \text{dog}(x) \end{array} \right]$$

The 'x' in 'dog(x)' does not represent a path in the flattened record type. Rather the flattened type should be, in abbreviatory notation:

$$\left[\begin{array}{c} x.x : Ind \\ x.e : \text{dog}(x.x) \end{array} \right]$$

This is a notational complication. Recall that 'dog(x)' in the first type uses 'x' to represent the path 'x.x', that is, the path to the 'x' in the immediate record containing 'dog(x)'.

Let us consider how this would work in the official notation with full paths represented in dependent fields. In order to define flattening and unflattening we would need to convert:

$$\left[\begin{array}{c} x \\ \vdots \end{array} \left[\begin{array}{c} x : Ind \\ e : \langle \lambda v:Ind . \text{dog}(v), \langle x.x \rangle \rangle \end{array} \right] \right]$$

to:

$$\left[\begin{array}{c} x.x : Ind \\ x.e : \langle \lambda v:Ind . \text{dog}(v), \langle x.x \rangle \rangle \end{array} \right]$$

and back again.

Relabelling can require adjustment of dependent fields A relabelling such as:

$$x.x \rightsquigarrow y$$

applied to the flattened type leads to a different problem involved in expansion. The relabelled flattened record would be:

$$\left[\begin{array}{ll} y & : \text{Ind} \\ x.e & : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle y \rangle \rangle \end{array} \right]$$

Expanding the relabelled flattened type in the way we did for records will result in:

$$\left[\begin{array}{ll} y & : \text{Ind} \\ x & : \left[\begin{array}{ll} e & : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle y \rangle \rangle \end{array} \right] \end{array} \right]$$

This is incorrect because the label ‘y’ is not present in

$$\left[\begin{array}{ll} e & : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle y \rangle \rangle \end{array} \right]$$

and this is not a type as require in Clause 7 of the definition of non-dependent record types on p. 285. Rather the correct result should be:

$$\left[\begin{array}{ll} y & : \text{Ind} \\ x & : \langle \lambda v:\text{Ind} . \left[\begin{array}{ll} e & : \text{dog}(v) \end{array} \right], \langle y \rangle \rangle \end{array} \right]$$

In abbreviatory notation this is:

$$\left[\begin{array}{ll} y & : \text{Ind} \\ x & : \left[\begin{array}{ll} e & : \text{dog}(\uparrow y) \end{array} \right] \end{array} \right]$$

This, then, is a problem for defining relabelling of record types in terms of a notion of flattening and unflattening as we did for records. This problem is not insurmountable given the techniques that we will develop below. However, there is another problem which seems to strongly suggest that we should not use flattening in the same way that we did for records.

Dependencies involving paths which are not total In the examples above the dependencies involve objects at the end of paths. It is also possible to have dependencies on records corresponding to the value of a partial path. Consider the following example:

$$\left[\begin{array}{l} s : \left[\begin{array}{l} x : Ind \\ c : \text{dog}(x) \\ e : \text{run}(x) \end{array} \right] \\ e : \text{see}(\text{sam}, s) \end{array} \right]$$

which might be a way of representing the type which is the content of *Sam saw a dog run*. The official notation for this type (with unabbreviated paths) is:

$$\left[\begin{array}{l} s : \left[\begin{array}{l} x : Ind \\ c : \langle \lambda v:Ind . \text{dog}(v), \langle s.x \rangle \rangle \\ e : \langle \lambda v:Ind . \text{run}(v), \langle s.x \rangle \rangle \end{array} \right] \\ e : \langle \lambda v:Rec . \text{see}(\text{sam}, v), \langle s \rangle \rangle \end{array} \right]$$

We might think that the flattening of this type is:

$$\left[\begin{array}{l} s.x : Ind \\ s.c : \langle \lambda v:Ind . \text{dog}(v), \langle s.x \rangle \rangle \\ s.e : \langle \lambda v:Ind . \text{run}(v), \langle s.x \rangle \rangle \\ e : \langle \lambda v:Rec . \text{see}(\text{sam}, v), \langle s \rangle \rangle \end{array} \right]$$

But this is, of course, not correct since the path ‘s’ in the ‘e’-field is not a path in the flattened type. Our strategy will be not to define relabelling in terms of flattening as we did for records, but to define relabelling of types in terms of direct replacement of paths in unflattened types. In order to do this we will, however, find it convenient to use a “flattened” representation of the unflattened type which explicitly lists its paths.

A.14.2 Relabellings of record types

Here we will consider how, for a given record type, T , an η -relabelling of T , $[T]_\eta$ can be defined.

We start by defining the notion of paths for record types. This can be defined in terms of the paths which are common to all witnesses for the types. If T is a record type, then the *set of paths of T* , $\text{paths}(T)$, is

$$\text{cpaths}(\llbracket T \rrbracket)$$

where cpaths is as defined on p. 280 and the *set of total paths of T* , $\text{tpaths}(T)$, is

$$\text{ctpaths}([\sim T])$$

where ctpaths is as defined on p. 280.

It may look from these definitions as if you have to compute all the witnesses for a record type before you can compute its paths, but in fact this is not true. For any type, we can compute what paths are witness of that type must have. For example, any witness of the type

$$\left[\begin{array}{lcl} x & : & \text{Ind} \\ c & : & \langle \lambda v:\text{Ind} . [\text{e}:\text{dog}(v)], \langle x \rangle \rangle \end{array} \right]$$

must be a record of the form

$$\left[\begin{array}{lcl} x & = & a \\ c & = & \left[\begin{array}{lcl} e & = & s \\ \dots & & \end{array} \right] \\ \dots & & \end{array} \right]$$

where:

$$\begin{array}{l} a : \text{Ind} \\ s : \text{dog}(a) \end{array}$$

Hence, any record of this type must have ‘ x ’ and ‘ $c.e$ ’ among its total paths and these will be the only total paths that are common to all witnesses for the type.

In a similar way, we can say that tree relations holding in all witnesses of a type hold for the types (see p. 278). Thus, for example, we can say that ‘ c ’ immediately dominates ‘ e ’ in the type in virtue of the fact that the relation holds in all witnesses of the type, despite the fact the ‘ e ’ is embedded in a function below ‘ c ’ in the type.

It will be convenient to define the computation of relabellings of types in terms of a notation which gives unique identifiers to the occurrences of labels and which eschews representing the functions in dependent fields explicitly, rather representing the dependencies by using the unique identifiers. To increase readability, we will only represent the unique identifiers on those labels which are referenced in some dependent field. As an example, consider the following types, presented in official notation.

$$\left[\begin{array}{lcl} x & : & Ind \\ c & : & \langle \lambda v:Ind . \text{dog}(v), \langle x \rangle \rangle \\ e & : & \langle \lambda v:Ind . \text{appear}([e:\text{approach}(v)]), \langle x \rangle \rangle \end{array} \right]$$

$$\left[\begin{array}{lcl} x & : & Ind \\ c & : & \langle \lambda v:Ind . [e:\text{dog}(v)], \langle x \rangle \rangle \end{array} \right]$$

In the unique identifier notation these will be:

$$\left[\begin{array}{lcl} x\{0\} & : & Ind \\ c & : & \text{dog}(\{0\}:Ind) \\ e & : & \text{appear}([e:\text{approach}(\{0\}:Ind)]) \end{array} \right]$$

$$\left[\begin{array}{lcl} x\{0\} & : & Ind \\ c & : & [e:\text{dog}(\{0\}:Ind)] \end{array} \right]$$

The unique identifier notation can be obtained from the official notation by

1. adding a unique identifier $\{i\}$ (where i is a natural number) to each occurrence of a label in the type, other than in the paths following a function in dependent fields. In practice we only add identifiers to those labels which are addressed in some dependent field.
2. replacing any path in a dependent field with the unique identifier associated with the final label occurrence on the path which it references.
3. replacing any pair in a dependent field of the form

$$\langle \lambda v_1 : T_1 \dots \lambda v_n : T_n . \varphi[v_1, \dots, v_n], \langle \{i_1\}, \dots, \{i_n\} \rangle \rangle$$

with

$$\varphi[\{i_1\} : T_1, \dots, \{i_n\} : T_n]$$

The final step involves a variant of β -conversion. It is important to represent the type restriction associated with the unique identifier by the domain type of the function since the domain type of the function may be distinct from the type associated with the identifier in the field labelled by the identifier.

It is important to note that we are manipulating the representations of types here, not the types themselves. The unique identifier notation does not introduce additional types.

We also need to show how to convert back from unique identifier notation for a type, T^* , to official notation. We can do this by carrying out the following in order:

1. replace all the dependencies on unique identifiers with appropriate pairs of functions and sequences of paths
2. remove the unique identifiers

We replace all the dependencies on unique identifiers by carrying out the following for each occurrence of a label, ℓ , in T^* associated with an identifier, $\{i\}$, which is referenced elsewhere in T^* :

1. for each sister, ℓ' , of $\ell\{i\}$, if ℓ' immediately dominates a subtree, τ , containing $\{i\}:T$ for some type, T , then:
 - if τ is not an ordered pair consisting of a function and a sequence of paths:
 - replace τ with $\langle \lambda v_i:T . \tau[\{i\}:T \rightsquigarrow v_i], \langle \ell \rangle \rangle$
 - otherwise, if τ is $\langle f, \langle \pi_1, \dots, \pi_n \rangle \rangle$:
 - replace τ with $\lambda v_i:T . f[\{i\}:T \rightsquigarrow v_i], \langle \ell, \pi_1, \dots, \pi_n \rangle \rangle$
2. Let $\ell_0 \dots \ell_n.\ell\{i\}$ be the path ending in $\ell\{i\}$ in T^* . For each ℓ_k in order in $\langle \ell_n, \dots, \ell_0 \rangle$, where π^* is the path starting with ℓ_k and ending with $\ell\{i\}$, and for each sister, ℓ' , of ℓ_k , if ℓ' immediately dominates a subtree, τ , containing $\{i\}:T$ for some type, T , then:
 - if τ is not an ordered pair consisting of a function and a sequence of paths:
 - replace τ with $\langle \lambda v_i:T . \tau[\{i\}:T \rightsquigarrow v_i], \langle \pi^* \rangle \rangle$
 - otherwise, if τ is $\langle f, \langle \pi_1, \dots, \pi_n \rangle \rangle$:
 - replace τ with $\langle \lambda v_i:T . f[\{i\}:T \rightsquigarrow v_i], \langle \pi^*, \pi_1, \dots, \pi_n \rangle \rangle$

Consider as an example the type which in abbreviatory notation is written (with abbreviatory representation of the path ‘ $\text{id}_0.x$ ’ in the ‘ $\text{id}_0.e$ ’-field) as:

$$\left[\begin{array}{l} \text{id}_0 \\ \text{id}_1 \end{array} : \left[\begin{array}{l} x : \text{Ind} \\ e : \text{dog}(x) \\ e : \text{run}(\uparrow \text{id}_0.x) \end{array} \right] \right]$$


This is a type which relates to our discussion of long term memory in Chapters 4–6. In official notation this type is represented as:

$$\left[\begin{array}{l} \text{id}_0 \\ \text{id}_1 \end{array} : \left[\begin{array}{l} x : \text{Ind} \\ e : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle x \rangle \rangle \\ e : \langle \lambda v:\text{Ind} . \left[e : \text{run}(v) \right], \langle \text{id}_0.x \rangle \rangle \end{array} \right] \right]$$


In the unique identifier notation this would be

$$\left[\begin{array}{l} \text{id}_0 : \left[\begin{array}{l} x\{0\} : \text{Ind} \\ e : \text{dog}(\{0\}:\text{Ind}) \end{array} \right] \\ \text{id}_1 : \left[\begin{array}{l} e : \text{run}(\{0\}:\text{Ind}) \end{array} \right] \end{array} \right]$$


This in turn can be represented in tree-notation to make the sisterhood relations clear.




To convert this back to official notation, we first look at the sister of ‘ $x\{0\}$ ’, that is, ‘ e ’, and replace the subtree it immediately dominates with the corresponding dependent field.



Next we go up one node from ‘ $x\{0\}$ ’ to ‘ id_0 ’ and look at its sister, ‘ id_1 ’. We then form an appropriate dependent field from its immediate subtree.



Finally, we remove the unique identifier ' $\{0\}$ '.



This is now the tree variant of the official notation for the type.

The flattened variant of the unique identifier notation,

$$\left[\begin{array}{l} \text{id}_0 \\ \text{id}_1 \end{array} : \left[\begin{array}{ll} x\{0\} & : \text{Ind} \\ e & : \text{dog}(\{0\}:Ind) \\ e & : \text{run}(\{0\}:Ind) \end{array} \right] \right]$$

is:

$$\left[\begin{array}{ll} \text{id}_0.x\{0\} & : \text{Ind} \\ \text{id}_0.e & : \text{dog}(\{0\}:Ind) \\ \text{id}_1.e & : \text{run}(\{0\}:Ind) \end{array} \right]$$

where we have made the paths in the original type explicit. Note, that this again is not introducing a new type. It is merely a variant notation for the original unflattened type. If we think of the notation itself as a record, that is, a set of ordered pairs, we can obtain the flattened representation of the type by applying the same procedure as we used for records. Thus the unflattened representation corresponds to the set

$$\{ \langle \text{id}_0, \{ \langle x\{0\}, Ind \rangle, \langle e, \text{dog}(\{0\}:Ind) \rangle \} \rangle, \langle \text{id}_1, \{ \langle e, \text{run}(\{0\}:Ind) \rangle \} \rangle \}$$

We can now apply the flattening algorithm defined on p. 279 to this set, yielding the set:

$$\begin{aligned} & \{ \langle \text{id}_0.x\{0\}, \text{Ind} \rangle \\ & \quad \langle \text{id}_0.e, \text{dog}(\{0\}:\text{Ind}) \rangle \\ & \quad \langle \text{id}_1.e, \text{run}(\{0\}:\text{Ind}) \rangle \end{aligned}$$

Thus we are not flattening the type, but we *are* flattening the *representation* of the type which can itself be conceived of as a record.

Now consider the relabelling:

$$\begin{aligned} \text{id}_0.x &\rightsquigarrow x \\ \text{id}_0.e &\rightsquigarrow e_1 \\ \text{id}_1.e &\rightsquigarrow e_2 \end{aligned}$$

This relabelling can be decorated with the unique identifiers corresponding to the flattened representation of the type. On the left hand side of the relabelling we copy whatever unique identifiers occur associated with the label in the type. We then place a copy of the unique identifier occurring on the left hand side also on the right hand side. Thus the decorated version of this relabelling corresponding to this type will be:

$$\begin{aligned} \text{id}_0.x\{0\} &\rightsquigarrow x\{0\} \\ \text{id}_0.e &\rightsquigarrow e_1 \\ \text{id}_1.e &\rightsquigarrow e_2 \end{aligned}$$

Now, again considering the flattened representation of the type as a record, we can apply the relabelling to the representation in the manner characterized on pp. 281ff. Thus for the above type this relabelling will yield:

$$\left[\begin{array}{ll} x\{0\} & : \text{Ind} \\ e_1 & : \text{dog}(\{0\}:\text{Ind}) \\ e_2 & : \text{run}(\{0\}:\text{Ind}) \end{array} \right]$$

which in official notation is

$$\left[\begin{array}{ll} x & : \text{Ind} \\ e_1 & : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle x \rangle \rangle \\ e_2 & : \langle \lambda v:\text{Ind} . \text{run}(v), \langle x \rangle \rangle \end{array} \right]$$

In this example the unique identifier occurred at the end of a path, but this is not always the case. Consider the example discussed on p. 297. In unique identifier notation this is

$$\left[\begin{array}{lcl} s\{1\} & : & \left[\begin{array}{lcl} x\{0\} & : & Ind \\ c & : & \text{dog}(\{0\}:Ind) \\ e & : & \text{run}(\{0\}:Ind) \end{array} \right] \\ e & : & \text{see}(\text{sam}, \{1\}:Rec) \end{array} \right]$$

The flattened variant of the unique identifier notation for this type is below, where the unique identifier ‘{1}’ occurs on three copies of the label ‘s’:

$$\left[\begin{array}{lcl} s\{1\}.x\{0\} & : & Ind \\ s\{1\}.c & : & \text{dog}(\{0\}:Ind) \\ s\{1\}.e & : & \text{run}(\{0\}:Ind) \\ e & : & \text{see}(\text{sam}, \{1\}:Rec) \end{array} \right]$$

We require that any relabelling for a type must preserve partial paths terminated with a label associated with a unique identifier in the sense that while the partial path itself may be changed, it must be changed to the same partial path in each path where the unique identifier occurs. Thus the following would be an admissible relabelling for this type:

$$\begin{aligned} s.x &\rightsquigarrow id_0.id_1.x \\ s.c &\rightsquigarrow id_0.id_1.e \\ s.e &\rightsquigarrow id_0.e \\ e &\rightsquigarrow id_2.e \end{aligned}$$

We can decorate this relabelling with unique identifiers corresponding to the type:

$$\begin{aligned} s\{1\}.x\{0\} &\rightsquigarrow id_0\{1\}.id_1.x\{0\} \\ s\{1\}.c &\rightsquigarrow id_0\{1\}.id_1.e \\ s\{1\}.e &\rightsquigarrow id_0\{1\}.e \\ e &\rightsquigarrow id_2.e \end{aligned}$$

How do we know where to place the identifier ‘{1}’ on the right-hand side of the annotated relabelling? It goes after the longest partial path which the three paths corresponding to left-hand side paths containing {1} have in common. So, for example, if the relabelling had been:

$$\begin{aligned}
s.x &\rightsquigarrow \text{id}_0.\text{id}_1.x \\
s.c &\rightsquigarrow \text{id}_0.\text{id}_1.e_1 \\
s.e &\rightsquigarrow \text{id}_0.\text{id}_1.e_2 \\
e &\rightsquigarrow \text{id}_2.e
\end{aligned}$$

the decorated version would be:

$$\begin{aligned}
s\{1\}.x\{0\} &\rightsquigarrow \text{id}_0.\text{id}_1\{1\}.x\{0\} \\
s\{1\}.c &\rightsquigarrow \text{id}_0.\text{id}_1\{1\}.e_1 \\
s\{1\}.e &\rightsquigarrow \text{id}_0.\text{id}_1\{1\}.e_2 \\
e &\rightsquigarrow \text{id}_2.e
\end{aligned}$$

According to the first relabelling the relabelled type would be:

$$\left[\begin{array}{ll} \text{id}_0\{1\}.\text{id}_1.x\{0\} & : \text{Ind} \\ \text{id}_0\{1\}.\text{id}_1.e & : \text{dog}(\{0\}:\text{Ind}) \\ \text{id}_0\{1\}.e & : \text{run}(\{0\}:\text{Ind}) \\ \text{id}_2.e & : \text{see}(\text{sam},\{1\}:\text{Rec}) \end{array} \right]$$

In standard notation with unabbreviated paths this type is represented as

$$\left[\begin{array}{l} \text{id}_0 : \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle \text{id}_0.\text{id}_1.x \rangle \rangle \end{array} \right] \\ e : \langle \lambda v:\text{Ind} . \text{run}(v), \langle \text{id}_0.\text{id}_1.x \rangle \rangle \end{array} \right] \\ \text{id}_2 : \langle \lambda v:\text{Rec} . [\text{e}:\text{see}(\text{sam},v)], \langle \text{id}_0 \rangle \rangle \end{array} \right]$$

Using abbreviated paths, this is

$$\left[\begin{array}{l} \text{id}_0 : \left[\begin{array}{l} \text{id}_1 : \left[\begin{array}{l} x : \text{Ind} \\ e : \langle \lambda v:\text{Ind} . \text{dog}(v), \langle x \rangle \rangle \end{array} \right] \\ e : \langle \lambda v:\text{Ind} . \text{run}(v), \langle \text{id}_1.x \rangle \rangle \end{array} \right] \\ \text{id}_2 : \langle \lambda v:\text{Rec} . [\text{e}:\text{see}(\text{sam},v)], \langle \text{id}_0 \rangle \rangle \end{array} \right]$$

Suppose we had tried to use the following relabelling instead.

$$\begin{aligned}
s.x &\rightsquigarrow \text{id}_0.\text{id}_1.x \\
s.c &\rightsquigarrow \text{id}_0.\text{id}_1.e
\end{aligned}$$

$$\begin{aligned} s.e &\rightsquigarrow \text{id}_1.e \\ e &\rightsquigarrow \text{id}_2.e \end{aligned}$$

This relabelling cannot be decorated according to the unique identifiers of the type because the three occurrences of ‘s’ do not have a corresponding partial path on the right-hand side. This relabelling is thus inadmissible for this type. We thus define admissibility of a relabelling for a type as follows:

A relabelling, η , is *admissible* for a record type, T , just in case η can be decorated according to the flattened unique identifier representation of T .

We can now define the notion of relabelling for types using the notions of flattening and relabelling we used for records:

If T is a record type with unique identifier representation, r , and η is an admissible relabelling for T , then the *result of relabelling T with η* , $[T]_\eta$, is that T' whose unique identifier representation is $[r]_\eta$.

A.15 Using records to restrict and specify record types

(These definitions were not included in Cooper, 2012b.)

If T is a type and r is a record, then $T \upharpoonright r$ is a type. $a : T \upharpoonright r$ iff $a \in r$ (see Appendix A.12) and $a : T$.

If T is a record type and r is a record, then $T \parallel r$, the *restriction* of T by r is the result of replacing each field, $\langle \ell, T' \rangle$, in T such that ℓ is a label in r , with

1. $\langle \ell, T' \upharpoonright r.\ell \rangle$ ¹¹, if T' is a type
2. $\langle \ell, \langle f', \Pi \rangle \rangle$, if $T' = \langle f, \Pi \rangle$ where f is a function and Π is a sequence of paths of length n and for any a_1, \dots, a_n , $f'(a_1) \dots (a_n)$ is defined iff $f(a_1) \dots (a_n)$ is defined and $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n) \upharpoonright r.\ell$

A variant of this notion of restriction is default restriction which will only require restriction of fields which are not already restricted. If T is a record type and r is a record, then $T \diagup r$, the *default restriction* of T by r is the result of replacing each field, $\langle \ell, T' \rangle$, in T such that ℓ is a label in r , with

¹¹That is, in an adaptation of our graphical notation for manifest fields, $[\ell:T']$ is replaced by $[\ell \in r.\ell:T']$

1. $\langle \ell, T' \upharpoonright r.\ell \rangle$, if T' is a type but not a restricted type. If T' is a singleton type then $\langle \ell, T' \rangle$ is replaced by itself.
2. $\langle \ell, \langle f', \Pi \rangle \rangle$, if $T' = \langle f, \Pi \rangle$ where f is a function and Π is a sequence of paths of length n and for any a_1, \dots, a_n , $f'(a_1) \dots (a_n)$ is defined iff $f(a_1) \dots (a_n)$ is defined and $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n) \upharpoonright r.\ell$ if $f(a_1) \dots (a_n)$ is not a restricted type. Otherwise, $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n)$.

If T is a record type and r is a record, then $T \parallel r$, the *specification* (or *anchoring*) of T by r ¹² is the result of replacing each field, $\langle \ell, T' \rangle$, in T such that ℓ is a label in r , with

1. $\langle \ell, T'_{r.\ell} \rangle$ ¹³, if T' is a type
2. $\langle \ell, \langle f', \Pi \rangle \rangle$, if $T' = \langle f, \Pi \rangle$ where f is a function and Π is a sequence of paths of length n and for any a_1, \dots, a_n , $f'(a_1) \dots (a_n)$ is defined iff $f(a_1) \dots (a_n)$ is defined and $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n)_{r.\ell}$

A variant of this notion of specification is default specification which will only require specification of fields which are not already specified. If T is a record type and r is a record, then $T // r$, the *default specification* (or *default anchoring*) of T by r is the result of replacing each field, $\langle \ell, T' \rangle$, in T such that ℓ is a label in r , with

1. $\langle \ell, T'_{r.\ell} \rangle$, if T' is a type but not a singleton type.¹⁴ If T' is a singleton type then $\langle \ell, T' \rangle$ is replaced by itself.
2. $\langle \ell, \langle f', \Pi \rangle \rangle$, if $T' = \langle f, \Pi \rangle$ where f is a function and Π is a sequence of paths of length n and for any a_1, \dots, a_n , $f'(a_1) \dots (a_n)$ is defined iff $f(a_1) \dots (a_n)$ is defined and $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n)_{r.\ell}$ if $f(a_1) \dots (a_n)$ is not a singleton type. Otherwise, $f'(a_1) \dots (a_n) = f(a_1) \dots (a_n)$.

Types can also be specified by records which have different labels to the type by using a relabelling. Thus $T \parallel_\eta r$ is the result of replacing each field in T , $\langle \ell, T' \rangle$, such that $\eta(\ell)$ is a label in r and $r.\eta(\ell) : T'$, with a manifest field $\langle \ell, T'_{r.\eta(\ell)} \rangle$. More exactly we will define $T \parallel_\eta r$ in terms of flattening, relabelling and specification by a record:

$$T \parallel_\eta r = \varphi^-([\varphi(T)]_\eta \parallel \varphi(r)]_{\eta^-}$$

¹² r could also be referred to as a *partial specifier* or *partial assignment* for T

¹³That is, in our standard graphical notation, $[\ell:T']$ is replaced by $[\ell=r.\ell:T']$

¹⁴That is, $\langle \ell, T' \rangle$ is not already a manifest field.

Here are two examples: suppose that η is a function with domain $\{\ell_1, \ell_2\}$ such that $\eta(\ell_1) = \ell_3$ and $\eta(\ell_2) = \ell_4$. Then:

$$\left[\begin{array}{c} \ell_1 \\ \ell_2 \end{array} : \begin{array}{c} T_1 \\ T_2 \end{array} \right] \parallel_{\eta} \left[\begin{array}{c} \ell_3 = a \\ \ell_4 = b \end{array} \right] = \left[\begin{array}{c} \ell_1=a \\ \ell_2=b \end{array} : \begin{array}{c} T_1 \\ T_2 \end{array} \right]$$

Suppose now that η is a function with domain $\{\ell_5.\ell_1, \ell_5.\ell_2, \ell_6\}$ (where $\ell_5.\ell_1$ and $\ell_5.\ell_2$ are complex labels) such that $\eta(\ell_5.\ell_1) = \ell_7.\ell_3$, $\eta(\ell_5.\ell_2) = \ell_8.\ell_4$ and $\eta(\ell_6) = \ell_6$. Then:

$$\left[\begin{array}{c} \ell_5 \\ \ell_6 \end{array} : \left[\begin{array}{c} \ell_1 : T_1 \\ \ell_2 : T_2 \end{array} \right] \right] \parallel_{\eta} \left[\begin{array}{c} \ell_7 = \left[\begin{array}{c} \ell_3 = a \\ \ell_4 = b \end{array} \right] \\ \ell_8 = \left[\begin{array}{c} \ell_4 = b \\ \ell_9 = c \end{array} \right] \end{array} \right] =$$

$$\left[\begin{array}{c} \ell_5 \\ \ell_6 \end{array} : \left[\begin{array}{c} \ell_1=a : T_1 \\ \ell_2=b : T_2 \end{array} \right] \right]$$

A.16 Strings and regular types

A *string algebra* over a set of objects O is a pair $\langle S, \frown \rangle$ where:

1. S is the closure of $O \cup \{\varepsilon\}$ (ε is the empty string) under the binary operation ‘ \frown ’ (“concatenation”)
2. for any s in S , $\varepsilon \frown s = s \frown \varepsilon = s$
3. for any s_1, s_2, s_3 in S , $(s_1 \frown s_2) \frown s_3 = s_1 \frown (s_2 \frown s_3)$. For this reason we normally write $s_1 \frown s_2 \frown s_3$ or more simply $s_1 s_2 s_3$.

The objects in S are called strings. Strings have length. ε has length 0. If s is a string in S with length n and a is an object in O then $s \frown a$ has length $n + 1$. We use $s[n]$ to represent the n th element of string s .

In TTR strings are records¹⁵ with fields labelled by a distinguished ordered countably infinite set of labels (corresponding to the natural numbers): t_0, t_1, \dots . ε is the empty record, that is the empty set. This has length 0. (Recall that records are sets of ordered pairs.) A string with one element a (of some type) is the record $[t_0=a]$. If s is a string whose highest label in the order is t_n ($n \geq 0$) and a is an object of some type then $s \frown a$ is $s \cup \{\langle t_{n+1}, a \rangle\}$. If s is a string whose highest label in the order is t_n then the length of s is $n + 1$. $s[n]$ is defined to be $s.t_n$. Concatenation (‘ \frown ’)

¹⁵This is new since Cooper (2012b).

can be extended to include concatenation of strings of arbitrary length. If s is a string and s_n is a string of length n , then $s \frown s_n$ is $s \frown s_n[0] \frown \dots \frown s_n[n-1]$.

If s is a string of length n of records such that for each i , $0 \leq i < n$, $s[i].\pi$ is a defined path, $\text{concat}_{0 \leq i < n}(s[i].\pi)$ denotes $s[0].\pi \frown \dots \frown s[n-1].\pi$. We use $\text{concat}_i(s[i].\pi)$ to represent $\text{concat}_{0 \leq i < \text{length}(s)}(s[i].\pi)$.

A system of complex types $\mathbf{TYPE}_S = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ with record types based on $\langle L, \mathbf{RType} \rangle$ has string types if

1. for each natural number i , $t_i \in L$
2. $String \in \mathbf{BType}$
3. $\emptyset \in String$
4. if $T \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_S} T$ then $[t_0=a] : String$
5. if $s :_{\mathbf{TYPE}_S} String$, t_n is a label in s such that there is no $i > n$ where t_i is a label in s , $T \in \mathbf{Type}$ and $a :_{\mathbf{TYPE}_S} T$ then $s \cup \{ \langle t_{n+1}, a \rangle \} :_{\mathbf{TYPE}_S} String$
6. Nothing is of type $String$ except as required above.

We can define types whose elements are strings. Such types correspond to regular expressions and we will call them *regular types*. Here we will define just two kinds of such types: concatenation types and Kleene-+ types.

A system of complex types with string types $\mathbf{TYPE}_S = \langle \mathbf{Type}, \mathbf{BType}, \langle \mathbf{PType}, \mathbf{Pred}, \mathbf{ArgIndices}, \mathbf{Arity} \rangle, \langle A, F \rangle \rangle$ has concatenation types if

1. a) for any $T_1, T_2 \in \mathbf{Type}$, $(T_1 \frown T_2) \in \mathbf{Type}$
b) for any $T_1, T_2, T_3 \in \mathbf{Type}$, $(T_1 \frown T_2) \frown T_3 = T_1 \frown (T_2 \frown T_3)$ ¹⁶
2. $a :_{\mathbf{TYPE}_S} T_1 \frown T_2$ iff $a = x \frown y$, $x :_{\mathbf{TYPE}_S} T_1$ and $y :_{\mathbf{TYPE}_S} T_2$

\mathbf{TYPE}_S has Kleene-+ types if

1. for any $T \in \mathbf{Type}$, $T^+ \in \mathbf{Type}$
2. $a :_{\mathbf{TYPE}_S} T^+$ iff $a = x_1 \frown \dots \frown x_n$, $n > 0$ and for i , $1 \leq i \leq n$, $x_i :_{\mathbf{TYPE}_S} T$

¹⁶This has been added to the definition in Cooper (2012b) to make associativity explicit.

\mathbf{TYPE}_S has Kleene-* types if

1. for any $T \in \mathbf{Type}$, $T^* \in \mathbf{Type}$
2. $a : \mathbf{TYPE}_S T^*$ iff $a = x_1 \hat{\ } \dots \hat{\ } x_n$, $n \geq 0$ and for i , $1 \leq i \leq n$, $x_i : \mathbf{TYPE}_S T$

Note that this definition distinguishes an object a and the unit string consisting of a , that is, $[t_0=a]$. We will use \hat{a} to represent this.

Strings are used standardly in formal language theory where strings of symbols or strings of words are normally considered. Following important insights by Tim Fernando Fernando (2004, 2006, 2008, 2009) we shall be concerned rather with strings of events. We use informal notations like ‘ “sam” ’ and ‘ “ran” ’ to represent phonological types of speech events (utterances of *Sam* and *ran*). Thus ‘ “sam” $\hat{\ }$ “ran” ’ is the type of speech events which are concatenations of an utterance of *Sam* and an utterance of *ran*.

Predicates which relate strings

We introduce a number of distinguished predicates which are used to relate strings. The following predicates all have arity $[String, String]$: *init*, *final*, *final_align*

init “ s_1 is an initial substring of s_2 ”

If s_1 is a string of length n and s_2 is a string of any length, then $s : \text{init}(s_1, s_2)$ iff the length of s_2 is greater than or equal to n and for each i , $0 \leq i < n$, $s_1[i] = s_2[i]$ and $s = s_2$.

final “ s_1 is a final substring of s_2 ”

If s_1 is a string of length n and s_2 is a string of length m , then $s : \text{final}(s_1, s_2)$ iff m is greater than or equal to n and for each i , $0 \leq i < n$, $s_1[i] = s_2[(m - n) + i]$ and $s = s_2$.

final_align “ s_1 is aligned with a final substring of s_2 ”

If $s_1 : \text{Rec}^+$ is a string of length n and $s_2 : \text{Rec}^+$ is a string of length m , then $s : \text{final_align}(s_1, s_2)$ iff

1. m is greater than or equal to n
2. s is a string of length m
3. for each i , $0 \leq i < n$,
 - a) $s[(m - n) + i] : \begin{bmatrix} \mathbf{e}_1 : \text{Rec} \\ \mathbf{e}_2 : \text{Rec} \end{bmatrix}$
 - b) $s[(m - n) + i].\mathbf{e}_1 = s_1[i]$
 - c) $s[(m - n) + i].\mathbf{e}_2 = s_2[(m - n) + i]$
4. otherwise for each i , $0 \leq i < m$, $s[i] = s_2[i]$

Appendix B

Grammar rules

B.1 Universal resources

B.1.1 Frames

AmbTempFrame (Chapter 5)

$$\left[\begin{array}{ll} x & : \textit{Real} \\ \textit{loc} & : \textit{Loc} \\ e & : \textit{temp}(\textit{loc}, x) \end{array} \right]$$

AgeFrame (Chapter 5)

$$\left[\begin{array}{l} x:\textit{Ind} \\ \textit{age}:\textit{Real} \\ c_{\textit{age}}:\textit{age_of}(x,\textit{age}) \end{array} \right]$$

DogFrame (Chapter 5)

$$\left[\begin{array}{l} x:\textit{Ind} \\ e:\textit{dog}(x) \\ \textit{age}:\textit{Real} \\ c_{\textit{age}}:\textit{age_of}(x,\textit{age}) \end{array} \right]$$

B.1.2 Scales

$\zeta_{\textit{temp}}$ (Chapter 5)

$\lambda r:\textit{AmbTempFrame} . r.x$

ζ_{age} (Chapter 5)

$\lambda r:\text{AgeFrame} . r.\text{age}$

B.1.3 Signs

Sign (Chapter 2)

$$\left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{cont} & : \text{Cont} \end{array} \right]$$

Sign (Chapter 3)

a recursive type

$$\sigma : \text{Sign} \text{ iff } \sigma : \left[\begin{array}{ll} \text{s-event} & : \text{SEvent} \\ \text{syn} & : \text{Syn} \\ \text{cont} & : \text{Cont} \end{array} \right]$$

SEvent (Chapter 2)

$$\left[\begin{array}{ll} \text{e-loc} & : \text{Loc} \\ \text{sp} & : \text{Ind} \\ \text{au} & : \text{Ind} \\ \text{e} & : \text{Phon} \\ \text{c}_{\text{loc}} & : \text{loc}(\text{e}, \text{e-loc}) \\ \text{c}_{\text{sp}} & : \text{speaker}(\text{e}, \text{sp}) \\ \text{c}_{\text{au}} & : \text{audience}(\text{e}, \text{au}) \end{array} \right]$$

Phon (Chapter 2)

Word^+

Cont (Chapter 2)

RecType

Cont (Chapter 3)

$\text{RecType} \vee \text{Ppty} \vee \text{Quant} \vee (\text{Ppty} \rightarrow \text{Quant})$

Ppty (Chapter 3)

$([x:Ind] \rightarrow RecType)$

Ppty (Chapter 5)

$$\left[\begin{array}{ll} \text{bg} & : \text{Type} \\ \text{fg} & : ([x:\text{bg}] \rightarrow RecType) \end{array} \right]$$

Ppty(T), where T is a type (Chapter 5)

$$\left[\begin{array}{ll} \text{bg}=T & : \text{Type} \\ \text{fg} & : ([x:\text{bg}] \rightarrow RecType) \end{array} \right]$$

PPpty (Chapter 4)

$$\left[\begin{array}{ll} \text{bg} & : RecType \\ \text{fg} & : (\text{bg} \rightarrow Ppty) \end{array} \right]$$

Abbreviations for properties (Chapter 5)

If p is a predicate with arity $\langle T \rangle$ then p' represents the property

$$\left[\begin{array}{ll} \text{bg} & = T \\ \text{fg} & = \lambda r: [x:T] . [e : p(r.x)] \end{array} \right]$$

If P is the property

$$\left[\begin{array}{ll} \text{bg} & = T_1 \\ \text{fg} & = \lambda r: [x:T_1] . T_2(r) \end{array} \right]$$

then $P \upharpoonright s$ represents

$$\left[\begin{array}{ll} \text{bg} & = T_1 \\ \text{fg} & = \lambda r: [x:T_1] . T_2(r) // [e=s] \end{array} \right]$$

This uses the definition of $T // r$ in Appendix A.15.

Quant (Chapter 3)

$(Ppty \rightarrow RecType)$

$PQuant$ (Chapter 4)

$$\left[\begin{array}{ll} bg & : RecType \\ fg & : (bg \rightarrow Quant) \end{array} \right]$$

Syn (Chapter 3)

$$\left[\begin{array}{ll} cat & : Cat \\ daughters & : Sign^* \end{array} \right]$$

Cat (Chapter 3)

$s, np, det, n, v, vp : Cat$

Category sign types:

S (Chapter 3)

$Sign \wedge [syn: [cat=s:Cat]]$

NP (Chapter 3)

$Sign \wedge [syn: [cat=np:Cat]]$

Det (Chapter 3)

$Sign \wedge [syn: [cat=det:Cat]]$

N (Chapter 3)

$Sign \wedge [syn: [cat=n:Cat]]$

V (Chapter 3)

$Sign \wedge [syn: [cat=v:Cat]]$

VP (Chapter 3)

$Sign \wedge [syn: [cat=vp:Cat]]$

$NoDaughters$ (Chapter 3)

$[syn: [daughters=\varepsilon:Sign^*]]$

B.1.4 Sign type construction operations

B.1.4.1 Lexicon

sign (Chapter 2)

If σ is a type of speech event and κ is a type (of situation) then

$$\text{sign}(\sigma, \kappa) = \left[\begin{array}{l} \text{s-event: } [e:\sigma] \\ \text{cont} = \left[\begin{array}{l} e:\kappa \\ \text{c}_{\text{tns}}:\text{final_align}(\uparrow\text{s-event.e}, e) \end{array} \right] : \text{RecType} \end{array} \right]$$

sign_{uc} (Chapter 2)

If σ is a type of speech event then

$$\text{sign}_{uc}(\sigma) = \left[\begin{array}{l} \text{s-event: } [e:\sigma] \\ \text{cont: RecType} \end{array} \right]$$

Lex (Chapter 3)

$$\begin{aligned} &\lambda T_1: \text{Type} \\ &\lambda T_2: \text{Type} . \\ &T_1 \wedge [\text{s-event: } [e:T_2]] \wedge \text{NoDaughters} \end{aligned}$$

Licensing condition associated with lexical resources (Chapter 3)

If $\text{Lex}(T, C)$ is a resource available to agent A , then for any $u, u :_A T$ licenses $:_A \text{Lex}(T, C)$
 $\wedge [\text{s-event: } [e=u:T]]$

Universal resources for lexical content construction

SemCommonNoun(p), where p is a predicate with arity $\langle \text{Ind} \rangle$ (Chapter 3)

$$\lambda r: [x:\text{Ind}] . [e : p(r.x)]$$

SemCommonNoun($p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}$), where p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type representing the background requirements (Chapter 5)

$$\left[\begin{array}{l} \text{bg} = T_{\text{bg}} \\ \text{fg} = \lambda c: T_{\text{bg}} . \left[\begin{array}{l} \text{bg} = T_{\text{restr}} \\ \text{fg} = \lambda r: [x:T_{\text{restr}}] . [e:p(r.x)] \end{array} \right] \end{array} \right]$$

SemIntransVerb(T_{bg}, p), where T_{bg} , the “background” or “presupposition” type, is a record type and p is a predicate with arity $\langle \text{Ind} \rangle$ (Chapter 4)

$$\left[\begin{array}{l} \text{bg} = T_{\text{bg}} \\ \text{fg} = \lambda r_1:T_{\text{bg}} . \lambda r_2:[x:\text{Ind}] . [e : p(r_2.x)] \end{array} \right]$$

$\text{SemIntransVerb}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$ where p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} (Chapter 5)

$$\left[\begin{array}{l} \text{bg} = T_{\text{bg}} \\ \text{fg} = \lambda c:T_{\text{bg}} . \left[\begin{array}{l} \text{bg} = T_{\text{restr}} \\ \text{fg} = \lambda r:[x:T_{\text{restr}}] . [e:p(r.x)] \end{array} \right] \end{array} \right]$$

$\text{SemPropName}(a)$, where $a:\text{Ind}$ (Chapter 3)

$$\lambda P:\text{Ppty} . P([x=a])$$

$\text{SemPropName}(T)$, where T is a phonological type (Chapter 4)

$$\left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, T) \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} x:\text{Ind} \\ e:\text{named}(x, T) \end{array} \right] . \lambda P:\text{Ppty} . P(r) \end{array} \right]$$

$\text{SemNumeral}(n)$, where $n:\text{Real}$ (Chapter 5)

$$\begin{array}{l} \lambda r:\text{Rec} . \\ \lambda P:\text{Ppty}(\text{Real}) . \\ P.\text{fg}([x=n]) \end{array}$$

SemIndefArt (Chapter 3)

$$\lambda Q:\text{Ppty} . \left[\begin{array}{l} \text{restr}=Q : \text{Ppty} \\ \text{scope}=P : \text{Ppty} \\ e : \text{exist}(\text{restr}, \text{scope}) \end{array} \right]$$

SemIndefArt (Chapter 5)

$$\lambda Q:PP\text{pty} . \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f:\text{Rec} \\ a:Q.\text{bg} \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} f:\text{Rec} \\ a:Q.\text{bg} \end{array} \right] . \lambda P:\text{Ppty} . \left[\begin{array}{l} \text{restr}=Q.\text{fg}(r.a):\text{Ppty} \\ \text{scope}=P:\text{Ppty} \\ e:\text{exist}(\text{restr}, \text{scope}) \end{array} \right] \end{array} \right]$$

SemDefArt (Chapter 5)

$$\lambda Q:PPpty . \left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} f: \left[\begin{array}{l} s:Rec \\ e:unique(Q.fg(\uparrow a),s) \end{array} \right] \\ a:Q.bg \end{array} \right] \\ fg = \lambda r: \left[\begin{array}{l} f: \left[\begin{array}{l} s:Rec \\ e:unique(Q.fg(\uparrow a),s) \end{array} \right] \\ a:Q.bg \end{array} \right] \end{array} \right] . \lambda P:Ppty . \left[\begin{array}{l} \text{restr} = Q.fg(r.a):Ppty \\ \text{scope} = P:Ppty \\ e:every(\text{restr}, \text{scope}) \end{array} \right]$$

SemBe (Chapter 3)

$$\lambda Q:Quant . \\ \lambda r_1:[x:Ind] . \\ Q(\lambda r_2:[x:Ind] . \left[\begin{array}{ll} x=r_2.X, r_1.X & : Ind \\ e & : be(x) \end{array} \right])$$

SemBe(T_{arg}, T_{bg}) where T_{arg} and T_{bg} are types (Chapter 5)

If $T_{bg} \sqsubseteq [sc:(T_{arg} \rightarrow Real)]$ then SemBe(T_{arg}, T_{bg}) is

$$\lambda r:T_{bg} . \\ \lambda Q:Quant . \\ \left[\begin{array}{l} \text{bg} = T_{arg} \\ \text{fg} = \lambda r_1:[x:T_{arg}] . \\ \quad Q \left(\left[\begin{array}{l} \text{bg} = [x:Real] \\ \text{fg} = \lambda r_2:[x:Real] . \left[\begin{array}{ll} x=r.sc(r_1.X), r_2.X & : Real \\ e & : be(x) \end{array} \right] \end{array} \right] \right) \end{array} \right]$$

Otherwise, SemBe(T_{arg}, T_{bg}) is

$$\lambda r:T_{bg} . \\ \lambda Q:Quant . \\ \left[\begin{array}{l} \text{bg} = T_{arg} \\ \text{fg} = \lambda r_1:[x:T_{arg}] . \\ \quad Q \left(\left[\begin{array}{l} \text{bg} = T_{arg} \\ \text{fg} = \lambda r_2:[x:T_{arg}] . \left[\begin{array}{ll} x=r_1.X, r_2.X & : T_{arg} \\ e & : be(x) \end{array} \right] \end{array} \right] \right) \end{array} \right]$$

Universal resources for associating lexical content with phonological types

$\text{LexCommonNoun}(T_{\text{phon}}, p)$, where T_{phon} is a phonological type and p is a predicate with arity $\langle Ind \rangle$ (Chapter 3)
is defined as

$\text{Lex}(T_{\text{phon}}, N) \wedge [\text{cont}=\text{SemCommonNoun}(p):P\text{pty}]$

$\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where T_{phon} is a phonological type, p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type (Chapter 5) is defined as

$\text{Lex}(T_{\text{phon}}, N) \wedge [\text{cont}=\text{SemCommonNoun}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}):PP\text{pty}]$

$\text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, T_{\text{bg}}, p)$, where T_{phon} is a phonological type and p is a predicate with arity $\langle \text{Ind} \rangle$ (Chapter 4) is defined as

$\text{Lex}(T_{\text{phon}}, VP) \wedge [\text{cont}=\text{SemIntransVerb}(T_{\text{bg}}, p):PP\text{pty}]$

$\text{Lex}_{\text{IntransVerb}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}})$, where T_{phon} is a phonological type, p is a predicate with arity $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$ and T_{bg} is a record type (Chapter 5) is defined as

$\text{Lex}(T_{\text{phon}}, VP) \wedge [\text{cont}=\text{SemIntransVerb}(p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}):PP\text{pty}]$

$\text{Lex}_{\text{PropName}}(T_{\text{Phon}}, a)$, where T_{Phon} is a phonological type and $a:\text{Ind}$ (Chapter 3) is defined as

$\text{Lex}(T_{\text{Phon}}, NP) \wedge [\text{cont}=\text{SemPropName}(a):Quant]$

$\text{Lex}_{\text{PropName}}(T_{\text{Phon}})$, where T_{Phon} is a phonological type (Chapter 4) is defined as

$\text{Lex}(T_{\text{Phon}}, NP) \wedge [\text{cont}=\text{SemPropName}(T_{\text{Phon}}):PQuant]$

$\text{Lex}_{\text{numeral}}$, where T_{phon} is a phonological type and n is a (real) number (Chapter 5) is defined as

$\text{Lex}(T_{\text{phon}}, NP) \wedge [\text{cont}=\text{SemNumeral}(n):PQuant]$

$\text{Lex}_{\text{IndefArt}}(T_{\text{Phon}})$, where T_{Phon} is a phonological type (Chapter 3) is defined as

$\text{Lex}(T_{\text{Phon}}, Det) \wedge [\text{cont}=\text{SemIndefArt}:(P\text{pty} \rightarrow Quant)]$

$\text{Lex}_{\text{IndefArt}}(T_{\text{phon}})$, where T_{phon} is a phonological type (Chapter 5) is defined as

$\text{Lex}(T_{\text{phon}}, Det) \wedge [\text{cont}=\text{SemIndefArt}:(PP\text{pty} \rightarrow PQuant)]$

$\text{Lex}_{\text{DefArt}}(T_{\text{phon}})$, where T_{phon} is a phonological type (Chapter 5) is defined as

$\text{Lex}(T_{\text{phon}}, Det) \wedge [\text{cont}=\text{SemDefArt}:(PP\text{pty} \rightarrow PQuant)]$

$\text{Lex}_{\text{be}}(T_{\text{Phon}})$, where T_{Phon} is a phonological type (Chapter 3) is defined as

$\text{Lex}(T_{\text{Phon}}, V) \wedge [\text{cont}=\text{SemBe}:(\text{Quant} \rightarrow \text{Ppty})]$

$\text{Lex}_{\text{be}}(T_{\text{Phon}}, T_{\text{arg}}, T_{\text{bg}})$, where T_{Phon} is a phonological type and T_{arg} and T_{bg} are types (Chapter 5) is defined as

$\text{Lex}(T_{\text{Phon}}, V) \wedge [\text{cont}=\text{SemBe}(T_{\text{arg}}, T_{\text{bg}}):(\text{Quant} \rightarrow \text{PPpty})]$

Universal resources for coercing lexical sign types to new lexical sign types

$\text{CommonNounIndToFrame}$ (Chapter 5)

If T_{phon} is a phonological type, p is a predicate and T_{bg} is a record type (the “background type” or “presupposition”) then

$\text{CommonNounIndToFrame}(\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, \text{Ind}, \text{Ind}, T_{\text{bg}})) =$
 $\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p_{\text{frame}}, \text{Rec}, \text{Rec}, T_{\text{bg}})$

where if p is a predicate with arity $\langle \text{Ind} \rangle$, then for any e and r ,

$$e : p_{\text{frame}}(r) \text{ implies } r : \begin{bmatrix} x:\text{Ind} \\ e:p(x) \end{bmatrix}$$

$\text{RestrictCommonNoun}$ (Chapter 5)

If T_{phon} is a phonological type, p is a predicate, T_{arg} is a type and that arity of p is $\langle T_{\text{arg}} \rangle$, $T_{\text{restr}} \sqsubseteq T_{\text{arg}}$, T_{bg} is a record type and $T_{\text{mod}} \sqsubseteq T_{\text{restr}}$ then

$\text{RestrictCommonNoun}(\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{restr}}, T_{\text{bg}}), T_{\text{mod}}) =$
 $\text{Lex}_{\text{CommonNoun}}(T_{\text{phon}}, p, T_{\text{arg}}, T_{\text{mod}}, T_{\text{bg}})$

B.1.4.2 Operations which construct sign combination functions

Licensing condition associated with sign combination functions (Chapter 3)

If $f : (T_1 \rightarrow \text{Type})$ is a sign combination function available to agent A , then for any u , $u :_A T_1$ licenses $:_A f(u)$

RuleDaughters (Chapter 3)

RuleDaughters maps two types to a sign combination function

$\lambda T_1 : \text{Type}$
 $\lambda T_2 : \text{Type} .$
 $\lambda u : T_1 . T_2 \wedge [\text{syn}:[\text{daughters}=u:T_1]]$

ConcatPhon (Chapter 3)

$$\lambda u: [s\text{-event}: [e: Phon]]^+ .$$

$$[\text{ s-event} \quad : \quad [\text{ e=concat}_i(u[i].s\text{-event}.e) \quad : \quad Phon \quad] \quad]$$

Phrase structure rule notation (Chapter 3)

If C, C_1, \dots, C_n are category sign types then,

$$C \longrightarrow C_1 \dots C_n \text{ represents } \text{RuleDaughters}(C, C_1 \frown \dots \frown C_n) \frown \text{ConcatPhon}$$

Combination of parametric contents (Chapter 4)

If $\alpha : \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow (T_1 \rightarrow T_2)) \end{array} \right]$ and $\beta : \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow T_1) \end{array} \right]$ then the *combination of α and β based on functional application*, $\alpha @ \beta$, is

$$\left[\begin{array}{l} \text{bg} = \left[\begin{array}{l} \text{f:}[\alpha.\text{bg}]^f. \\ \text{a:}[\beta.\text{bg}]^a. \end{array} \right] \\ \text{fg} = \lambda r: \left[\begin{array}{l} \text{f:}[\alpha.\text{bg}]^f. \\ \text{a:}[\beta.\text{bg}]^a. \end{array} \right] . \alpha.\text{fg}(r.f)(\beta.\text{fg}(r.a)) \end{array} \right]$$

where $[T]^\pi$ represents the result of prefixing each path-name occurring as an argument to a predicate in T with π .

ContForwardApp (Chapter 3)

$$\lambda T_1:Type \lambda T_2:Type .$$

$$\lambda u: [\text{cont:}(T_2 \rightarrow T_1)] \frown [\text{cont:}T_2] .$$

$$[\text{cont}=u[0].\text{cont}(u[1].\text{cont}):T_1]$$

ContForwardApp (Chapter 4)

$$\lambda T_1:Type \lambda T_2:Type .$$

$$\lambda u: \left[\text{cont:} \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow (T_2 \rightarrow T_1)) \end{array} \right] \right] \frown \left[\text{cont:} \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow T_2) \end{array} \right] \right] .$$

$$\left[\text{cont}=u[0].\text{cont}@u[1].\text{cont:} \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow T_1) \end{array} \right] \right]$$

ContForwardApp (Chapter 5)

$$\lambda T_1:Type \lambda T_2:Type .$$

$$\lambda u: [\text{cont}:(T_2 \rightarrow T_1)] \cap [\text{cont}:T_2] .$$

$$[\text{cont}=u[0].\text{cont}(u[1].\text{cont}):T_1]$$

ContSForwardApp (Chapter 5)

$$\lambda T_1:Type \lambda T_2:Type .$$

$$\lambda u: \left[\text{cont}: \left[\text{bg}:RecType \right. \right. \left. \left. \text{fg}:(\text{bg} \rightarrow (T_2 \rightarrow T_1)) \right] \right] \cap \left[\text{cont}: \left[\text{bg}:RecType \right. \right. \left. \left. \text{fg}:(\text{bg} \rightarrow T_2) \right] \right] .$$

$$\left[\text{cont}=u[0].\text{cont}@u[1].\text{cont}: \left[\text{bg}:RecType \right. \right. \left. \left. \text{fg}:(\text{bg} \rightarrow T_1) \right] \right]$$

B.2 English resources

B.2.1 Lexicon

(Chapter 2)

sign(“Dudamel is a conductor”, conductor(dudamel)),
 sign(“Beethoven is a composer”, composer(beethoven)),
 sign(“Uchida is a pianist”, pianist(uchida)),
 sign_{uc}(“ok”),
 sign_{uc}(“aha”)

(Chapter 3)

Lex(“Dudamel”, *NP*)
 Lex(“Beethoven”, *NP*)
 Lex(“a”, *Det*)
 Lex(“composer”, *N*)
 Lex(“conductor”, *N*)
 Lex(“is”, *V*)
 Lex(“ok”, *S*)
 Lex(“aha”, *S*)

Lex_{PropName}(“Dudamel”, *d*), where *d:Ind*
 Lex_{PropName}(“Beethoven”, *b*), where *b:Ind*
 Lex_{CommonNoun}(“composer”, composer), where ‘composer’ is a predicate with arity $\langle [x:Ind] \rangle$
 Lex_{CommonNoun}(“conductor”, conductor), where ‘conductor’ is a predicate with arity $\langle [x:Ind] \rangle$
 Lex_{IndefArt}(“a”)
 Lex_{be}(“is”)

(Chapter 4)

Lex_{PropName}("Sam")
 Lex_{IntransVerb}("leave", *Rec*, leave)

(Chapter 5)

Lex_{DefArt}("the")
 Lex_{IndefArt}("a")
 Lex_{CommonNoun}("dog", dog, *Ind*, *Ind*, *Rec*)
 Lex_{CommonNoun}("dog", dog_frame, *Rec*, *Rec*, *Rec*) (derived by CommonNounIndToFrame)
 Lex_{CommonNoun}("dog", dog_frame, *Rec*, *DogFrame*, *Rec*) (derived by RestrictCommonNoun)
 Lex_{CommonNoun}("temperature", temperature, *Rec*, *Rec*, *Rec*)
 Lex_{CommonNoun}("temperature", temperature, *Rec*, *AmbTempFrame*, *Rec*) (derived by RestrictCommonNoun)
 Lex_{IntransVerb}("runs", run, *Ind*, *Ind*, *Rec*)
 Lex_{IntransVerb}("rises", rise, *Rec*, *Rec*, *Rec*)
 Lex_{be}("is", *Ind*, *Rec*)
 Lex_{be}("is", *AgeFrame*, [sc:(*AgeFrame*→*Real*)])
 Lex_{be}("is", *AmbTempFrame*, [sc:(*AmbTempFrame*→*Real*)])
 Lex_{numeral}("nine", 9)
 Lex_{numeral}("ninety", 90)

B.2.2 Phrase structure

(Chapter 3)

$S \longrightarrow NP VP$
 $NP \longrightarrow Det N$
 $VP \longrightarrow V NP$

B.2.3 Non-compositional Constructions

CnstrIsA (Chapter 3)

$$\lambda u: V \wedge [s\text{-event}: [e: "is"]] \frown NP \wedge \left[\text{syn}: \left[\begin{array}{l} \text{daughters: } Det \wedge [s\text{-event}: [e: "a"]] \\ \frown N \wedge [cont: Ppty] \end{array} \right] \right]$$

$$VP \wedge [cont = u[2].\text{syn}.\text{daughters}[2].\text{cont}: Ppty]$$

B.2.4 Interpreted phrase structure

(Chapter 3)

$S \longrightarrow NP VP \frown \text{ContForwardApp}(Ppty, RecType)$

$$NP \longrightarrow Det\ N \underset{\cdot}{\wedge} ContForwardApp(Ppty, Quant)$$

$$VP \longrightarrow V\ NP \underset{\cdot}{\wedge} CnstrIsA$$

A more readable abbreviatory notation for these rules is:

$$S \longrightarrow NP\ VP \mid NP'(VP')$$

$$NP \longrightarrow Det\ N \mid Det'(N')$$

$$VP \longrightarrow [{}_V\ \text{“is”}] [{}_{NP}\ [{}_{Det}\ \text{“a”}]\ N] \mid N'$$

Note that this last rule does not correspond to a context-free phrase-structure rule.

(Chapter 5)

$$S \longrightarrow NP\ VP \underset{\cdot}{\wedge} ContSForwardApp(Ppty, RecType)$$

$$NP \longrightarrow Det\ N \underset{\cdot}{\wedge} ContForwardApp(PPpty, PQuant)$$

$$VP \longrightarrow V\ NP \underset{\cdot}{\wedge} ContSForwardApp(Quant, Ppty)$$

A more readable abbreviatory notation for these rules is:

$$S \longrightarrow NP\ VP \mid NP' @ VP'$$

$$NP \longrightarrow Det\ N \mid Det'(N')$$

$$VP \longrightarrow V\ NP \mid V' @ NP'$$

Appendix C

Dialogue rules

C.1 Universal resources

C.1.1 Types of Information States

InfoState (Chapter 2)

$$\left[\begin{array}{l} \text{private:} \left[\text{agenda:} [MoveType(SELF)] \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move:} Move(SELF) \\ \text{chart:} Chart \\ \text{e:m-interp(chart,move)} \end{array} \right] \vee ERec \\ \text{commitments:} RecType \end{array} \right] \end{array} \right]$$

[??We need other options than *SELF*]

InitInfoState (Chapter 2)

The type of initial or empty information states

$$\left[\begin{array}{l} \text{private:} \left[\text{agenda=} [] : [RecType] \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} ERec \\ \text{commitments=} Rec : RecType \end{array} \right] \end{array} \right]$$

GameBoard (Chapter 4)

$T : GameBoard$ iff $T \sqsubseteq InfoState$

TotalInfoState (Chapter 4)

$$\left[\begin{array}{ll} \text{ltm} & : \text{RecType} \\ \text{gb} & : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right]$$

C.1.2 Action functions

Licensing conditions on type acts If $f : (T \rightarrow \text{Type})$ is an action function then for any object a and agent A , $a :_A T$ licenses $:_A f(a)!$

de se: If $f : (T \rightarrow (\text{Ind} \rightarrow \text{Type}))$ is an action function then for any object a and agent A , $a :_A T$ licenses $:_A f(a)(A)!$

ExecTopAgenda (Chapter 2)

$$\lambda r: \left[\begin{array}{ll} \text{private} & : \left[\begin{array}{ll} \text{agenda} & : \text{ne}[\text{RecType}] \end{array} \right] \\ \left[\begin{array}{ll} \text{move} & : \text{fst}(r.\text{private}.\text{agenda}) \\ \text{chart} & : \text{Chart} \\ \text{e} & : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \end{array} \right] .$$

C.1.3 Perception functions (type shifts)

Licensing conditions on type acts If $f : (T \rightarrow \text{Type})$ is a perception function then for any object o and agent A , $o :_A T$ licenses $o :_A f(o)$

de se: If $f : (T \rightarrow (\text{Ind} \rightarrow \text{Type}))$ is a perception function then for any object o and agent A , $o :_A T$ licenses $o :_A f(o)(A)$

PerceiveSpeechAct(T), $T \sqsubseteq \text{Phon}$ (Chapter 2)

$$\lambda e: \left[\begin{array}{ll} \text{e} : T \\ \text{au} = \text{SELF} : \text{Ind} \end{array} \right] . \left[\begin{array}{ll} \text{move} & : \left[\begin{array}{ll} \text{e} & : \text{SpeechAct} \wedge [\text{au} = \text{SELF} : \text{Ind}] \\ \text{cnt} & : \text{Cnt} \\ \text{c}_{\text{cnt}} & : \text{content}(\text{e}, \text{cnt}) \end{array} \right] \\ \text{chart} & : \mathfrak{C}_T \\ \text{e} & : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right]$$

where \mathfrak{C}_T is the type of charts assigned to utterances of type T (as a result of parsing). In Chapter 2 \mathfrak{C}_T is equated with Σ_T , the type of signs associated with utterances of type T .

C.1.4 Update functions

Licensing conditions on type acts If $f : (T_1 \rightarrow (T_2 \rightarrow \text{Type}))$ is an update function, A is an agent, s_i is A 's current information state, $s_i :_A T_i$, $T_i \sqsubseteq T_1$ (and $s_i : T_1$), then an event $e :_A T_2$

(and $e : T_2$) licenses $s_{i+1} :_A T_i \boxed{\wedge} f(s_i)(e)$.

IntegrateOwnAssertion (Chapter 2)

$$\lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : {}_{ne} [MoveType(SELF)] \end{array} \right] \\ \lambda u: \left[\begin{array}{l} \text{move} : \text{fst}(r.\text{private}.\text{agenda}) \wedge \left[\begin{array}{l} \text{e:} \left[\begin{array}{l} \text{sp=SELF:Ind} \\ \text{au:Ind} \end{array} \right] \\ \text{e:Assertion} \end{array} \right] \wedge \left[\begin{array}{l} \text{e:Acknowledgement} \wedge \left[\begin{array}{l} \text{sp=u.move.e.au:Ind} \\ \text{au=SELF:Ind} \end{array} \right] \\ \text{cnt=u.move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] : [MoveType(SELF)] \\ \text{chart} : Chart \\ \text{e} : \text{m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} \text{e:Acknowledgement} \wedge \left[\begin{array}{l} \text{sp=u.move.e.au:Ind} \\ \text{au=SELF:Ind} \end{array} \right] \\ \text{cnt=u.move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{rst}(r.\text{private}.\text{agenda}) \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move=u.move:Move(SELF)} \\ \text{chart=u.chart:Chart} \\ \text{e=u.e:m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right] \end{array} \right]$$

IntegrateOtherAssertion (Chapter 2)

$$\lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : [RecType] \end{array} \right] \\ \lambda u: \left[\begin{array}{l} \text{move:} \left[\begin{array}{l} \text{e:Assertion} \wedge \left[\begin{array}{l} \text{sp:Ind} \\ \text{au=SELF:Ind} \end{array} \right] \\ \text{cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{chart:Chart} \\ \text{e:m-interp}(\text{chart},\text{move}) \end{array} \right] \\ \left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda=} \left[\begin{array}{l} \text{e:Acknowledgement} \wedge \left[\begin{array}{l} \text{sp=SELF:Ind} \\ \text{au=u.move.e.sp:Ind} \end{array} \right] \\ \text{cnt=u.move.cnt:RecType} \\ \text{c}_{\text{cnt}}:\text{content}(e,\text{cnt}) \end{array} \right] \\ \text{r.private.agenda} \end{array} \right] : [RecType] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move=u.move:Move} \\ \text{chart=u.chart:Chart} \\ \text{e=u.e:m-interp}(\text{chart},\text{move}) \end{array} \right] \end{array} \right] \end{array} \right] \end{array} \right]$$

IntegrateOwnAcknowledgement (Chapter 2)

$$\lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : {}_{ne} [RecType] \end{array} \right] \\ \text{shared} : \left[\begin{array}{l} \text{latest-utterance} : \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} \text{content} : RecType \end{array} \right] \end{array} \right] \\ \text{commitments} : RecType \end{array} \right] \end{array} \right]$$

$$\lambda u: \left[\begin{array}{l} \text{move} \quad : \quad \text{fst}(r.\text{private}.\text{agenda}) \wedge [e:\text{Acknowledgement}] \wedge [e:[\text{sp}=\text{SELF}:\text{Ind}]] \\ \text{chart} \quad : \quad \text{Chart} \\ e \quad : \quad \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] .$$

$$\left[\begin{array}{l} \text{private:} \left[\text{agenda} = \text{rst}(r.\text{private}.\text{agenda}):[\text{RecType}] \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move} = u.\text{move}:\text{Move} \\ \text{chart} = u.\text{chart}:\text{Chart} \\ e = u.e:\text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \\ \text{commitments} = [\text{prev}:r.\text{commitments}] \wedge u.\text{move}.\text{cnt}:\text{RecType} \end{array} \right] \end{array} \right]$$

IntegrateOtherAcknowledgement (Chapter 2)

$$\begin{array}{l}
\lambda r: \left[\begin{array}{l} \text{private} : \left[\begin{array}{l} \text{agenda} : {}_{ne} [RecType] \end{array} \right] \\ \text{shared} : \left[\begin{array}{l} \text{latest-utterance} : \left[\begin{array}{l} \text{move} : \left[\begin{array}{l} \text{content} : RecType \end{array} \right] \end{array} \right] \\ \text{commitments} : RecType \end{array} \right] \end{array} \right] \\
\lambda u: \left[\begin{array}{l} \text{move} : \text{fst}(r.\text{private}.\text{agenda}) \wedge [e: Acknowledgement] \wedge [e: [au=SELF:Ind]] \\ \text{chart} : Chart \\ e : \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \end{array} \right] . \\
\left[\begin{array}{l} \text{private:} \left[\begin{array}{l} \text{agenda} = \text{rst}(r.\text{private}.\text{agenda}): [RecType] \end{array} \right] \\ \text{shared:} \left[\begin{array}{l} \text{latest-utterance:} \left[\begin{array}{l} \text{move} = u.\text{move}: Move \\ \text{chart} = u.\text{chart}: Chart \\ e = u.e: \text{m-interp}(\text{chart}, \text{move}) \end{array} \right] \\ \text{commitments} = [\text{prev}: r.\text{commitments}] \wedge u.\text{move}.\text{cnt}: RecType \end{array} \right] \end{array} \right]
\end{array}$$

Licensing conditions on accommodation updates (Chapter 4)

If A is an agent, s_i is A 's current information state, f is a parametric content of type T_f such that

$$T_f \sqsubseteq \left[\begin{array}{l} \text{bg} : RecType \\ \text{fg} : (\text{bg} \rightarrow RecType) \end{array} \right]$$

and $s_i :_A T_i$ for some T_i such that

$$T_i \sqsubseteq \left[\begin{array}{l} \text{ltm}: RecType \\ \text{gb:} \left[\begin{array}{l} \text{shared:} \left[\begin{array}{l} \text{commitments}: RecType \\ \text{latest-move:} [\text{cont} = f: T_f] \end{array} \right] \end{array} \right] \end{array} \right]$$

then

if there is some η which is a relabelling of $\varphi(f.\text{bg})$ such that

$$\varphi(s_i.\text{gb}.\text{shared}.\text{commitments}) \sqsubseteq [\varphi(f.\text{bg})]_\eta$$

then $s_{i+1} :_A T_i \sqcap \mathbf{AccGB}(\eta)(s_i)(f)$ is licensed

else if there is some η which is a relabelling of $\varphi(f.\text{bg})$ such that $\varphi(s_i.\text{ltm}) \sqsubseteq [\varphi(f.\text{bg})]_\eta$

then $s_{i+1} :_A T_i \sqcap \mathbf{AccLTM}(\eta)(s_i)(f)$ is licensed

else $s_{i+1} :_A T_i \sqcap \mathbf{AccNM}(s_i)(f)$ is licensed

AccLTM(η) (“accommodate match with long term memory”, Chapter 4)

$$\lambda r: \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] \\ \lambda f: \left[\begin{array}{l} \text{bg} : \text{RecType} \\ \text{fg} : (\text{bg} \rightarrow \text{RecType}) \end{array} \right] \cdot \\ \left[\begin{array}{l} \text{ltm} = r.\text{ltm} : \text{RecType} \\ \text{gb} = \lambda r_1 : \text{ltm} . ((r.\text{gb})(r_1)) \boxed{\wedge} \\ \left[\begin{array}{l} \text{shared:} \left[\begin{array}{l} \text{commitments} = \left[\begin{array}{l} \text{prev:}(r.\text{gb})(\uparrow^3 \text{ltm}).\text{shared.commitments} \\ \text{bg:} f.\text{bg} \parallel_{\eta} r_1 \\ \text{e:} f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \end{array} \right] \end{array} \right] : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right]$$

AccNM (“accommodate no match”, Chapter 4)

$$\lambda r: \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] \\ \lambda f: \left[\begin{array}{l} \text{bg} : \text{RecType} \\ \text{fg} : (\text{bg} \rightarrow \text{RecType}) \end{array} \right] \cdot \\ \left[\begin{array}{l} \text{ltm} = r.\text{ltm} : \text{RecType} \\ \text{gb} = \lambda r_1 : \text{ltm} . ((r.\text{gb})(r_1)) \boxed{\wedge} \\ \left[\begin{array}{l} \text{shared:} \left[\begin{array}{l} \text{commitments} = \left[\begin{array}{l} \text{prev:}(r.\text{gb})(\uparrow^3 \text{ltm}).\text{shared.commitments} \\ \text{bg:} f.\text{bg} \\ \text{e:} f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \end{array} \right] \end{array} \right] : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right]$$

AccGB(η) (“accommodate match on gameboard”, Chapter 4)

AccGB(η) =

$$\lambda r: \left[\begin{array}{l} \text{ltm} : \text{RecType} \\ \text{gb} : (\text{ltm} \rightarrow \text{GameBoard}) \end{array} \right] \cdot \\ \lambda f: \left[\begin{array}{l} \text{bg:RecType} \\ \text{fg:}(\text{bg} \rightarrow \text{RecType}) \end{array} \right] \cdot \\ \left[\begin{array}{l} \text{ltm} = r.\text{ltm} : \text{RecType} \\ \text{gb} = \lambda r_1 : \text{ltm} . r.\text{gb}(r_1) \boxed{\wedge} \left[\begin{array}{l} \text{shared:} \left[\begin{array}{l} \text{commitments} = \left[\begin{array}{l} \text{prev:} r.\text{gb}.\text{shared.commitments} \\ \text{bg:} f.\text{bg} \parallel_{\eta} \text{prev} \\ \text{fg:} f.\text{fg}(\text{bg}) \end{array} \right] : \text{RecType} \end{array} \right] \end{array} \right] : (\text{ltm} \rightarrow \text{RecType}) \end{array} \right]$$

C.2 English resources

Bibliography

- Artstein, Ron, Mark Core, David DeVault, Kallirroi Georgila, Elsi Kaiser and Amanda Stent, eds. (2011) *SemDial 2011* (Los Angeles): Proceedings of the 15th Workshop on the Semantics and Pragmatics of Dialogue.
- Austin, J. (1962) *How to Do Things with Words*, Oxford University Press, ed. by J. O. Urmson.
- Austin, J. L. (1961) Truth, in J. O. Urmson and G. J. Warnock (eds.), *J. L. Austin: Philosophical Papers*, Oxford University Press, Oxford.
- Barsalou, Lawrence W. (1992a) *Cognitive psychology. An overview for cognitive scientists*, Lawrence Erlbaum Associates, Hillsdale, NJ.
- Barsalou, Lawrence W. (1992b) Frames, concepts, and conceptual fields, in A. Lehrer and E. F. Kittay (eds.), *Frames, fields, and contrasts: New essays in semantic and lexical organization*, pp. 21–74, Lawrence Erlbaum Associates, Hillsdale, NJ.
- Barsalou, Lawrence W. (1999) Perceptual symbol systems, *Behavioral and Brain Sciences*, Vol. 22, pp. 577–660.
- Barwise, Jon (1989) *The Situation in Logic*, CSLI Publications, Stanford.
- Barwise, Jon and Robin Cooper (1981) Generalized quantifiers and natural language, *Linguistics and Philosophy*, Vol. 4, No. 2, pp. 159–219.
- Barwise, Jon and Robin Cooper (1993) Extended Kamp Notation: a Graphical Notation for Situation Theory, in P. Aczel, D. Israel, Y. Katagiri and S. Peters (eds.), *Situation Theory and its Applications*, Vol. 3, CSLI, Stanford.
- Barwise, Jon and John Perry (1983) *Situations and Attitudes*, Bradford Books, MIT Press, Cambridge, Mass.
- Bäuerle, Rainer, Urs Egli and Arnim von Stechow, eds. (1979) *Semantics from Different Points of View* (*Springer Series in Language and Communication* 6), Springer.
- Bennett, Michael Ruisdael (1974) *Some extensions of a Montague fragment of English*, PhD dissertation, UCLA. Distributed by Indiana University Linguistics Club.

- Blackburn, Patrick, Maarten de Rijke and Yde Venema (2001) *Modal logic* (*Cambridge Tracts in Theoretical Computer Science* 53), Cambridge University Press.
- Boas, Hans C. and Ivan A. Sag, eds. (2012) *Sign-Based Construction Grammar*, CSLI Publications.
- Bos, J. (1996) Predicate logic unplugged, in P. Dekker and M. Stokhof (eds.), *Proceedings of the Tenth Amsterdam Colloquium*, pp. 133–143, ILLC/Department of Philosophy, University of Amsterdam, Amsterdam.
- Botvinick, Matthew M. (2008) Hierarchical models of behavior and prefrontal function, *Trends in Cognitive Sciences*, Vol. 12, No. 5, pp. 201 – 208.
- Botvinick, Matthew M., Yael Niv and Andrew C. Barto (2009) Hierarchically organized behavior and its neural foundations: A reinforcement learning perspective, *Cognition*, Vol. 113, No. 3, pp. 262 – 280. Reinforcement learning and higher cognition.
- Breitholtz, Ellen (2010) Clarification Requests as Enthymeme Elicitors, in *Aspects of Semantics and Pragmatics of Dialogue. SemDial 2010, 14th Workshop on the Semantics and Pragmatics of Dialogue* .
- Breitholtz, Ellen (2014) *Enthymemes in Dialogue: A micro-rhetorical approach*, PhD dissertation, University of Gothenburg.
- Breitholtz, Ellen and Robin Cooper (2011) Enthymemes as Rhetorical Resources, in Artstein *et al.* (2011).
- Breitholtz, Ellen and Jessica Villing (2008) Can Aristotelian Enthymemes Decrease the Cognitive Load of a Dialogue System User?, in *Proceedings of LonDial 2008, the 12th SEMDIAL workshop*.
- Bullock, Barbara E. and Almeida Jacqueline Toribio, eds. (2009) *The Cambridge Handbook of Linguistic Code-Switching*, Cambridge University Press.
- Carlson, Gregory N. (1982) Generic Terms and Generic Sentences, *Journal of Philosophical Logic*, Vol. 11, pp. 145–81.
- Carnap, Rudolf (1956) *Meaning and Necessity: A Study in Semantics and Modal Logic*, second edition, University of Chicago Press.
- Chierchia, Gennaro (1995) *Dynamics of Meaning: Anaphora, Presupposition, and the Theory of Grammar*, University of Chicago Press, Chicago.
- Chierchia, Gennaro and Raymond Turner (1988) Semantics and property theory, *Linguistics and Philosophy*, Vol. 11, No. 3, pp. 261–302.

- Cooper, Robin (1982) Binding in wholewheat* syntax (*unenriched with inaudibilia), in P. Jacobson and G. K. Pullum (eds.), *The Nature of Syntactic Representation* (*Synthese Language Library* 15), Reidel Publishing Company.
- Cooper, Robin (1983) *Quantification and Syntactic Theory*, Reidel Publishing Company.
- Cooper, Robin (1991) Three lectures on situation theoretic grammar, in M. Filgueiras, L. Damas, N. Moreira and A. P. Tomás (eds.), *Natural Language Processing, EAIA 90, Proceedings, Lecture Notes in Artificial Intelligence* 476, pp. 101–140, Springer Verlag, Berlin.
- Cooper, Robin (1996) The Role of Situations in Generalized Quantifiers, in S. Lappin (ed.), *The Handbook of Contemporary Semantic Theory*, Blackwell, Oxford.
- Cooper, Robin (2005) Records and Record Types in Semantic Theory, *Journal of Logic and Computation*, Vol. 15, No. 2, pp. 99–112.
- Cooper, Robin (2010) Frames in formal semantics, in H. Loftsson, E. Rögnvaldsson and S. Helgadóttir (eds.), *IceTAL 2010*, Springer Verlag.
- Cooper, Robin (2011) Copredication, Quantification and Frames, in S. Pogodalla and J.-P. Prost (eds.), *Logical Aspects of Computational Linguistics: 6th International Conference, LACL 2011*, pp. 64–79, Springer.
- Cooper, Robin (2012a) Intensional quantifiers, in T. Graf, D. Paperno, A. Szabolcsi and J. Tellings (eds.), *Theories of Everything: In Honor of Ed Keenan*, UCLA Working Papers in Linguistics 17, pp. 69–71, Department of Linguistics, UCLA.
- Cooper, Robin (2012b) Type Theory and Semantics in Flux, in R. Kempson, N. Asher and T. Fernando (eds.), *Handbook of the Philosophy of Science*, Vol. 14: Philosophy of Linguistics, pp. 271–323, Elsevier BV. General editors: Dov M. Gabbay, Paul Thagard and John Woods.
- Cooper, Robin (2013a) Clarification and Generalized Quantifiers, *Dialogue and Discourse*, Vol. 4, No. 1, pp. 1–25.
- Cooper, Robin (2013b) Update conditions and intensionality in a type-theoretic approach to dialogue semantics, in R. Fernández and A. Isard (eds.), *Proceedings of the 17th Workshop on the Semantics and Pragmatics of Dialogue*, pp. 15–24, University of Amsterdam.
- Cooper, Robin (2015) Natural reasoning: truth or judgement based? presented at Co-CoNat’15, Bloomington, Ind., 19th-20th July, 2015 <https://sites.google.com/site/typetheorywithrecords/publications/natlogtruthjudge.pdf>.
- Cooper, Robin (fthc) Type Theory and Semantics in Flux, in m (ed.), *Handbook of the Philosophy of Science*, Vol. 14: Philosophy of Linguistics, Elsevier BV. General editors: Dov M. Gabbay, Paul Thagard and John Woods.

- Cooper, Robin, Simon Dobnik, Shalom Lappin and Staffan Larsson (2014a) A Probabilistic Rich Type Theory for Semantic Interpretation, in Cooper *et al.* (2014b), pp. 72–79, Association for Computational Linguistics.
- Cooper, Robin, Simon Dobnik, Shalom Lappin and Staffan Larsson, eds. (2014b) *Proceedings of the EACL 2014 Workshop on Type Theory and Natural Language Semantics (TTNLS)*, Association for Computational Linguistics, Gothenburg, Sweden.
- Cooper, Robin, Simon Dobnik, Shalom Lappin and Staffan Larsson (2015) Probabilistic Type Theory and Natural Language Semantics, *Linguistic Issues in Language Technology*, Vol. 10, No. 4, pp. 1–45.
- Cooper, Robin and Jonathan Ginzburg (2011a) Negation in Dialogue, in Artstein *et al.* (2011), pp. 130–139.
- Cooper, Robin and Jonathan Ginzburg (2011b) Negative inquisitiveness and alternatives-based negation, in *Proceedings of the Amsterdam Colloquium, 2011*.
- Cooper, Robin and Ruth Kempson, eds. (2008) *Language in Flux: Dialogue Coordination, Language Variation, Change and Evolution (Communication, Mind and Language 1)*, College Publications, London.
- Cooper, Robin and Aarne Ranta (2008) Natural Languages as Collections of Resources, in Cooper and Kempson (2008), pp. 109–120.
- Cresswell, M.J. (1985) *Structured Meanings: The Semantics of Propositional Attitudes*, MIT Press.
- Davidson, Donald (1967) The Logical Form of Action Sentences, in N. Rescher (ed.), *The Logic of Decision and Action*, University of Pittsburgh Press. Reprinted in Davidson (1980).
- Davidson, Donald (1980) *Essays on Actions and Events*, Oxford University Press, New edition 2001.
- Davidson, Donald and Gilbert Harman, eds. (1972) *Semantics of Natural Language*, Reidel Publishing Company.
- Dowty, David (1989) On the semantic content of the notion of ‘Thematic Role’, in G. Chierchia, B. H. Partee and R. Turner (eds.), *Properties, Types and Meanings*, Vol. II: Semantic Issues, pp. 69–130, Kluwer, Dordrecht.
- Dowty, David, Robert Wall and Stanley Peters (1981) *Introduction to Montague Semantics*, Reidel (Springer).
- van Eijck, Jan and Christina Unger (2010) *Computational Semantics with Functional Programming*, Cambridge University Press.
- Elbourne, Paul (2012) *Definite Descriptions*, Clarendon Press, Oxford.

- Fernando, Tim (2001) Conservative Generalized Quantifiers and Presupposition, in R. Hastings, B. Jackson and Z. Zvolenszky (eds.), *Proceedings of the 11th Semantics and Linguistic Theory Conference (held May 11-13, 2001, at New York University)*, pp. 172–191.
- Fernando, Tim (2004) A finite-state approach to events in natural language semantics, *Journal of Logic and Computation*, Vol. 14, No. 1, pp. 79–92.
- Fernando, Tim (2006) Situations as Strings, *Electronic Notes in Theoretical Computer Science*, Vol. 165, pp. 23–36.
- Fernando, Tim (2008) Finite-state descriptions for temporal semantics, in H. Bunt and R. Muskens (eds.), *Computing Meaning, Volume 3 (Studies in Linguistics and Philosophy 83)*, pp. 347–368, Springer.
- Fernando, Tim (2009) Situations in LTL as strings, *Information and Computation*, Vol. 207, No. 10, pp. 980–999.
- Fernando, Tim (2011) Constructing Situations and Time, *Journal of Philosophical Logic*, Vol. 40, pp. 371–396.
- Fillmore, Charles J. (1970) Subjects, Speakers, and Roles, *Synthese*, Vol. 21, pp. 251–274.
- Fillmore, Charles J. (1982) Frame semantics, in *Linguistics in the Morning Calm*, pp. 111–137, Hanshin Publishing Co., Seoul.
- Fillmore, Charles J. (1985) Frames and the semantics of understanding, *Quaderni di Semantica*, Vol. 6, No. 2, pp. 222–254.
- von Fintel, Kai and Irene Heim (2011) *Intensional Semantics*. MIT, Spring 2011 Edition, <http://web.mit.edu/fintel/fintel-heim-intensional.pdf>.
- Fodor, Janet Dean (1970) *The Linguistic Description of Opaque Contexts*, PhD dissertation, MIT.
- Fox, Chris and Shalom Lappin (2005) *Foundations of Intensional Semantics*, Blackwell Publishing.
- Frege, Gottlob (1892) Über Sinn und Bedeutung, *Zeitschrift für Philosophie und philosophische Kritik*, Vol. 100, pp. 25–50. Translated in Geach and Black (1980).
- Gawron, Jean Mark and Stanley Peters (1990) *Anaphora and Quantification in Situation Semantics*, CSLI Publications.
- Geach, P. and M. Black, eds. (1980) *Translations from the Philosophical Writings of Gottlob Frege*, third edition, Blackwell, Oxford.
- Gibson, James J. (1986) *The Ecological Approach to Visual Perception*, Lawrence Erlbaum Associates.

- Gil, David (2000) Syntactic categories, cross-linguistic variation and universal grammar, in P. M. Vogel and B. Comrie (eds.), *Approaches to the typology of word classes (Empirical approaches to language typology 23)*, Mouton de Gruyter, Berlin.
- Ginzburg, Jonathan (1994) An update semantics for dialogue, in H. Bunt (ed.), *Proceedings of the 1st International Workshop on Computational Semantics*, Tilburg University.
- Ginzburg, Jonathan (2010) Relevance for Dialogue, in Łupkowski and Purver (2010), pp. 121–129, Polish Society for Cognitive Science.
- Ginzburg, Jonathan (2012) *The Interactive Stance: Meaning for Conversation*, Oxford University Press, Oxford.
- Ginzburg, Jonathan and Robin Cooper (2004) Clarification, ellipsis, and the nature of contextual updates in dialogue, *Linguistics and Philosophy*, Vol. 27, No. 3, pp. 297–365.
- Ginzburg, Jonathan and Robin Cooper (2014) Quotation via Dialogical Interaction, *Journal of Logic, Language and Information*, Vol. 23, No. 3, pp. 287–311.
- Ginzburg, Jonathan, Robin Cooper and Tim Fernando (2014) Propositions, Questions, and Adjectives: a rich type theoretic approach, in Cooper *et al.* (2014b), pp. 89–96, Association for Computational Linguistics.
- Ginzburg, Jonathan and Raquel Fernández (2010) Computational Models of Dialogue, in A. Clark, C. Fox and S. Lappin (eds.), *The Handbook of Computational Linguistics and Natural Language Processing*, Wiley-Blackwell.
- Ginzburg, Jonathan and Ivan A. Sag (2000) *Interrogative Investigations: The Form, Meaning, and Use of English Interrogatives*, CSLI Lecture Notes 123, CSLI Publications, Stanford, California.
- Globus, Gordon G. (1995) *The Postmodern Brain (Advances in Consciousness Research 1)*, John Benjamins Publishing Company.
- Groenendijk, Jeroen and Floris Roelofsen (2012) Course Notes on Inquisitive Semantics, NASSLLI 2012. Available at <https://sites.google.com/site/inquisitivesemantics/documents/NASSLLI-2012-inquisitive-semantics-lecture-notes.pdf>.
- de Groote, Philippe and Ekaterina Lebedeva (2010) Presupposition Accommodation as Exception Handling, in *Proceedings of SIGDIAL 2010: the 11th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pp. 71–74.
- Grosz, Barbara J., Aravind K. Joshi and Scott Weinstein (1983) Providing a unified account of definite noun phrases in discourse, in *Proceedings of ACL-83*, pp. 44–50, Cambridge, MA.

- Grosz, Barbara J., Aravind K. Joshi and Scott Weinstein (1995) Centering: A framework for modeling the local coherence of discourse, *Computational Linguistics*, Vol. 21, No. 2, pp. 202–225.
- Gupta, Anil (1980) *The Logic of Common Nouns: An Investigation in Quantified Model Logic*, Yale University Press, New Haven.
- Halliday, M. A. K. (1977) Text as semantic choice in social contexts, in T. van Dijk and J. Petöfi (eds.), *Grammars and descriptions*, pp. 176–225, Walter de Gruyter, Berlin.
- Halpern, J. (2003) *Reasoning About Uncertainty*, MIT Press, Cambridge MA.
- Heim, Irene and Angelika Kratzer (1998) *Semantics in Generative Grammar*, Blackwell Publishing.
- Hughes, G.E. and M.J. Cresswell (1968) *Introduction to modal logic*, Methuen and Co., Ltd.
- Hughes, G.E. and M.J. Cresswell (1996) *A New Introduction to Modal Logic*, Routledge.
- Jackendoff, Ray (1979) How to Keep Ninety from Rising, *Linguistic Inquiry*, Vol. 10, No. 1, pp. 172–177.
- Jackendoff, Ray (2002) *Foundations of Language: Brain, Meaning, Grammar, Evolution*, Oxford University Press.
- Joshi, Aravind K. and Scott Weinstein (1981) Control of inference: Role of some aspects of discourse structure-centering, in *Proceedings of the IJCAI*, pp. 385–387, Vancouver, CA.
- Jurafsky, Daniel and James H. Martin (2009) *Speech and Language Processing*, second edition, Pearson Education.
- Kallmeyer, Laura and Rainer Osswald (2013) Syntax-driven semantic frame composition in Lexicalized Tree Adjoining Grammars, *Journal of Language Modelling*, Vol. 1, No. 2, pp. 267–330.
- Kamp, Hans (1979) Events, Instants and Temporal Reference, in Bäuerle *et al.* (1979), pp. 376–418.
- Kamp, Hans (1990) Prolegomena to a Structural Theory of Belief and other Attitudes, in C. A. Anderson and J. Owens (eds.), *Propositional Attitudes: the Role of Content in Logic, Language and Mind*, CSLI Publications, Stanford.
- Kamp, Hans, Josef van Genabith and Uwe Reyle (2011) Discourse Representation Theory, in D. Gabbay and F. Guenther (eds.), *Handbook of Philosophical Logic*, Vol. 15, Springer Science+Business Media B.V. .
- Kamp, Hans and Uwe Reyle (1993) *From Discourse to Logic*, Kluwer, Dordrecht.

- Kant, Immanuel (1781) *Critik der reinen Vernunft (Critique of Pure Reason)*, Johann Friedrich Hartknoch, Riga, second edition 1787.
- Kaplan, David (1978) On the Logic of Demonstratives, *Journal of Philosophical Logic*, Vol. 8, pp. 81–98.
- Keenan, E. L. and J. Stavi (1986) Natural Language Determiners, *Linguistics and Philosophy*, Vol. 9, pp. 253–326.
- Keller, William R. (1988) Nested Cooper Storage: The Proper Treatment of Quantification in Ordinary Noun Phrases, in U. Reyle and C. Rohrer (eds.), *Natural Language Parsing and Linguistic Theories*, pp. 432–447, Reidel, Dordrecht.
- King, Jeffrey C. (2014) Structured Propositions, in E. N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy*, spring 2014 edition, <http://plato.stanford.edu/archives/spr2014/entries/propositions-structured/>.
- King, Jeffrey C., Scott Soames and Jeff Speaks (2014) *New Thinking about Propositions*, Oxford University Press.
- Kracht, Marcus and Udo Klein (2014) The Grammar of Code Switching, *Journal of Logic, Language and Information*, Vol. 23, pp. 313–329.
- Kratzer, Angelika (1977) What ‘Must’ and ‘Can’ Must and Can Mean, *Linguistics and Philosophy*, Vol. 1, pp. 337–55.
- Kratzer, Angelika (1981) The Notional Category of Modality, in H. J. Eikmeyer and H. Rieser (eds.), *Words, Worlds, and Contexts*, pp. 38–74, de Gruyter, Berlin and New York.
- Kratzer, Angelika (2012) *Modals and Conditionals: New and Revised Perspectives*, Oxford University Press.
- Kratzer, Angelika (2014) Situations in Natural Language Semantics, in E. N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy*, spring 2014 edition, <http://plato.stanford.edu/archives/spr2014/entries/situations-semantics/>.
- Krifka, Manfred (1990) Four Thousand Ships Passed through the Lock: Object-induced Measure Functions on Events, *Linguistics and Philosophy*, Vol. 13, pp. 487–520.
- Kripke, Saul (1972) Naming and Necessity, in Davidson and Harman (1972), pp. 253–355.
- Kripke, Saul (1979) A Puzzle about Belief, in A. Margalit (ed.), *Meaning and Use*, Reidel.
- Kölbel, Max (2004) Faultless Disagreement, *Proceedings of the Aristotelian Society*, Vol. 104, pp. 53–73.
- Lappin, Shalom (2015) Curry Typing, Polymorphism, and Fine-Grained Intensionality, in Lappin and Fox (2015).

- Lappin, Shalom and Chris Fox, eds. (2015) *The Handbook of Contemporary Semantic Theory*, second edition, Wiley-Blackwell.
- Larson, Richard K. and Peter Ludlow (1993) Interpreted Logical Forms, *Synthese*, Vol. 96, pp. 305–55.
- Larsson, Staffan (2002) *Issue-based Dialogue Management*, PhD dissertation, University of Gothenburg.
- Larsson, Staffan (2010) Accommodating innovative meaning in dialogue, in Łupkowski and Purver (2010), pp. 83–90, Polish Society for Cognitive Science.
- Larsson, Staffan (2011) The TTR perceptron: Dynamic perceptual meanings and semantic coordination., in Artstein *et al.* (2011).
- Larsson, Staffan and Robin Cooper (2009) Towards a formal view of corrective feedback, in A. Alishahi, T. Poibeau and A. Villavicencio (eds.), *Proceedings of the Workshop on Cognitive Aspects of Computational Language Acquisition*, pp. 1–9.
- Larsson, Staffan and David R. Traum (2001) Information state and dialogue management in the TRINDI dialogue move engine toolkit, *Natural Language Engineering*, Vol. 6, No. 3&4, pp. 323–340.
- Lasersohn, Peter (2005) The Temperature Paradox as Evidence for a Presuppositional Analysis of Definite Descriptions, *Linguistic Inquiry*, Vol. 36, No. 1, pp. 127–134.
- Lewis, David (1972) General Semantics, in Davidson and Harman (1972), pp. 169–218.
- Lewis, David (1973) *Counterfactuals*, Harvard University Press, Revised printing 1986.
- Lewis, David (1979a) Attitudes de dicto and de se, *Philosophical Review*, Vol. 88, pp. 513–543. Reprinted in Lewis (1983).
- Lewis, David (1979b) Scorekeeping in a Language Game, *Journal of Philosophical Logic*, Vol. 8, pp. 339–359.
- Lewis, David (1983) *Philosophical Papers, Volume 1*, Oxford University Press.
- Lewis, David K. (1981) Ordering Semantics and Premise Semantics for Counterfactuals, *Journal of Philosophical Logic*, Vol. 10, pp. 217–234.
- Linell, Per (2009) *Rethinking Language, Mind, and World Dialogically: Interactional and contextual theories of human sense-making*, Advances in Cultural Psychology: Constructing Human Development, Information Age Publishing, Inc., Charlotte, N.C.
- Löbner, Sebastian (1979) *Intensionale Verben und Funktionalbegriffe. Untersuchung zur Syntax und Semantik von wechseln und den vergleichbaren Verben des Deutschen*, Narr, Tübingen.

- Löbner, Sebastian (1981) Intensional Verbs and Functional Concepts: More on the “Rising Temperature” Problem, *Linguistic Inquiry*, Vol. 12, No. 3, pp. 471–477.
- Löbner, Sebastian (2014) Evidence for frames from human language, in T. Gamerschlag, D. Gerland, W. Petersen and R. Osswald (eds.), *Frames and Concept Types* (*Studies in Linguistics and Philosophy* 94), pp. 23–68, Springer, Heidelberg, New York.
- Löbner, Sebastian (in prep) Functional Concepts and Frames. Available from http://semanticsarchive.net/Archive/jI1NGEwO/Loebner_Functional_Concepts_and_Frames.pdf.
- Ludlow, Peter (2014) *Living Words: Meaning Underdetermination and the Dynamic Lexicon*, Oxford University Press.
- Łupkowski, Paweł and Matthew Purver, eds. (2010) Aspects of Semantics and Pragmatics of Dialogue. SemDial 2010, 14th Workshop on the Semantics and Pragmatics of Dialogue. Poznań: Polish Society for Cognitive Science.
- Maier, Emar (2009) Proper names and indexicals trigger rigid presuppositions, *Journal of Semantics*, Vol. 26, pp. 253–315.
- Martin-Löf, Per (1984) *Intuitionistic Type Theory*, Bibliopolis, Naples.
- McCarthy, J. and P. J. Hayes (1969) Some philosophical problems from the standpoint of artificial intelligence, *Machine Intelligence*, Vol. 4, pp. 463–502.
- McCawley, James D. (1979) Presupposition and Discourse Structure, in C.-K. Oh and D. A. Dinneen (eds.), *Presupposition* (*Syntax and Semantics* 11), Academic Press.
- Menzel, Christopher (2015) Possible Worlds, in E. N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy*, summer 2015 edition, <http://plato.stanford.edu/archives/sum2015/entries/possible-worlds/>.
- Montague, Richard (1970) Universal Grammar, *Theoria*, Vol. 36, pp. 373–398.
- Montague, Richard (1973) The Proper Treatment of Quantification in Ordinary English, in J. Hintikka, J. Moravcsik and P. Suppes (eds.), *Approaches to Natural Language: Proceedings of the 1970 Stanford Workshop on Grammar and Semantics*, pp. 247–270, D. Reidel Publishing Company, Dordrecht.
- Montague, Richard (1974) *Formal Philosophy: Selected Papers of Richard Montague*, Yale University Press, New Haven, ed. and with an introduction by Richmond H. Thomason.
- Ninan, Dilip (2010) *De Se* Attitudes: Ascription and Communication, *Philosophy Compass*, Vol. 5, No. 7, pp. 551–567.

- Nordström, Bengt, Kent Petersson and Jan M. Smith (1990) *Programming in Martin-Löf's Type Theory* (*International Series of Monographs on Computer Science* 7), Clarendon Press, Oxford.
- Partee, Barbara H. (1977) Possible World Semantics and Linguistic Theory, *The Monist*, Vol. 60, No. 3, pp. 303–326.
- Partee, Barbara H. (1979) Semantics – Mathematics or Psychology?, in Bäuerle *et al.* (1979).
- Partee, Barbara H. (1986) Noun Phrase Interpretation and Type-Shifting Principles, in J. Groenendijk, D. de Jongh and M. Stokhof (eds.), *Studies in Discourse Representation Theory and the Theory of Generalized Quantifiers*, Foris Publications.
- Partee, Barbara H. (2014) The History of Formal Semantics: Changing Notions of Linguistic Competence. 9th Annual Joshua and Verona Whatmough Lecture, Harvard, <https://udrive.oit.umass.edu/partee/Partee2014Harvard.pdf>, <https://www.youtube.com/watch?v=0VV-1NDKMEc>.
- Partee, Barbara H. and Vladimir Borschev (2012) Sortal, Relational, and Functional Interpretations of Nouns and Russian Container Constructions, *Journal of Semantics*, Vol. 29, No. 4, pp. 445–486.
- Partee, Barbara H., Alice G.B. ter Meulen and Robert E. Wall (1990) *Mathematical Methods in Linguistics*, Springer.
- Perry, John (1979) The Problem of the Essential Indexical, *Noûs*, Vol. 13, No. 1, pp. 3–21. Reprinted in Perry (1993).
- Perry, John (1993) *The Problem of the Essential Indexical and Other Essays*, Oxford University Press.
- Peters, Stanley and Dag Westerståhl (2006) *Quantifiers in Language and Logics*, Oxford University Press.
- Poesio, Massimo, Rosemary Stevenson, Barbara Di Eugenio and Janet Hitzeman (2004) Centering: A Parametric Theory and Its Instantiations, *Computational Linguistics*, Vol. 30, No. 3, pp. 309–363.
- Portner, Paul (2009) *Modality* (*Oxford Surveys in Semantics and Pragmatics* 1), Oxford University Press.
- Prinz, Jesse J. and Lawrence W. Barsalou (2014) Steering a Course for Embodied Representation, in E. Dietrich and A. B. Markman (eds.), *Cognitive Dynamics: Conceptual and Representational Change in Humans and Machines*, pp. 51–77, Psychology Press. Previously published in 2000 by Lawrence Erlbaum.

- Pross, Tillmann (ms) Fodor's puzzle and the semantics of attitude reports. Available from <http://www.ims.uni-stuttgart.de/institut/mitarbeiter/prosstn/files/pross-attitudes.pdf>.
- Purver, Matthew, Eleni Gregoromichelaki, Wilfried Meyer-Viol and Ronnie Cann (2010) Splitting the *Is* and Crossing the *You*s: Context, Speech Acts and Grammar, in Łupkowski and Purver (2010), pp. 43–50, Polish Society for Cognitive Science.
- Quine, W. V. (1948) On what there is, *Review of Metaphysics*, Vol. ??, No. ??, . reprinted in *From a Logical Point of View* (Cambridge, Massachusetts; Harvard University Press: 1953).
- Ranta, Aarne (1994) *Type-Theoretical Grammar*, Clarendon Press, Oxford.
- Recanati, François (2010) *Truth-Conditional Pragmatics*, Clarendon Press Oxford.
- Rescher, Nicholas (1999) How Many Possible Worlds Are There?, *Philosophy and Phenomenological Research*, Vol. 59, No. 2, pp. 403–420.
- Rey, Georges (2015) The Analytic/Synthetic Distinction, in E. N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy*, fall 2015 edition, <http://plato.stanford.edu/archives/fall2015/entries/analytic-synthetic/>.
- Ribas-Fernandes, José J.F., Alec Solway, Carlos Diuk, Joseph T. McGuire, Andrew G. Barto, Yael Niv and Matthew M. Botvinick (2011) A Neural Signature of Hierarchical Reinforcement Learning, *Neuron*, Vol. 71, No. 2, pp. 370 – 379.
- Romero, Maribel (2008) The Temperature Paradox and Temporal Interpretation, *Linguistic Inquiry*, Vol. 39, No. 4, pp. 655–667.
- Romoli, Jacopo and Yasutada Sudo (2009) De Re / De Dicto Ambiguity and Presupposition Projection, in A. Riester and T. Solstad (eds.), *Proceedings of Sinn und Bedeutung 13*, pp. 425–438. <http://semanticsarchive.net/Archive/Dhh0TI2Z/sub13proc.pdf>.
- Ruppenhofer, Josef, Michael Ellsworth, Miriam R.L. Petruck, Christopher R. Johnson and Jan Scheffczyk (2006) FrameNet II: Extended Theory and Practice. Available from the FrameNet website.
- Russell, Bertrand (1903) *Principles of Mathematics*, Cambridge University Press.
- Sacks, H., E.A. Schegloff and G. Jefferson (1974) A simplest systematics for the organization of turn-taking for conversation, *Language*, Vol. 50, pp. 696–735.
- Sag, Ivan A., Thomas Wasow and Emily M. Bender (2003) *Syntactic Theory: A Formal Introduction*, 2nd edition, CSLI Publications, Stanford.
- de Saussure, Ferdinand (1916) *Cours de linguistique générale*, Payot, Lausanne and Paris, edited by Charles Bally and Albert Séchehaye.

- Schlenker, Philippe (2011) Indexicality and *De Se* Reports, in C. Maienborn, K. v. Heusinger and P. Portner (eds.), *Semantics: an international handbook of natural language meaning*, pp. 1561–1604, de Gruyter.
- Schubert, Lenhart K. (2000) The Situations We Talk about, in J. Minker (ed.), *Logic-Based Artificial Intelligence*, pp. 407–439, Kluwer Academic Publishers, Dordrecht.
- Schwager, Magdalena (2009) Speaking of qualities, in E. Cormany, S. Ito and D. Lutz (eds.), *Proceedings of SALT 19*.
- Searle, John R. (1969) *Speech Acts: an Essay in the Philosophy of Language*, Cambridge University Press.
- Seligman, Jerry and Larry Moss (1997) Situation Theory, in J. van Benthem and A. ter Meulen (eds.), *Handbook of Logic and Language*, North Holland and MIT Press.
- Shanahan, Murray (2009) The Frame Problem, in E. N. Zalta (ed.), *The Stanford Encyclopedia of Philosophy*, winter 2009 edition, <http://plato.stanford.edu/archives/win2009/entries/frame-problem/>.
- Shieber, Stuart (1986) *An Introduction to Unification-Based Approaches to Grammar*, CSLI Publications, Stanford.
- Steedman, Mark (2012) *Taking Scope: The Natural Semantics of Quantifiers*, The MIT Press.
- Suppes, Patrick (1960) *Axiomatic Set Theory*, The University Series in Undergraduate Mathematics, D. van Nostrand Company, Inc.
- Thomason, Richmond H. (1979) Home is where the heart is, in P. A. French, Uehling, Jr., Theodore E. and H. K. Wettstein (eds.), *Contemporary perspectives in the philosophy of language*, pp. 209–219, University of Minnesota Press, Minneapolis.
- Thomason, Richmond H. (1980) A model theory for propositional attitudes, *Linguistics and Philosophy*, Vol. 4, pp. 47–70.
- Tonhauser, Judith (2007) Nominal Tense? The Meaning of Guaraní Nominal Temporal Markers, *Language*, Vol. 83, No. 4, pp. 831–869.
- Traum, David R. (1994) *A Computational Theory of Grounding in Natural Language Conversation*, PhD dissertation, University of Rochester, Department of Computer Science.
- Turner, Raymond (2005) Semantics and Stratification, *Journal of Logic and Computation*, Vol. 15, No. 2, pp. 145–158.
- Wadler, Philip (2015) Propositions as Types, *Communications of the ACM*, Vol. 58, No. 12, pp. 75–84.

- Walker, Marilyn A., Aravind K. Joshi and Ellen F. Prince, eds. (1998) *Centering Theory in Discourse*, Oxford University Press, Oxford.
- Zweig, Eytan (2008) *Dependent Plurals and Plural Meaning*, PhD dissertation, New York University.
- Zweig, Eytan (2009) Number-neutral bare plurals and the multiplicity implicature, *Linguistics and Philosophy*, Vol. 32, pp. 353–407.